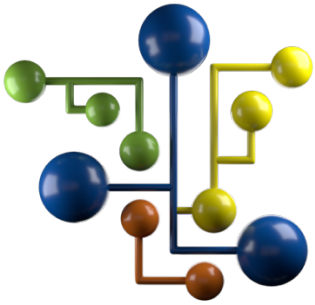

Data Science Academy

www.datascienceacademy.com.br

Engenharia de Dados com Hadoop e Spark

Instalação e Configuração do Ecossistema Hadoop

Sumário

1. Versão

Este documento foi criado pela equipe Data Science Academy e pode ser distribuído livremente, desde que se faça menção à fonte.

Versão	Ação	Data
1.0	Criação do documento	28/07/2019

2. Configuração do Ambiente

Item	Versão
Virtual Box	6.0.8
Sistema Operacional	CentOS 7.6 (64 bits) ou CentOS 6.8 (32 bits)
Interface Gráfica	Gnome
Firefox Web Browser	60.7
Java	1.8
Apache Hadoop	3.2.0
Apache Zookeeper	3.5.5
Apache Hbase	2.2.0
Apache Hive	3.1.1
Apache Pig	0.17.0
Apache Spark	2.4.3
Apache Sqoop	1.4.7
Apache Flume	1.9.0
Apache Ambari	2.4.1

2.1. Criação da Máquina Virtual no VirtualBox

O Oracle VM Virtual Box é gratuito e pode ser baixado em https://www.virtualbox.org. Existem versões disponíveis para Windows, MAC, Linux e Solaris. Aqui utilizaremos a versão 6.0 e o tutorial será o mesmo independente do sistema operacional do seu computador. O uso do VirtualBox não é obrigatório e você pode instalar em uma máquina física se desejar.

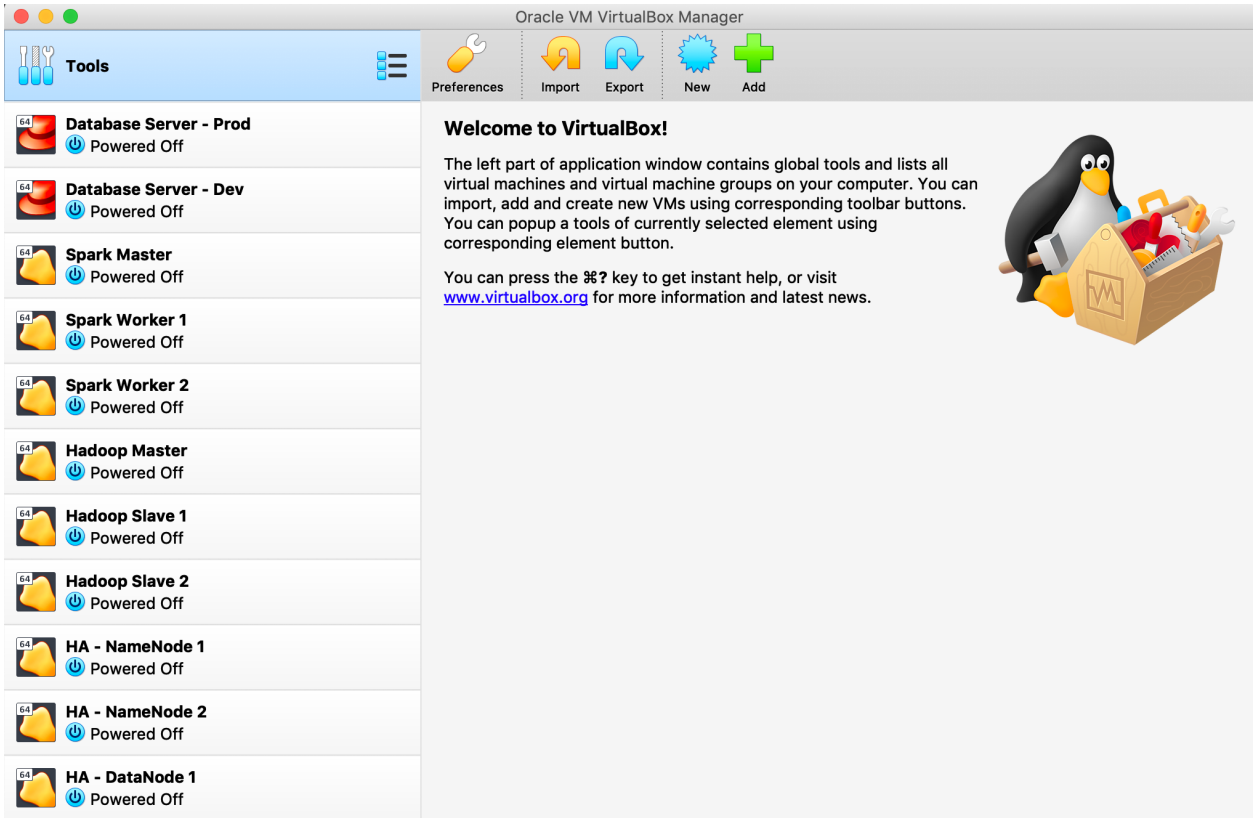

Abrindo o Gerenciador do Oracle Virtual Box

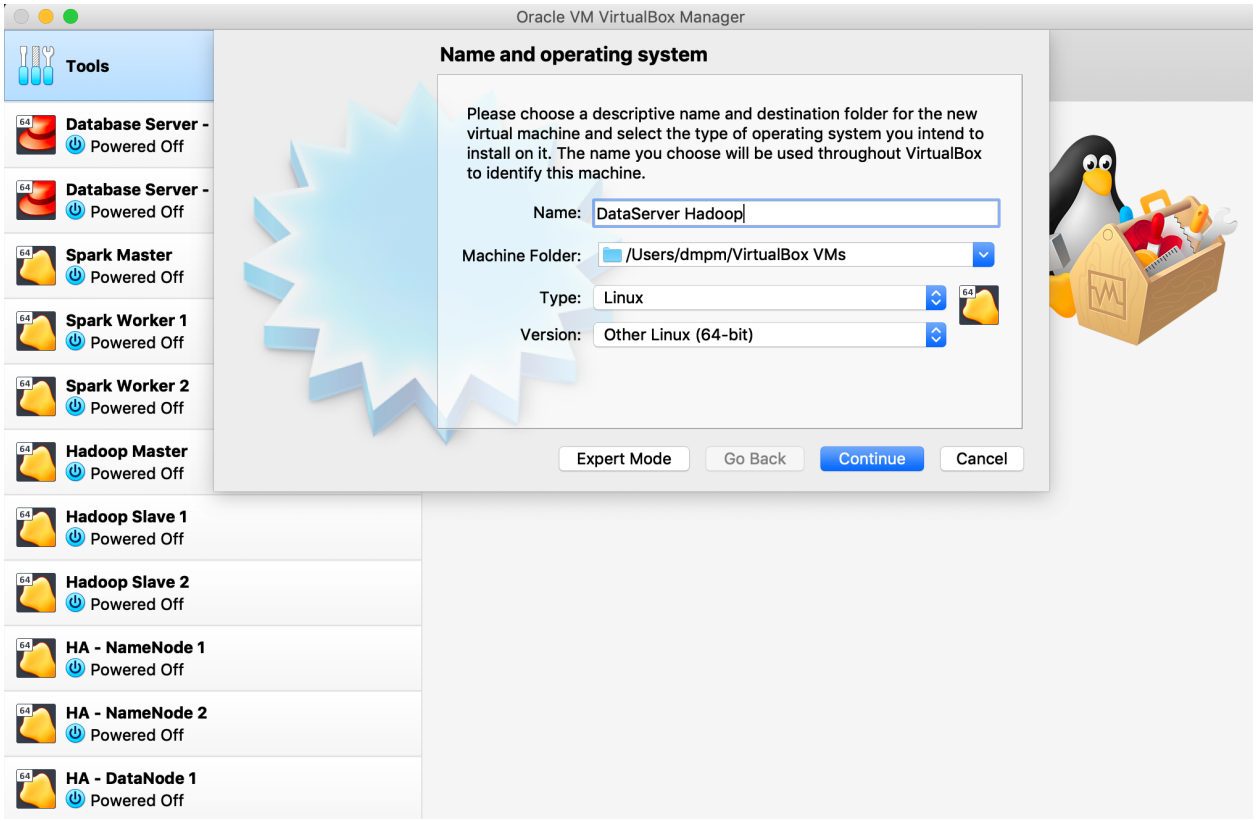

Definindo o nome da máquina virtual e a versão do sistema operacional

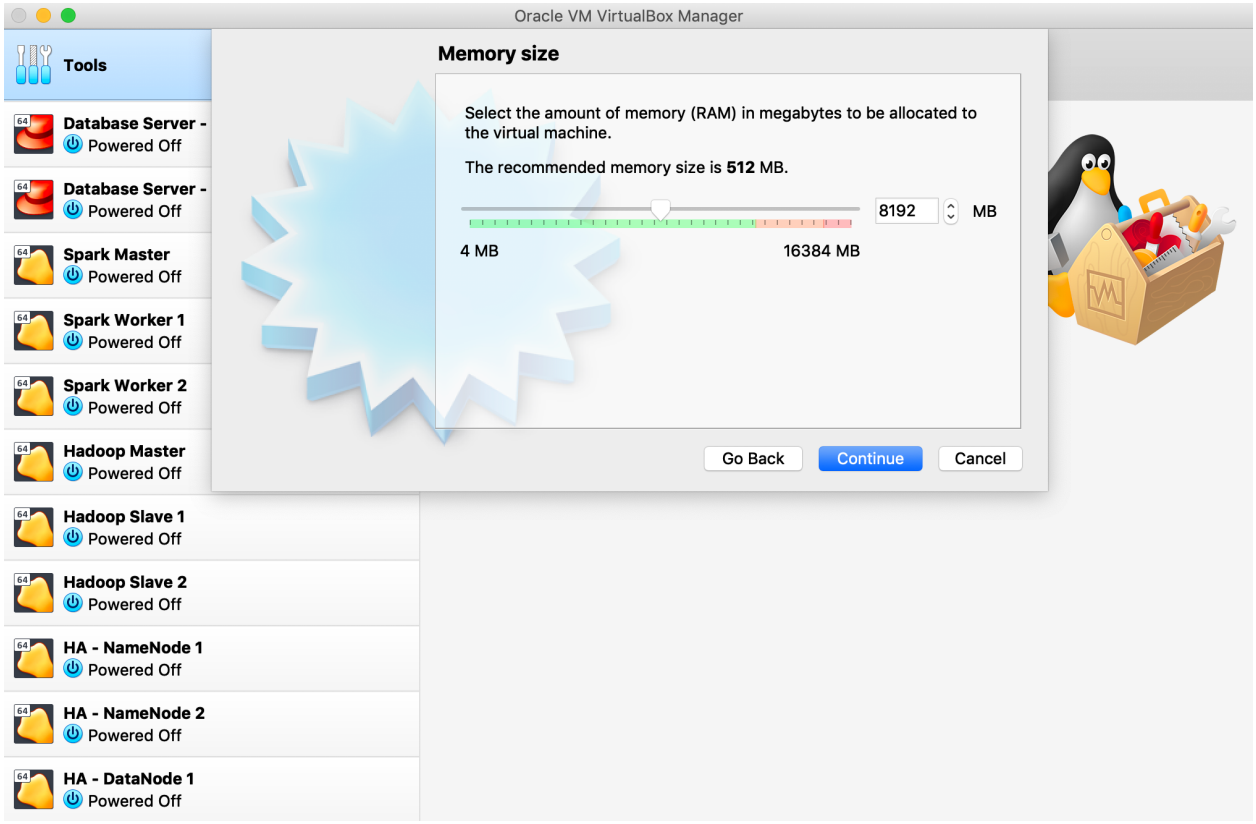

Configure metade da memória física do seu computador para a VM

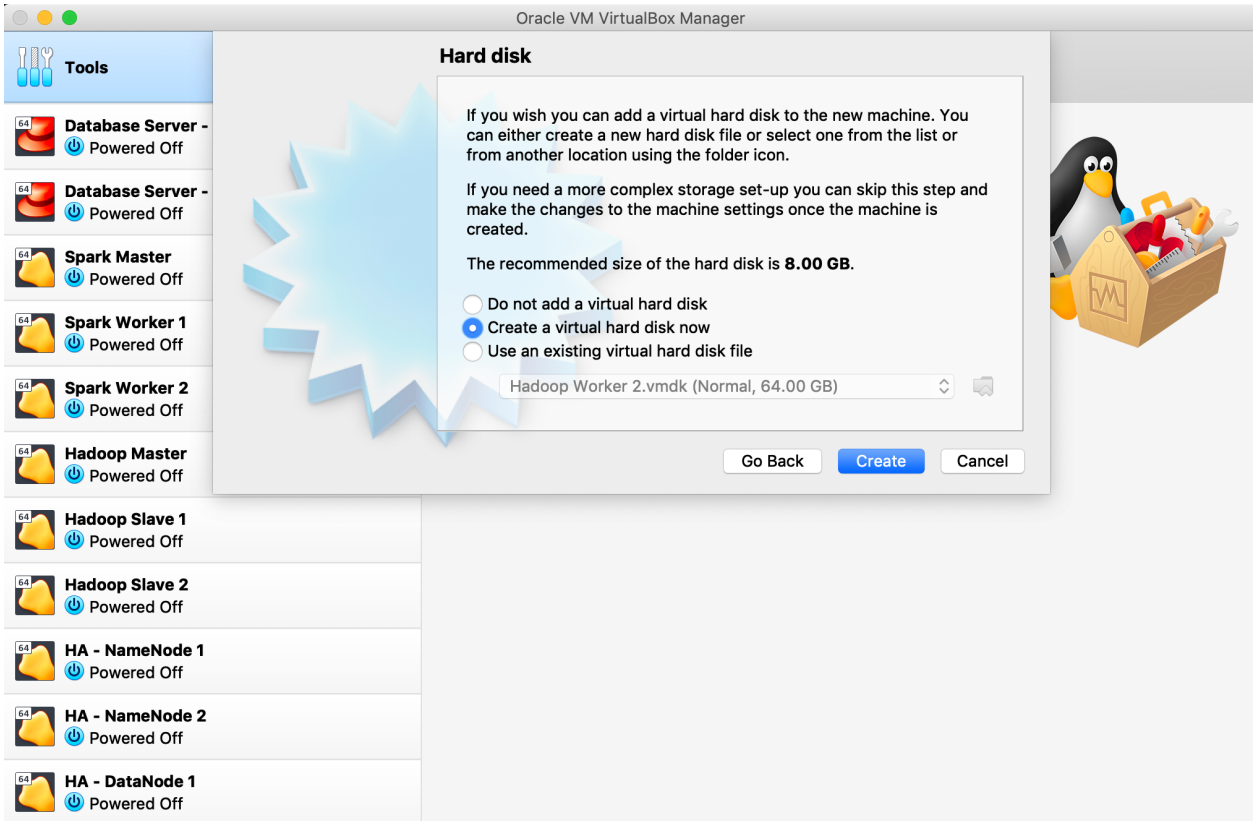

Criar um novo disco rígido virtual

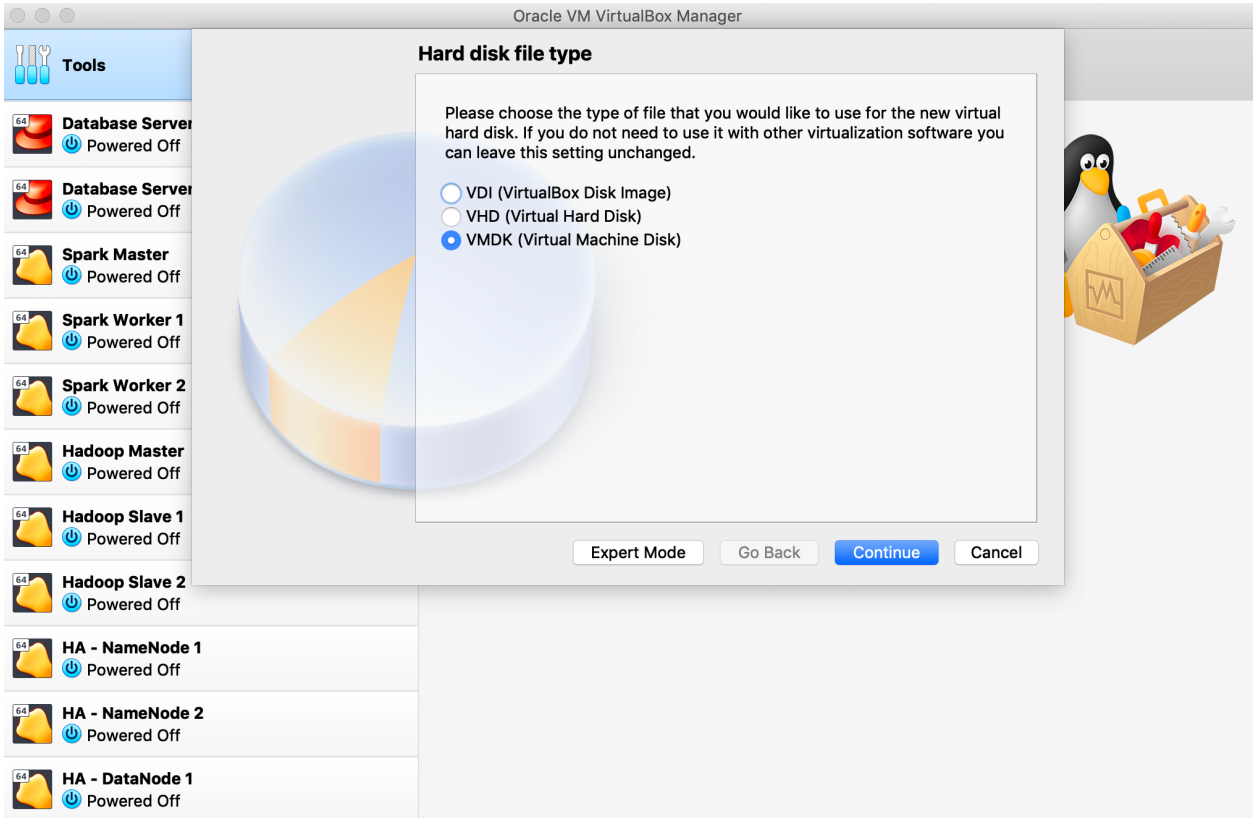

Selecione a opção VMDK

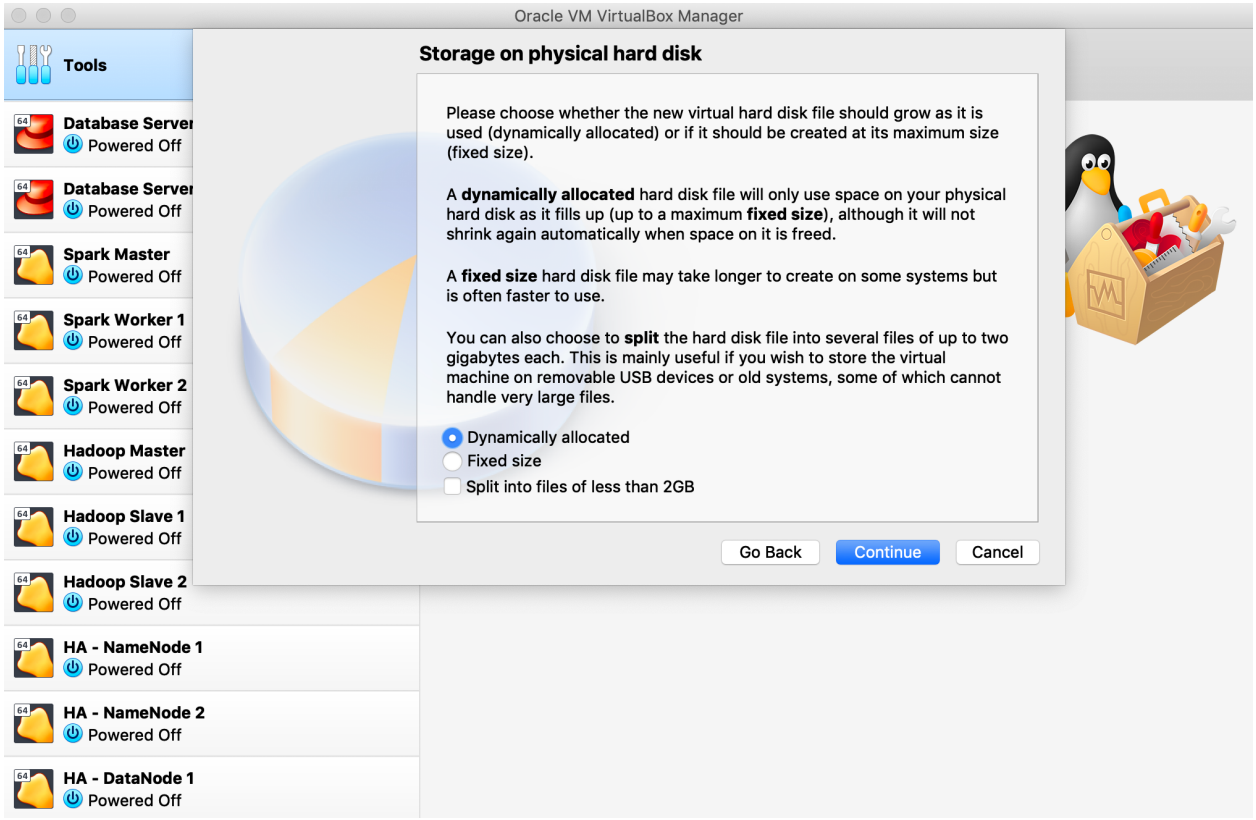

O disco deve ser alocado dinamicamente

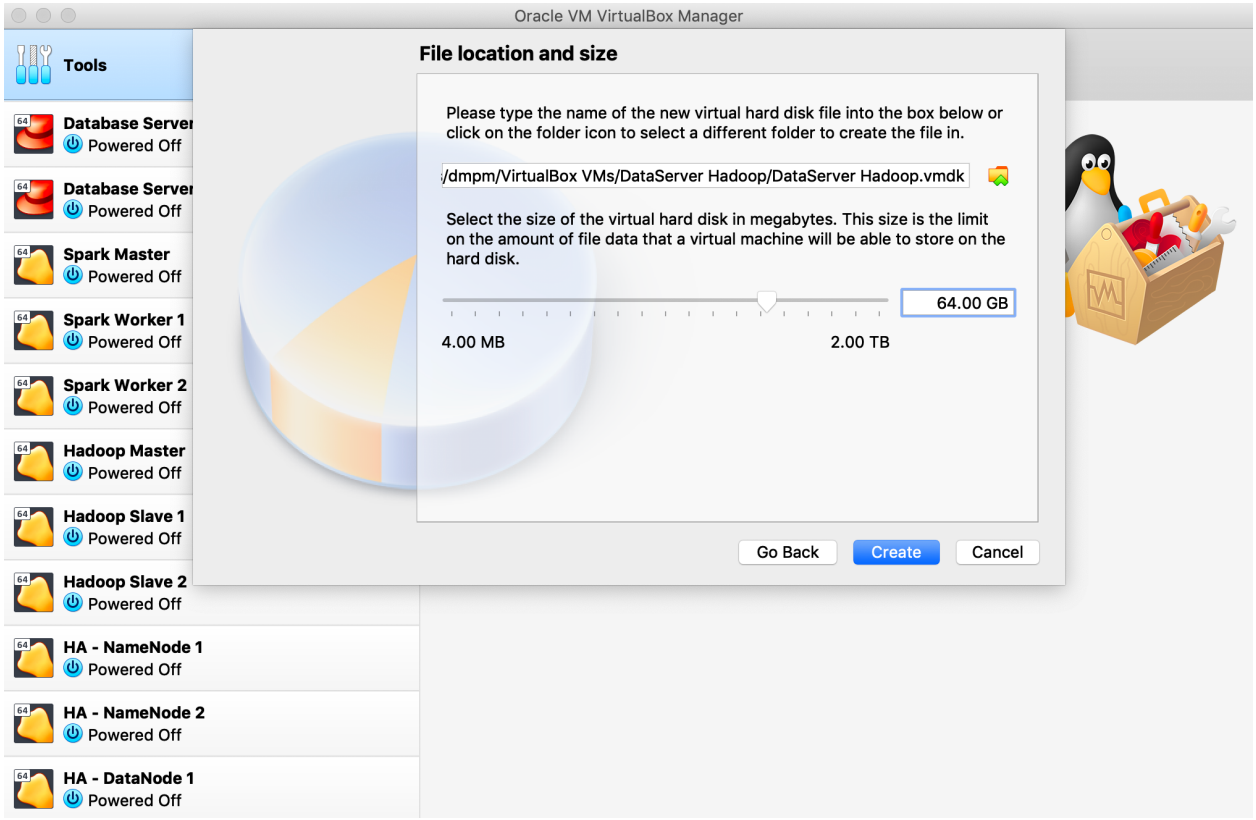

Selecione 64 GB para o disco virtual

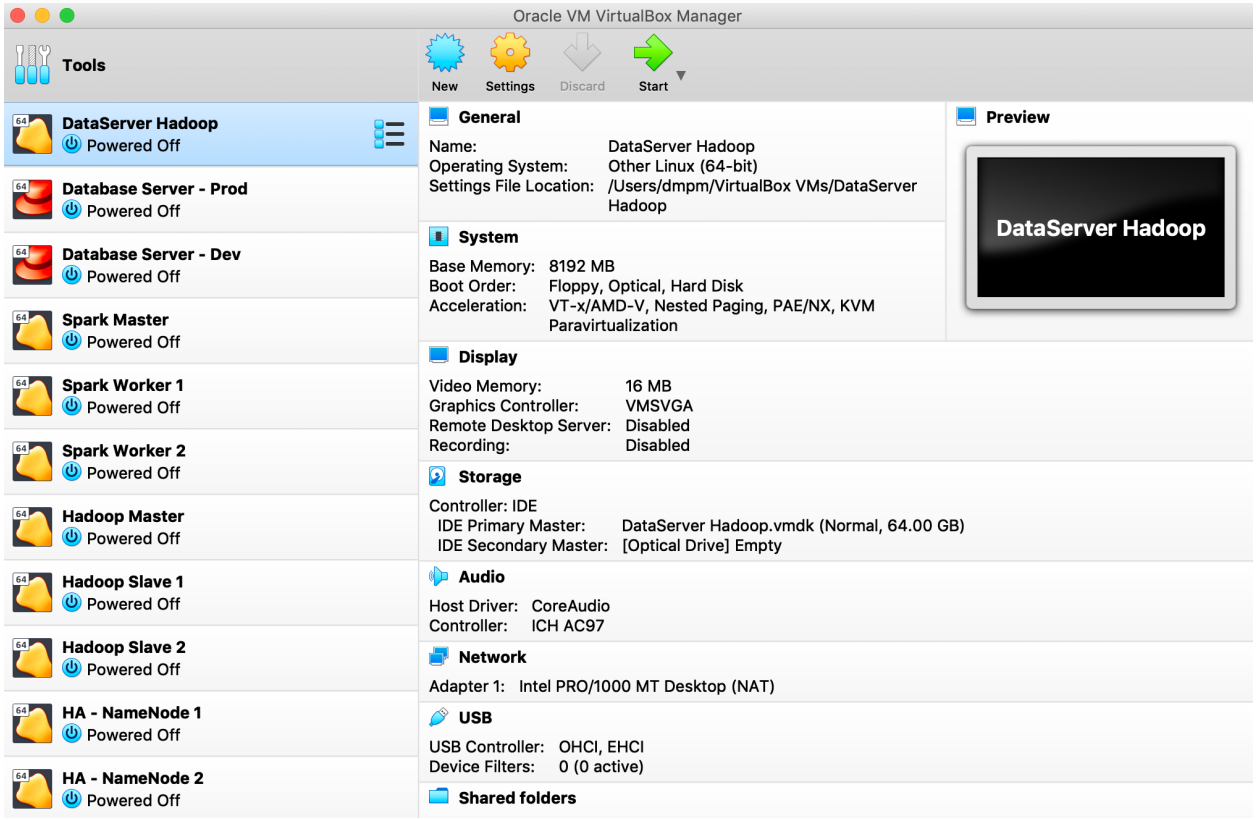

Máquina virtual criada. Selecione a VM e clique no botão Start para inicializar a VM.

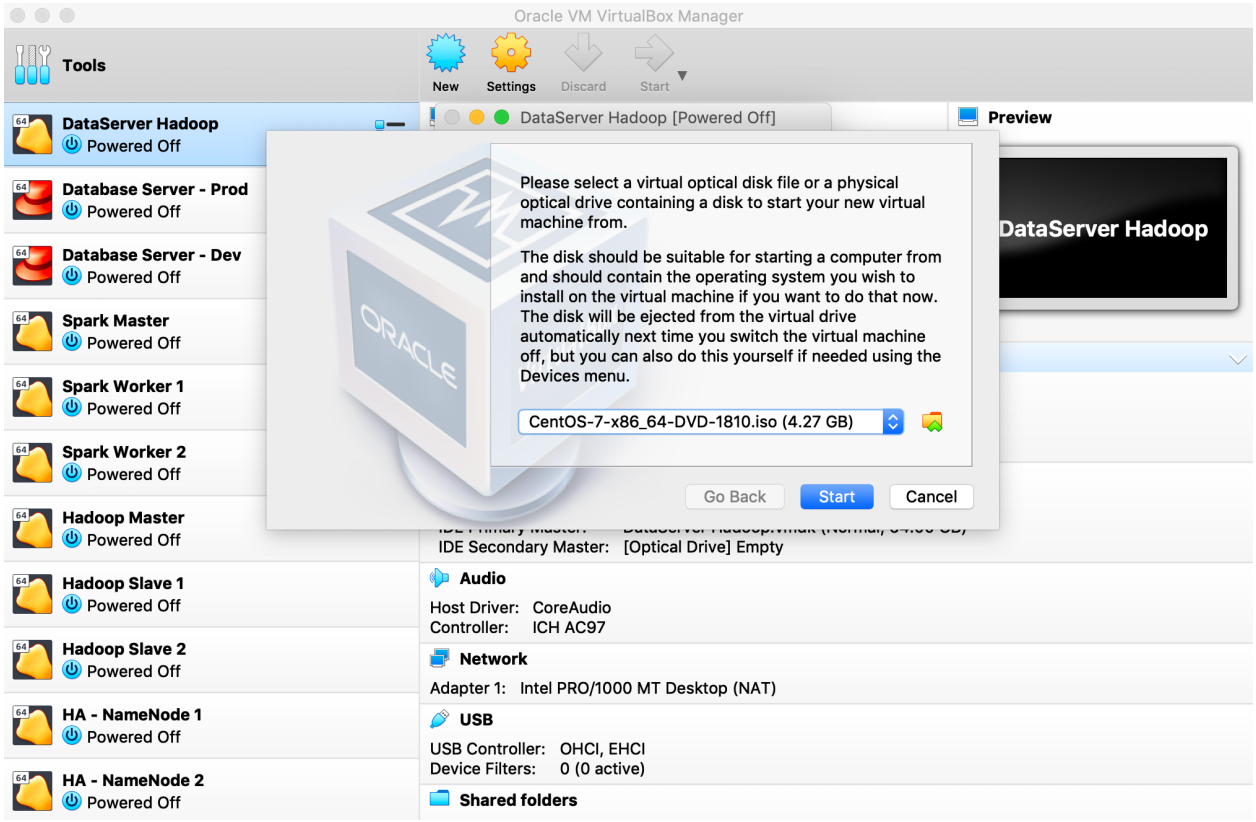

Selecione a mídia de instalação do sistema operacional

Utilizaremos o CentOS versão 7 64 bits. Caso sua máquina seja 32 bits você deve usar CentOS versão 6.8. Em ambos os casos faça o download do DVD de instalação como imagem .iso

CentOS 64 bits (versão 7):
http://isoredirect.centos.org/centos/7/isos/x86_64/

CentOS 32 bits (versão 6.8):
http://centos.mirror.netelligent.ca/centos/

2.2. Instalação do Sistema Operacional

Selecione a opção de Instalação do Sistema Operacional CentOS 7

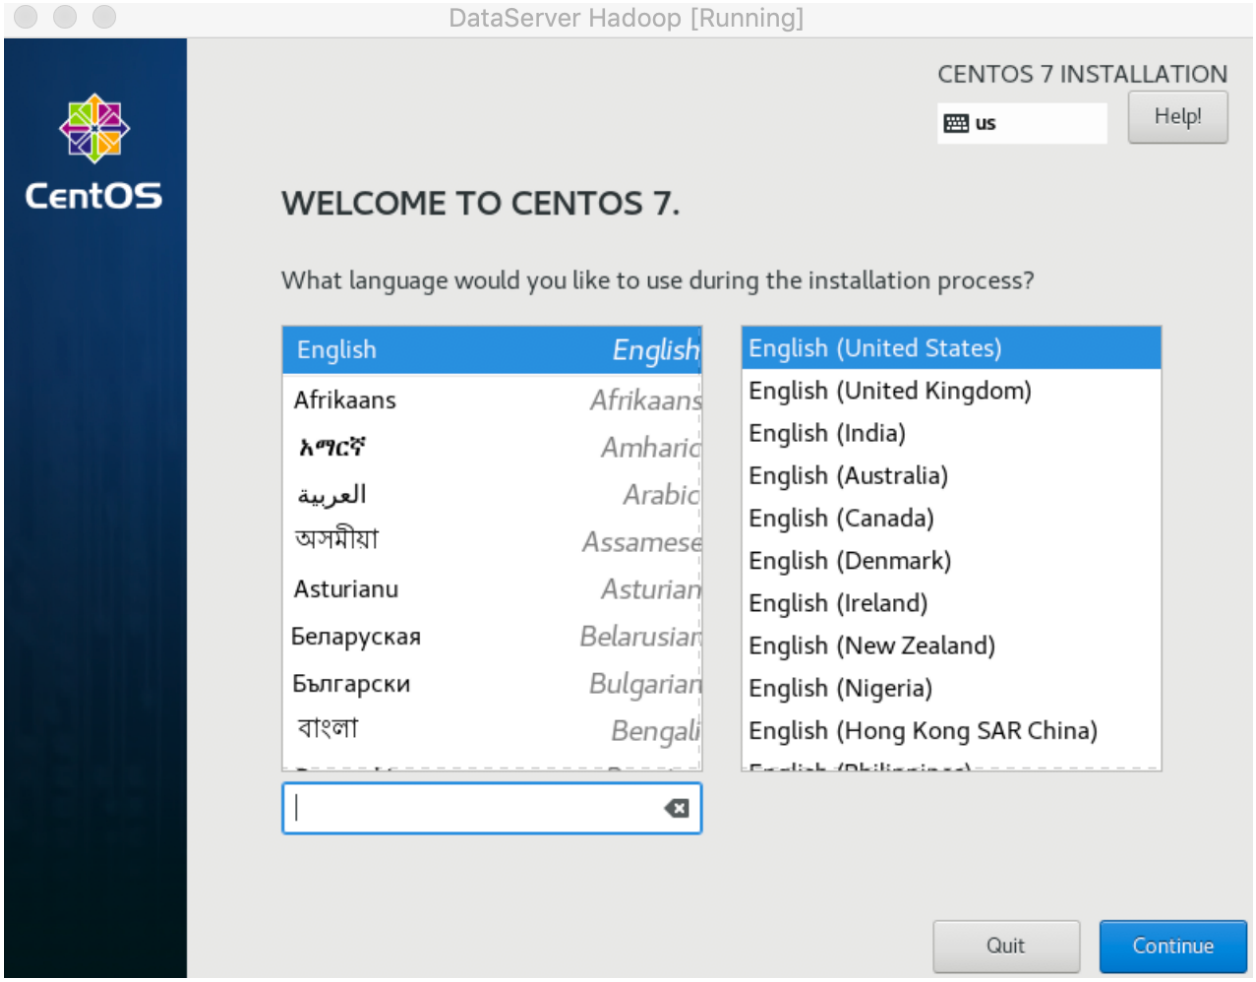

Seleção do idioma usado na instalação

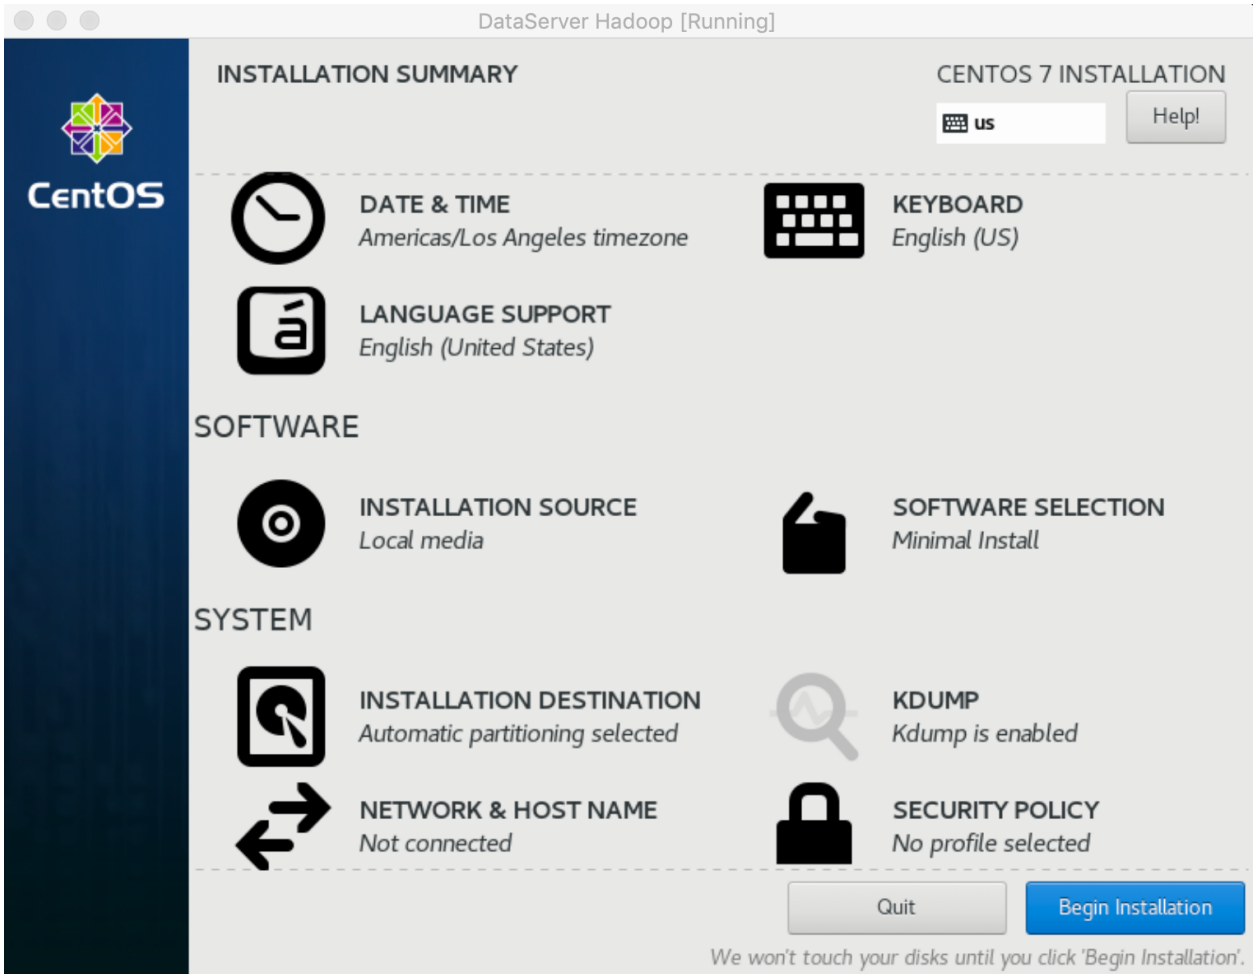

Opções de configuração


DataServer Hadoop [Running]
DATE & TIME
Done
CENTOS 7 INSTALLATION
us
Help!
Region: Americas
City: Sao Paulo
Network Time
OFF
23:08 PM
24-hour
AM/PM
06 / 21 / 2019
You need to set up networking first if you want to use NTP

Timezone

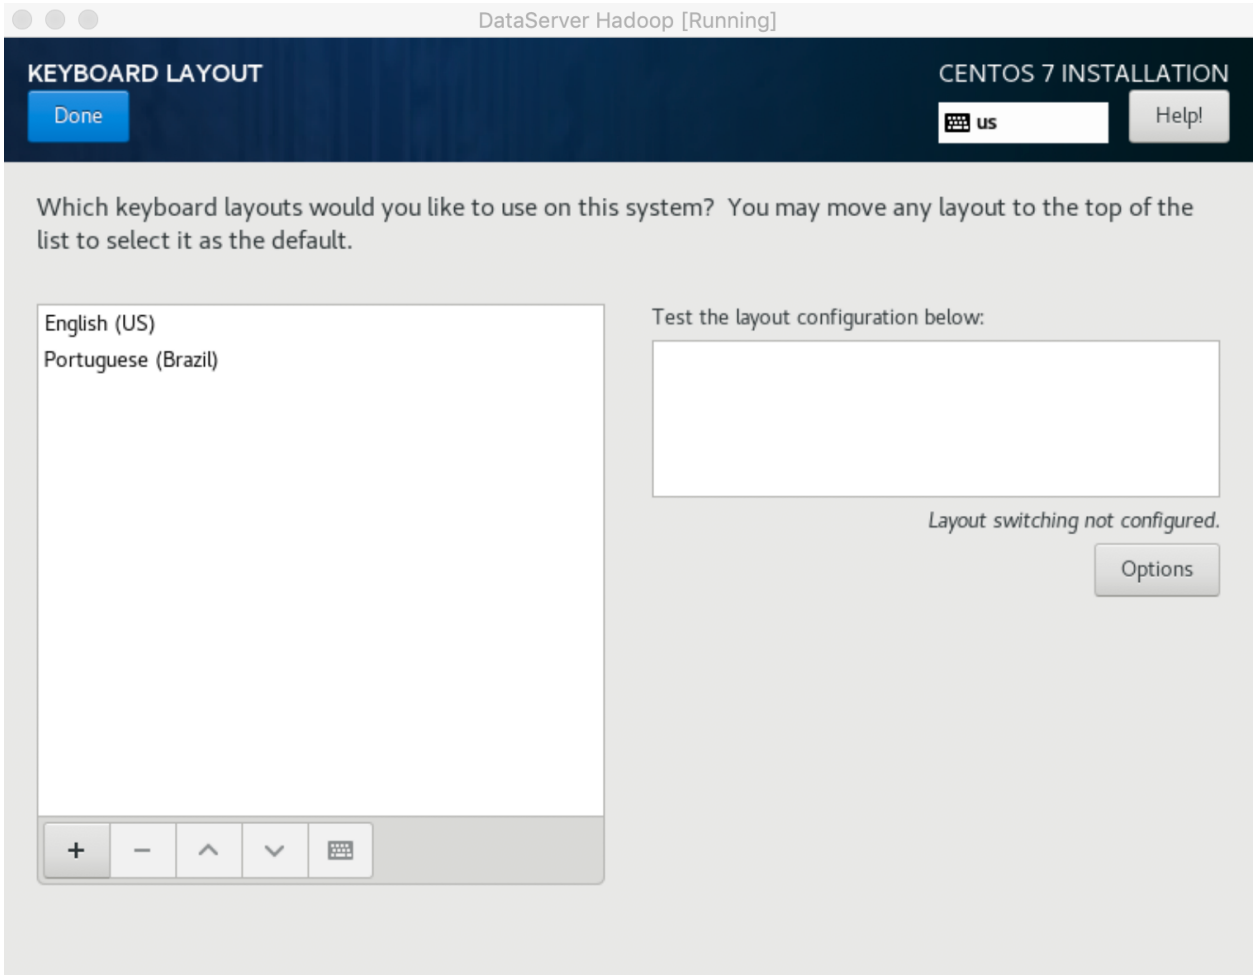

Layout do teclado

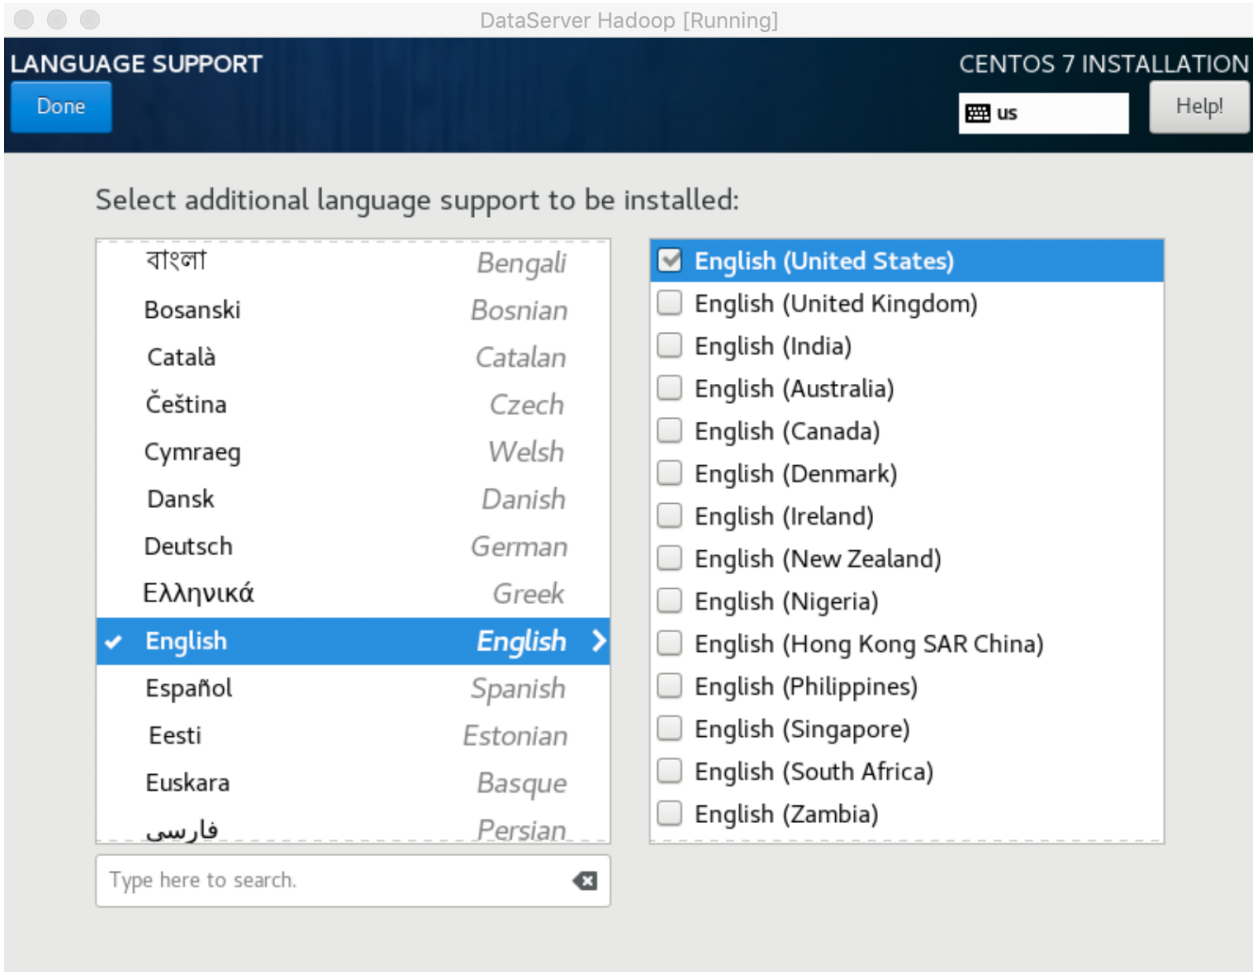

Idioma do sistema operacional

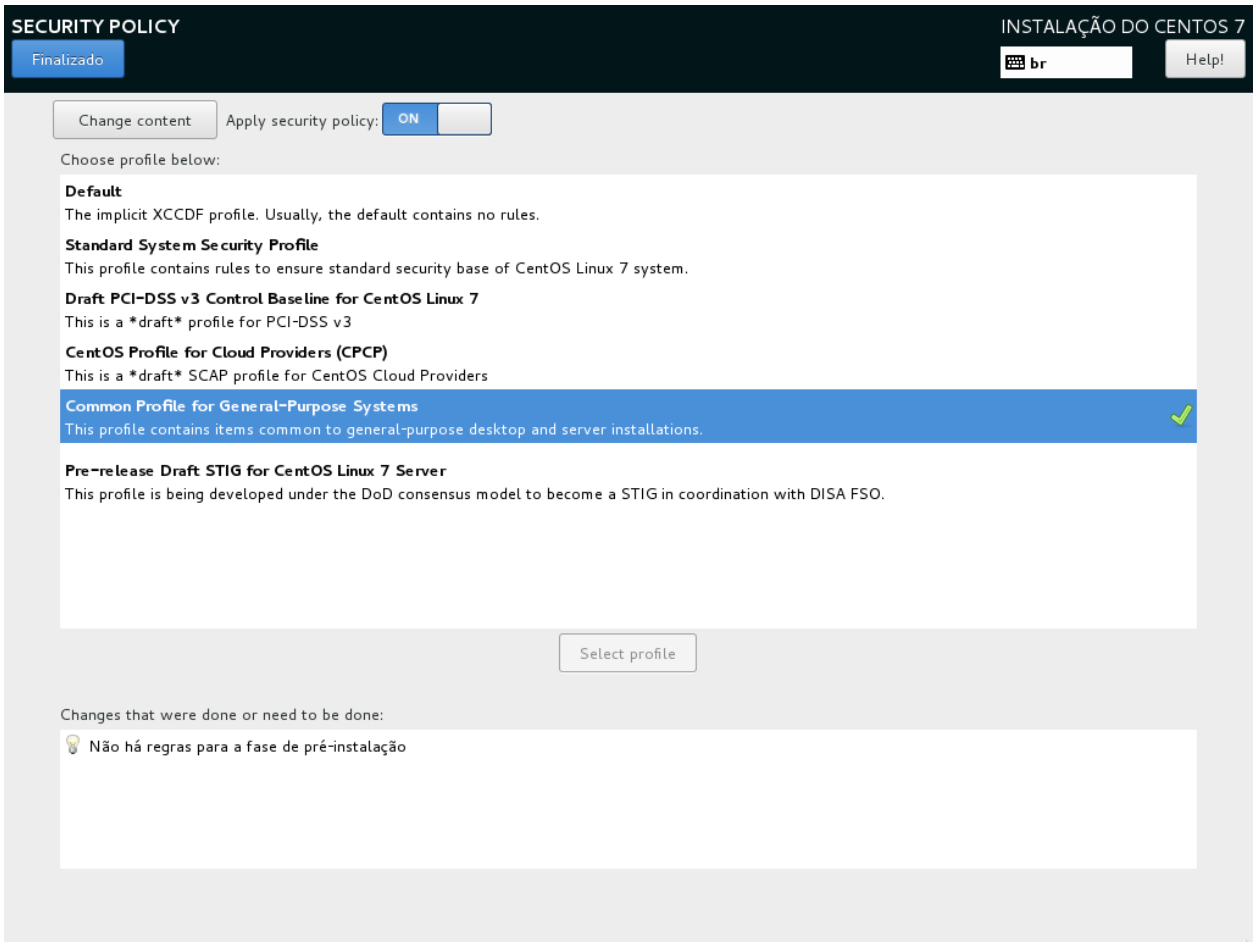

Política de segurança padrão

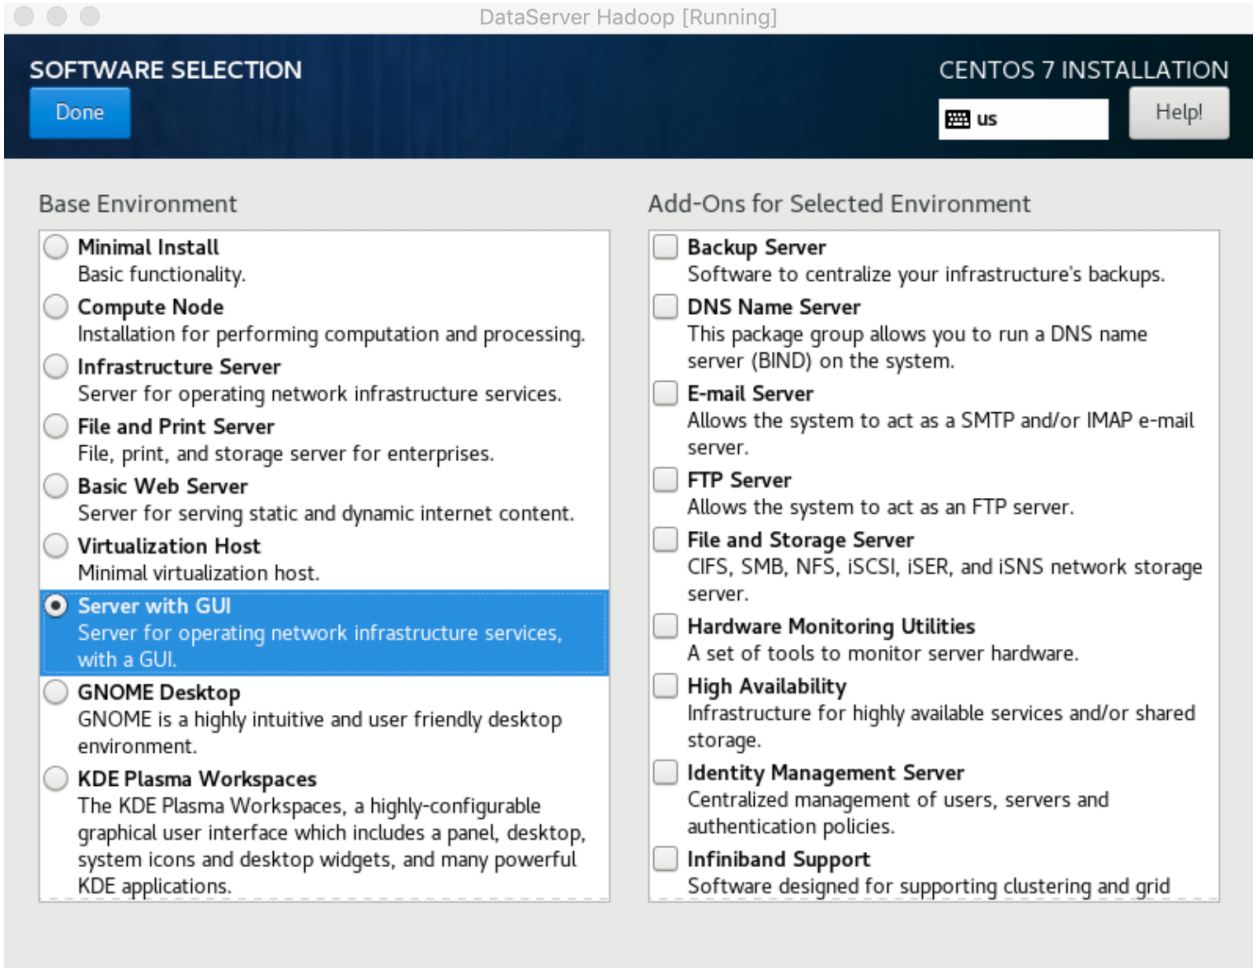

Selecione a opção Server with GUI

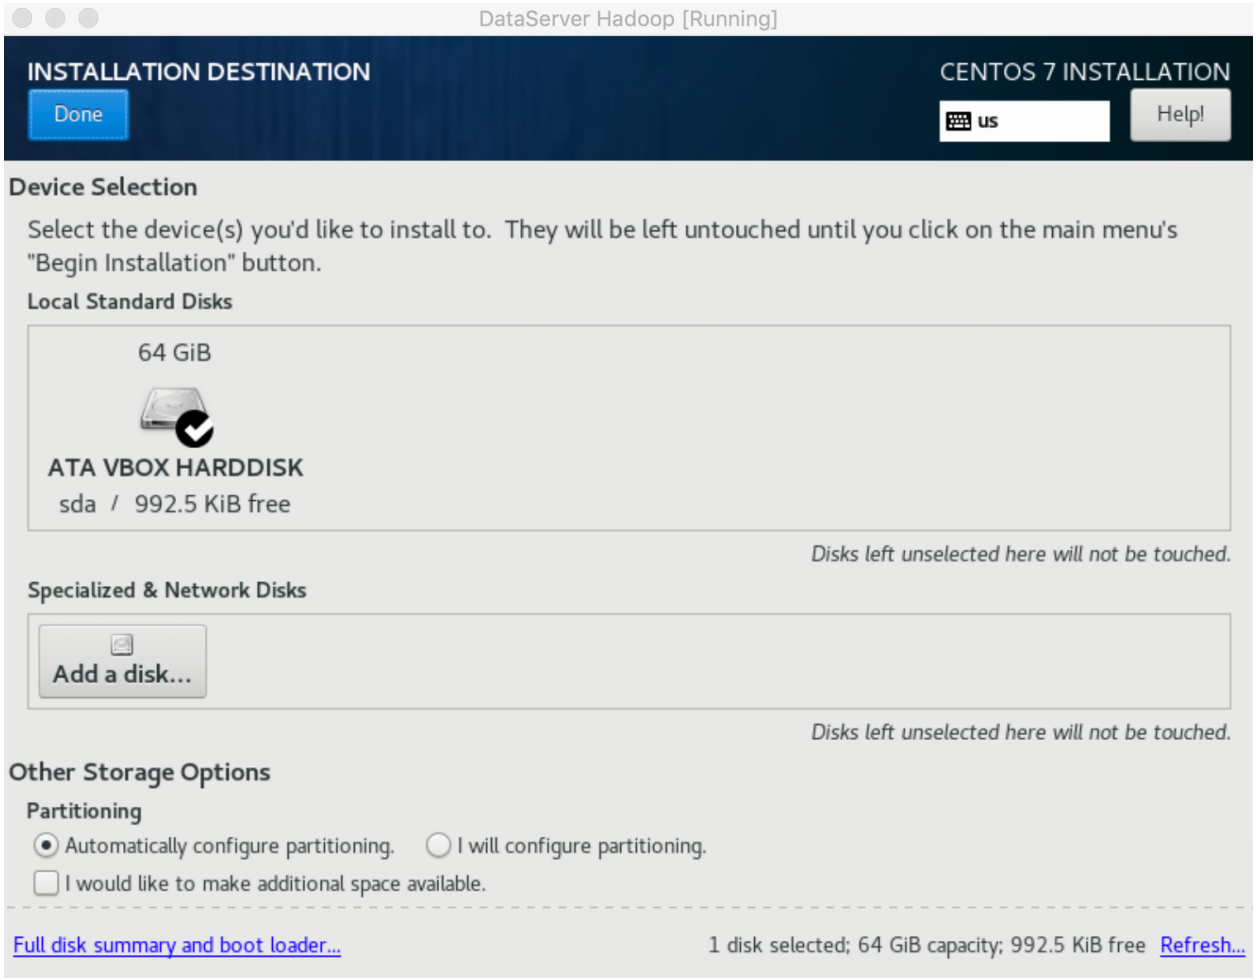
DataServer Hadoop [Running]
INSTALLATION DESTINATION
Done
CENTOS 7 INSTALLATION
us
Help!
Device Selection
Select the device(s) you'd like to install to. They will be left untouched until you click on the main menu's "Begin Installation" button.
Local Standard Disks
64 GiB
ATA VBOX HARDDISK
sda / 992.5 KiB free
Disks left unselected here will not be touched.
Specialized & Network Disks
Add a disk...
Disks left unselected here will not be touched.
Other Storage Options
Partitioning
Automatically configure partitioning.
I will configure partitioning.
I would like to make additional space available.
Full disk summary and boot loader...
1 disk selected; 64 GiB capacity; 992.5 KiB free
Refresh...

Disco

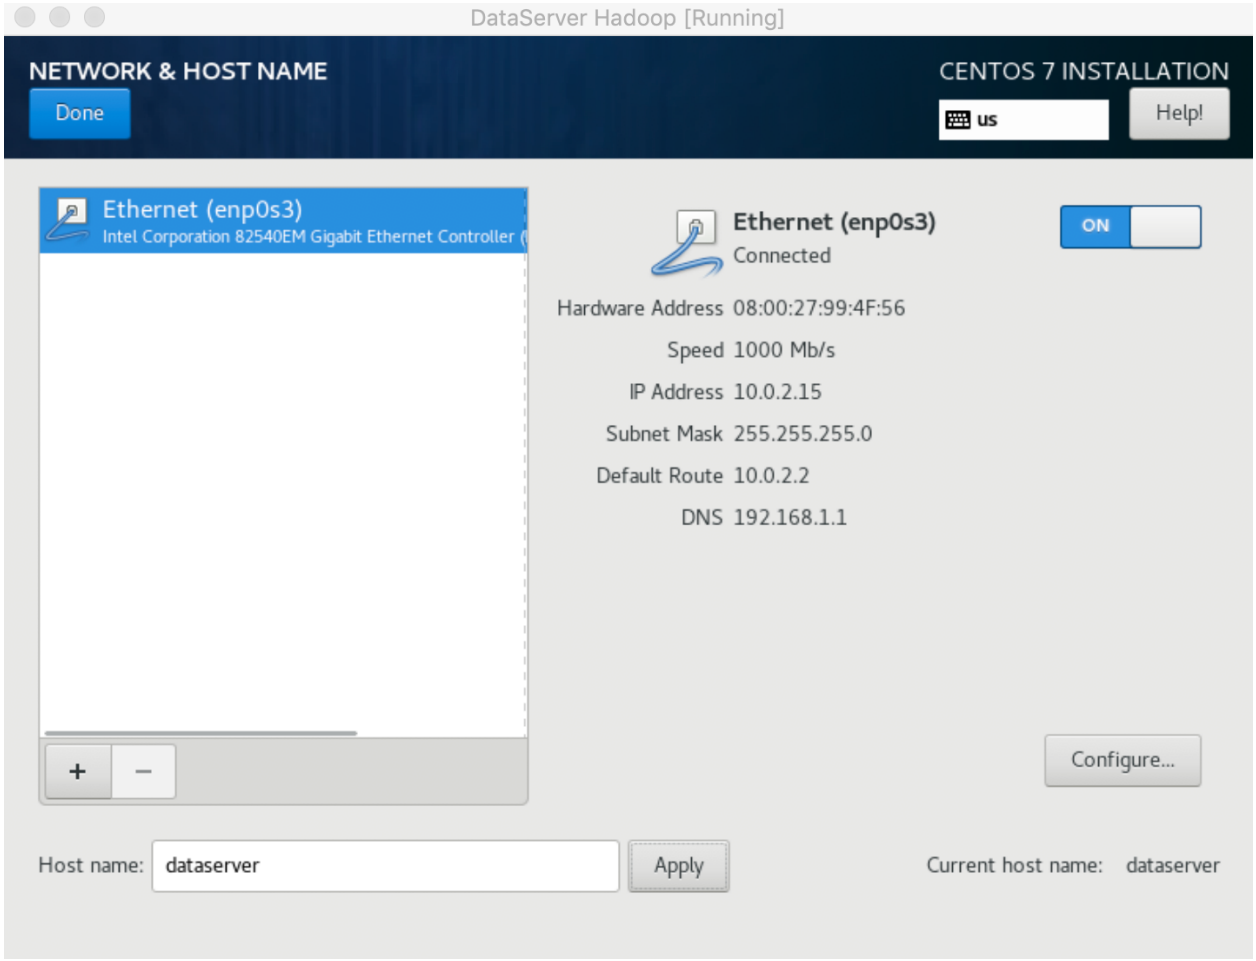

Configuração de rede e nome do servidor – **dataserver –** Clique em Apply
Certifique-se de habilitar a opção de ativar a Ethernet (botão on)

Definir senha do root e criar um novo usuário

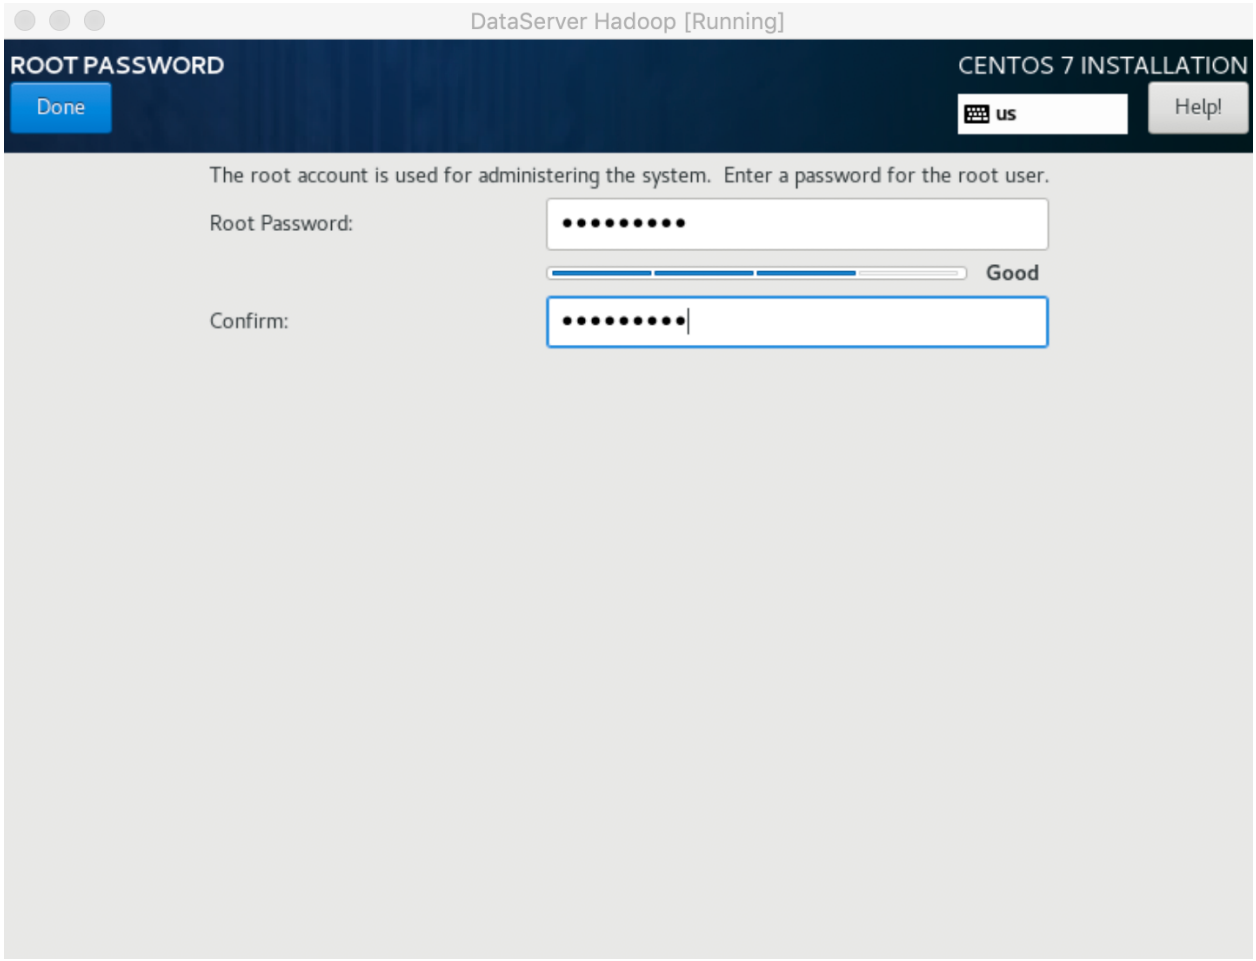

Definir senha do root – usuário administrador
Senha: **dsahadoop**

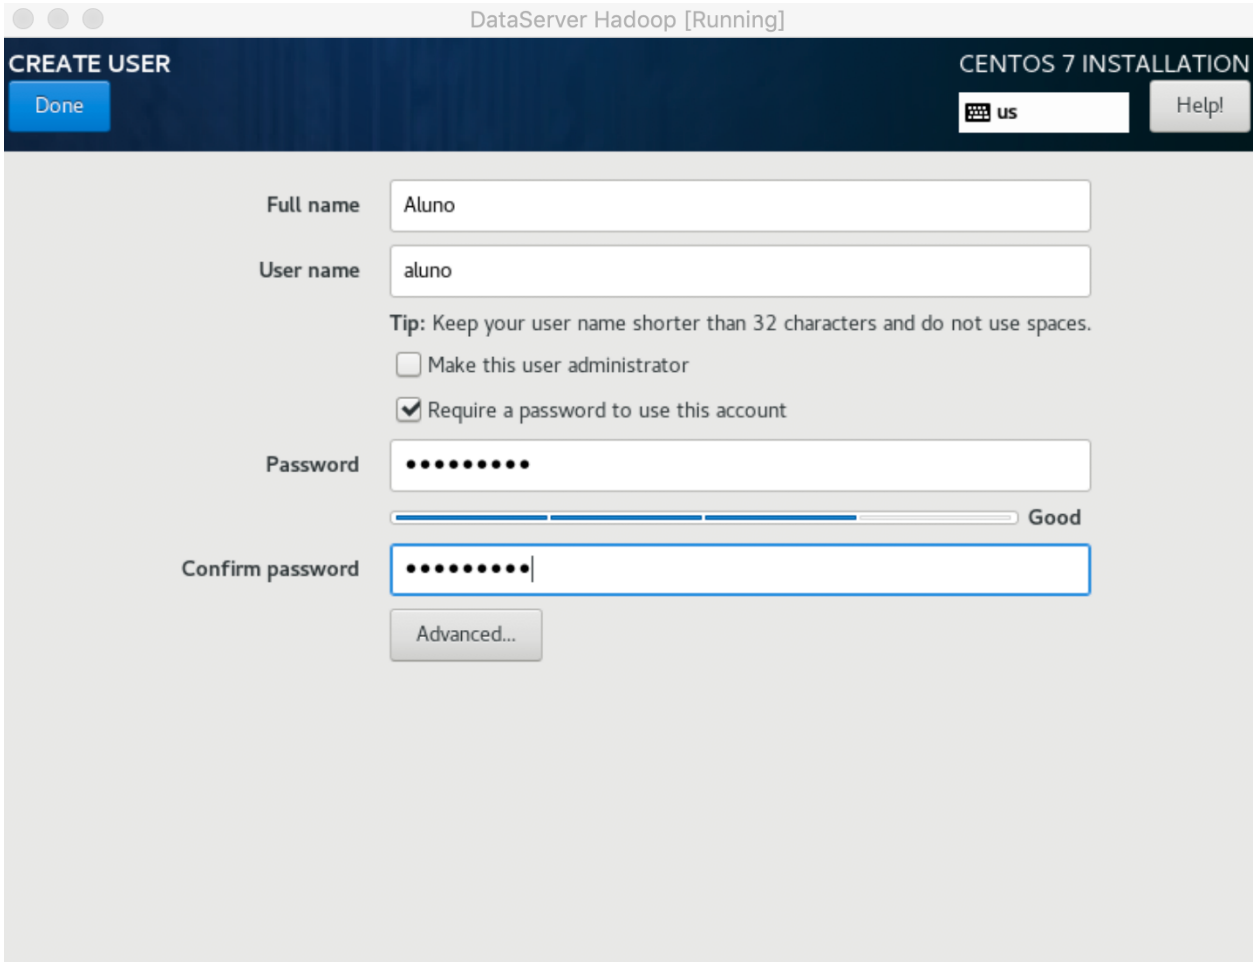

Criação de um usuário – Aluno
(username: aluno, senha: **dsahadoop**)

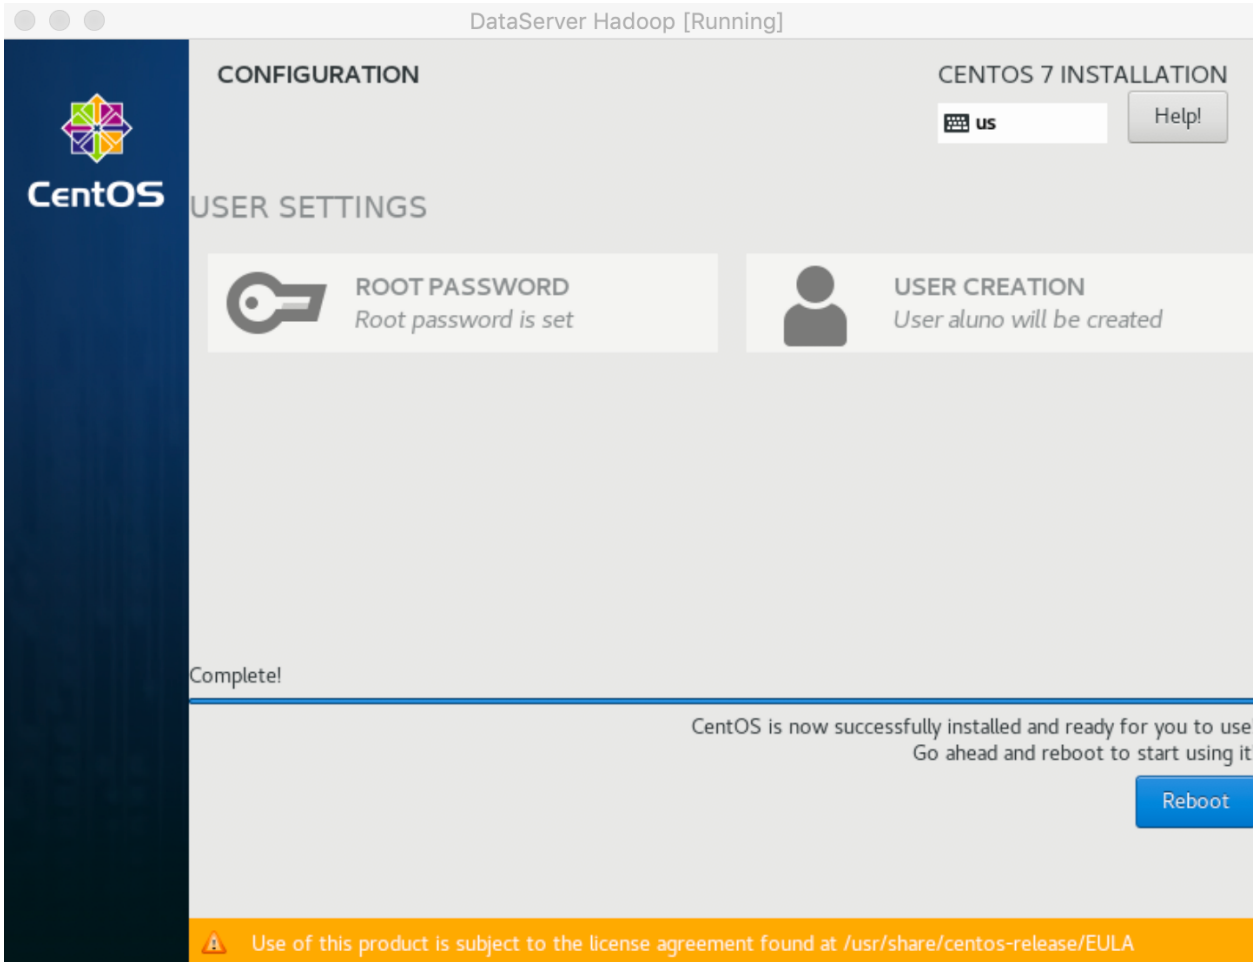

Conclusão da instalação. Clique no botão Reboot.

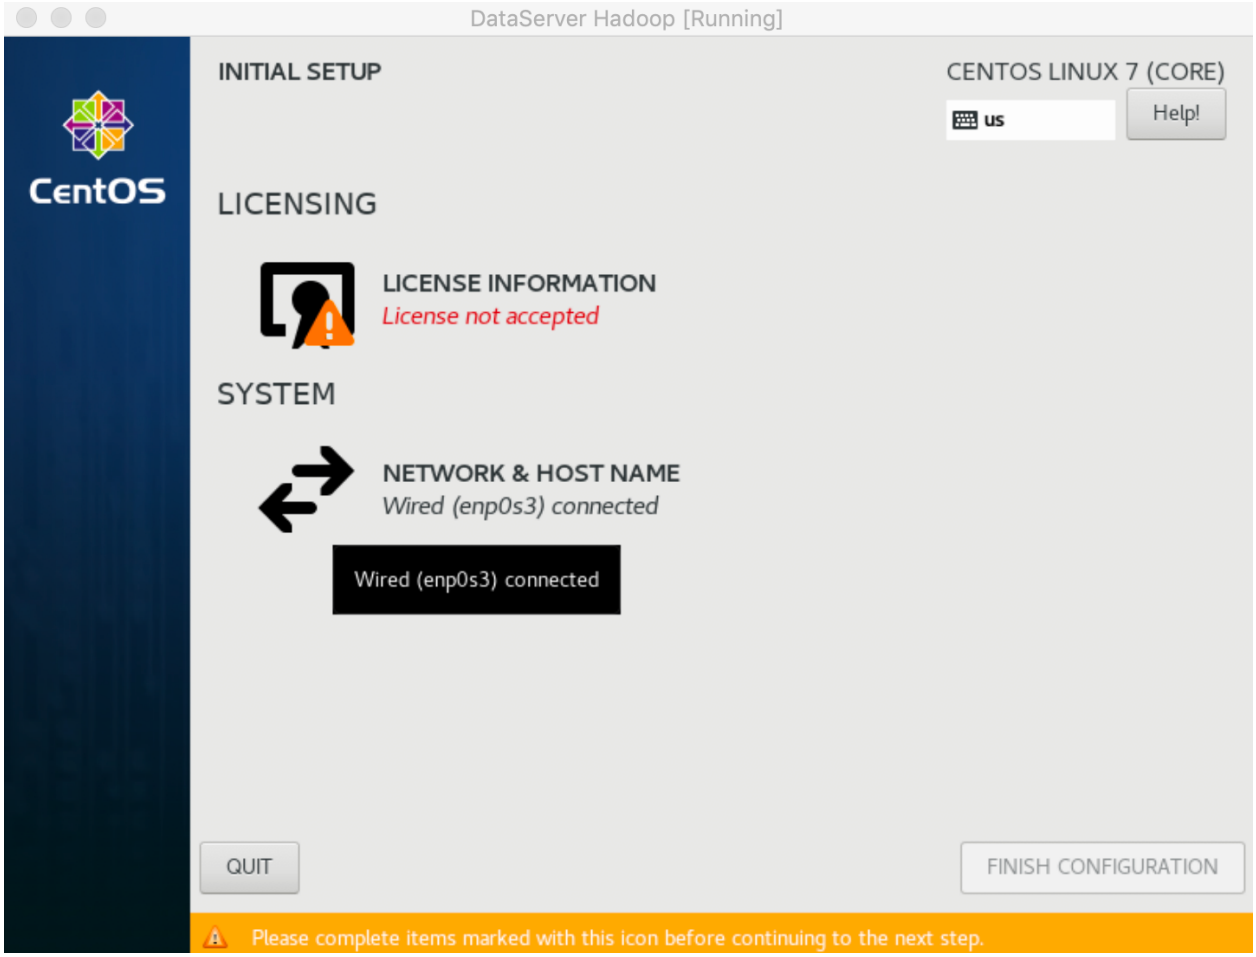

Clique em License Information para aceitar os termos de uso

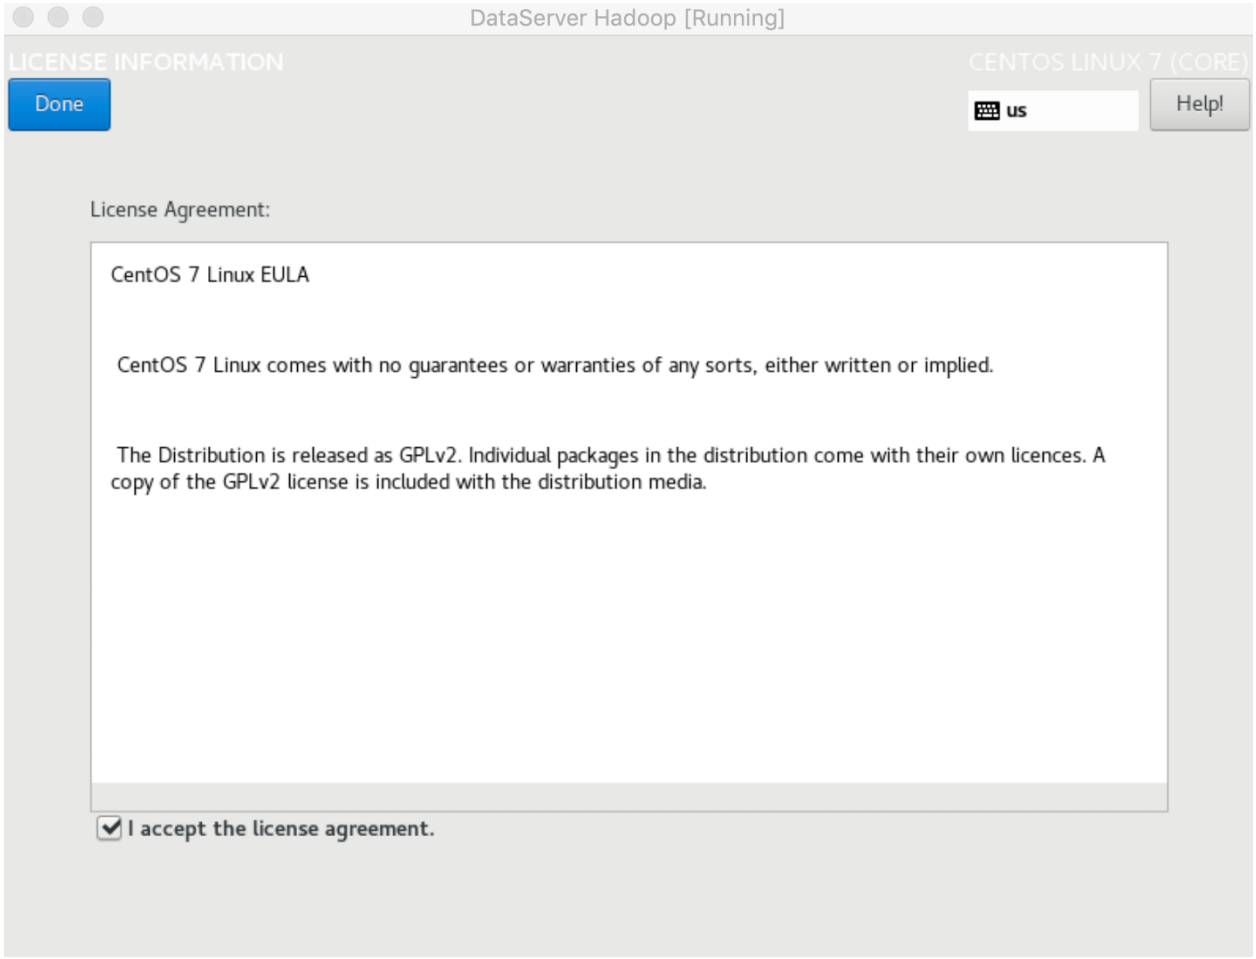

Marque a caixa e pressione Done.

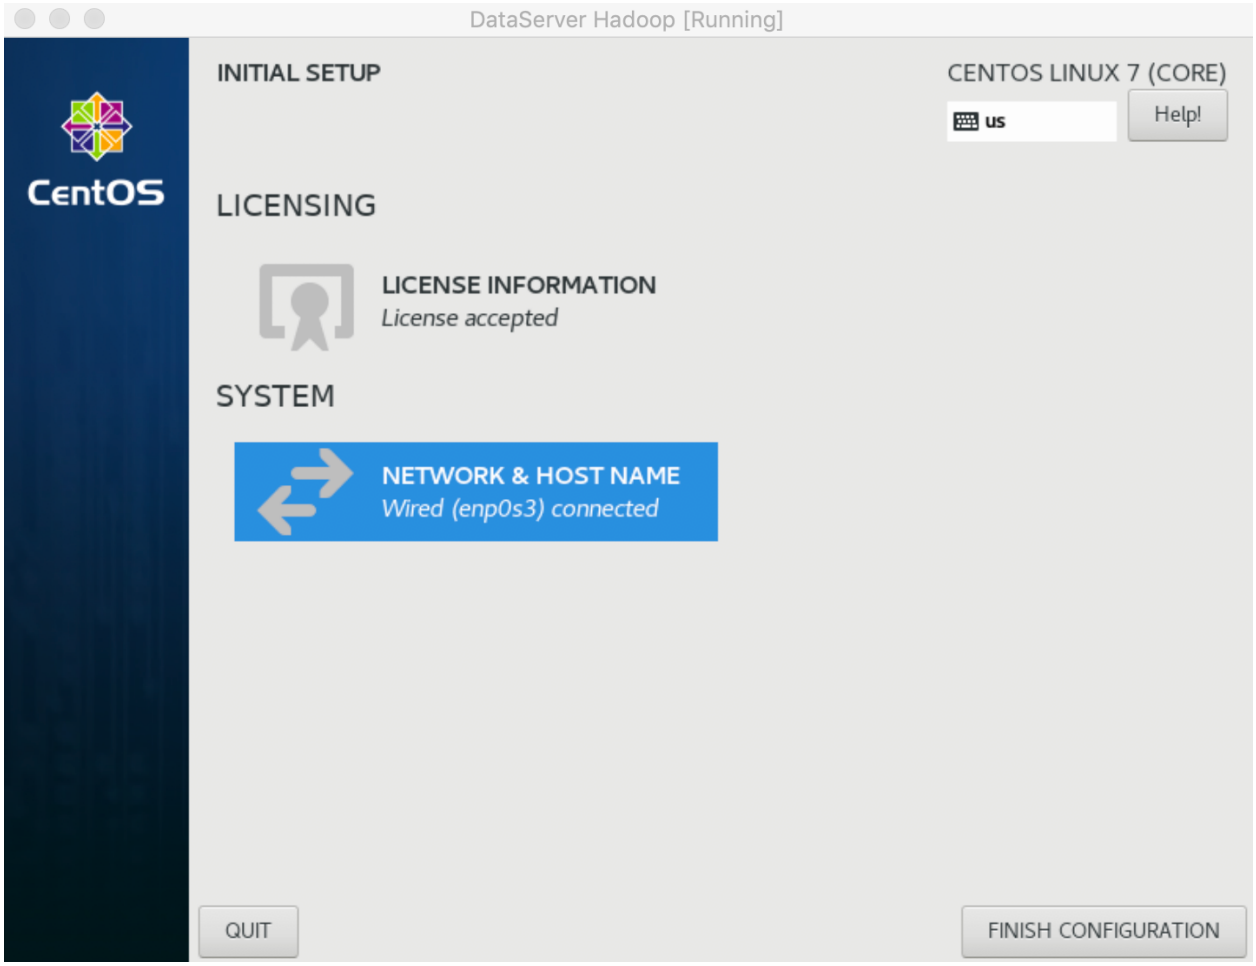

Se necessário, revise se o nome da máquina está correto e se a rede está ativada e pressione "Finish Configuration"

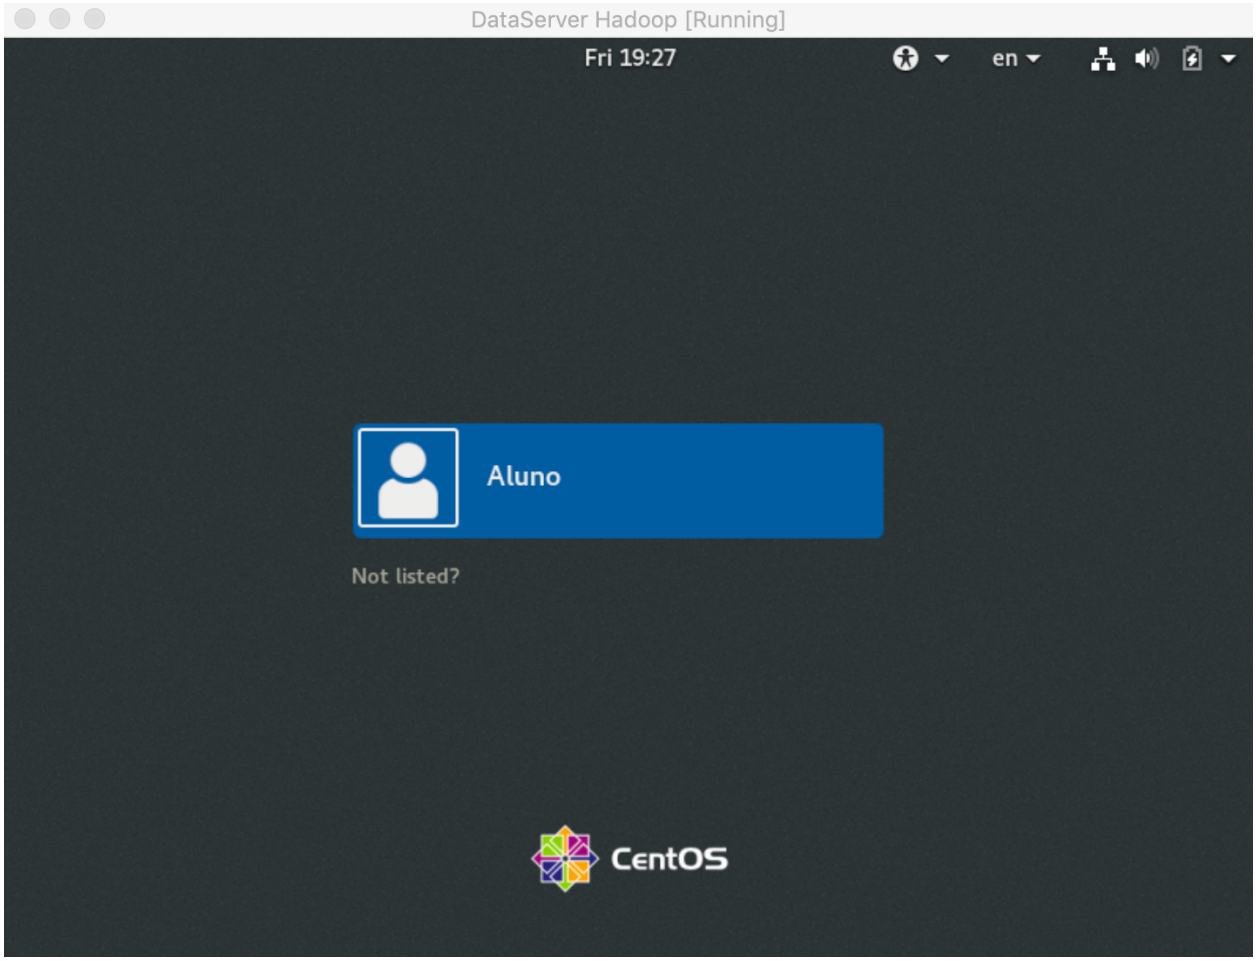

Clique no usuário Aluno

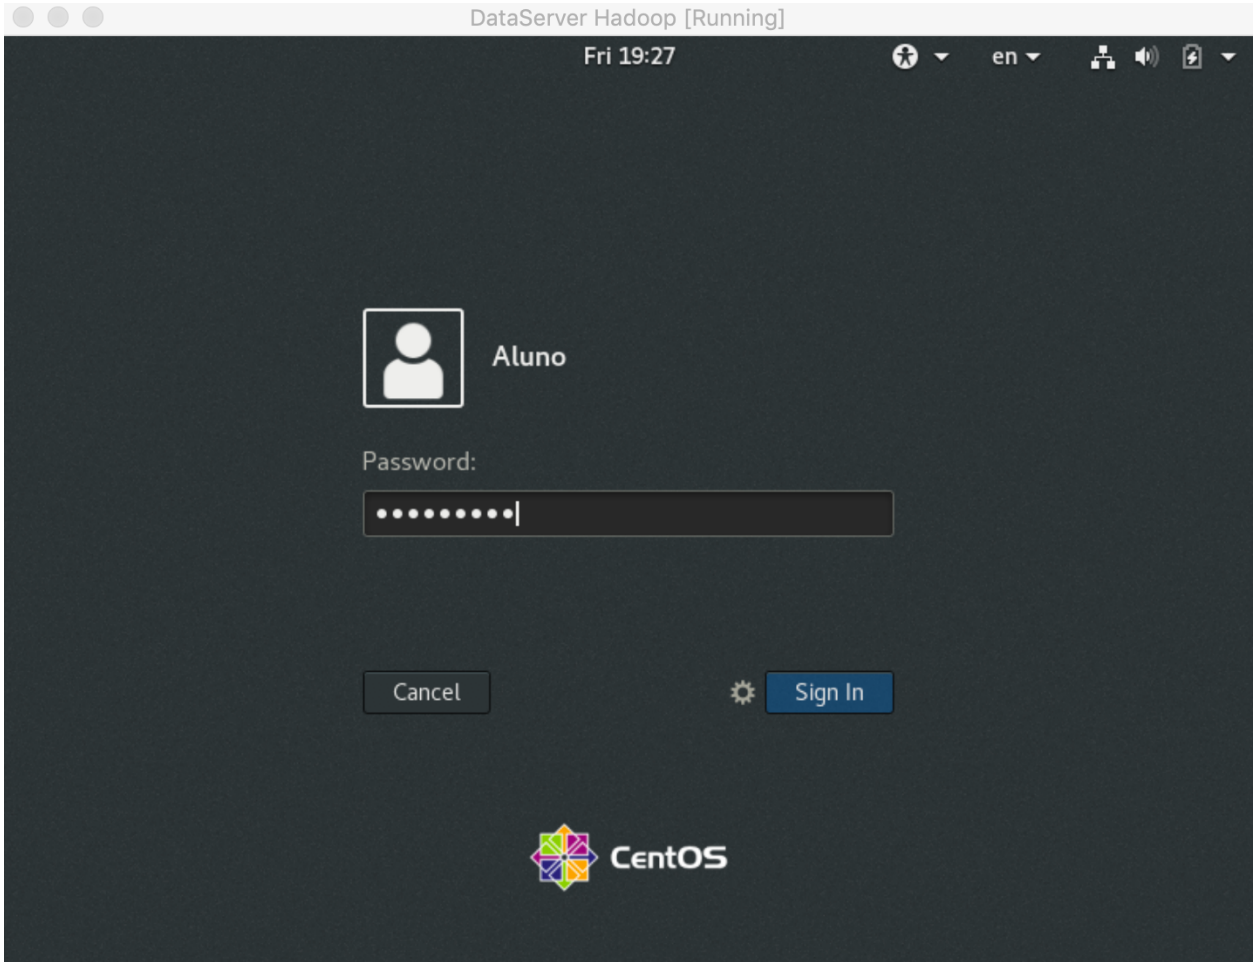

Usuário/Senha (**dsahadoop**)

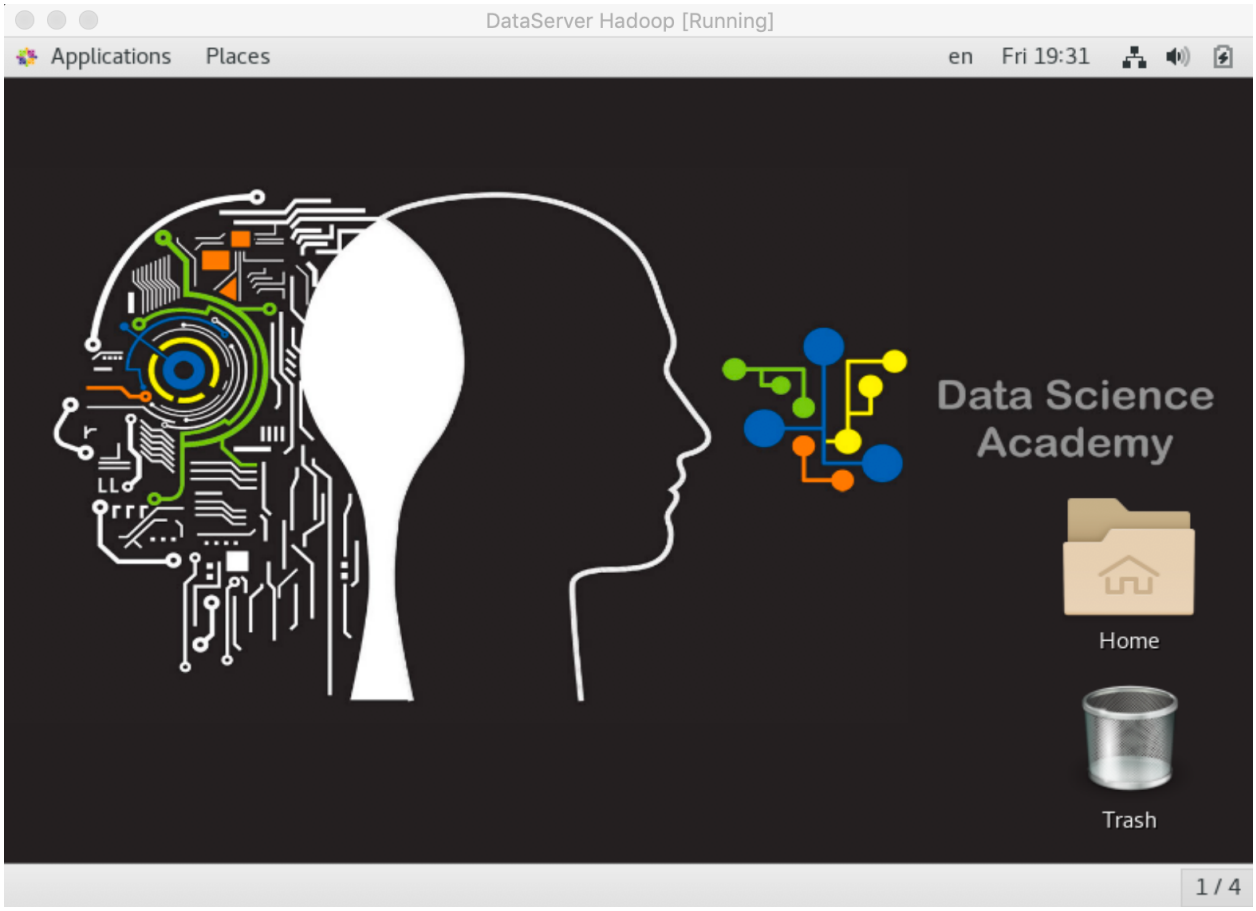

Instalação concluída com sucesso

Primeiro checkpoint:

Desligue a VM (clique no ícone da bateria e então em desligar).
Clique no menu File do VirtualBox e clique em Export Appliance.
Será gerada uma cópia de segurança da sua máquina virtual.

→ **VM: DataServer-Hadoop-v1.0.ova (Apenas SO)**

2.3. Instalação de Utilitários do Sistema Operacional

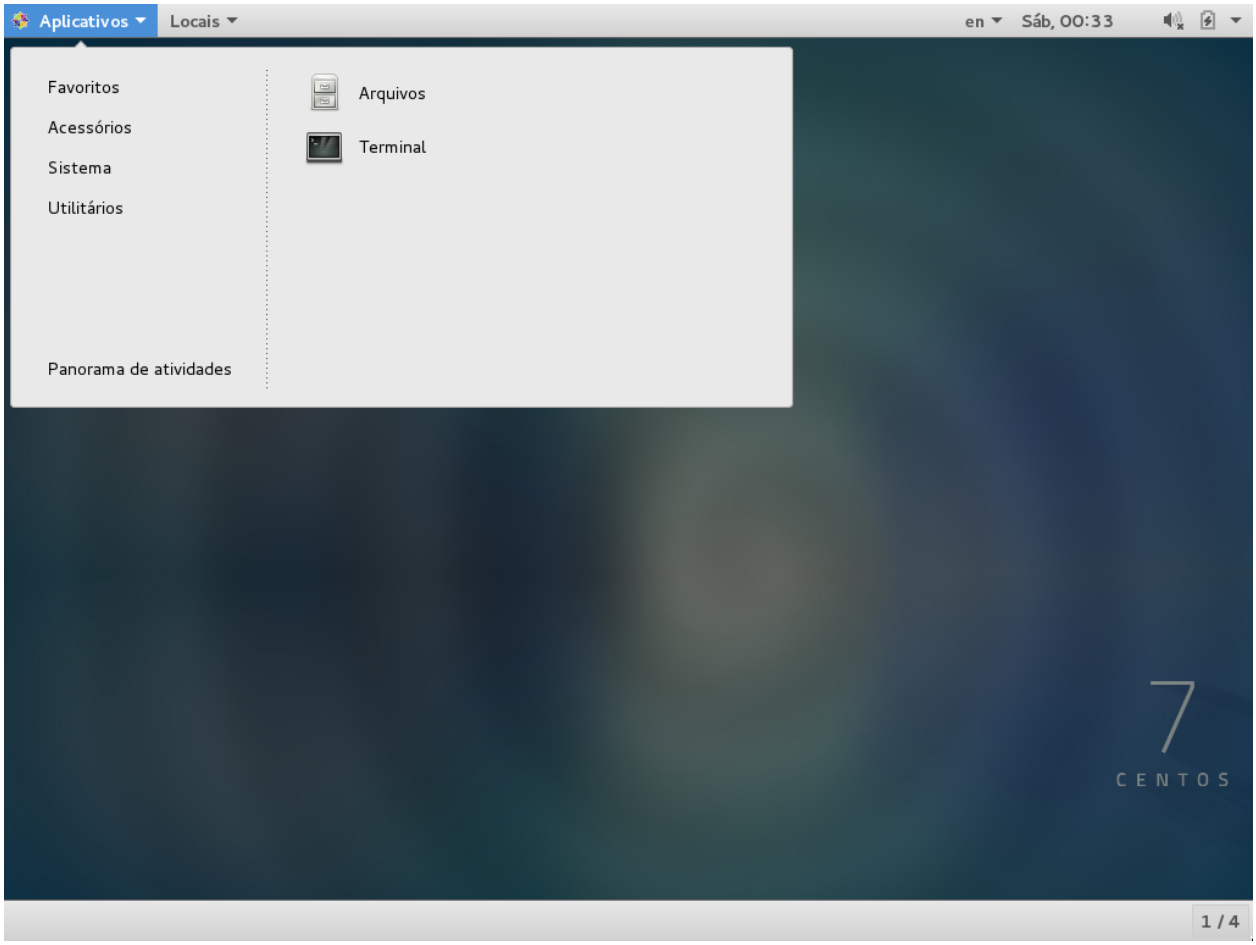

Abrindo o terminal

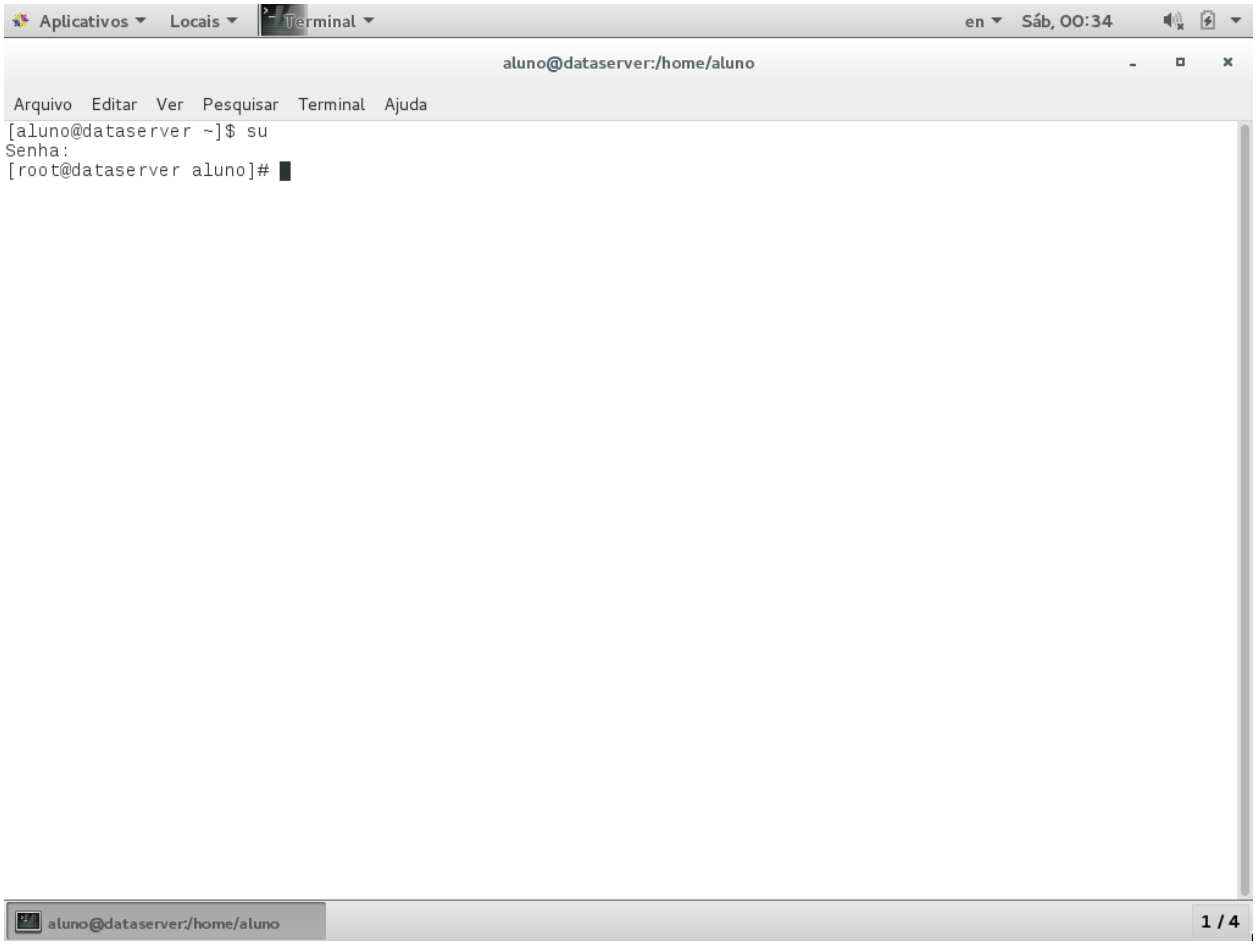

Efetuar login como root, usando o comando su. Senha: **dsahadoop**

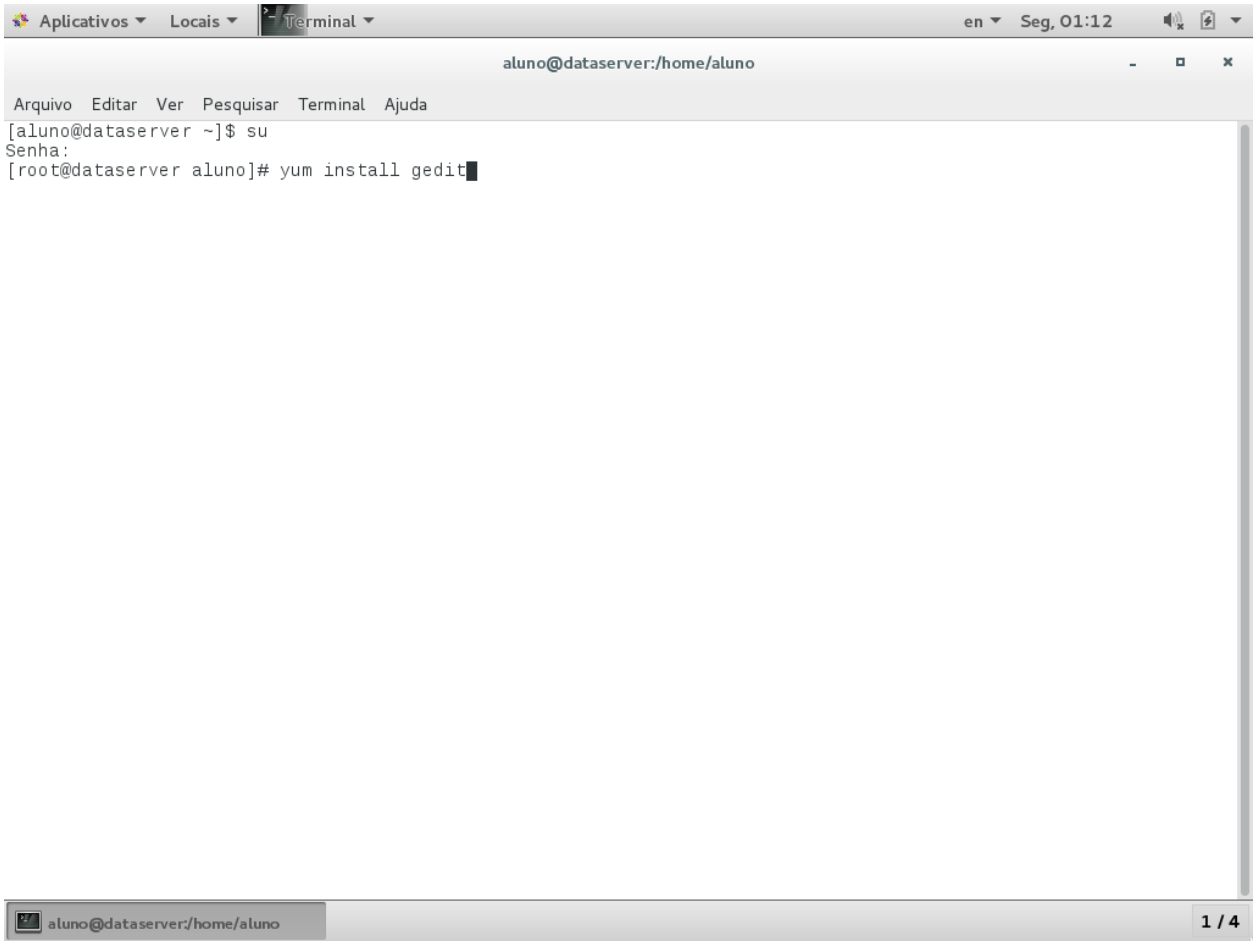

Instalar o editor de texto gedit, com o comando **yum install gedit**

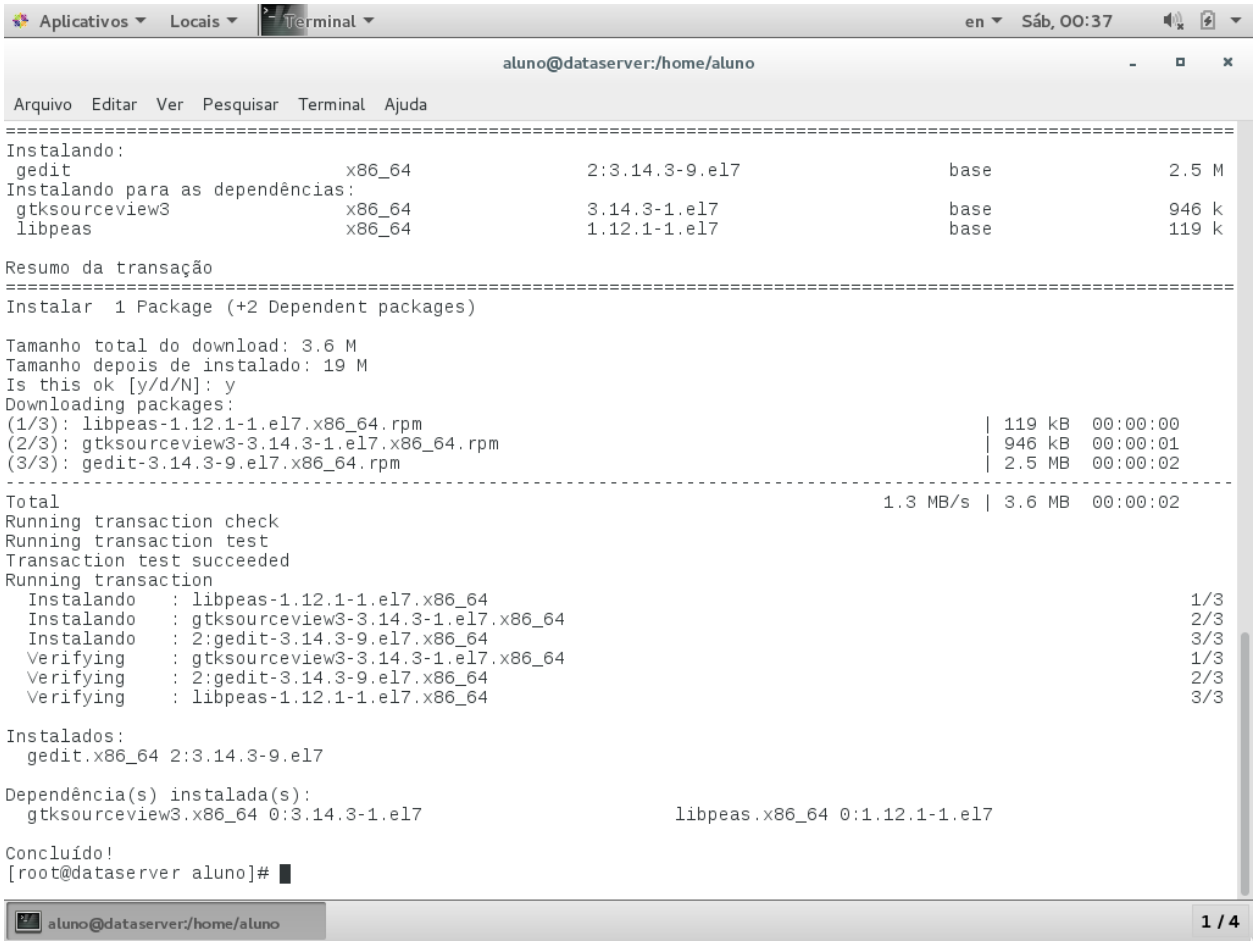

gedit instalado

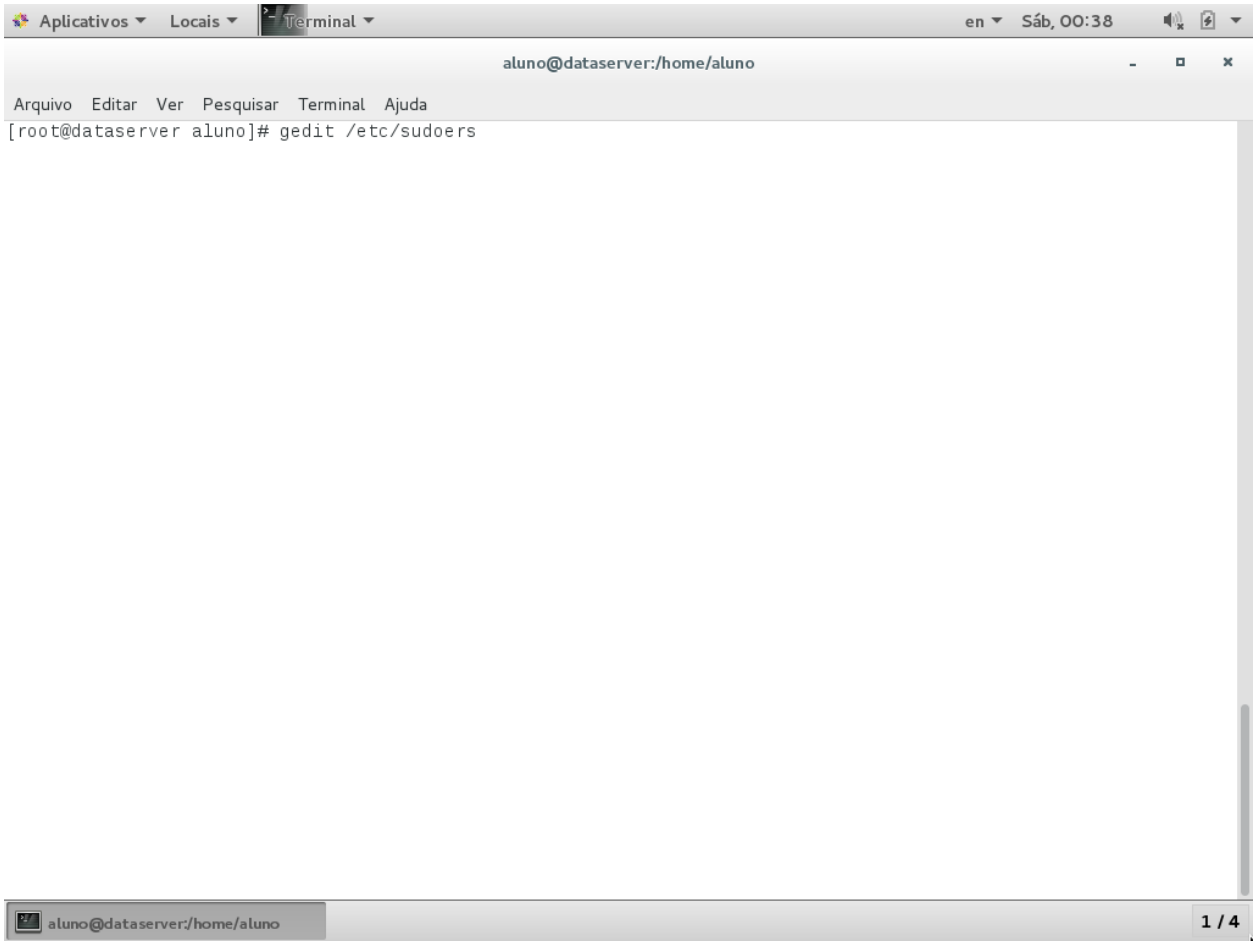

Editar o arquivo /etc/sudoers usando o gedit

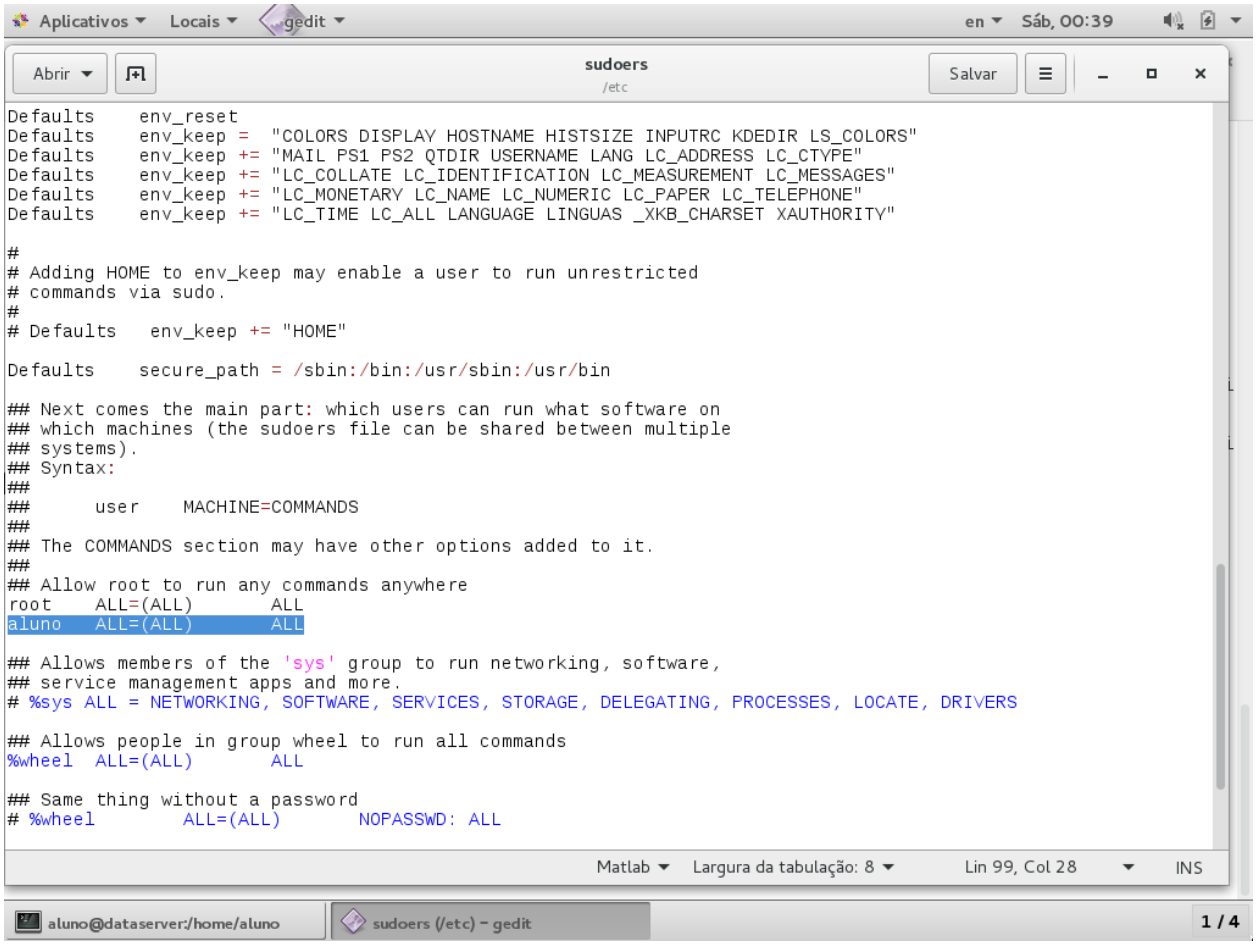

Incluir no arquivo a linha marcada acima e salvar o arquivo. Isso permitirá o usuário aluno executar comandos de administrador (root).

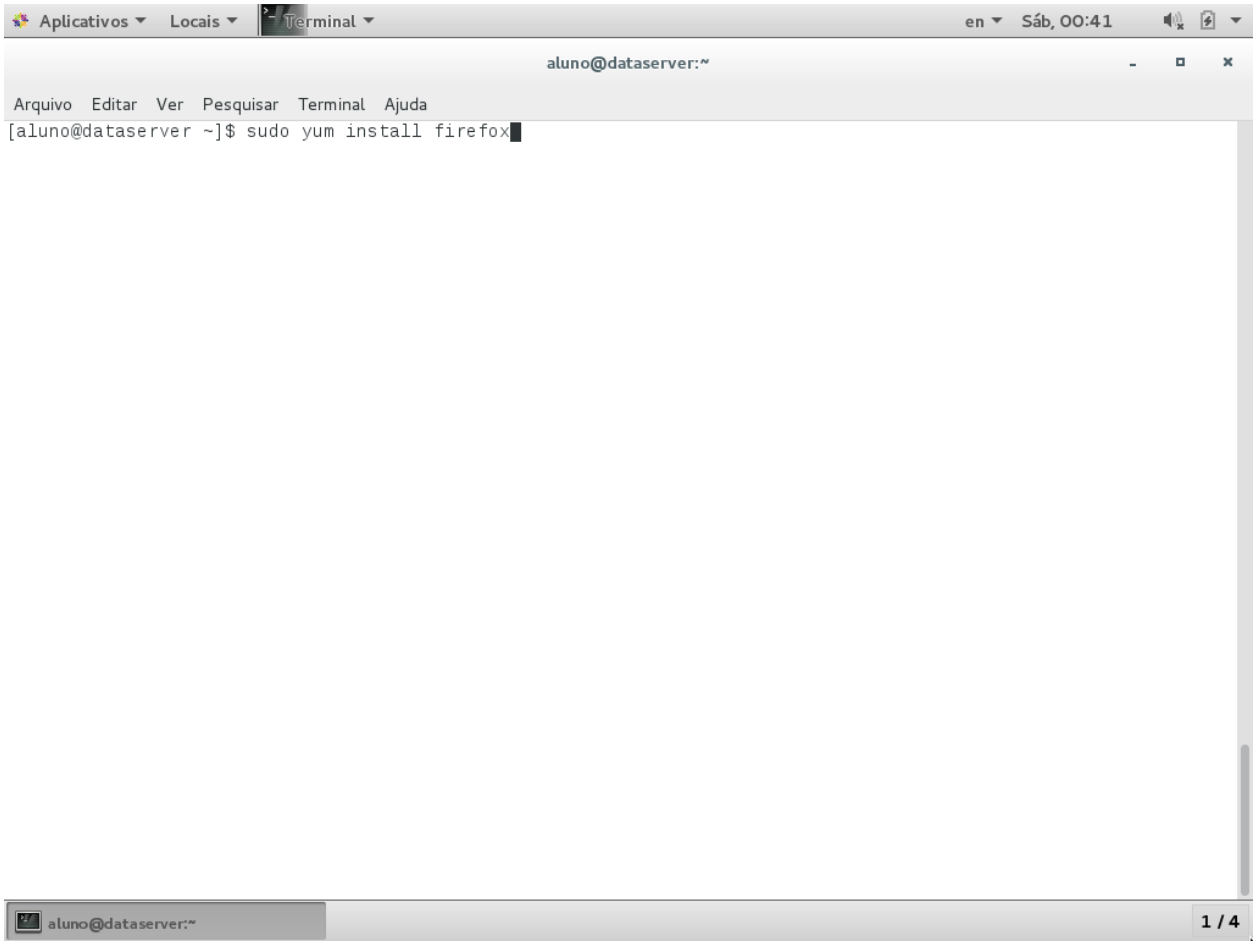

Conectado como usuário aluno, instalar o Firefox com o comando: **sudo yum install firefox**

Instalação concluída

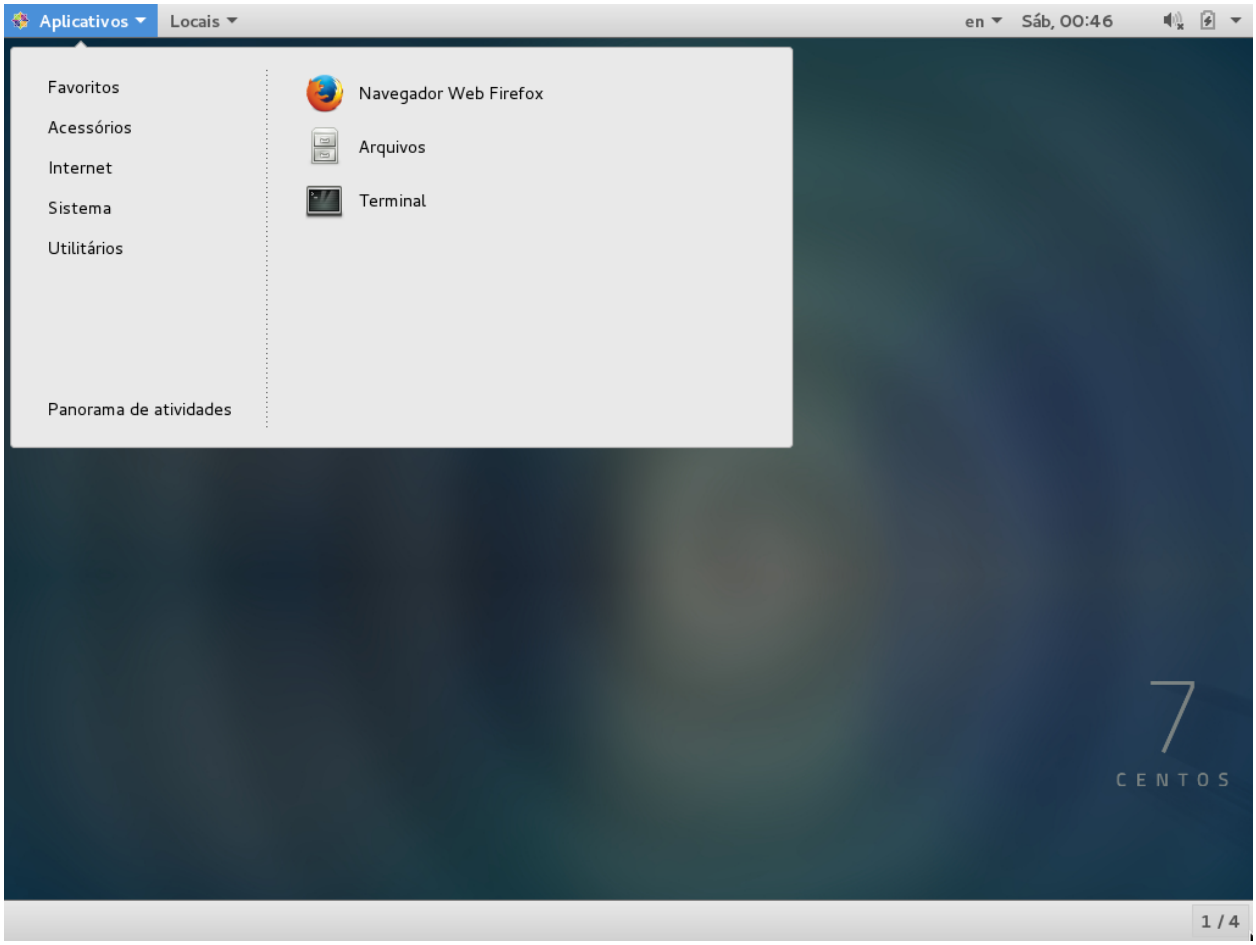

Firefox instalado

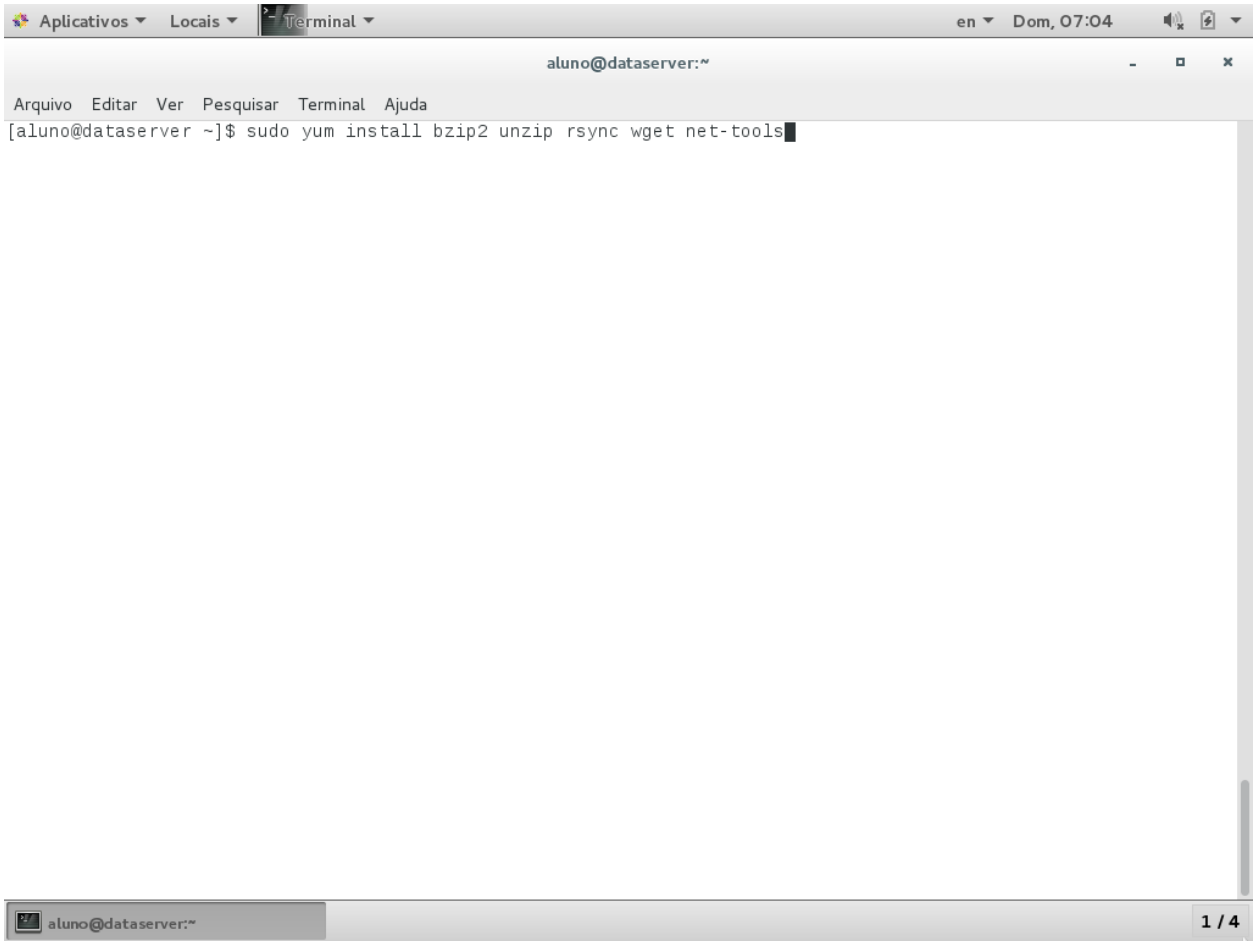

Instalar outros aplicativos: bzip2, unzip, rsync, wget e net-tools

Aplicativos instalados

2.4. Instalação do MySQL

A instalação do MySQL pode ser feita via linha de comando com o seguinte procedimento:

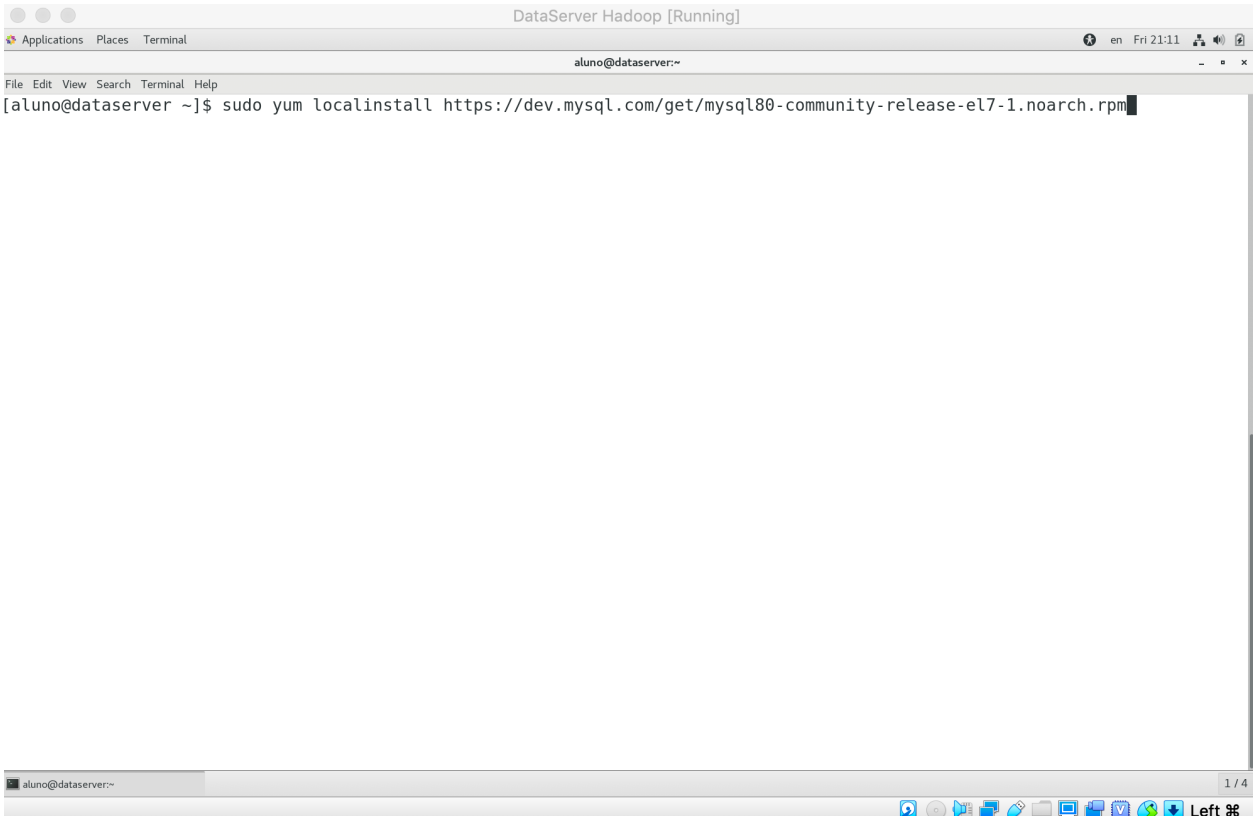

Execute o comando abaixo para baixar o pacote de instalação do MySQL para o CentOS:

sudo yum localinstall https://dev.mysql.com/get/mysql80-community-release-el7-1.noarch.rpm

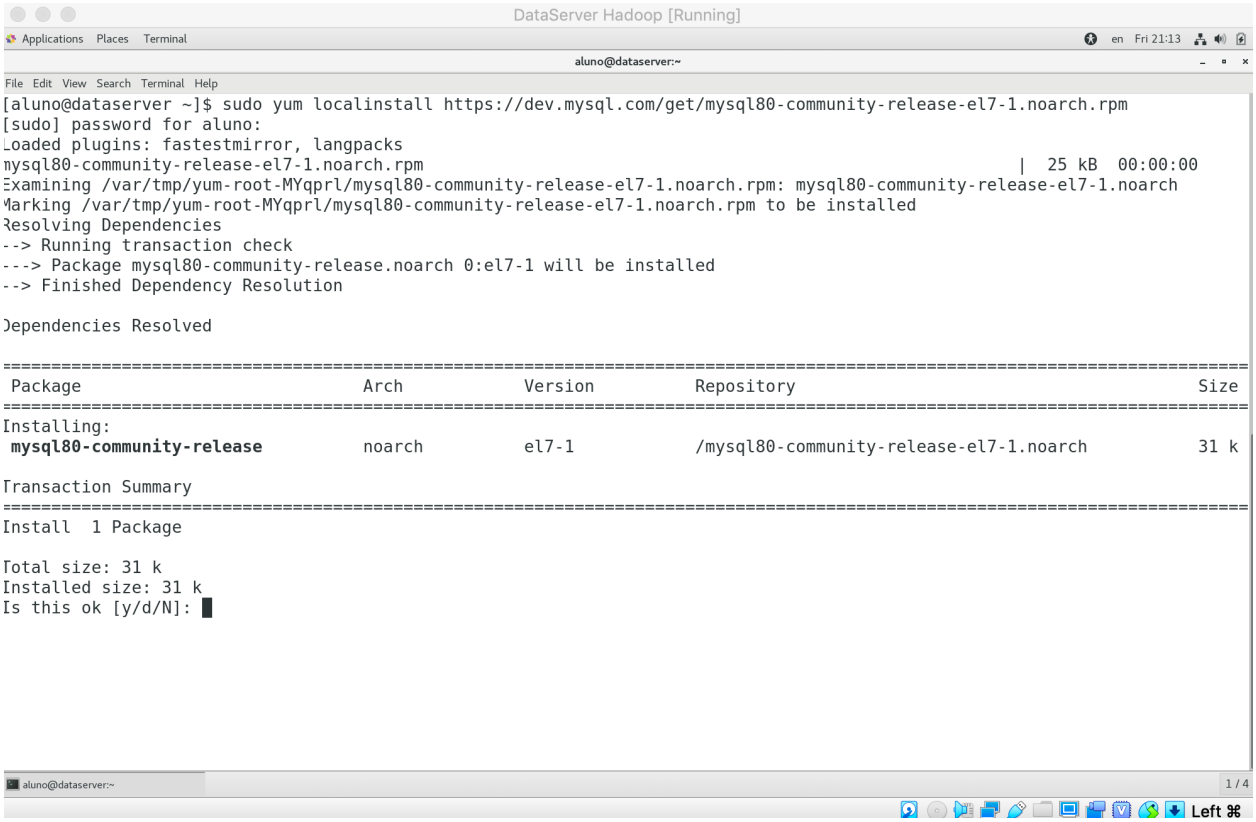
DataServer Hadoop [Running]
Applications Places Terminal
en Fri 21:13
aluno@dataserver:~
File Edit View Search Terminal Help
[aluno@dataserver ~]$ sudo yum localinstall https://dev.mysql.com/get/mysql80-community-release-el7-1.noarch.rpm
[sudo] password for aluno:
Loaded plugins: fastestmirror, langpacks
mysql80-community-release-el7-1.noarch.rpm | 25 kB 00:00:00
Examining /var/tmp/yum-root-MYqprl/mysql80-community-release-el7-1.noarch.rpm: mysql80-community-release-el7-1.noarch
Marking /var/tmp/yum-root-MYqprl/mysql80-community-release-el7-1.noarch.rpm to be installed
Resolving Dependencies
--> Running transaction check
---> Package mysql80-community-release.noarch 0:el7-1 will be installed
--> Finished Dependency Resolution
Dependencies Resolved
Package Arch Version Repository Size
Installing:
mysql80-community-release noarch el7-1 /mysql80-community-release-el7-1.noarch 31 k
Transaction Summary
Install 1 Package
Total size: 31 k
Installed size: 31 k
Is this ok [y/d/N]:
aluno@dataserver:~
1 / 4
Left ⌘

Pressione y

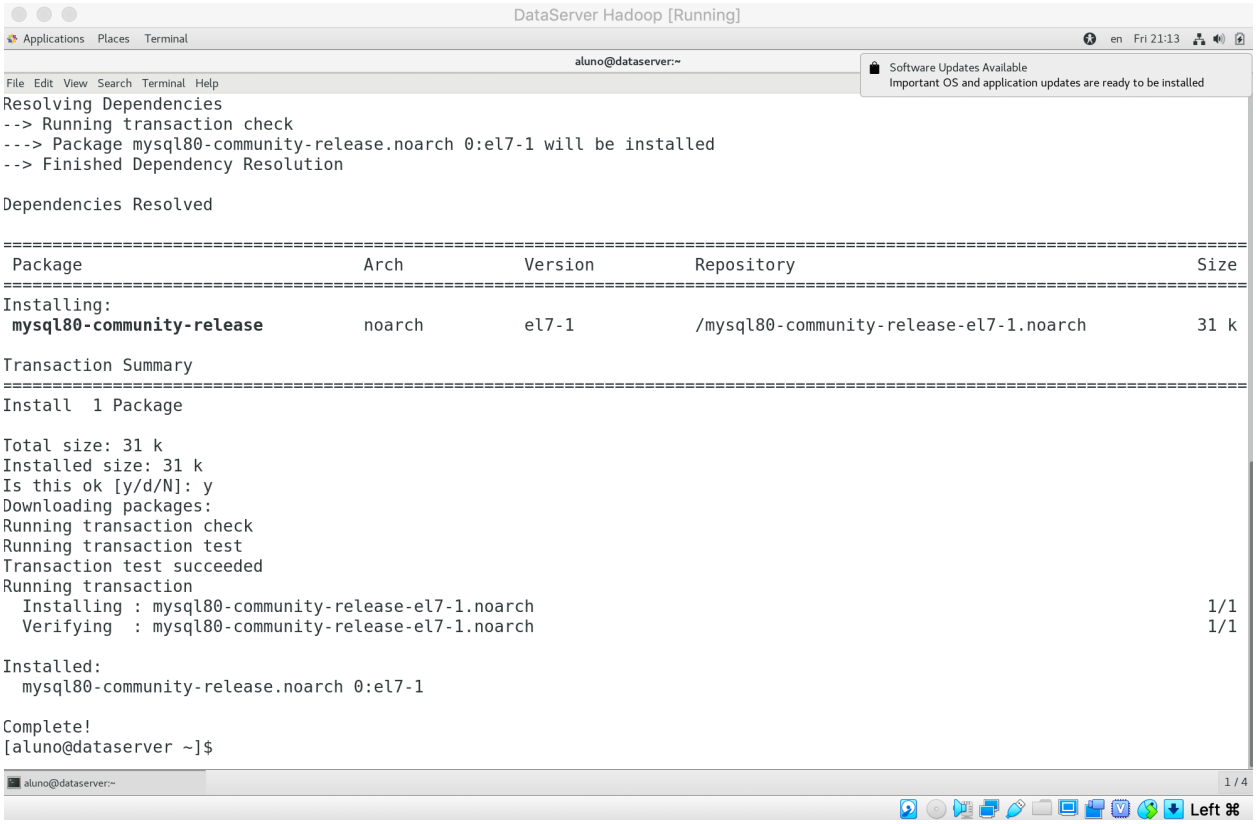

Download do pacote concluído

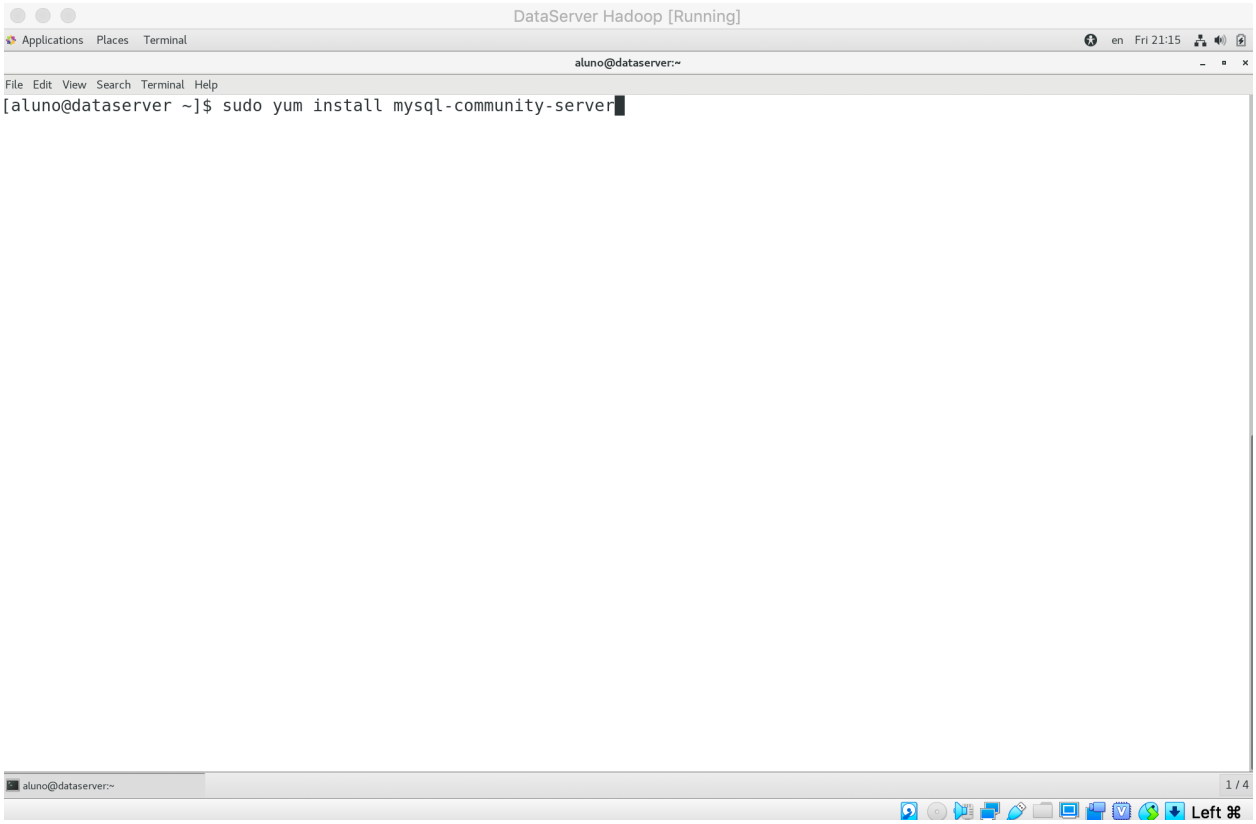

Esse comando inicia a instalação do MySQL:

sudo yum install mysql-community-server

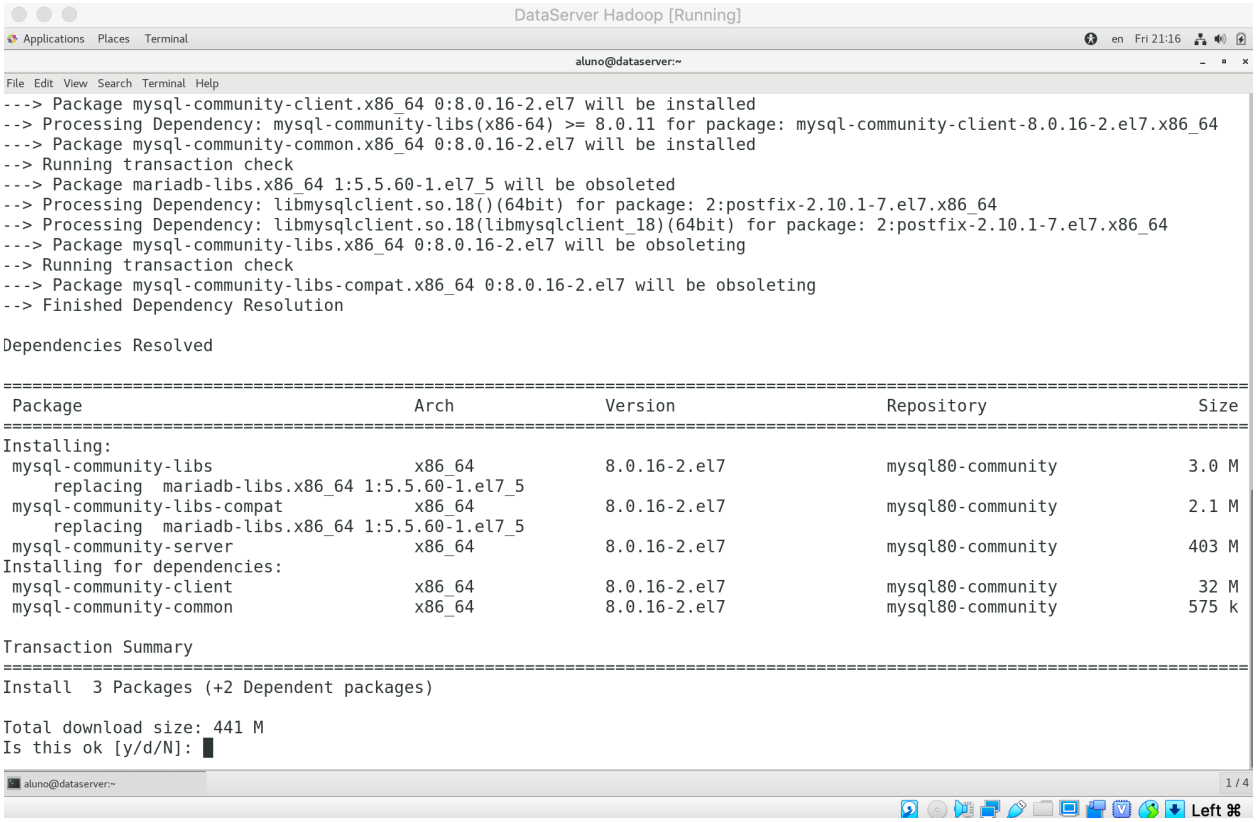
DataServer Hadoop [Running]
Applications Places Terminal
en Fri 21:16
aluno@dataserver:~
File Edit View Search Terminal Help
---> Package mysql-community-client.x86_64 0:8.0.16-2.el7 will be installed
--> Processing Dependency: mysql-community-libs(x86-64) >= 8.0.11 for package: mysql-community-client-8.0.16-2.el7.x86_64
---> Package mysql-community-common.x86_64 0:8.0.16-2.el7 will be installed
--> Running transaction check
---> Package mariadb-libs.x86_64 1:5.5.60-1.el7_5 will be obsoleted
--> Processing Dependency: libmysqlclient.so.18()(64bit) for package: 2:postfix-2.10.1-7.el7.x86_64
--> Processing Dependency: libmysqlclient.so.18(libmysqlclient_18)(64bit) for package: 2:postfix-2.10.1-7.el7.x86_64
---> Package mysql-community-libs.x86_64 0:8.0.16-2.el7 will be obsoleting
--> Running transaction check
---> Package mysql-community-libs-compat.x86_64 0:8.0.16-2.el7 will be obsoleting
--> Finished Dependency Resolution
Dependencies Resolved
Package Arch Version Repository Size
Installing:
mysql-community-libs x86_64 8.0.16-2.el7 mysql80-community 3.0 M
replacing mariadb-libs.x86_64 1:5.5.60-1.el7_5
mysql-community-libs-compat x86_64 8.0.16-2.el7 mysql80-community 2.1 M
replacing mariadb-libs.x86_64 1:5.5.60-1.el7_5
mysql-community-server x86_64 8.0.16-2.el7 mysql80-community 403 M
Installing for dependencies:
mysql-community-client x86_64 8.0.16-2.el7 mysql80-community 32 M
mysql-community-common x86_64 8.0.16-2.el7 mysql80-community 575 k
Transaction Summary
Install 3 Packages (+2 Dependent packages)
Total download size: 441 M
Is this ok [y/d/N]:
aluno@dataserver:~
1 / 4
Left ⌘

Pressione y

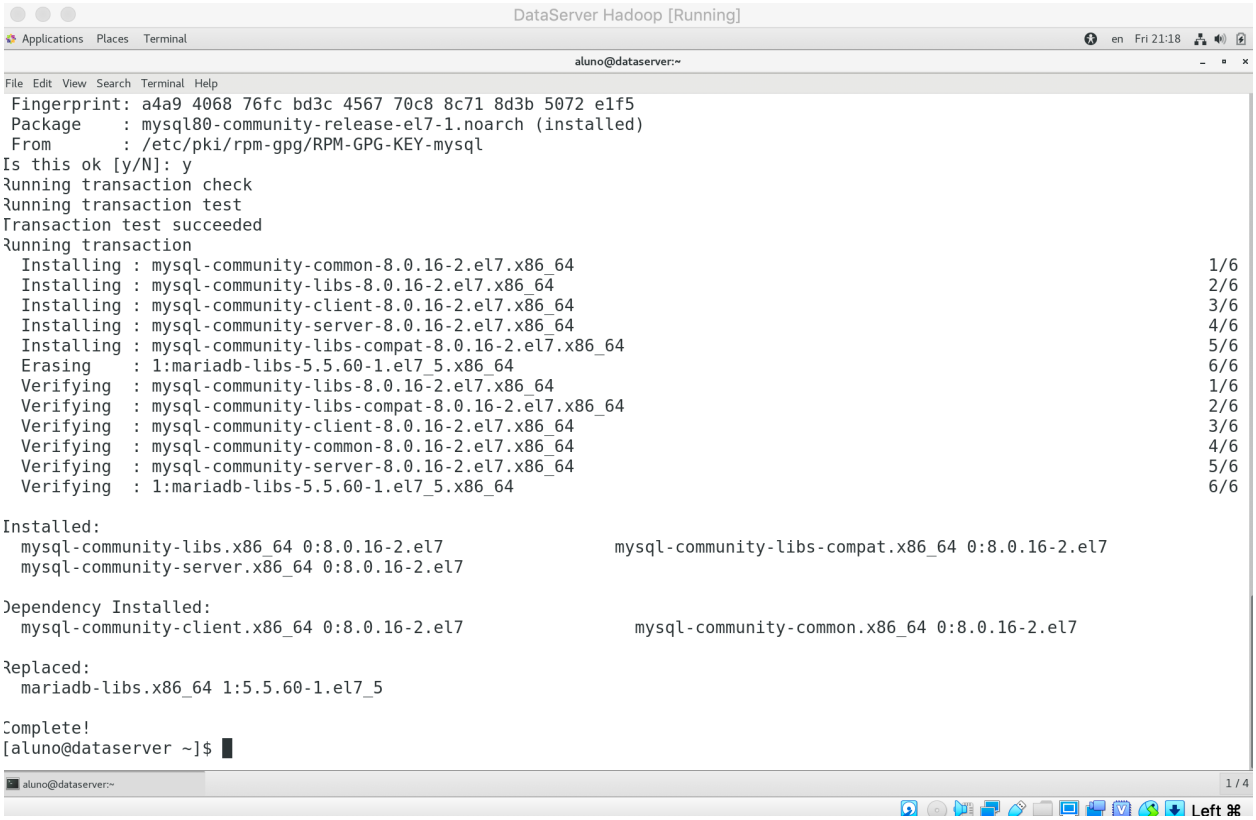

Instalação concluída

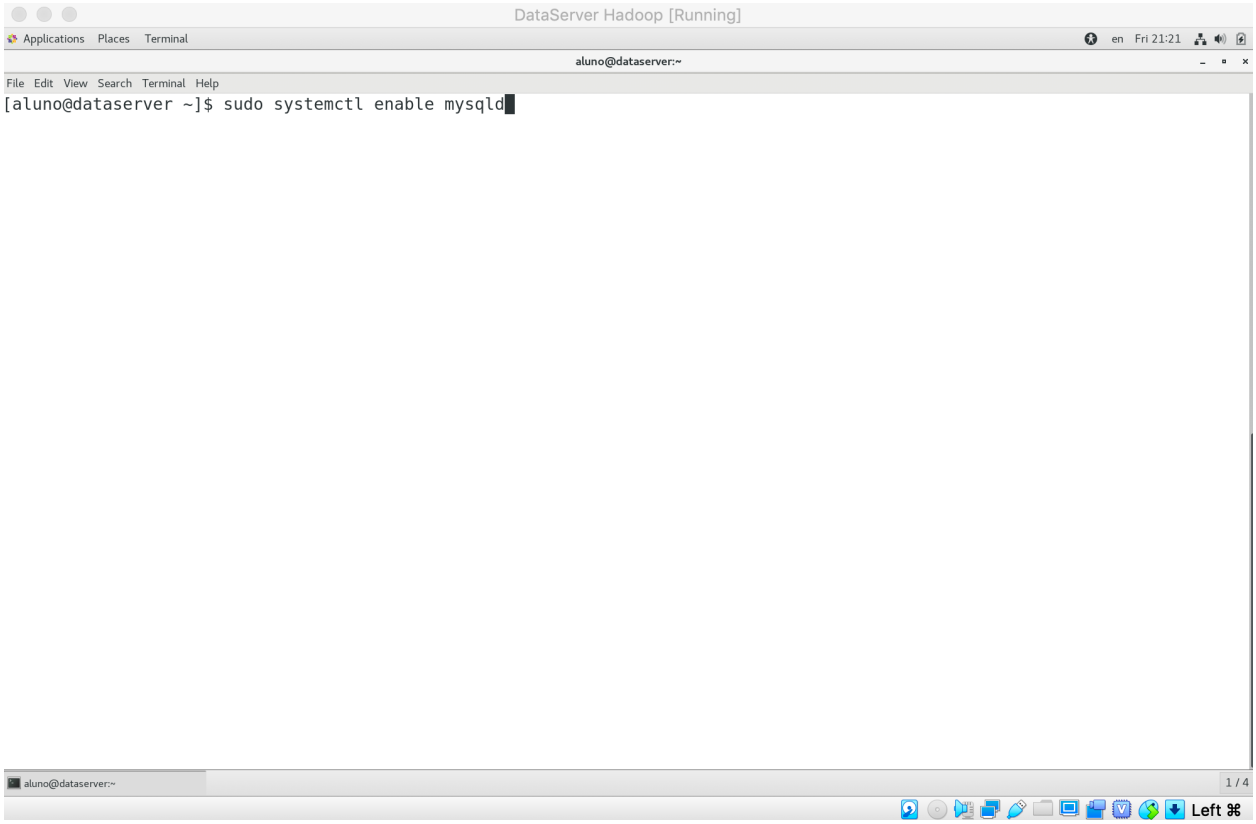

Ativando o serviço do MySQL

sudo systemctl enable mysqld

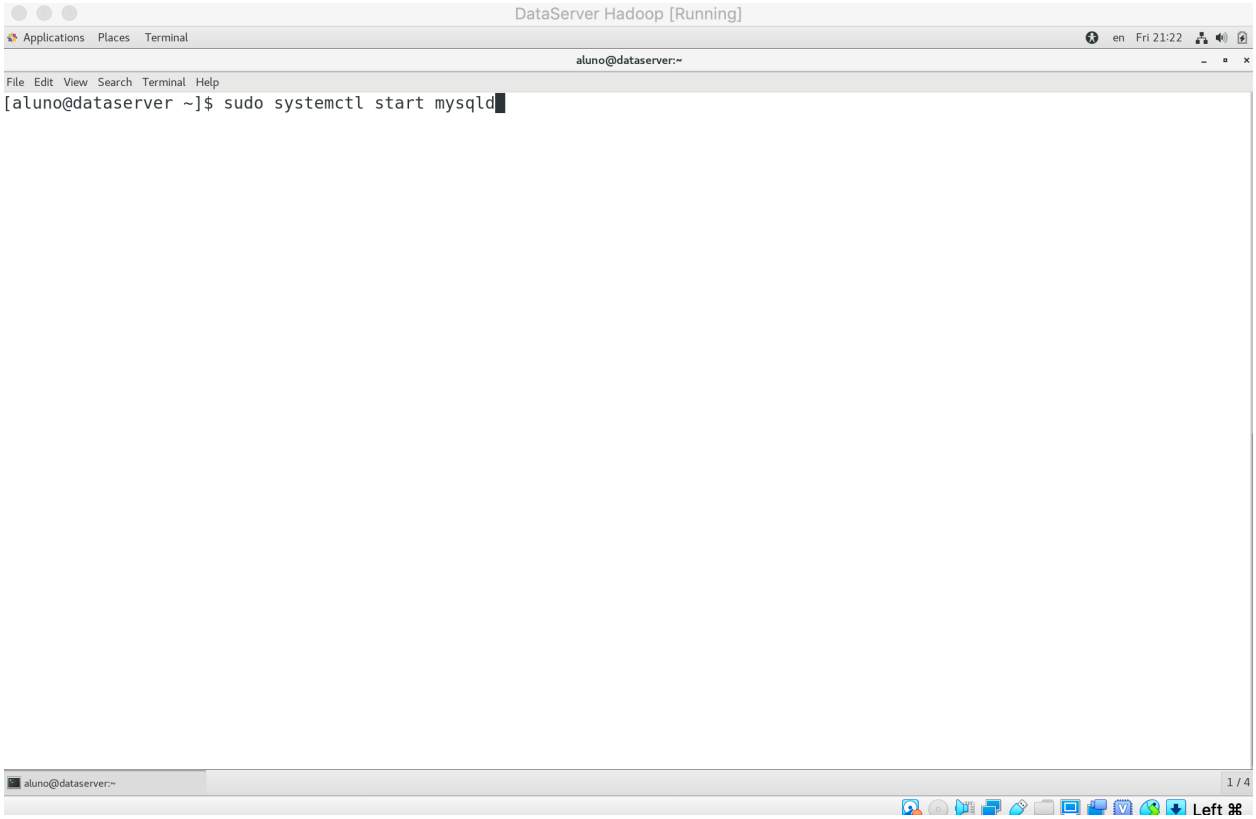

Inicia o MySQL

sudo systemctl start mysqld

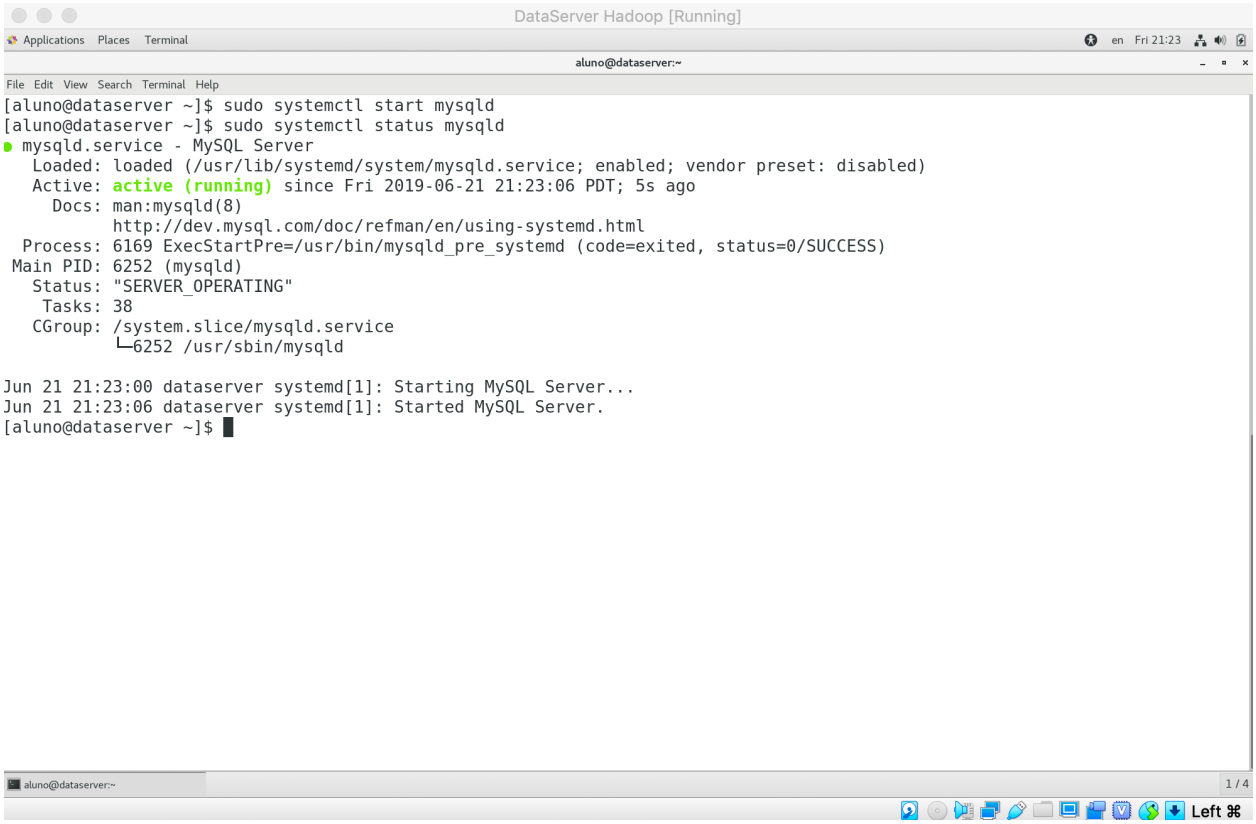

MySQL em execução

3. Instalação do servidor ssh

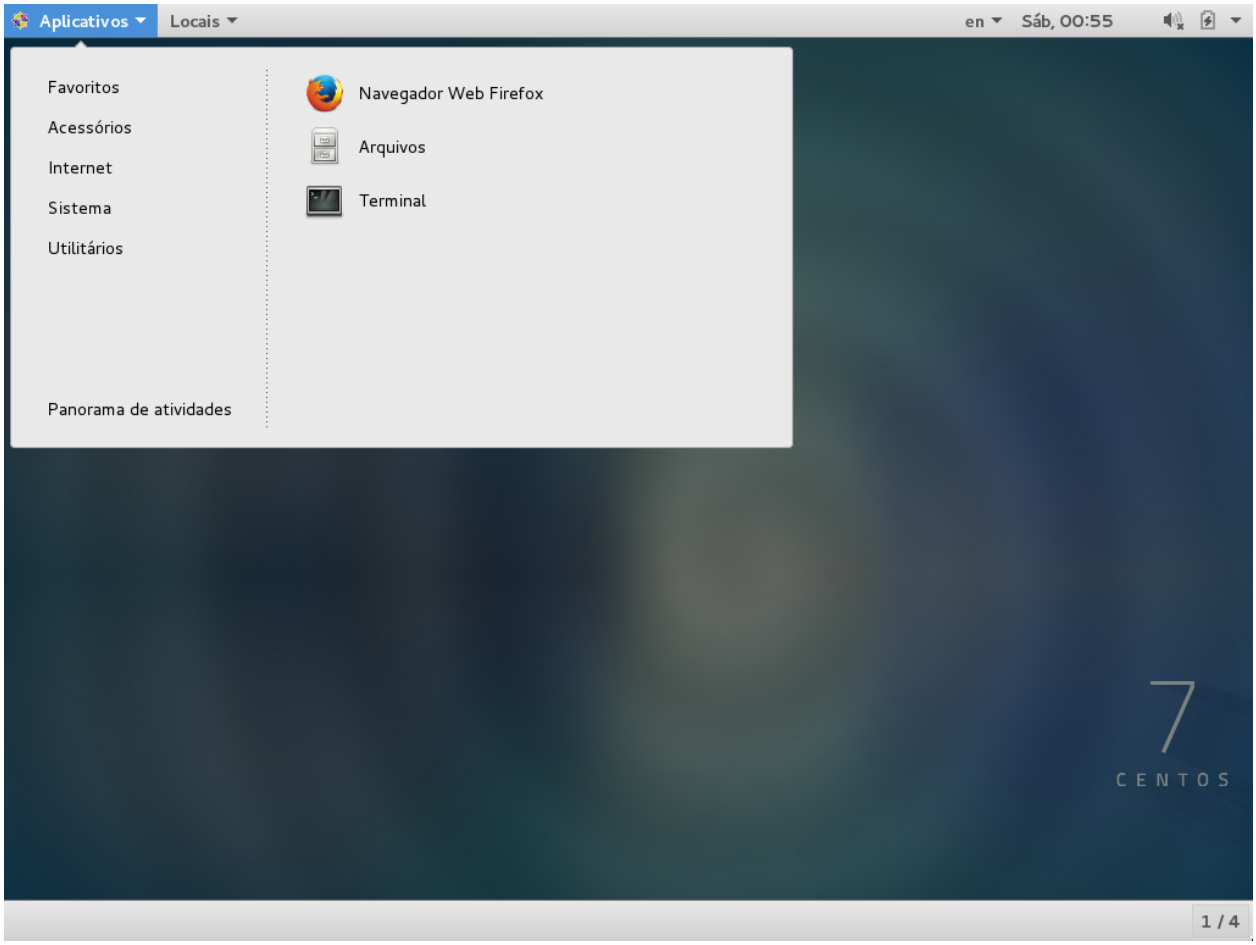

Abrindo o terminal

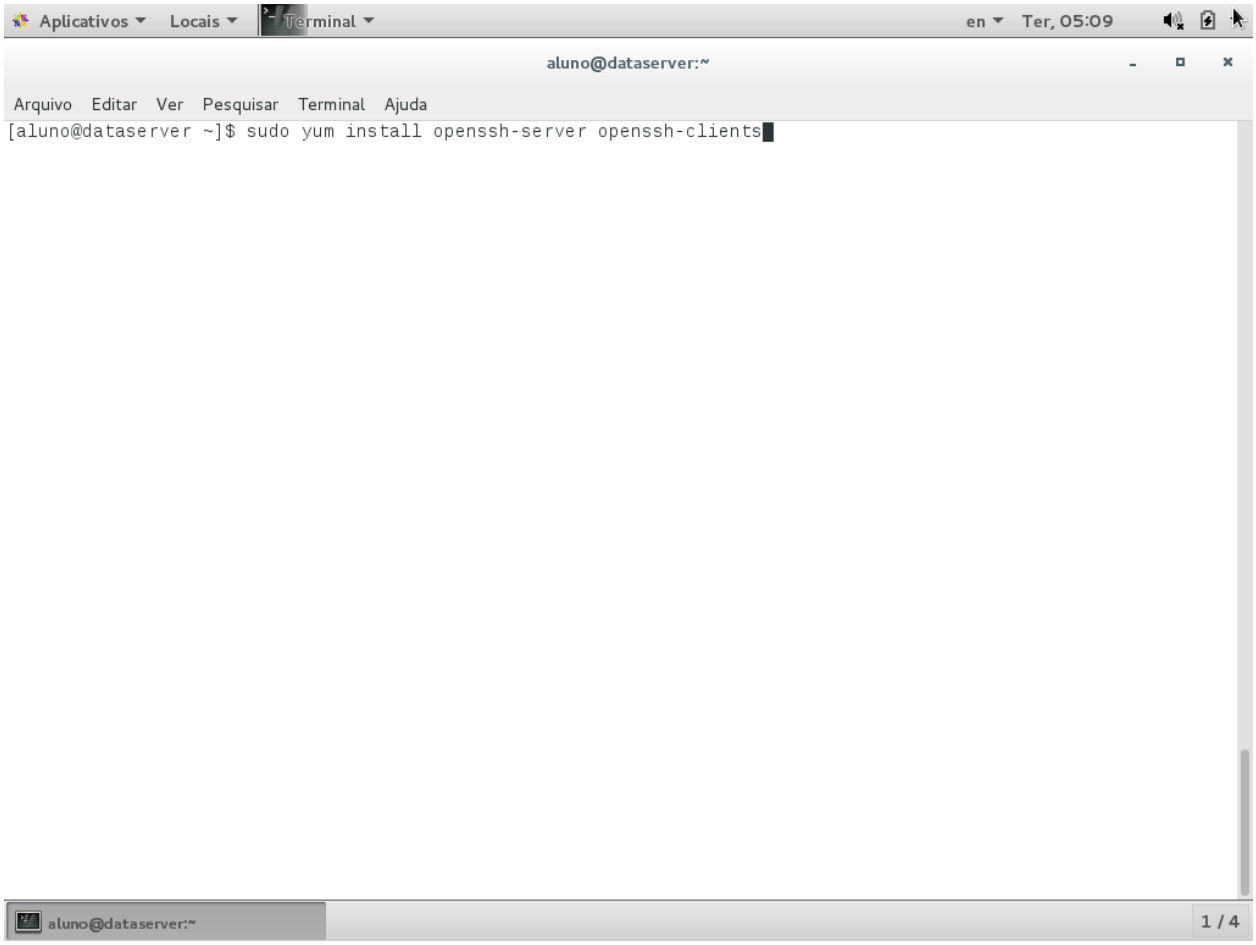

sudo yum install openssh-server openssh-clients

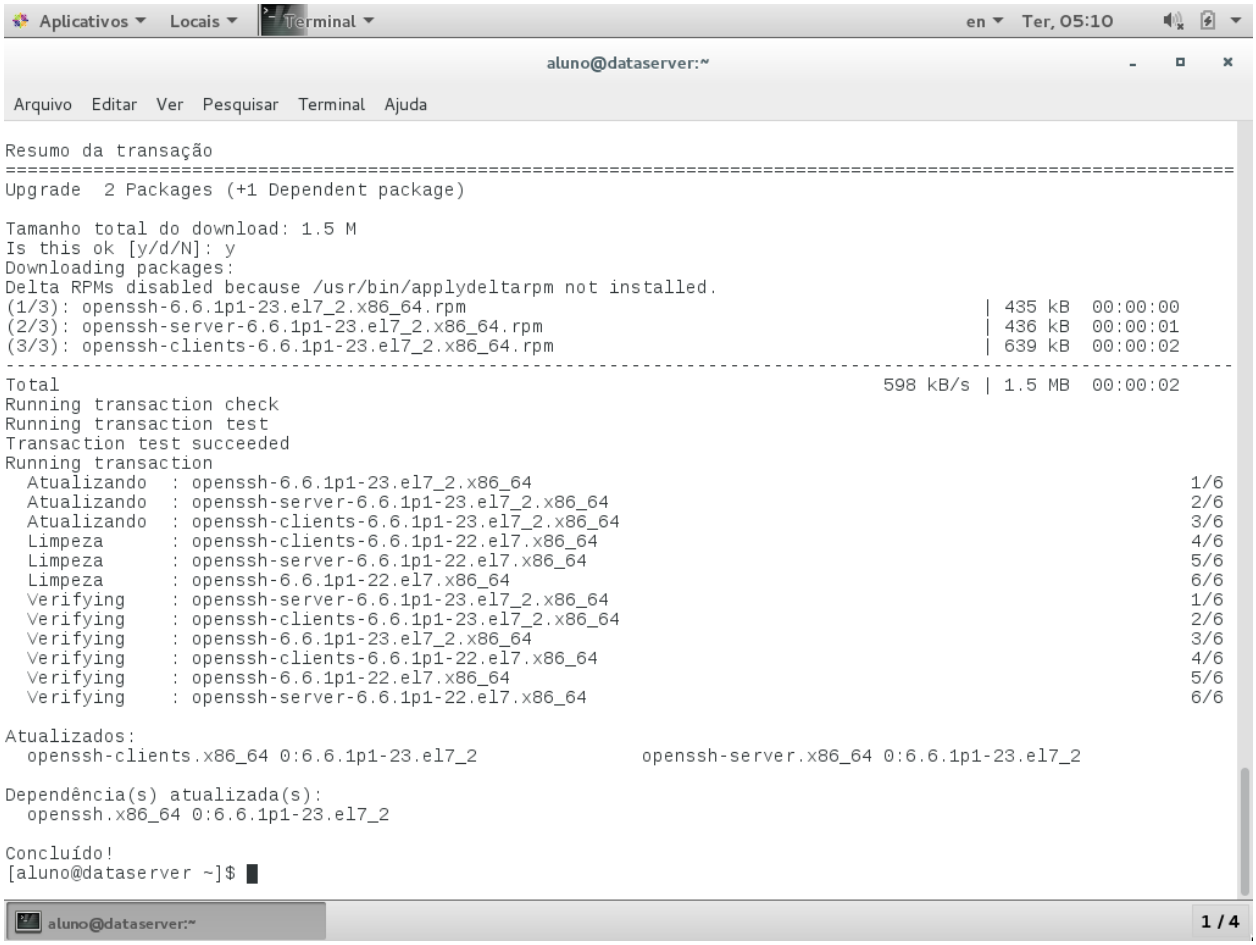
Aplicativos Locais Terminal en Ter, 05:10
aluno@dataserver:~
Arquivo Editar Ver Pesquisar Terminal Ajuda
Resumo da transação
Upgrade 2 Packages (+1 Dependent package)
Tamanho total do download: 1.5 M
Is this ok [y/d/N]: y
Downloading packages:
Delta RPMs disabled because /usr/bin/applydeltarpm not installed.
(1/3): openssh-6.6.1p1-23.el7_2.x86_64.rpm | 435 kB 00:00:00
(2/3): openssh-server-6.6.1p1-23.el7_2.x86_64.rpm | 436 kB 00:00:01
(3/3): openssh-clients-6.6.1p1-23.el7_2.x86_64.rpm | 639 kB 00:00:02
Total 598 kB/s | 1.5 MB 00:00:02
Running transaction check
Running transaction test
Transaction test succeeded
Running transaction
Atualizando : openssh-6.6.1p1-23.el7_2.x86_64 1/6
Atualizando : openssh-server-6.6.1p1-23.el7_2.x86_64 2/6
Atualizando : openssh-clients-6.6.1p1-23.el7_2.x86_64 3/6
Limpeza : openssh-clients-6.6.1p1-22.el7.x86_64 4/6
Limpeza : openssh-server-6.6.1p1-22.el7.x86_64 5/6
Limpeza : openssh-6.6.1p1-22.el7.x86_64 6/6
Verifying : openssh-server-6.6.1p1-23.el7_2.x86_64 1/6
Verifying : openssh-clients-6.6.1p1-23.el7_2.x86_64 2/6
Verifying : openssh-6.6.1p1-23.el7_2.x86_64 3/6
Verifying : openssh-clients-6.6.1p1-22.el7.x86_64 4/6
Verifying : openssh-6.6.1p1-22.el7.x86_64 5/6
Verifying : openssh-server-6.6.1p1-22.el7.x86_64 6/6
Atualizados:
openssh-clients.x86_64 0:6.6.1p1-23.el7_2 openssh-server.x86_64 0:6.6.1p1-23.el7_2
Dependência(s) atualizada(s):
openssh.x86_64 0:6.6.1p1-23.el7_2
Concluído!
[aluno@dataserver ~]$
aluno@dataserver:~
1 / 4

Concluído

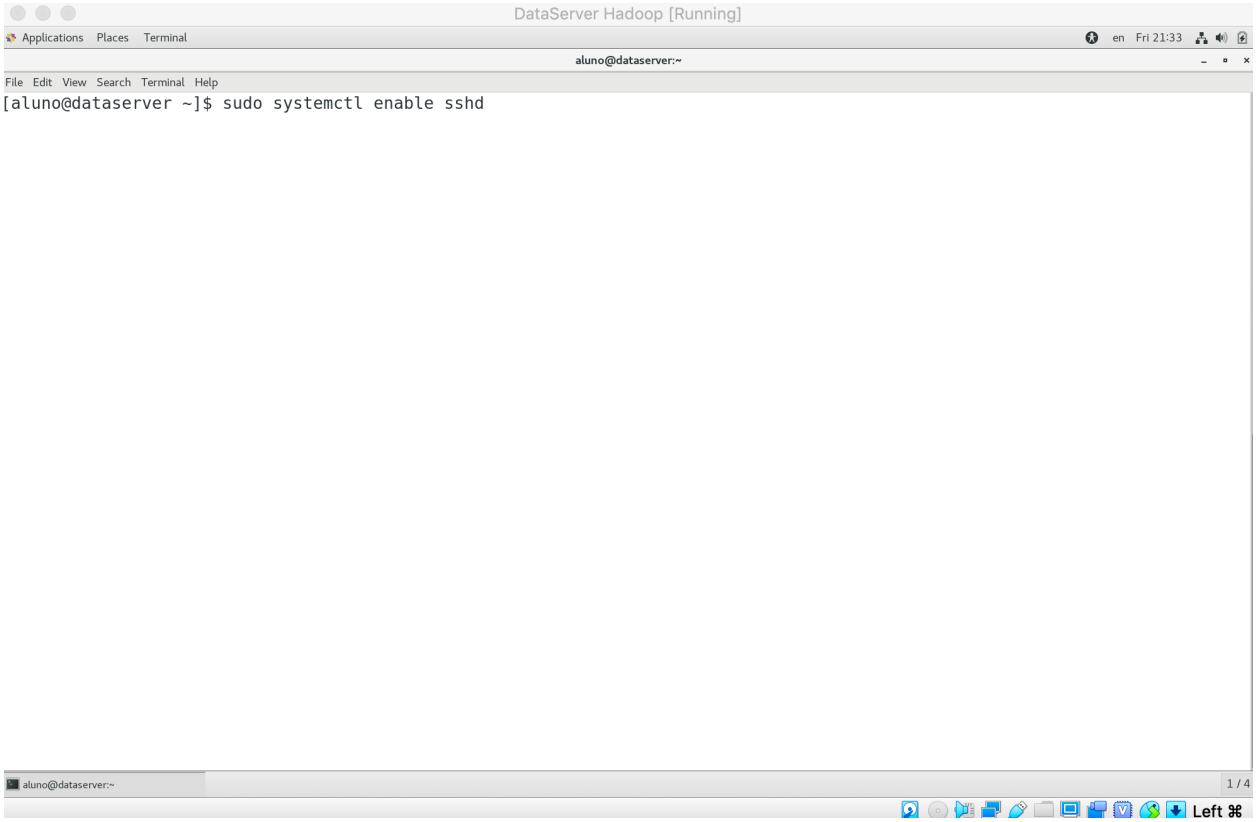

Habilitando o serviço

sudo systemctl enable sshd

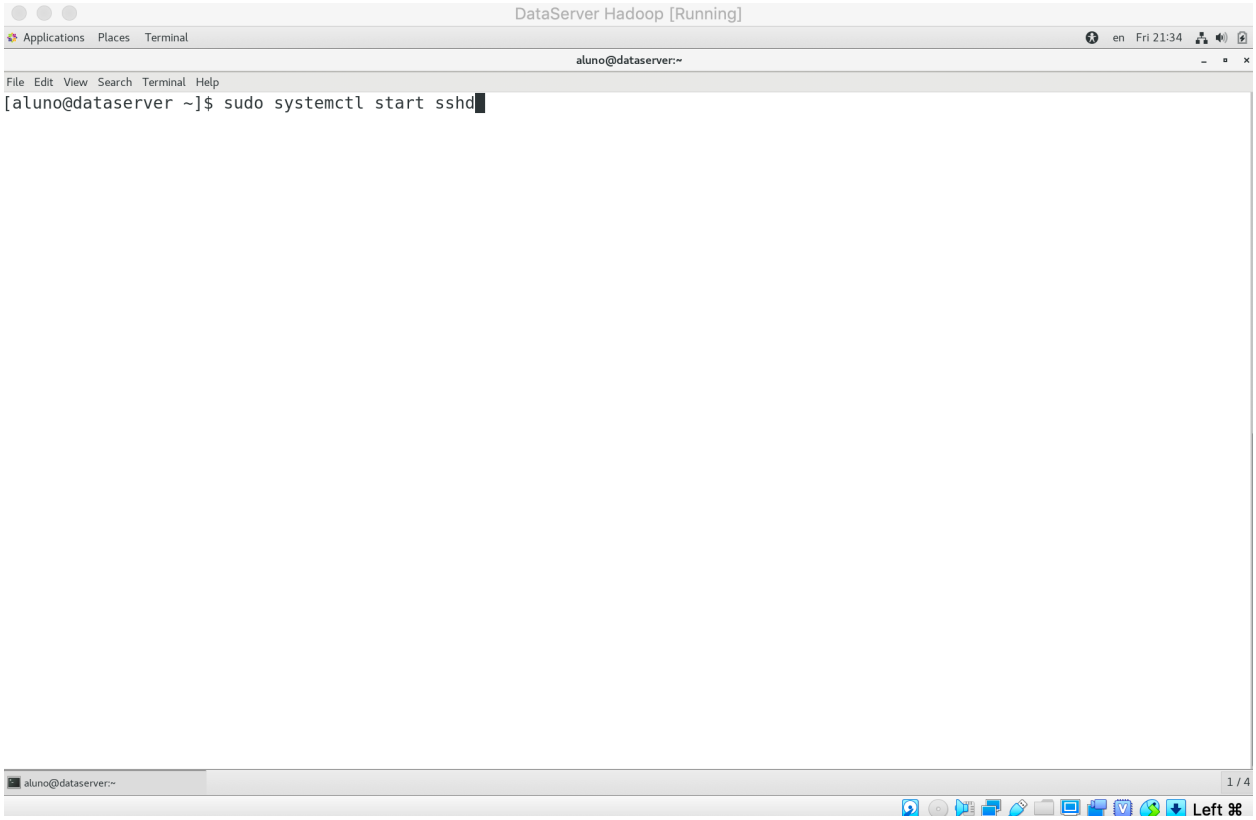

Iniciando o serviço

sudo systemctl start sshd

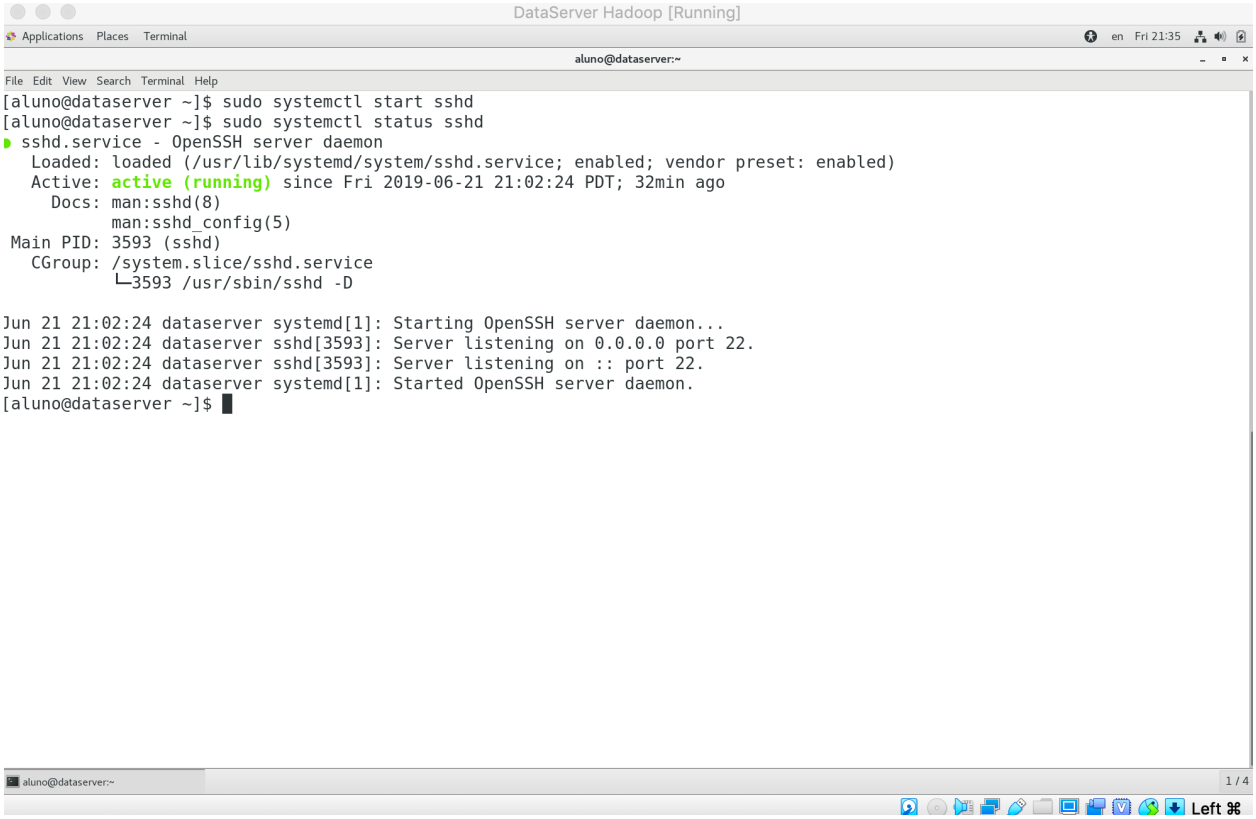

Serviço em execução

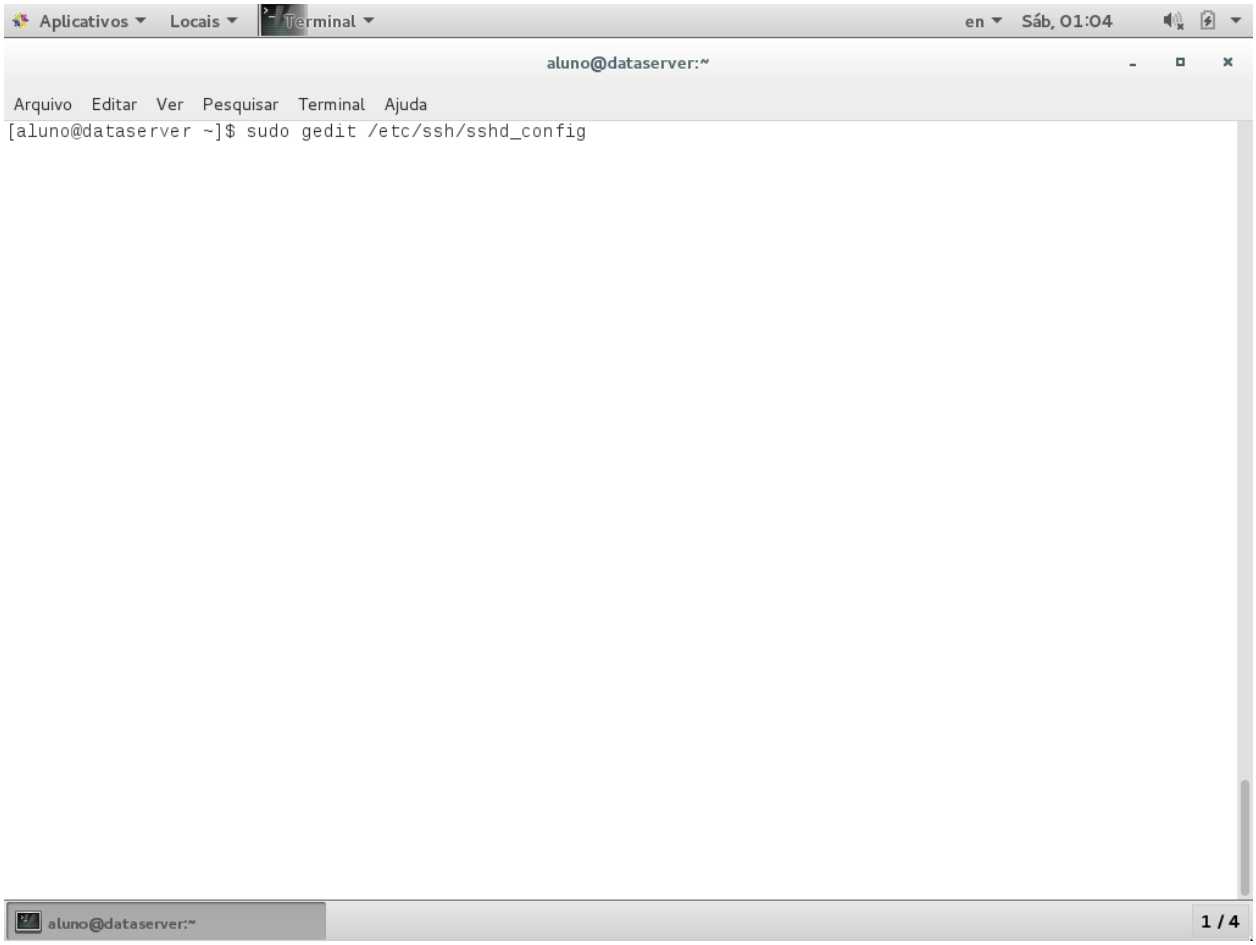

sudo gedit /etc/ssh/sshd_config

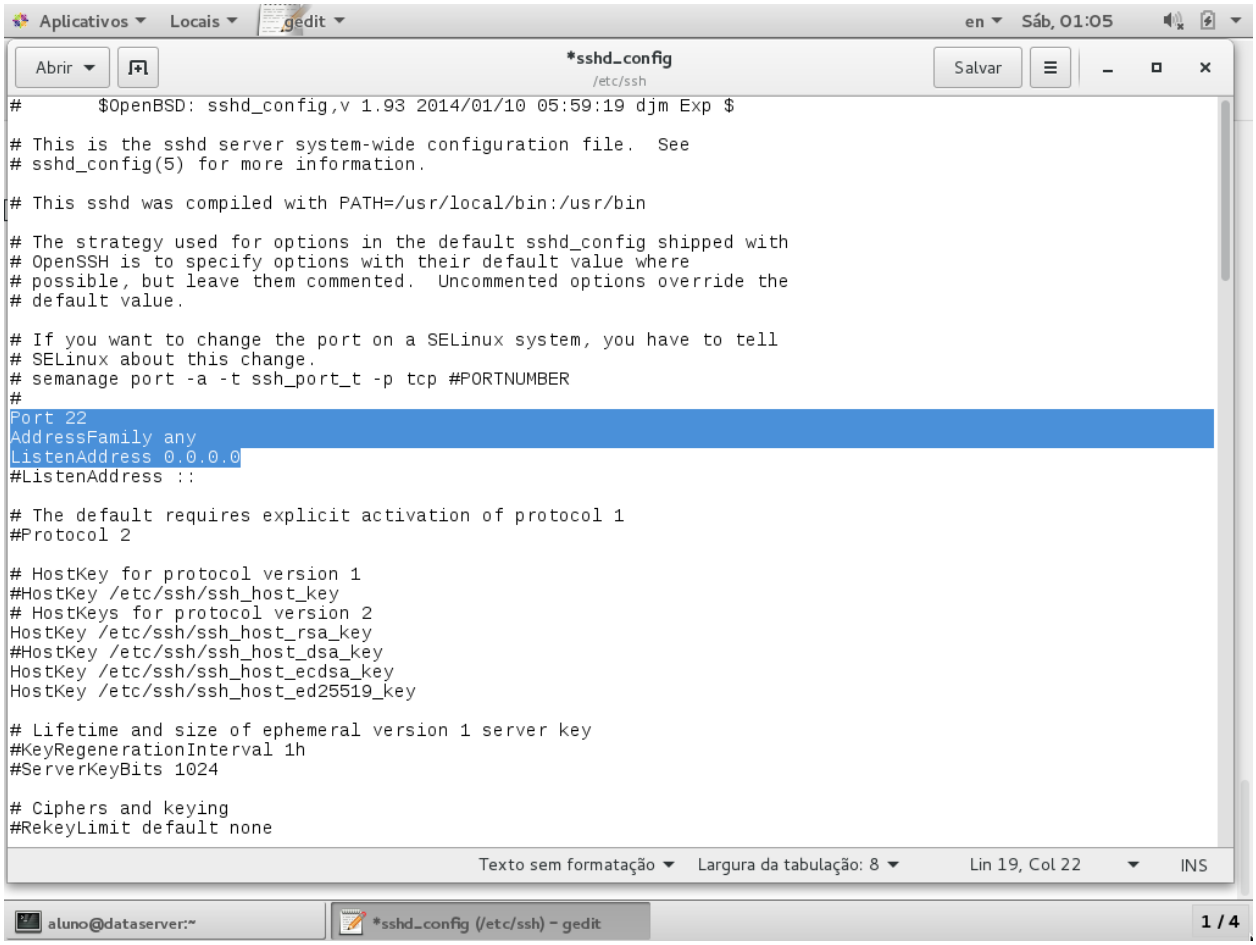

Primeira parte da configuração ssh.

Remova o símbolo (#) de comentário das 3 linhas marcadas acima.

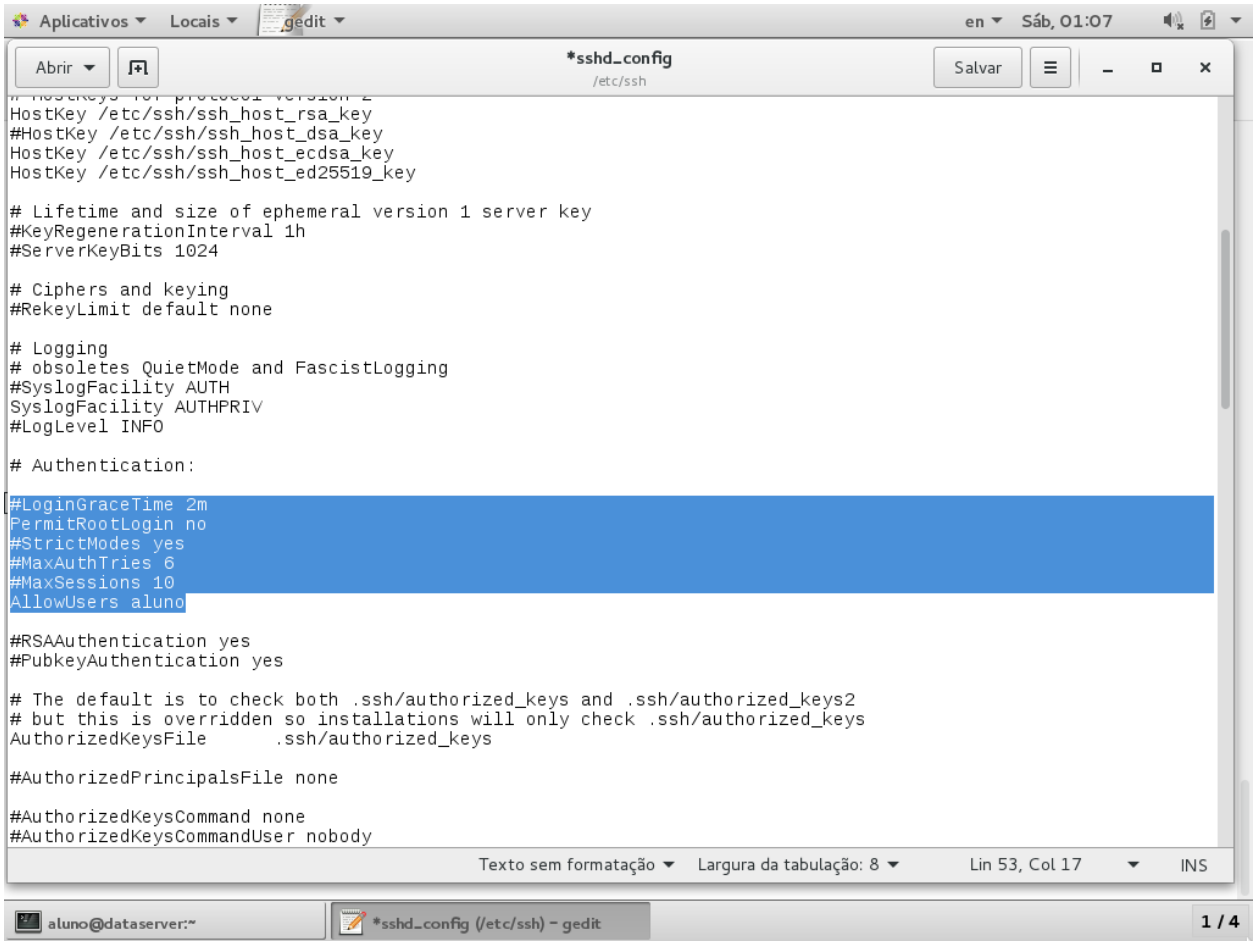

Segunda parte da configuração do ssh

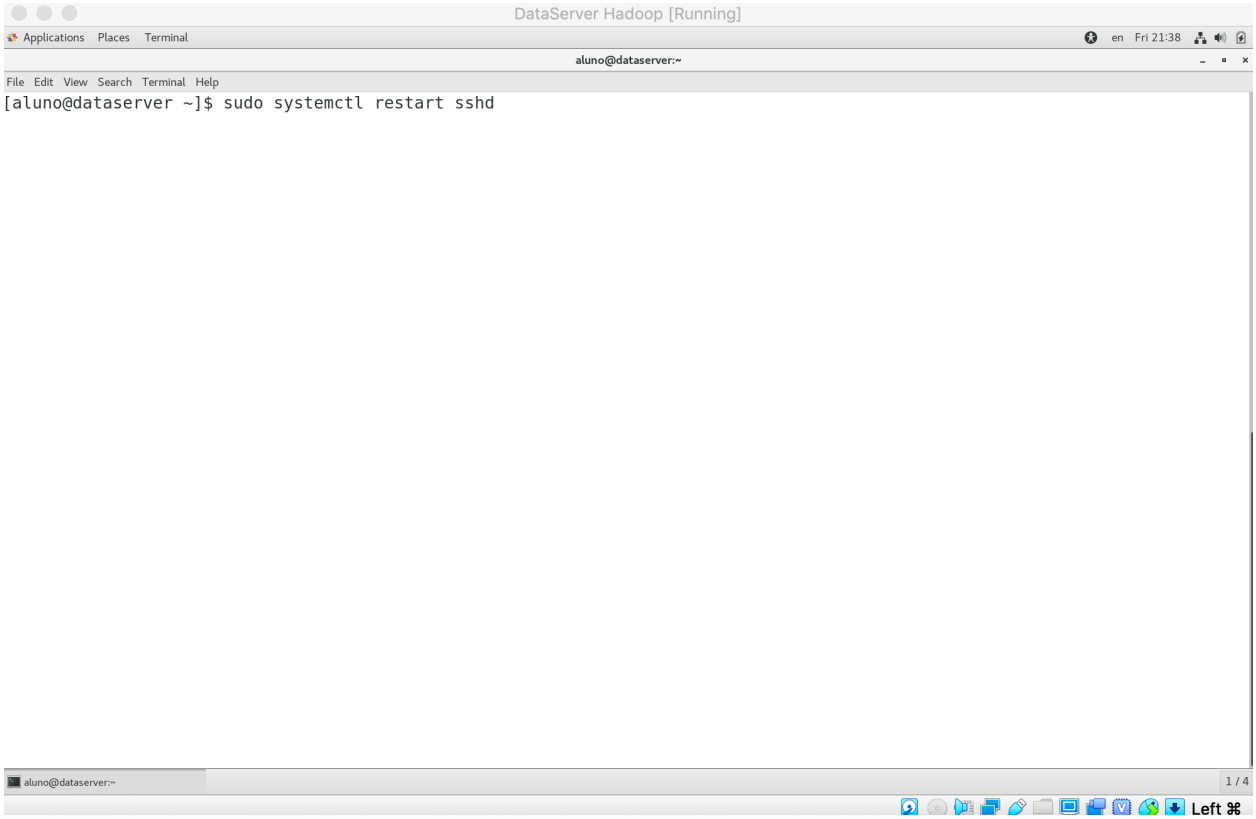

sudo systemctl restart sshd

4. Instalação do Java 8

4.1. Removendo o OpenJDK

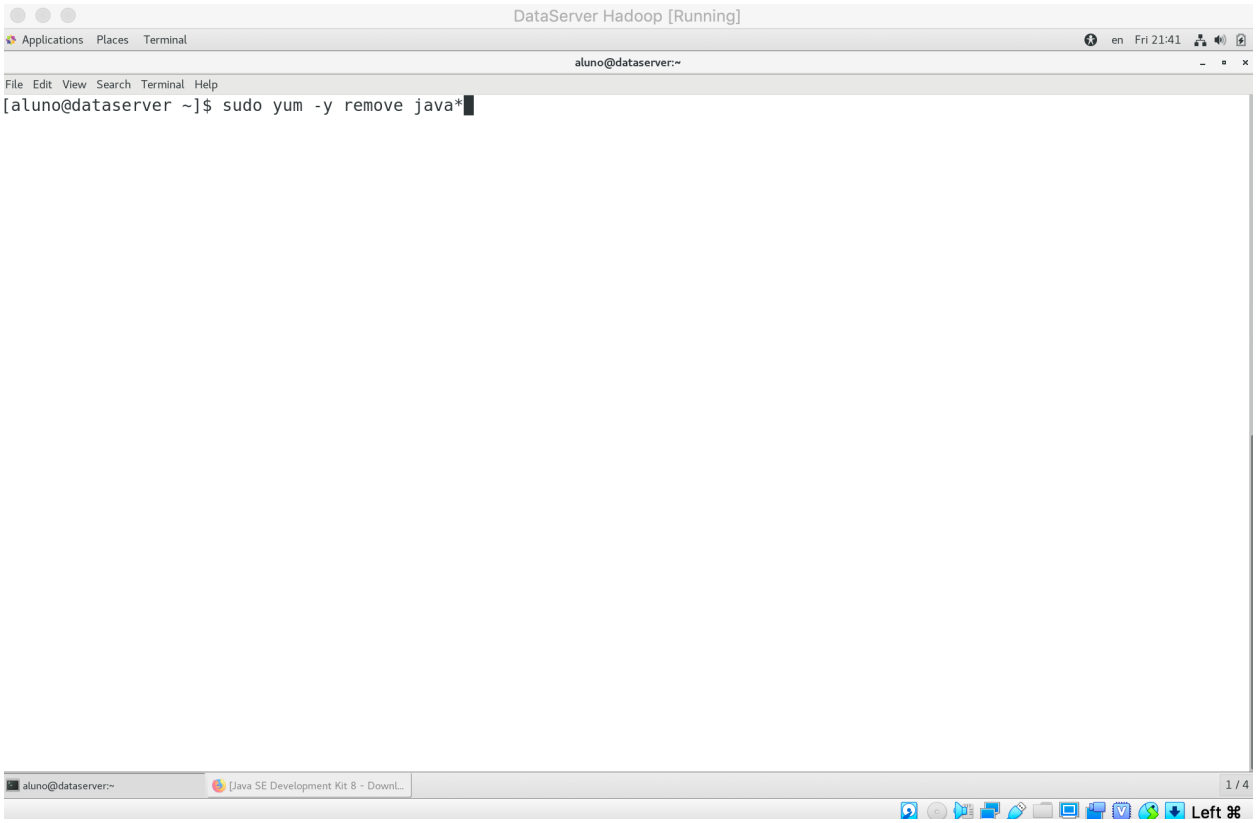

Removendo o OpenJDK

sudo yum -y remove java*

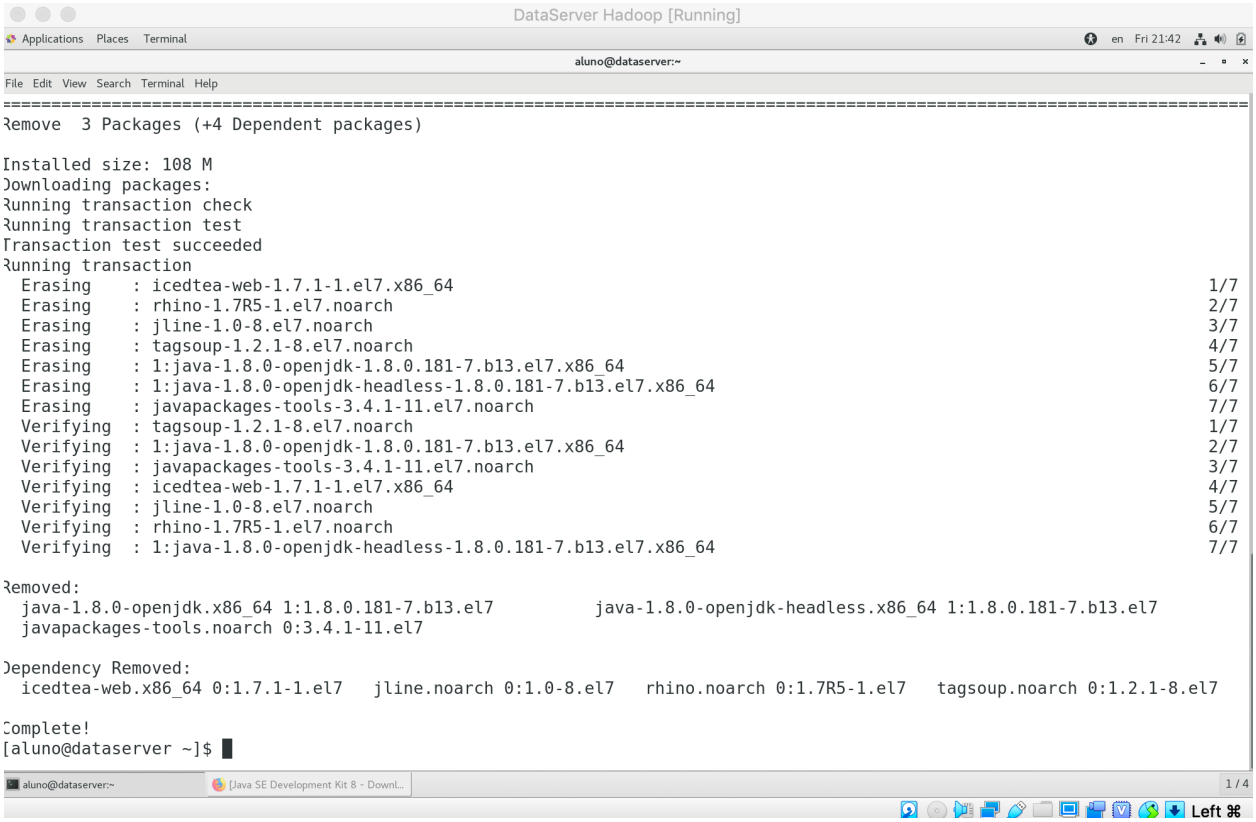

Concluído

Agora acesse o site da Oracle e faça download do Java JDK 1.8 para Linux

4.2. Instalação do JDK

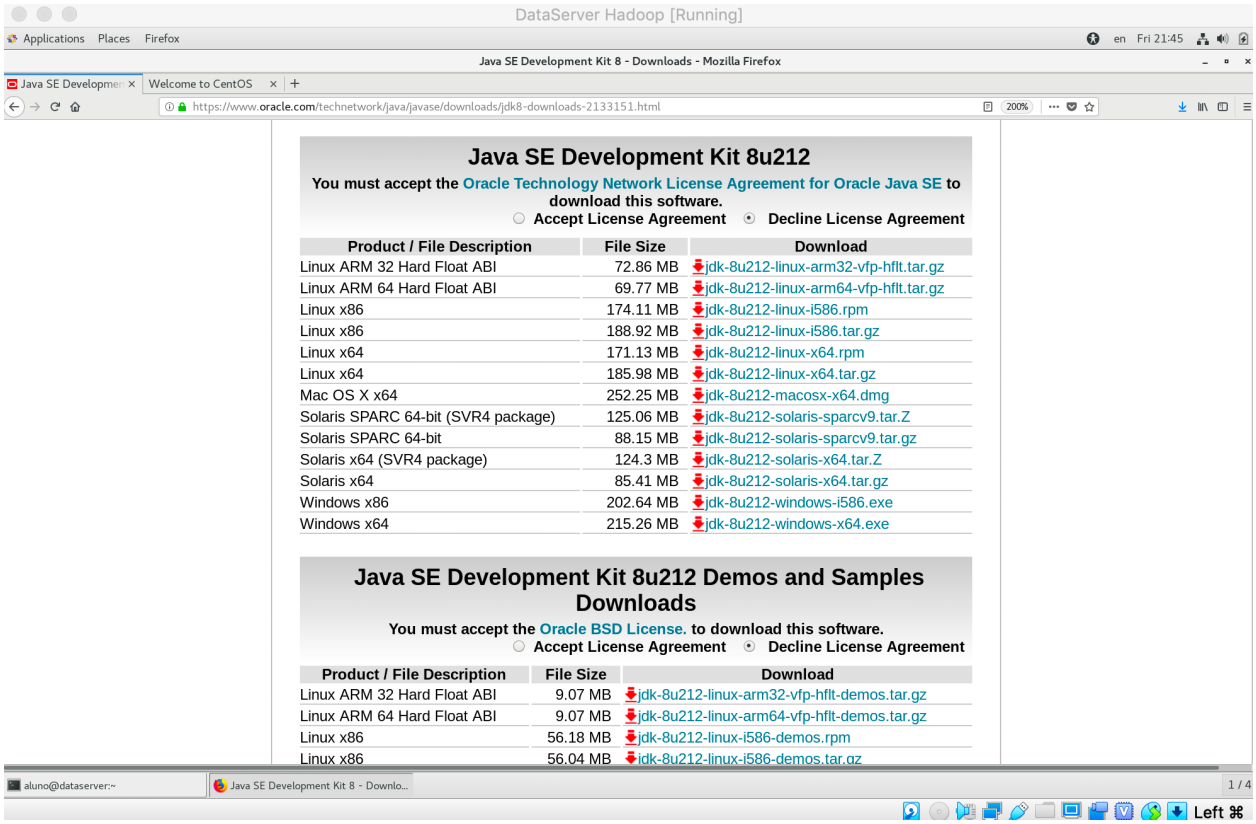

No site da Oracle, fazer o download do JDK

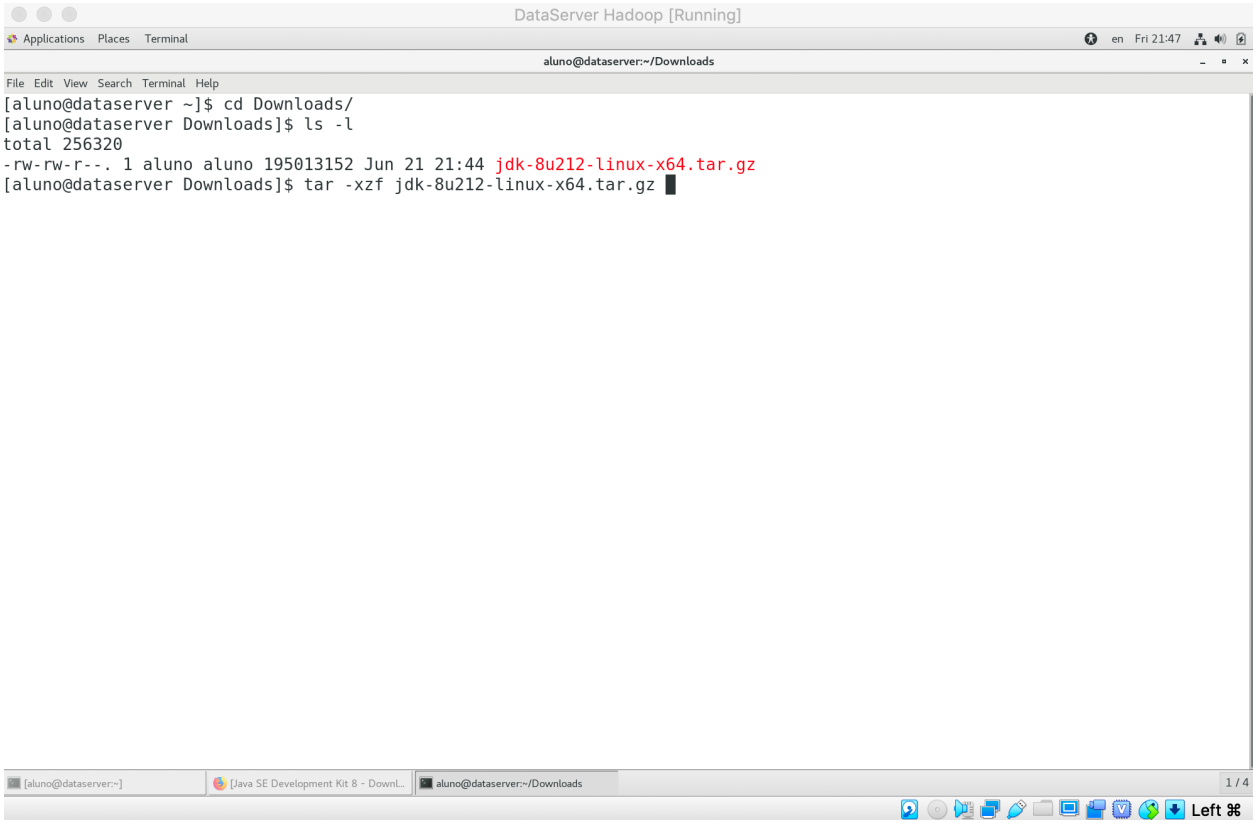

Executar o comando tar para descompactar o arquivo: **tar -xzf jdk-8u212-linux-x64.tar.gz**

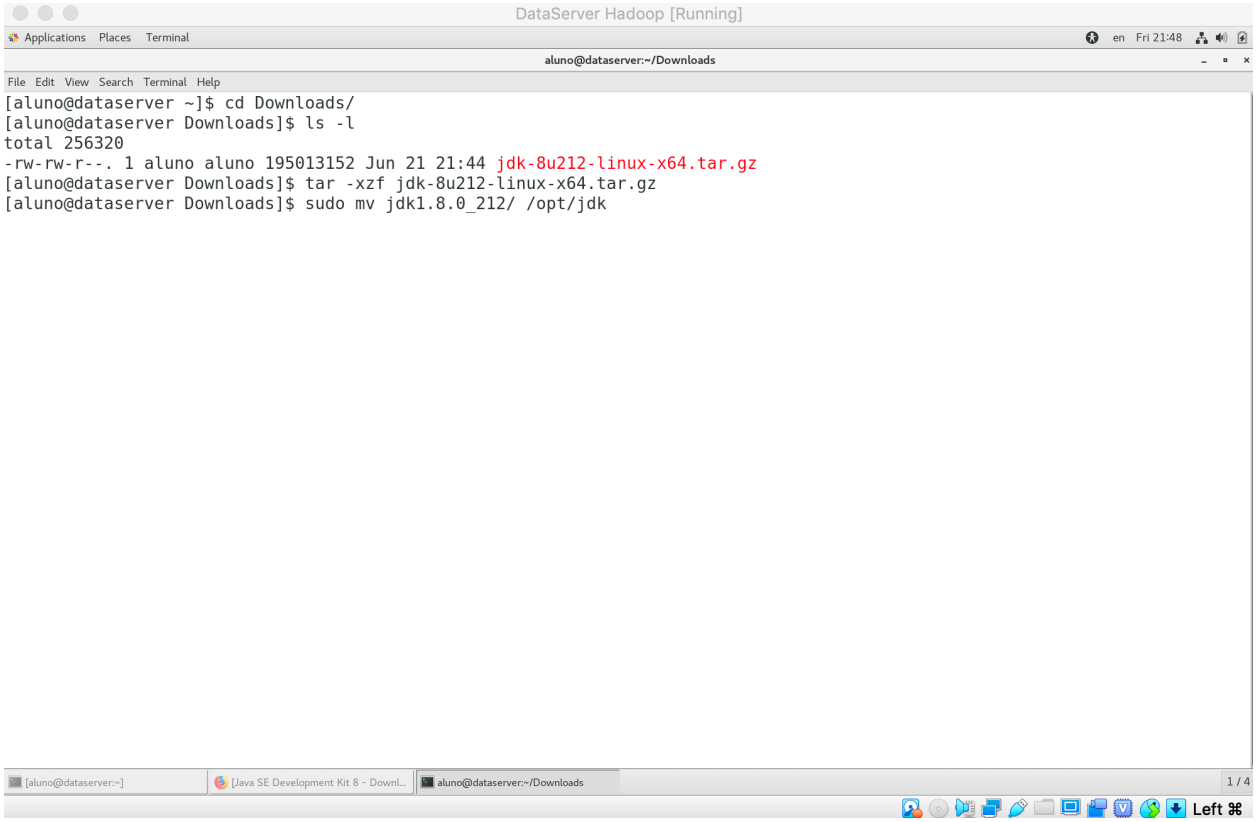

Mover o diretório do JDK

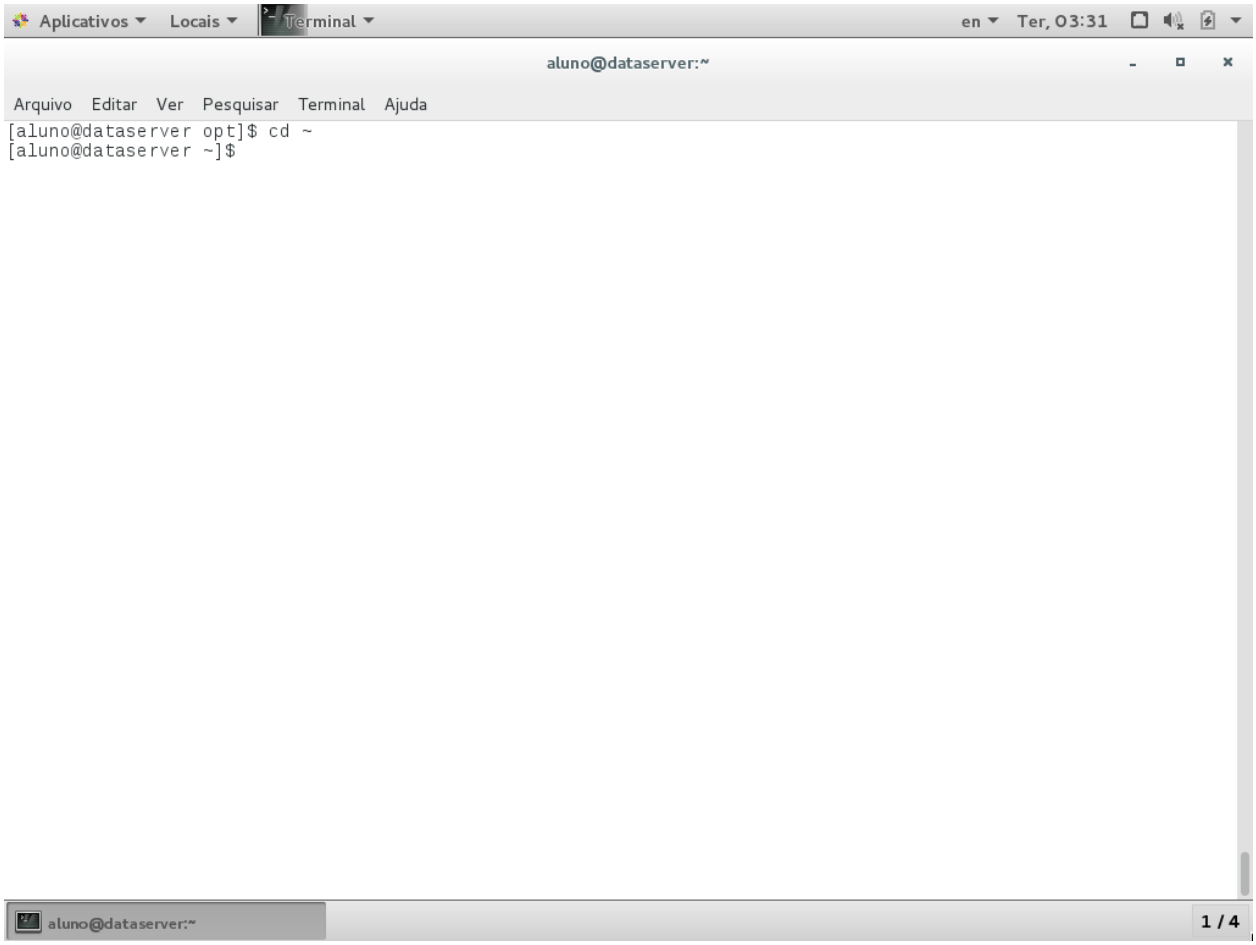
Aplicativos
Locais
Terminal
en
Ter, 03:31
aluno@dataserver:~
Arquivo Editar Ver Pesquisar Terminal Ajuda
[aluno@dataserver opt]$ cd ~
[aluno@dataserver ~]$
aluno@dataserver:~
1 / 4

cd ~

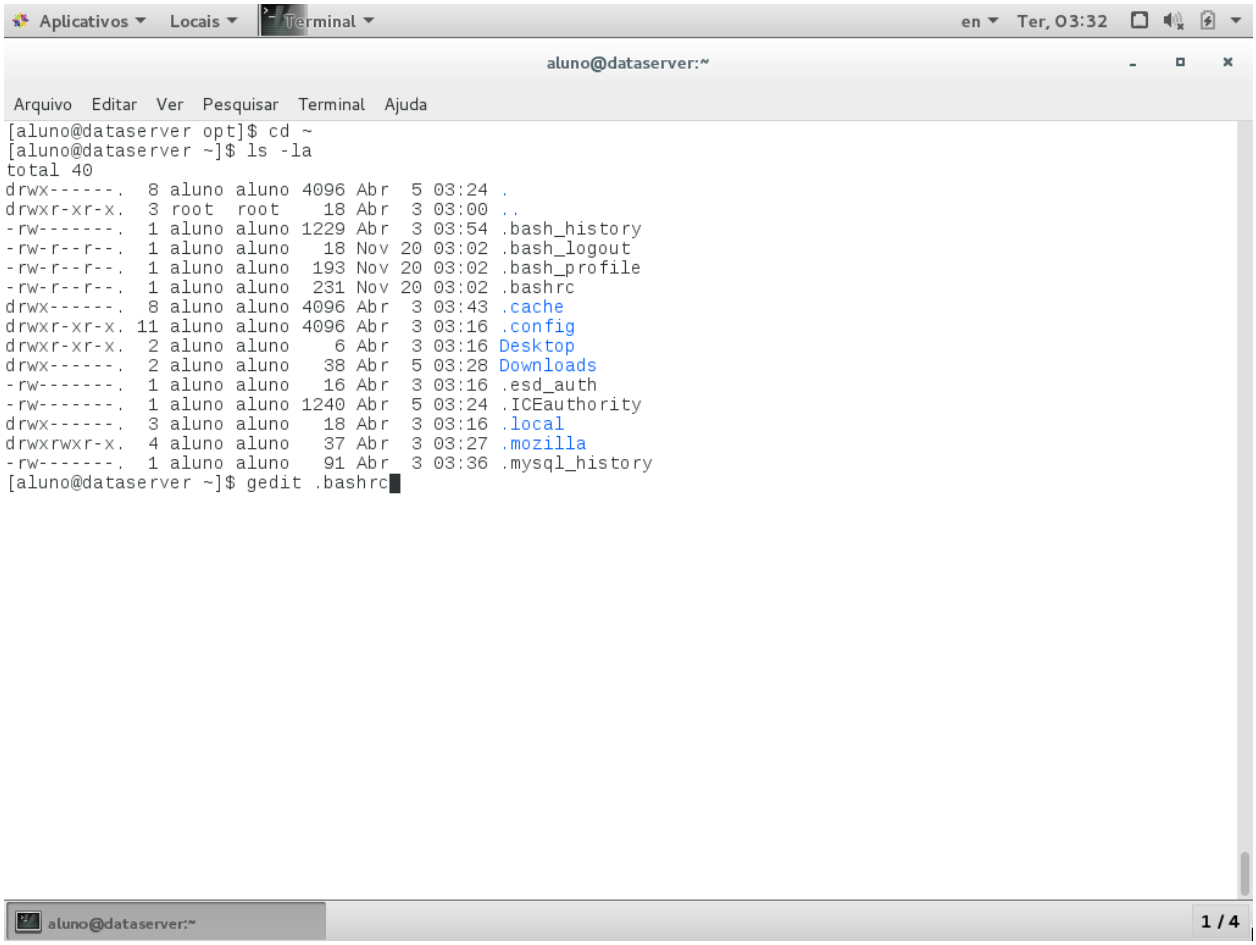

gedit .bashrc

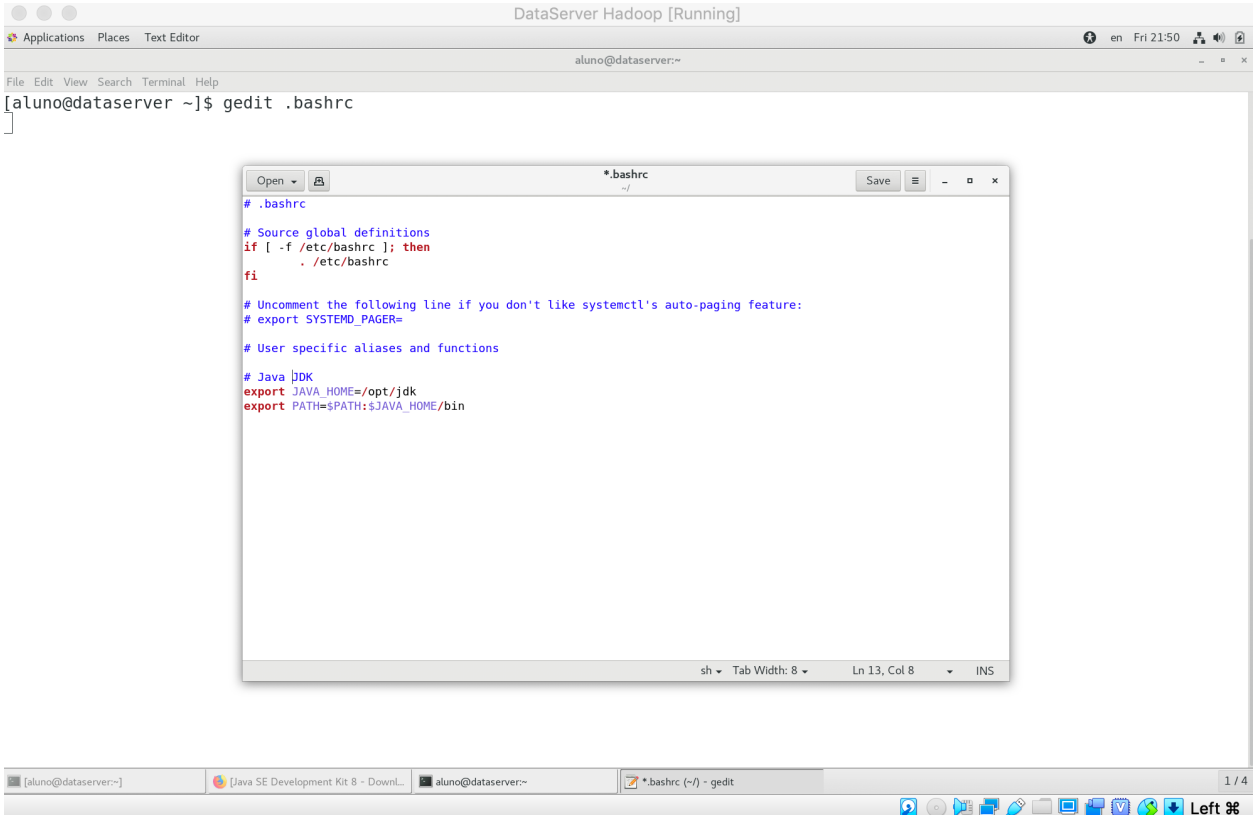

Editar as variáveis de ambiente conforme acima e salvar o arquivo

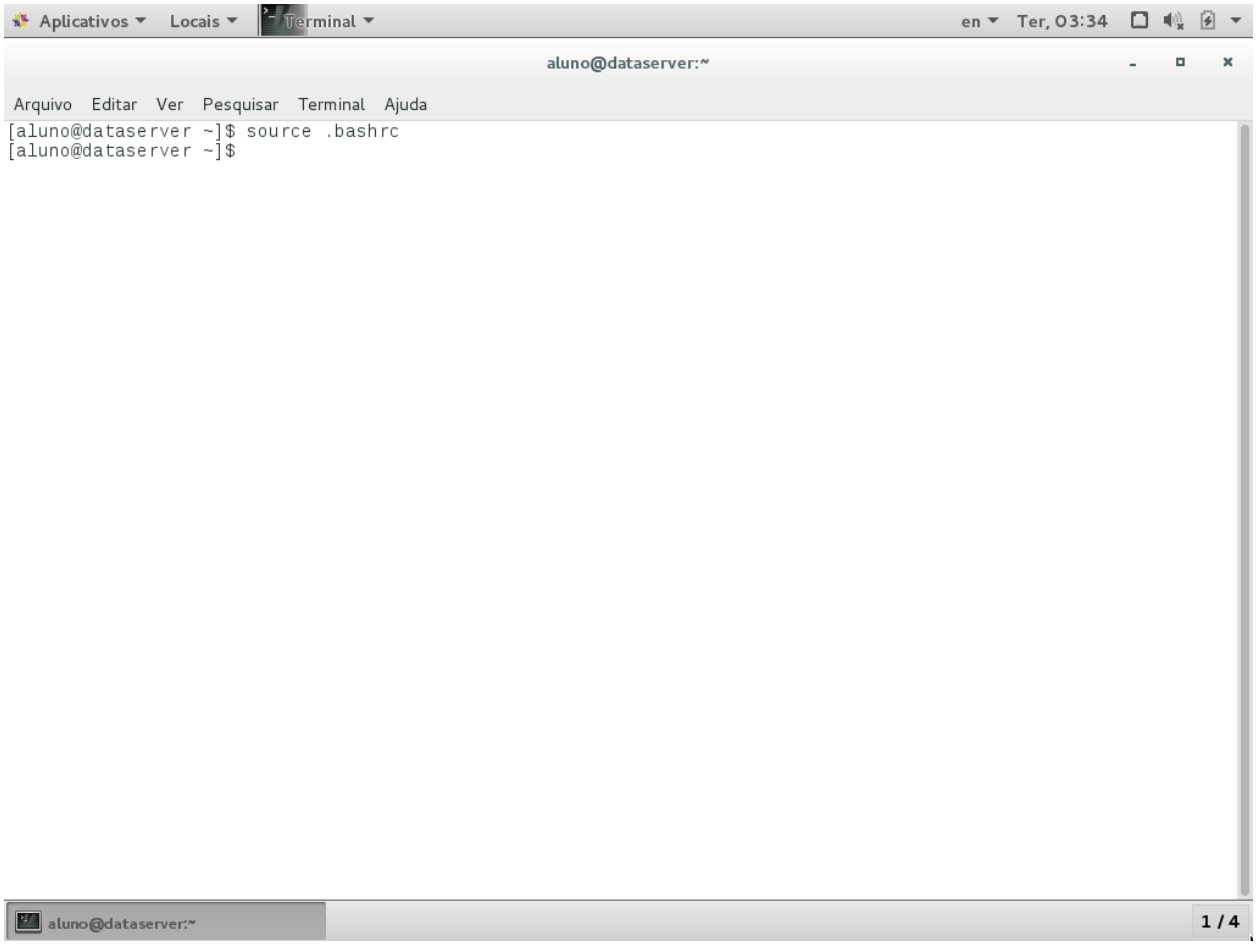

source .bashrc

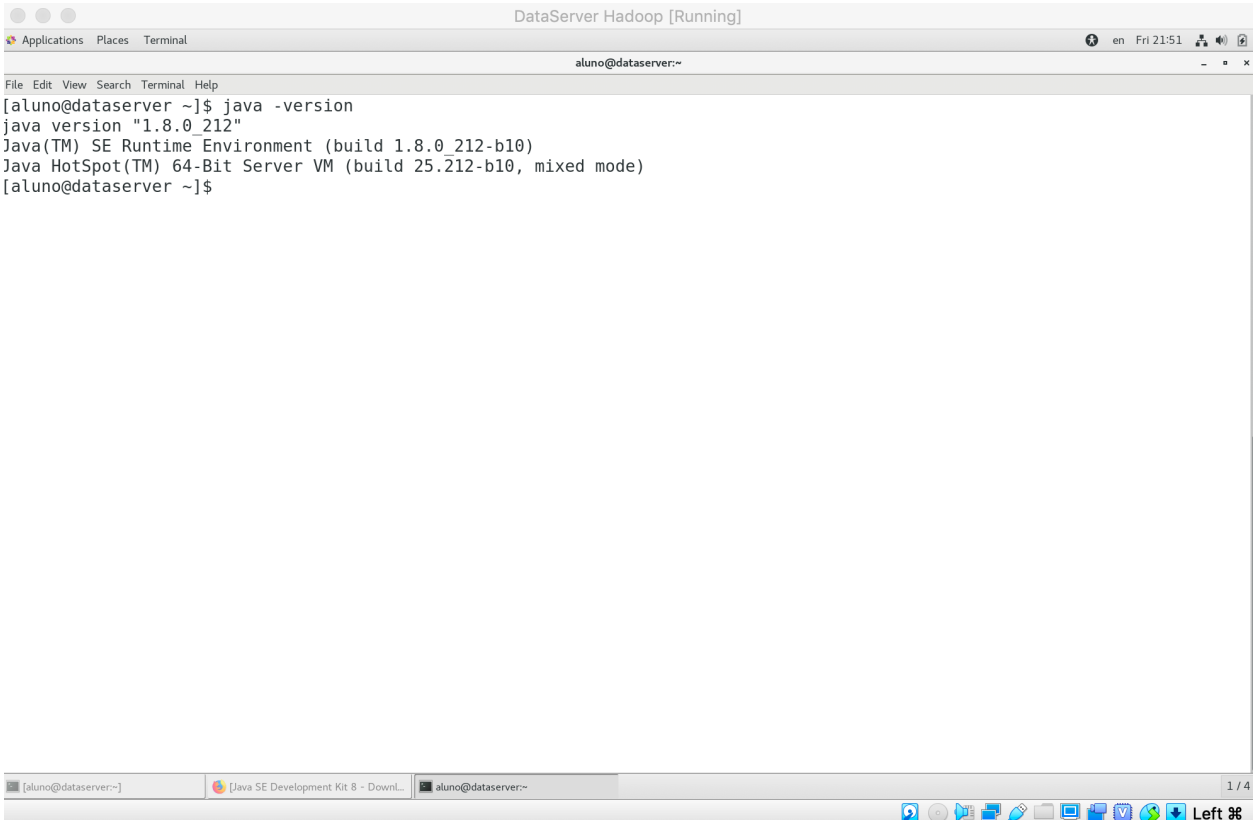

Checando a versão do Java JDK

Segundo checkpoint:

Desligue a VM (clique no ícone da bateria e então em desligar).
Clique no menu File do VirtualBox e clique em Export Appliance.
Será gerada uma cópia de segurança da sua máquina virtual.

→ VM: DataServer-Hadoop-v2.0.ova

5. Instalação e Configuração do Hadoop

5.1. Criando o usuário hadoop

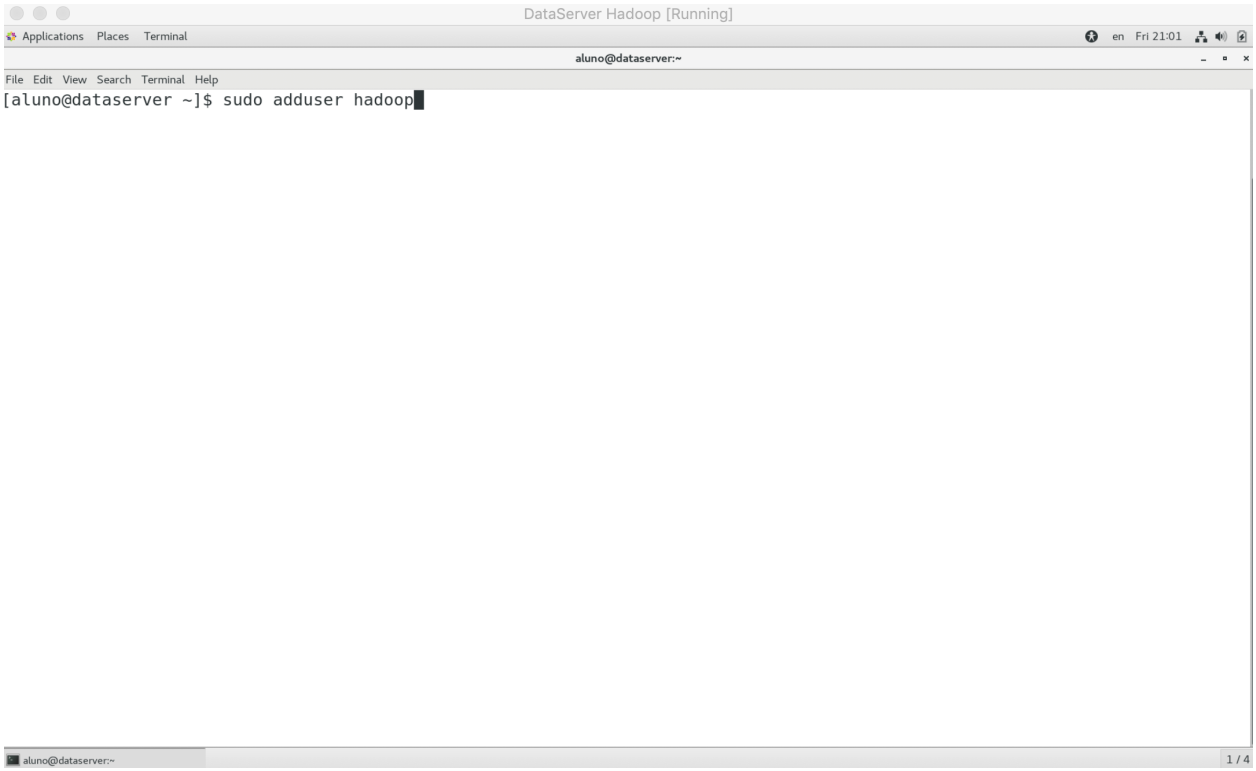

sudo adduser hadoop – para criar o usuário hadoop

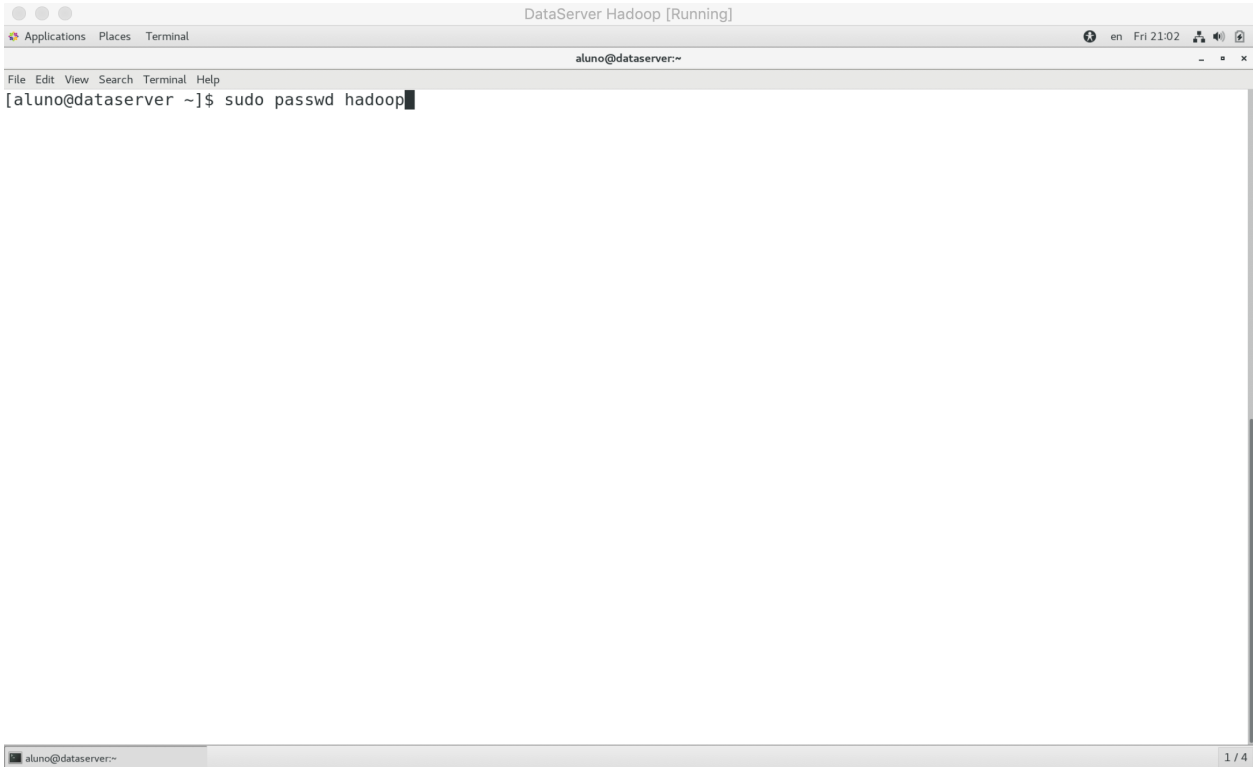

sudo passwd hadoop – para definir a senha do usuário hadoop (**dsahadoop**)

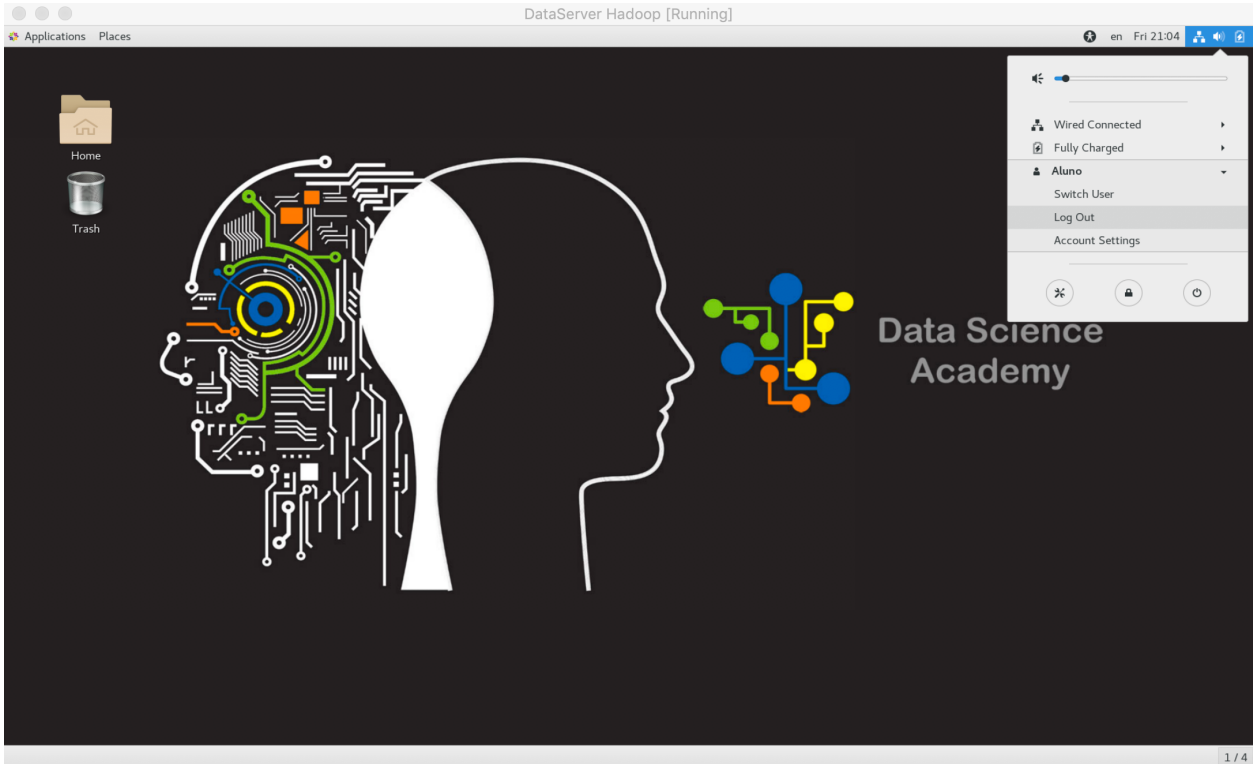

Efetue logout como usuário aluno

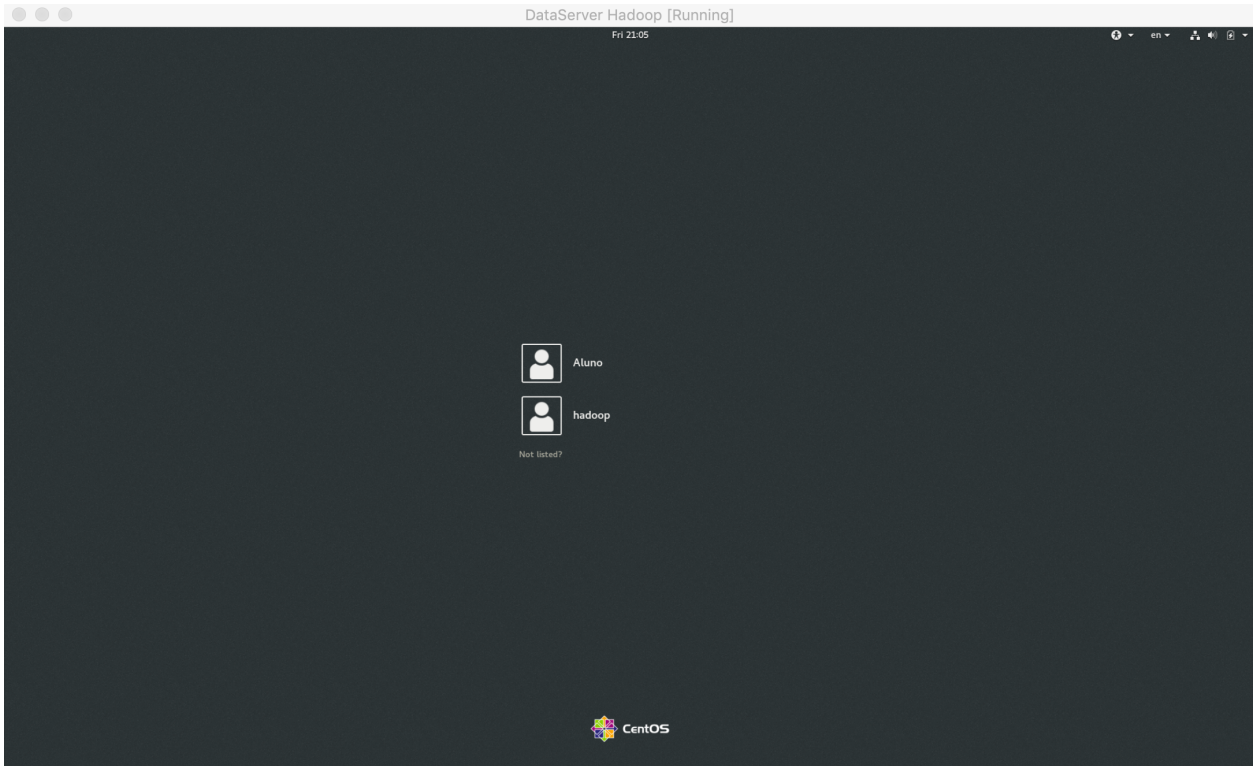

E efetue login como usuário hadoop

Adicione o usuário hadoop no arquivo /etc/sudoers conforme você fez com o usuário aluno

5.2. Configuração do ssh sem senha

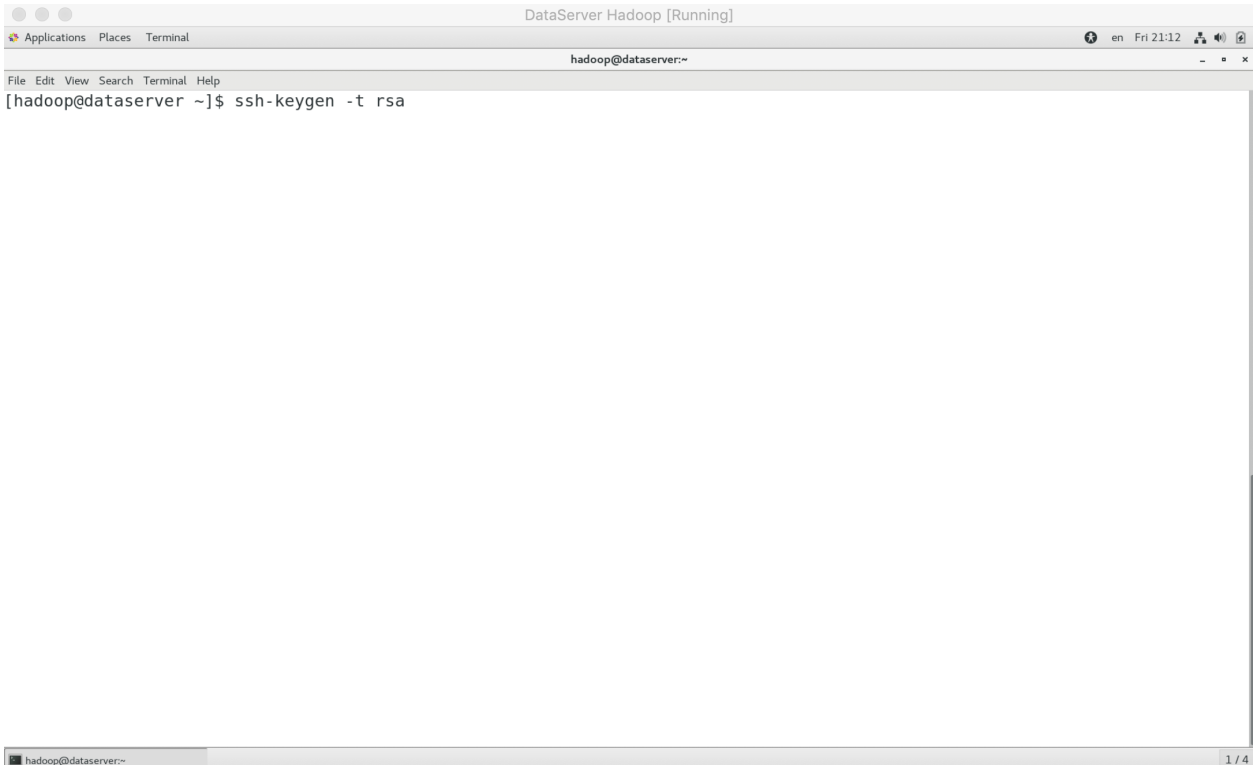

ssh-keygen –t rsa

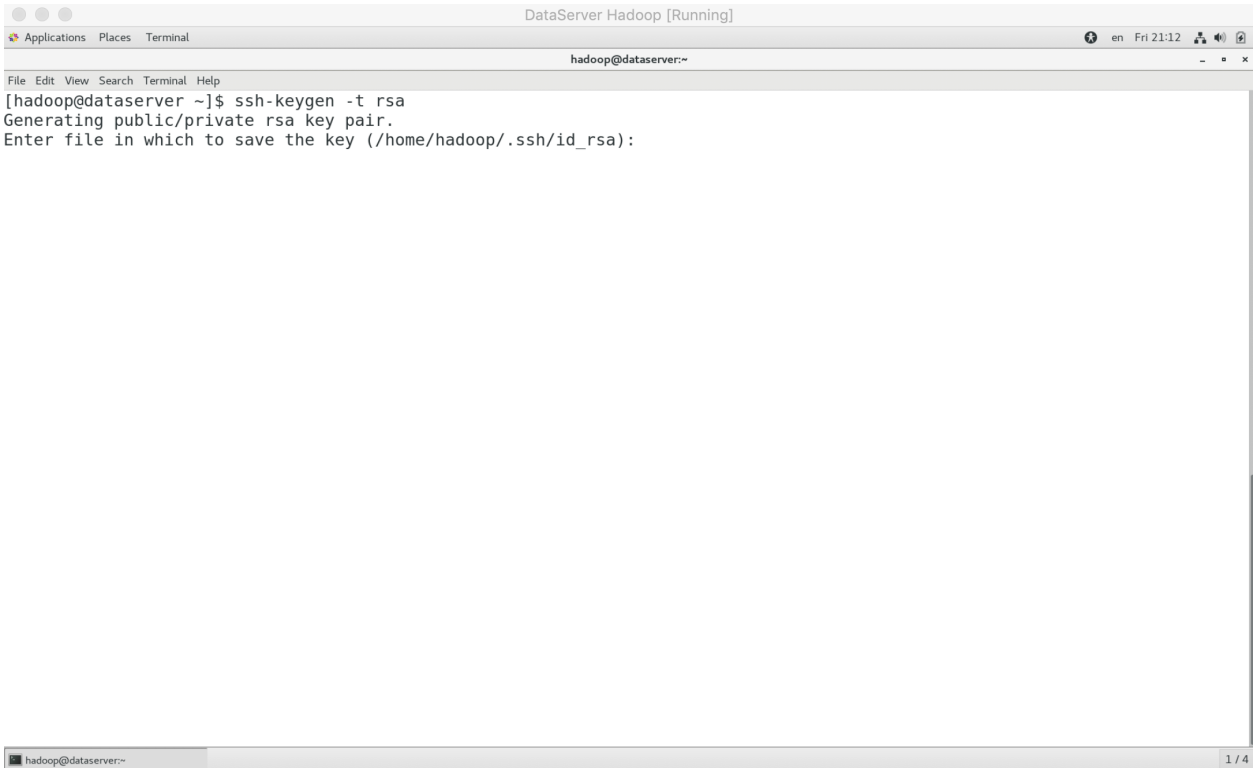

Pressionar Enter para confirmar o diretório onde as chaves serão geradas

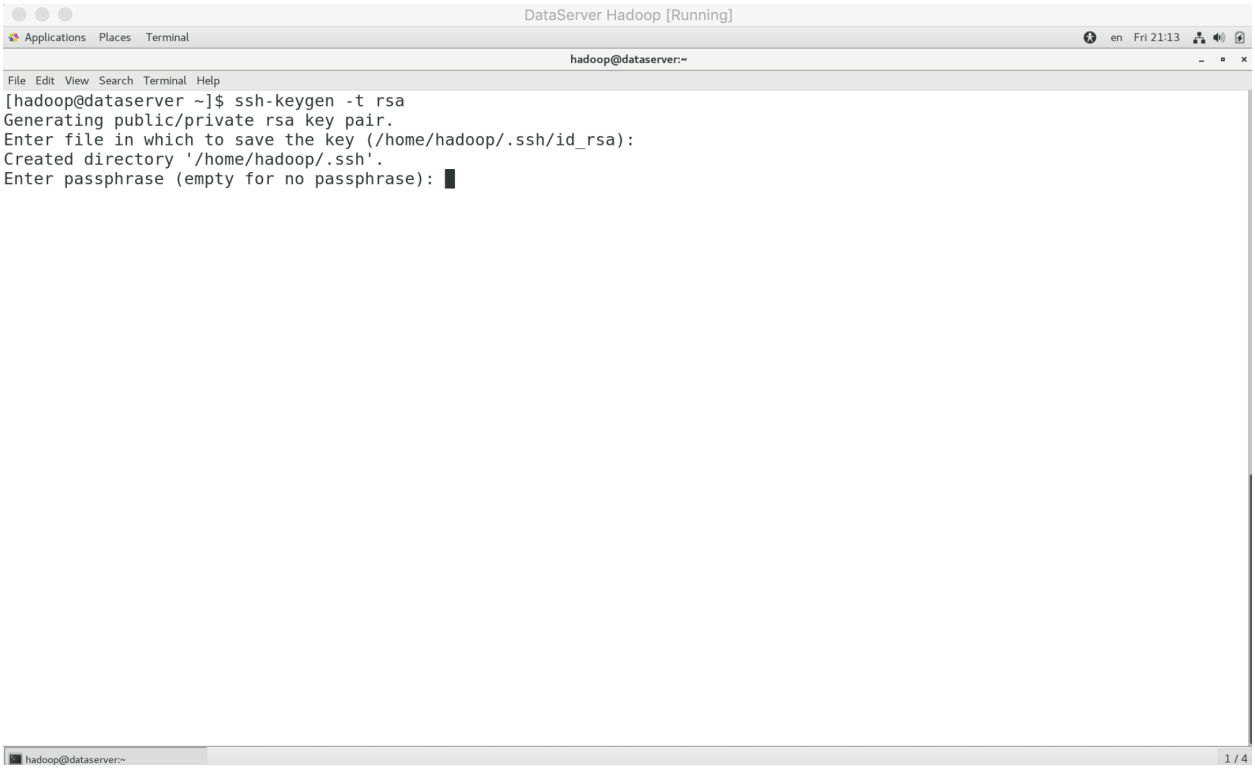

Pressionar Enter

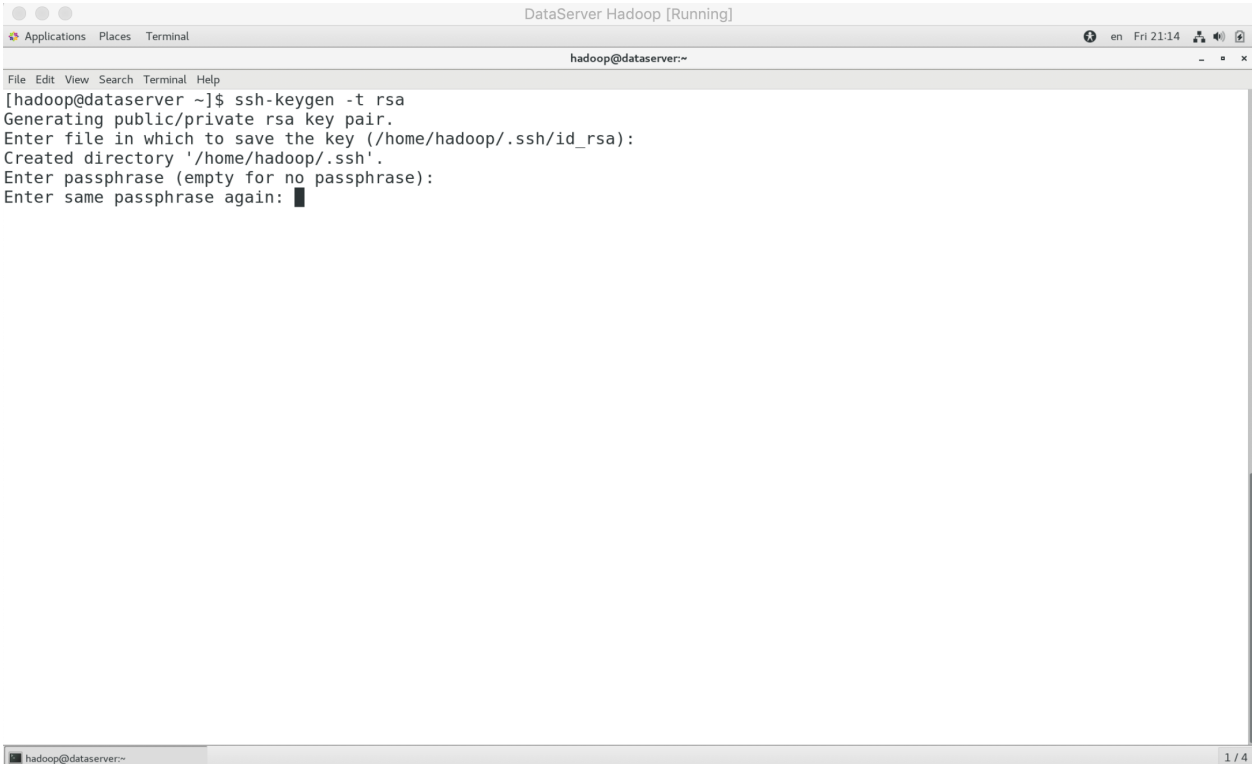

Pressionar Enter novamente

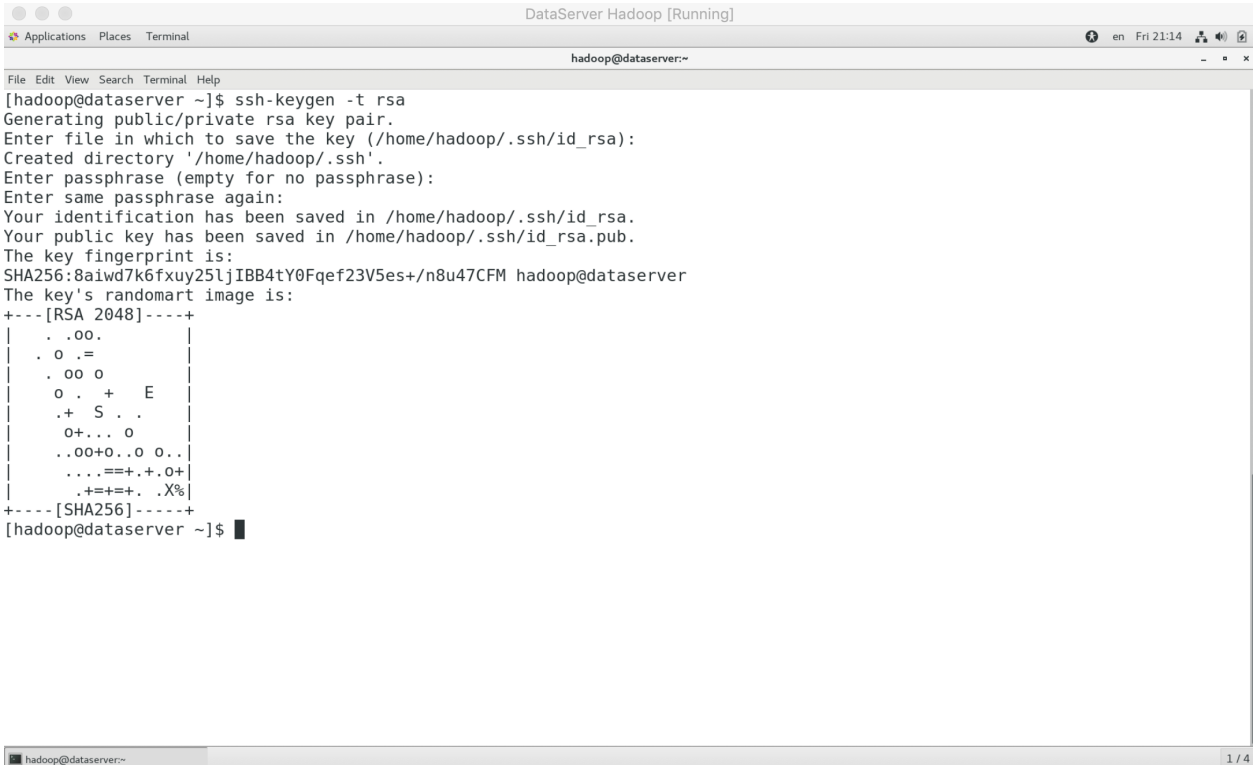

Chaves de segurança geradas

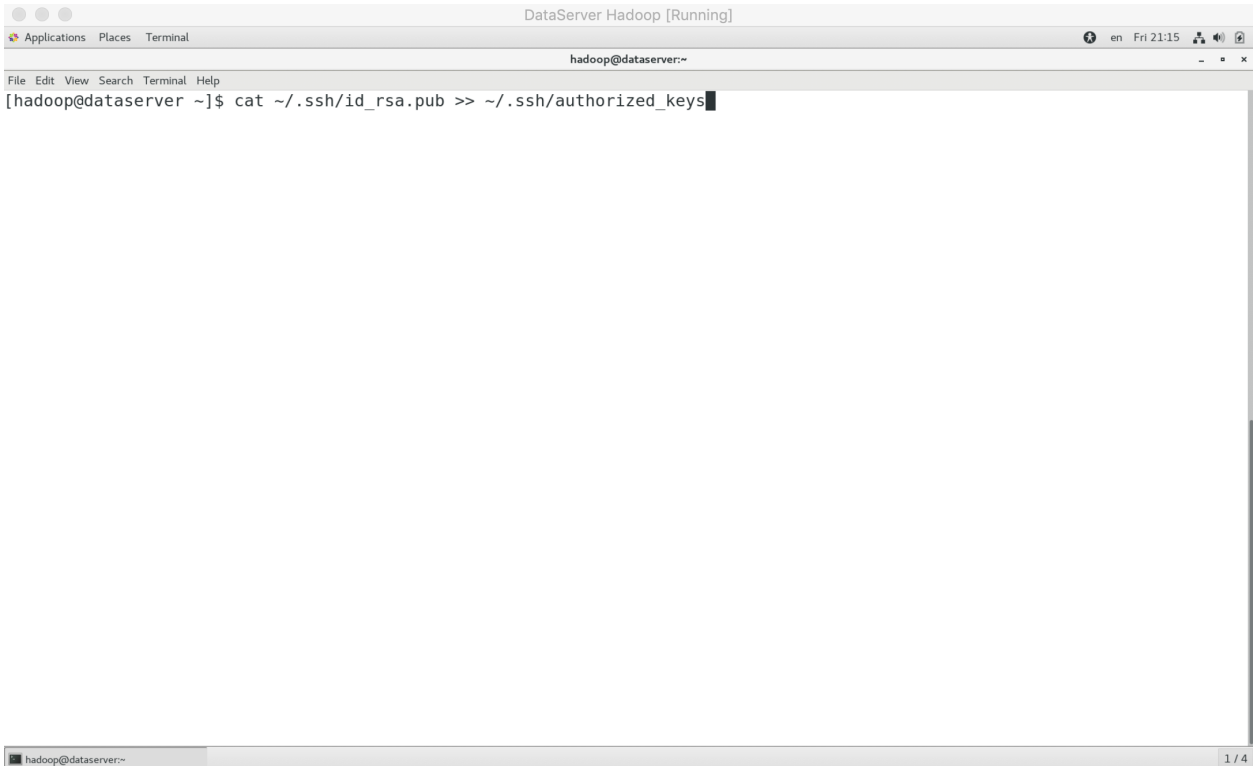

cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys

Esse comando cópia a chave pública para o arquivo authorized_keys do ssh

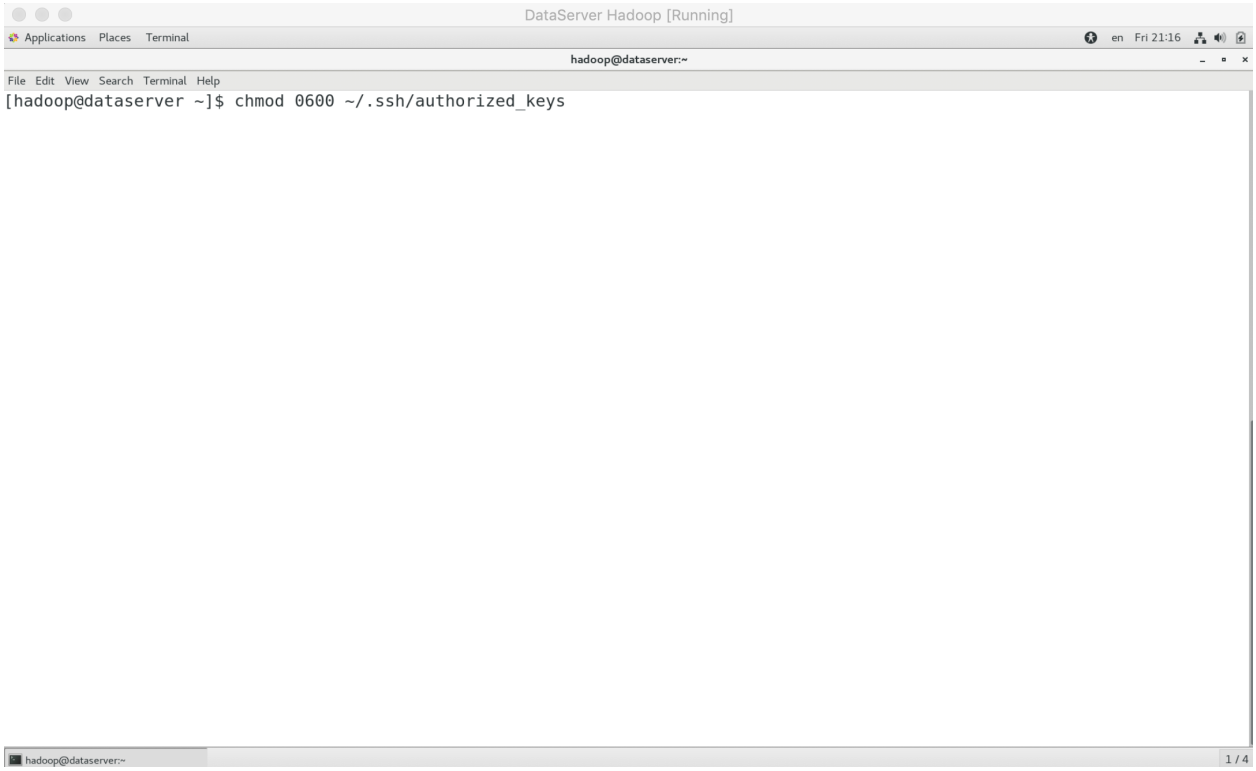

chmod 0600 ~/.ssh/authorized_keys

Esse comando define a permissão do arquivo authorized_keys

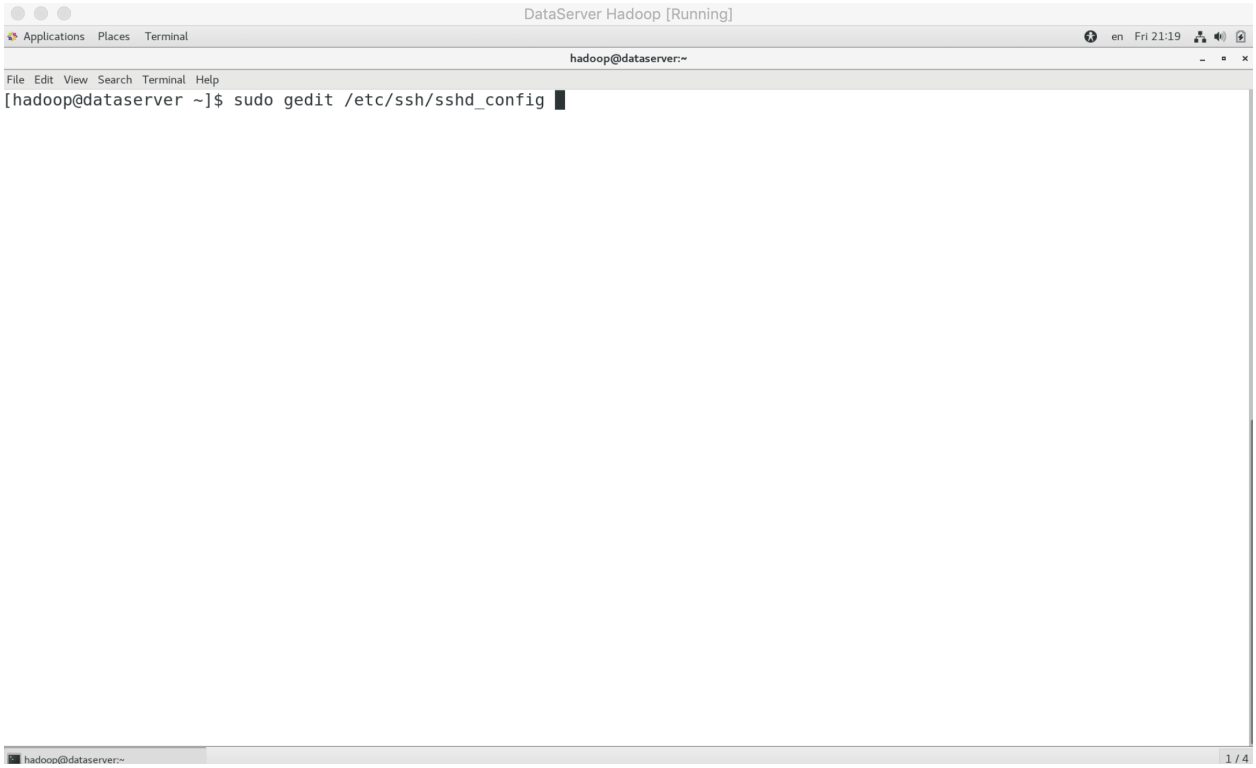

sudo gedit /etc/ssh/sshd_config

Edite o arquivo de configuração do ssh

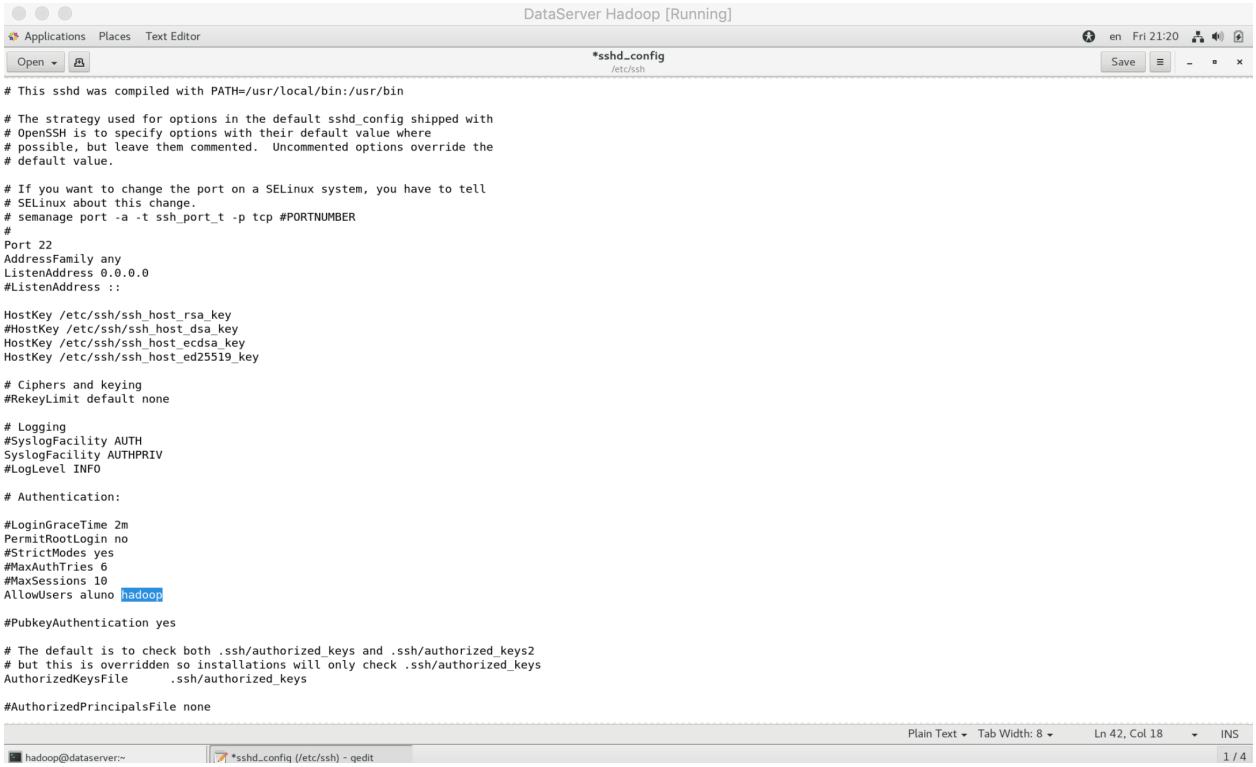

Inclua o usuário hadoop na linha AllowUsers, salve o arquivo e feche-o

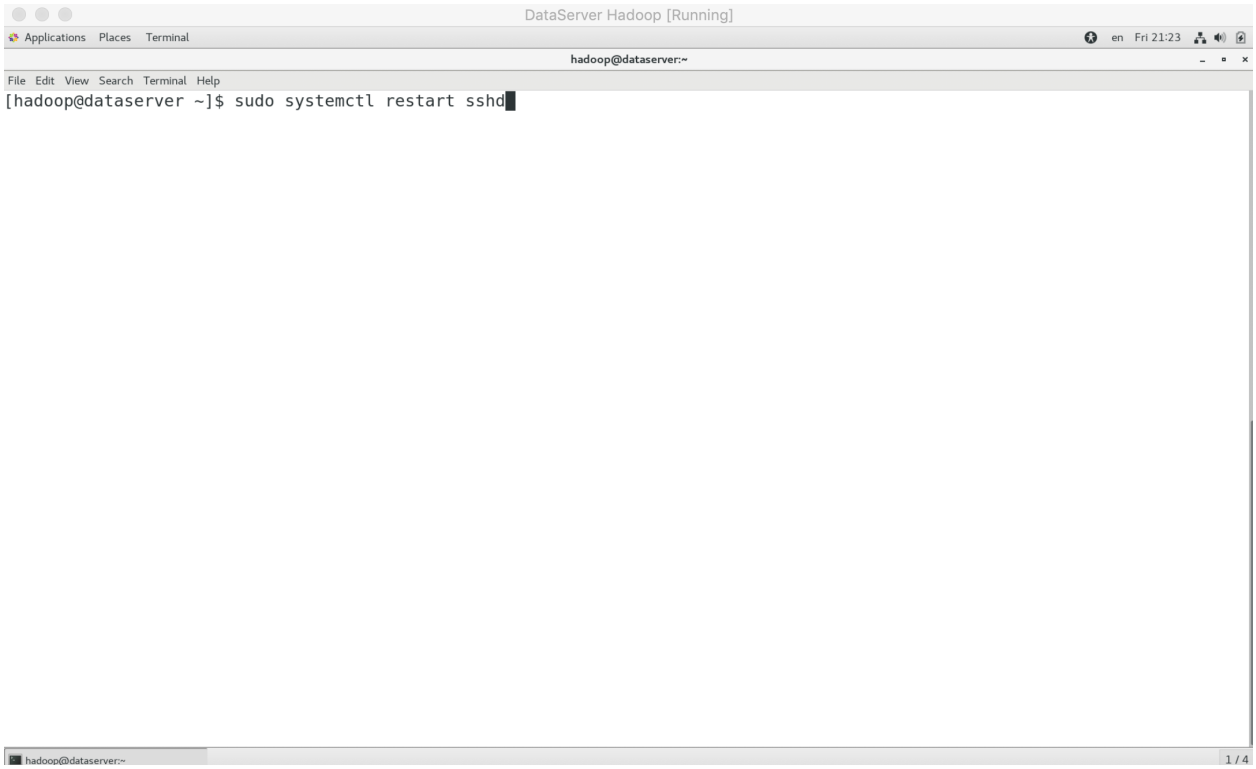

sudo systemctl restart sshd

Reinicie o serviço ssh

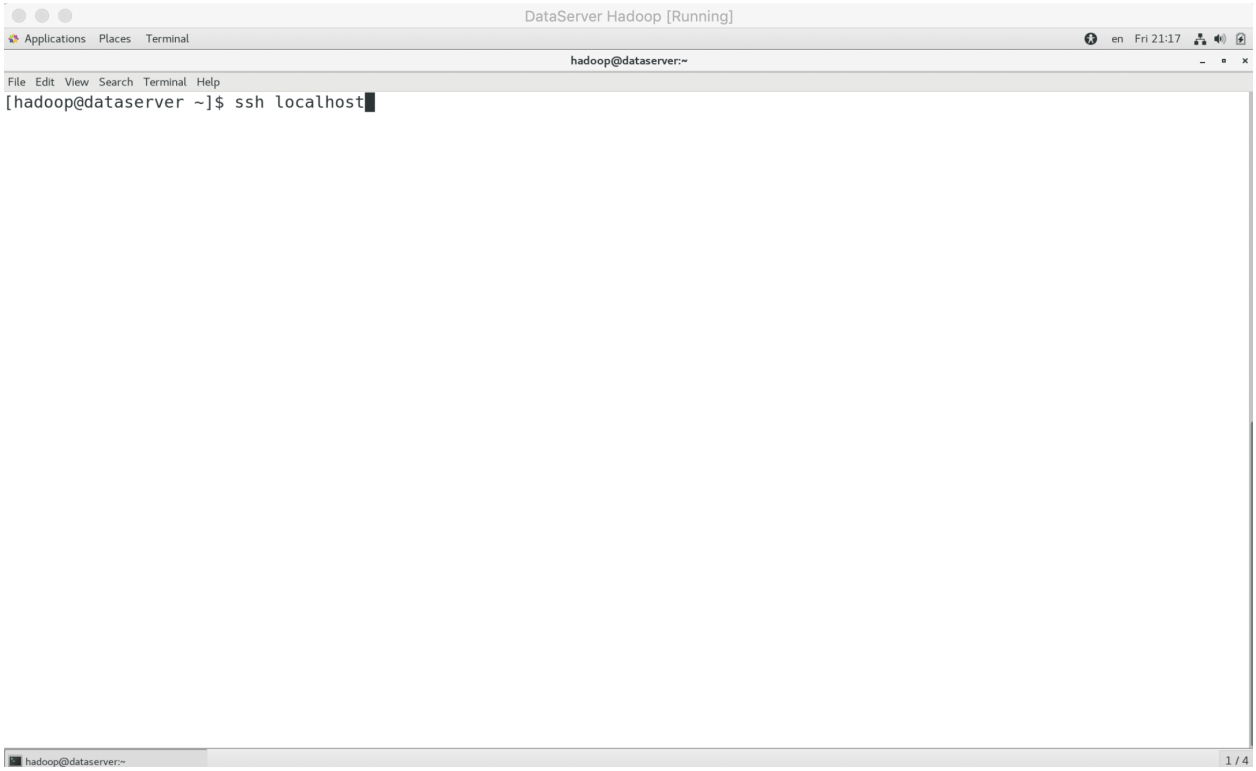

ssh localhost

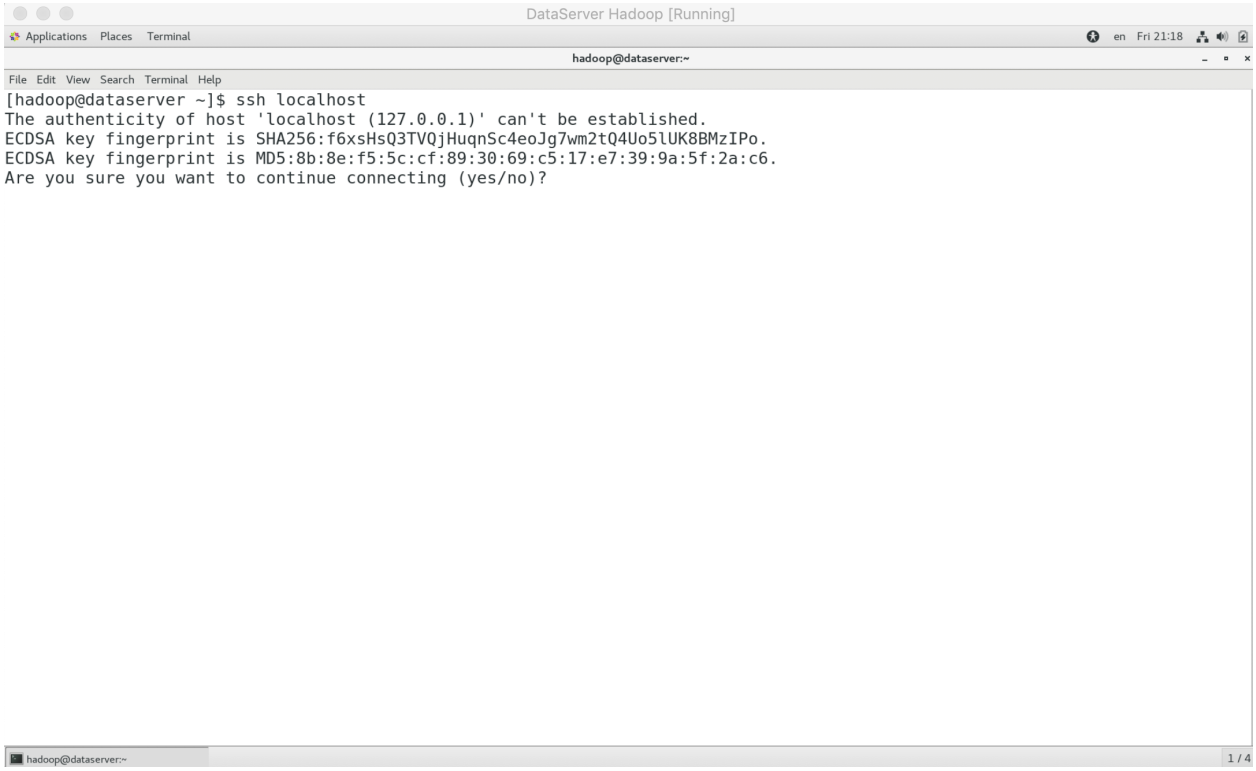

Yes

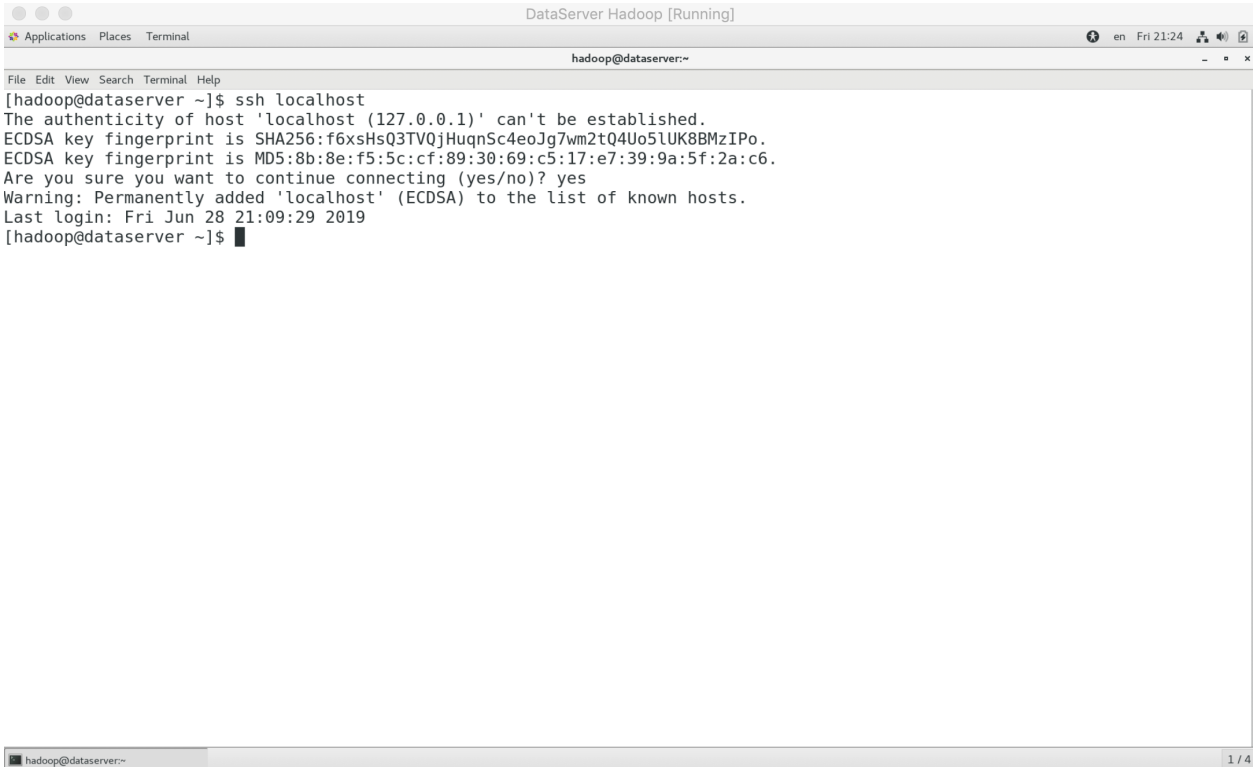

Conexão ssh sem senha efetuada com sucesso. Digite exit e pressione Enter.

Parabéns, seu ambiente está pronto para receber o Hadoop!!

5.3. Download e Instalação do Hadoop

5.3.1. Editando o arquivo hosts

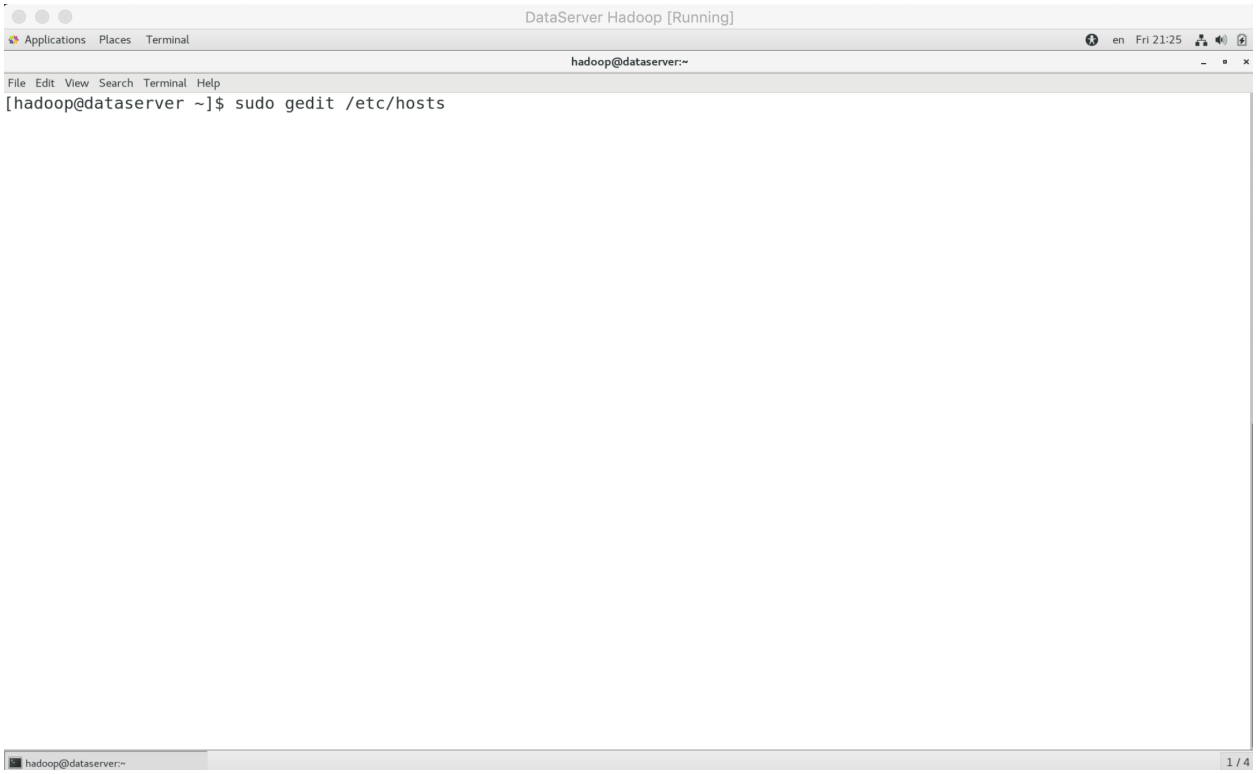

Editar o arquivo hosts

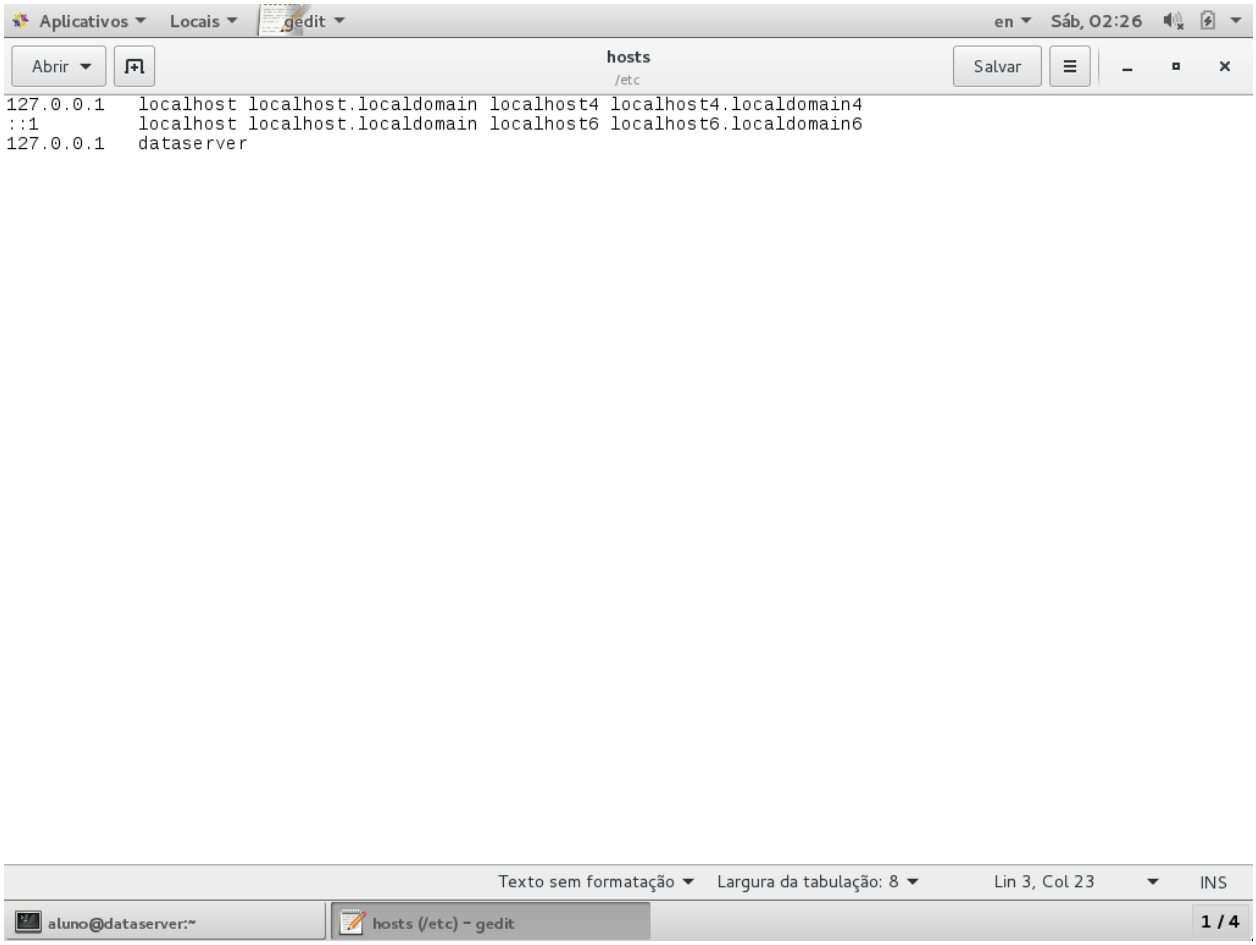

Incluir a última linha conforme acima

5.3.2. Download do Hadoop

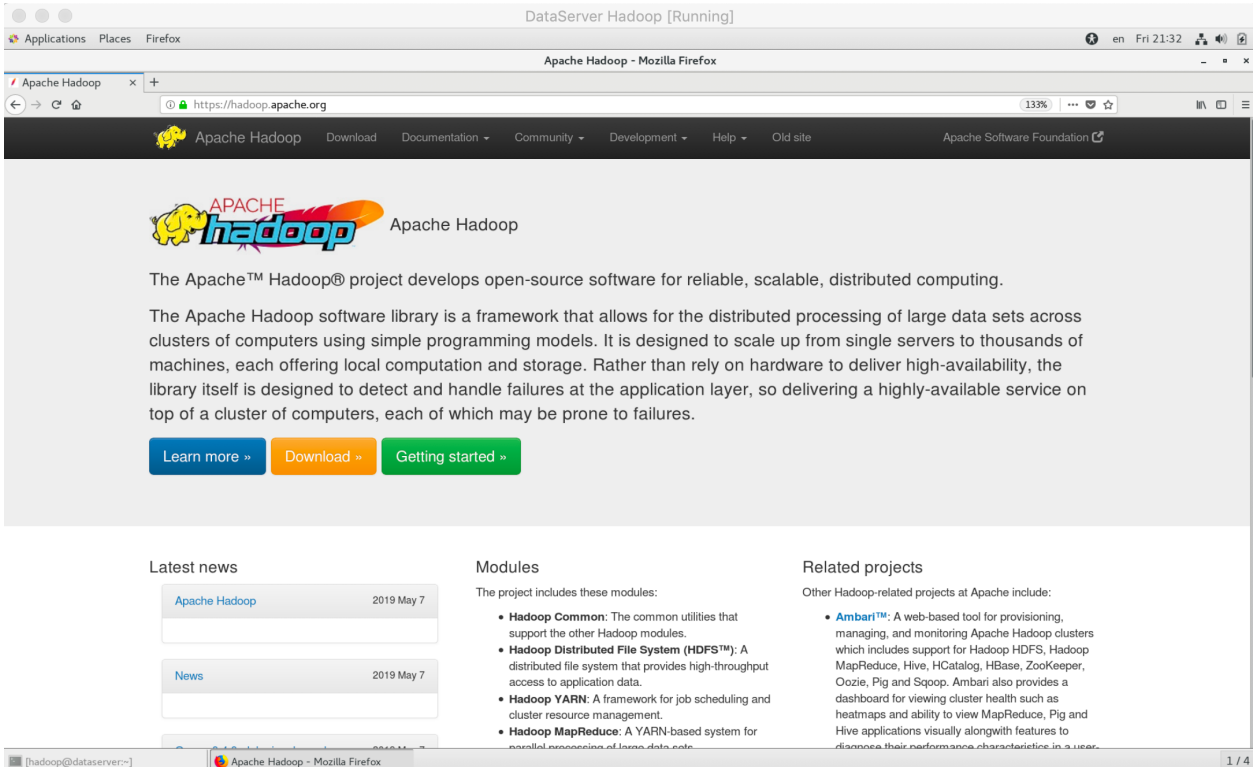

Acesse a página do Hadoop e clique em Download

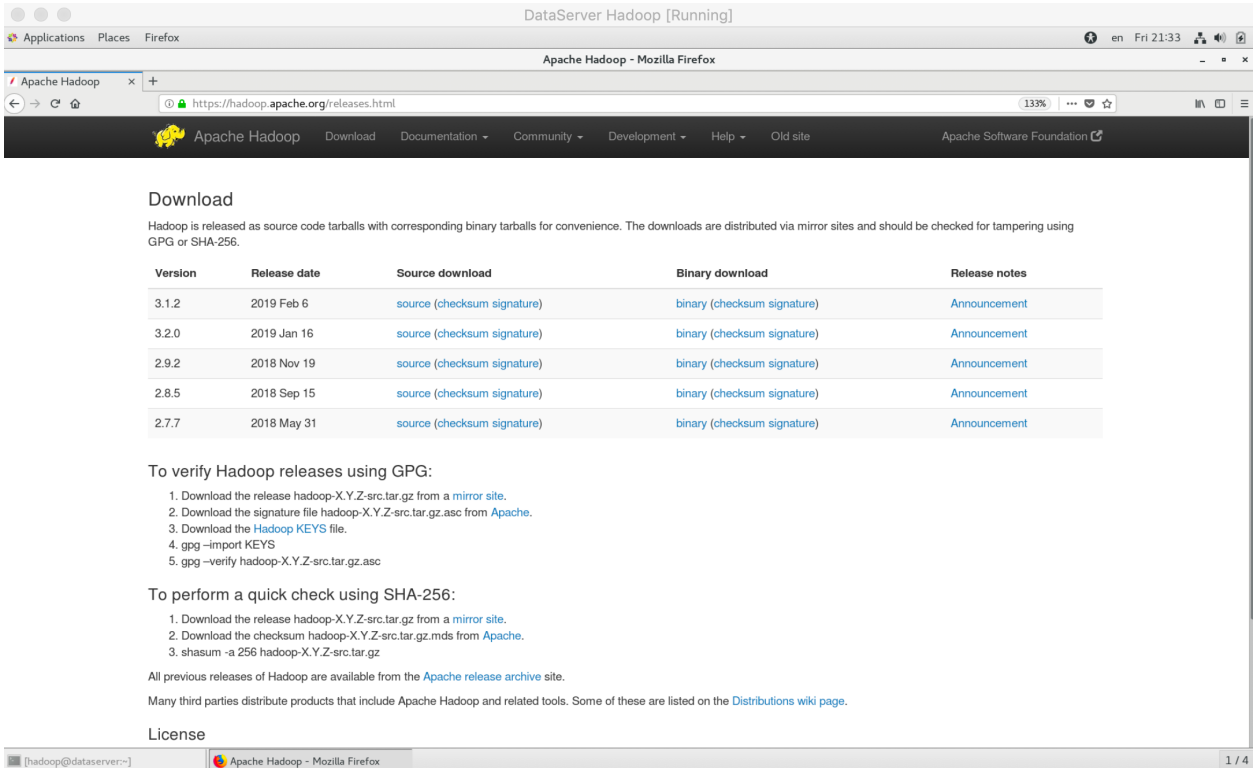

Faça o download da versão 3.2, opção binary.
O arquivo será baixado no diretório /home/hadoop/Downloads

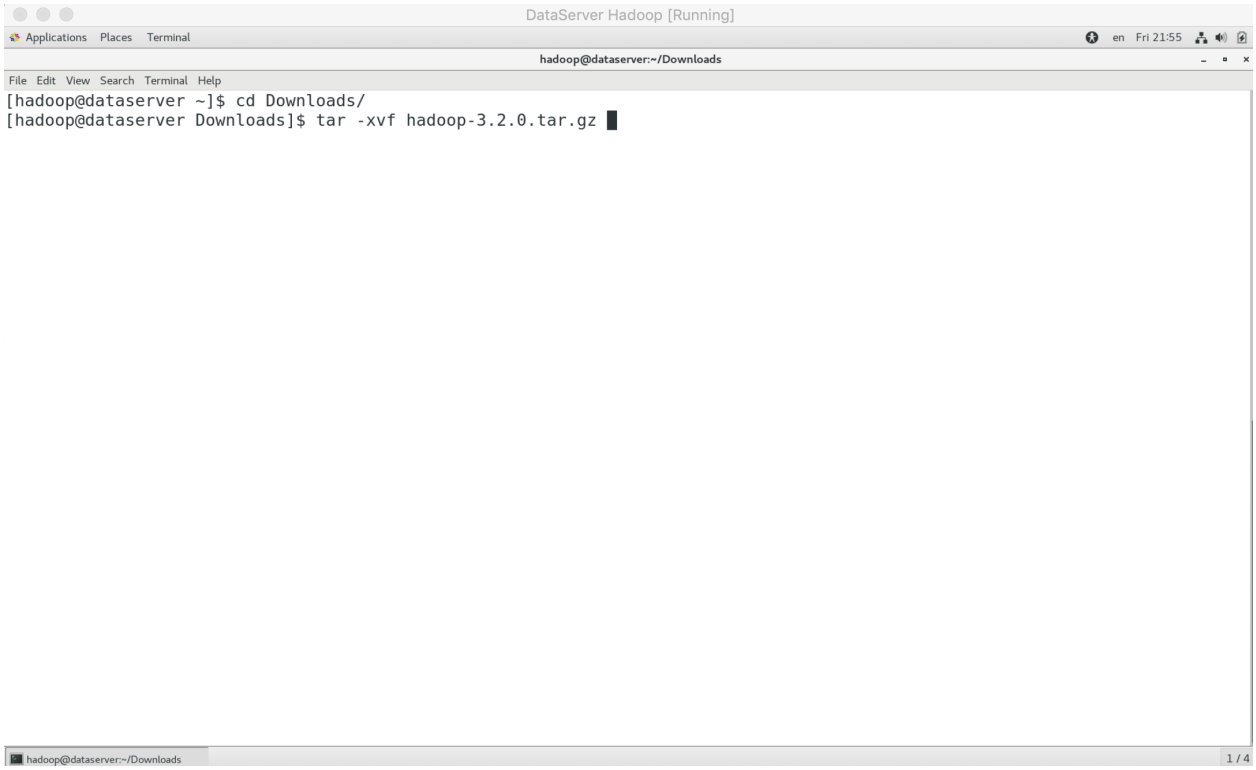

Descompacte o arquivo

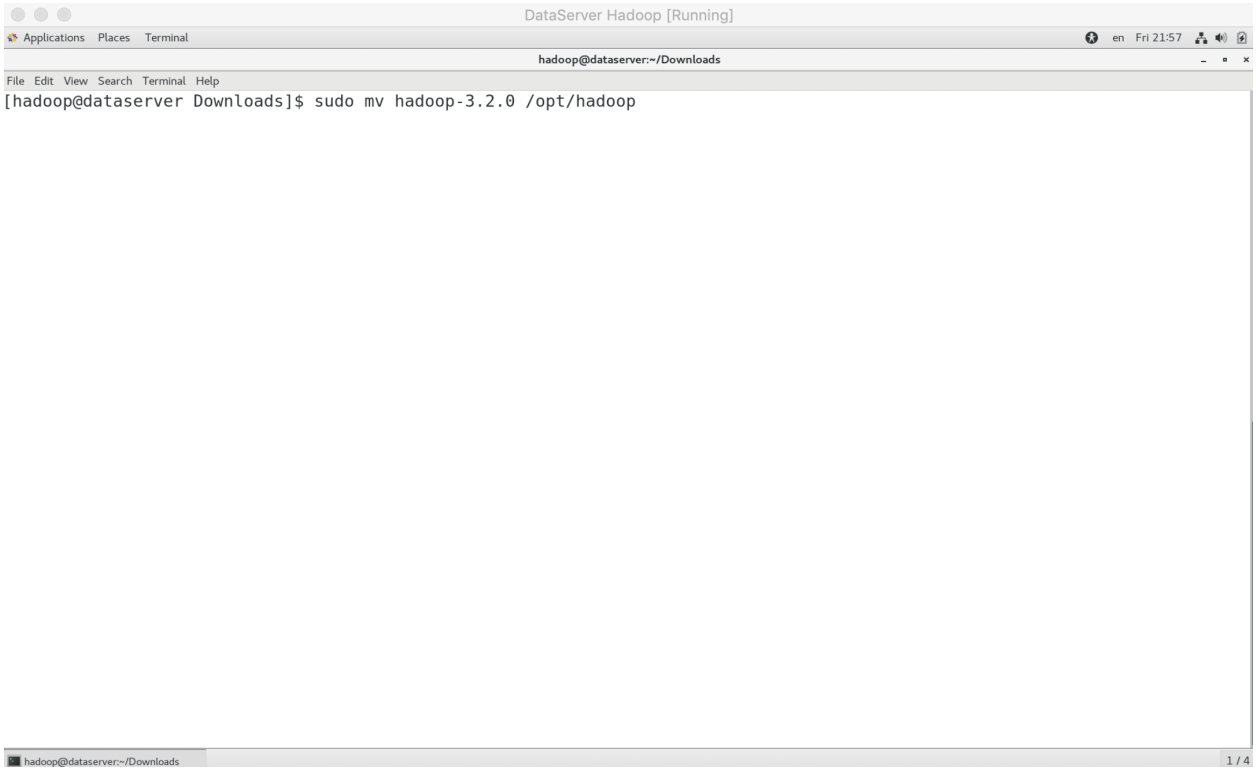

Mover o diretório para /opt/hadoop

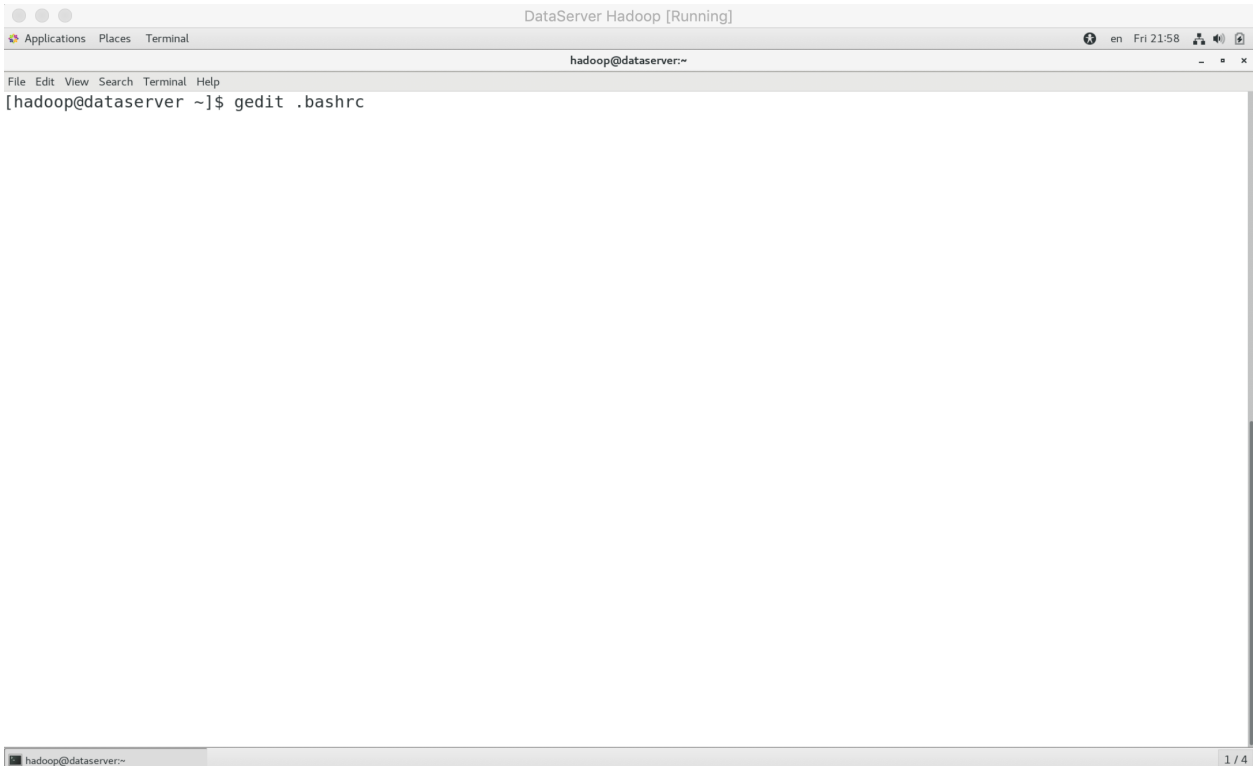

Abrir o arquivo de profile do usuário hadoop

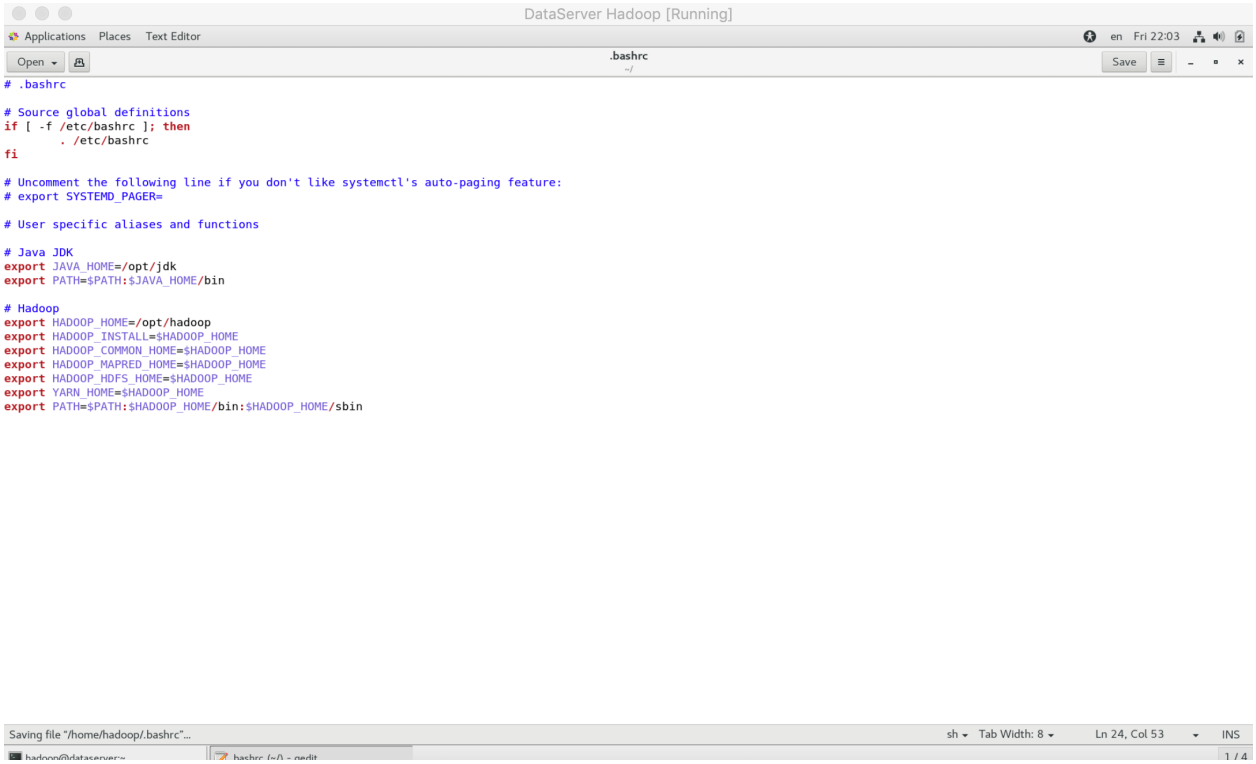

Configure as variáveis de ambiente conforme mostrado na imagem acima e salve o arquivo

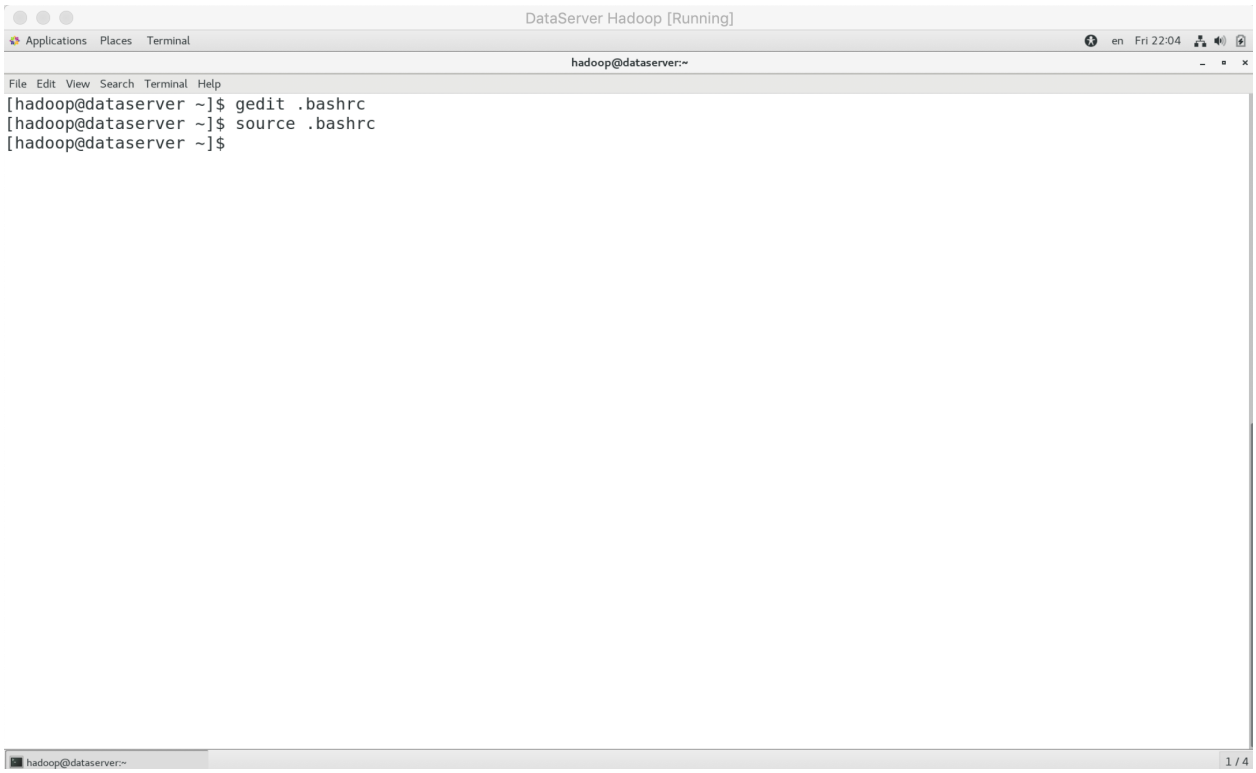

Execute source .bashrc para efetivar as mudanças das variáveis no SO

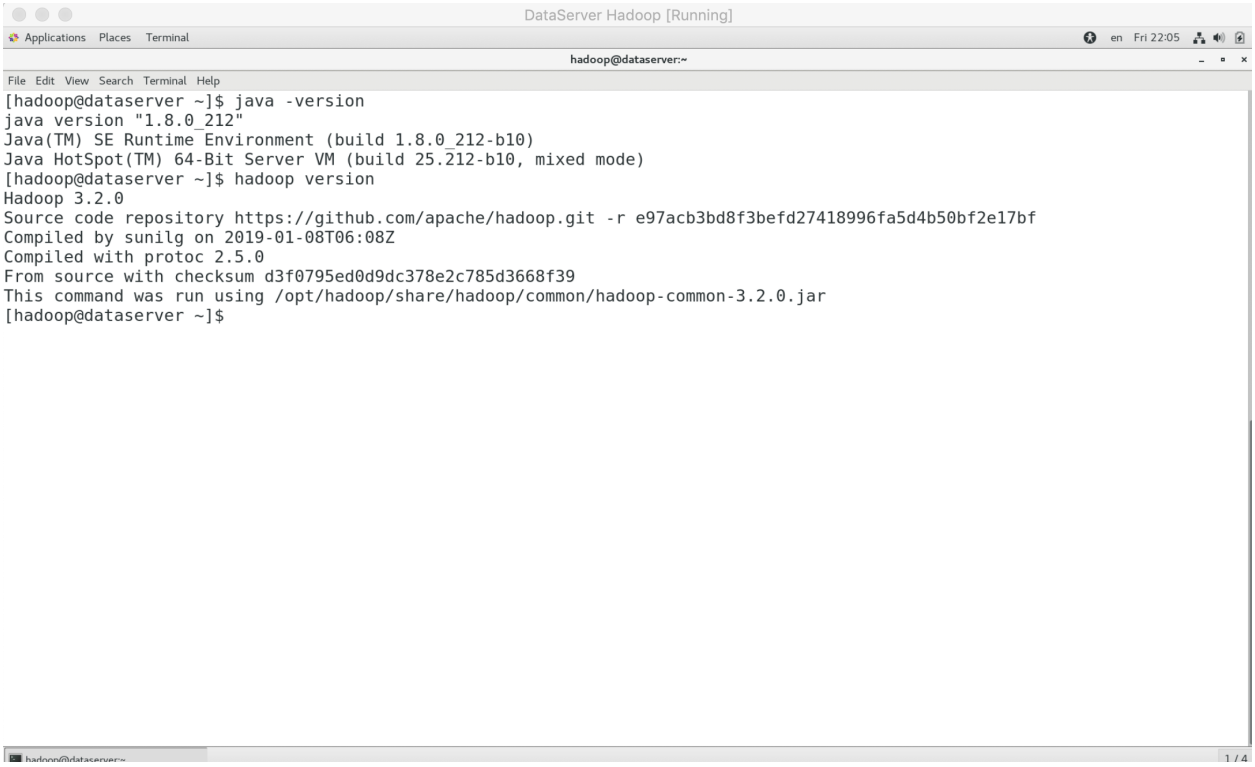

Java e Hadoop instalados e configurados com sucesso!!!

5.4. Configuração do Hadoop

5.4.1. Editar arquivos de configuração do Hadoop

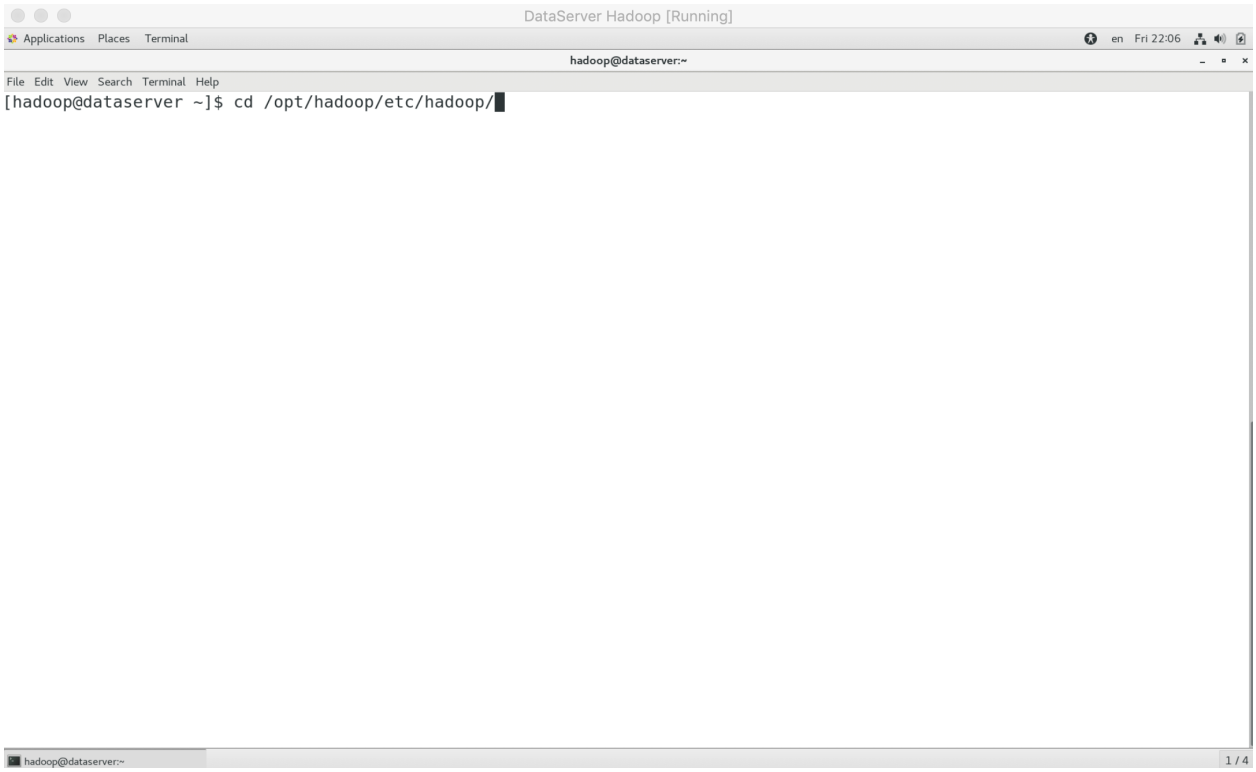

Os arquivos de configuração do Hadoop estão em
[Diretório de instalação do Hadoop]/etc/hadoop
Nesse caso: /opt/hadoop/etc/hadoop

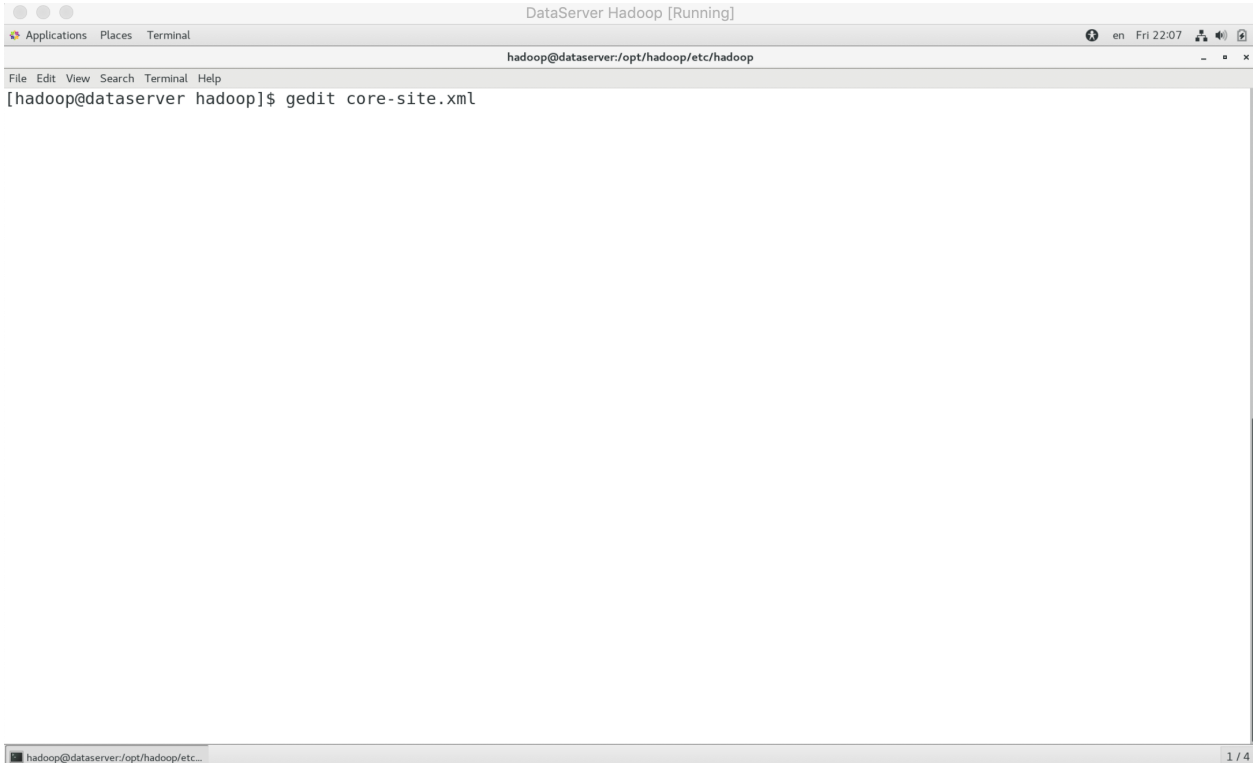

Editar o arquivo core-site.xml

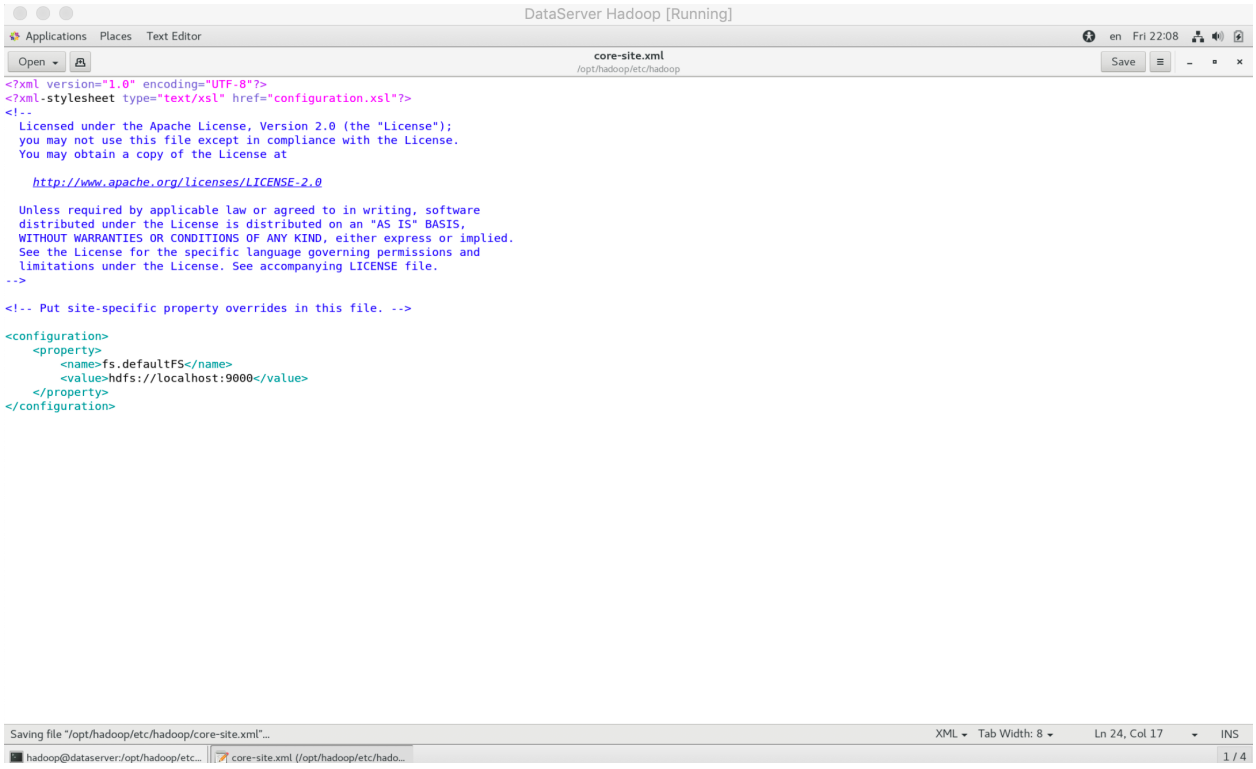

Acrescentar as propriedades conforme acima e salvar o arquivo.
Essa propriedade indica o endereço do HDFS.

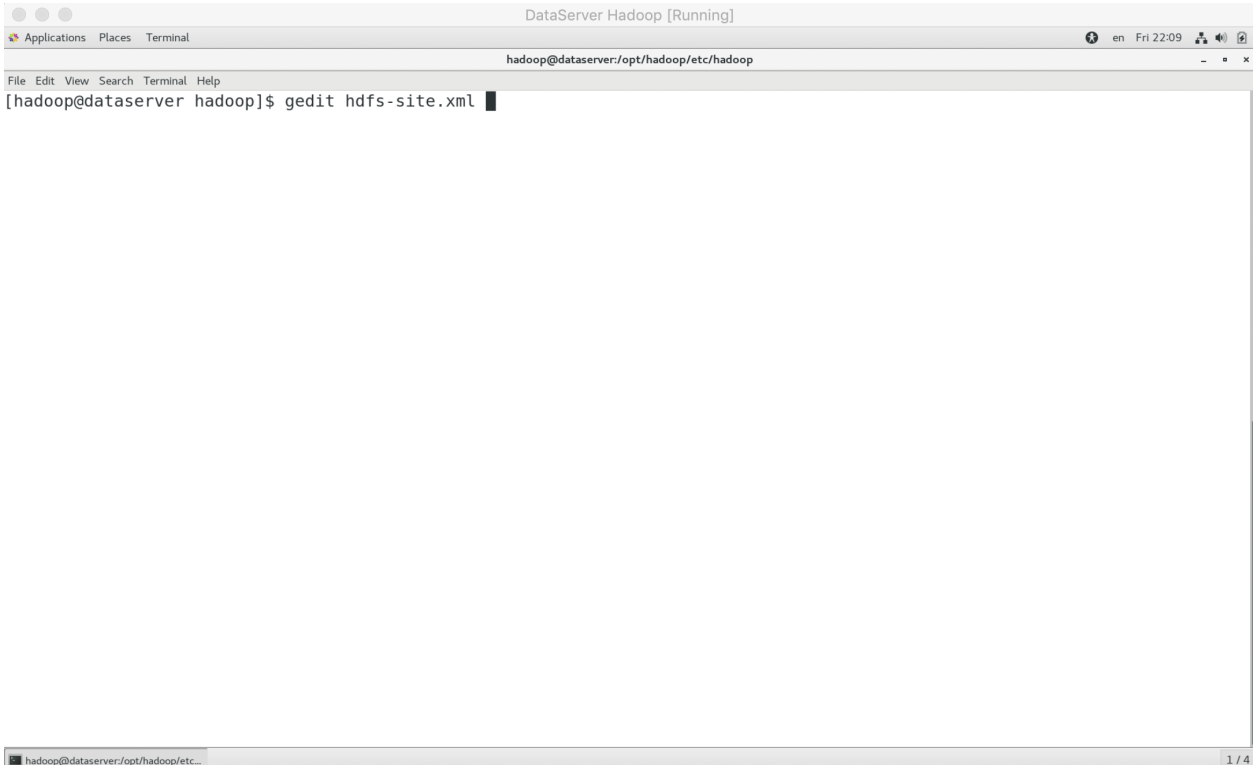

Editar o arquivo hdfs-site.xml

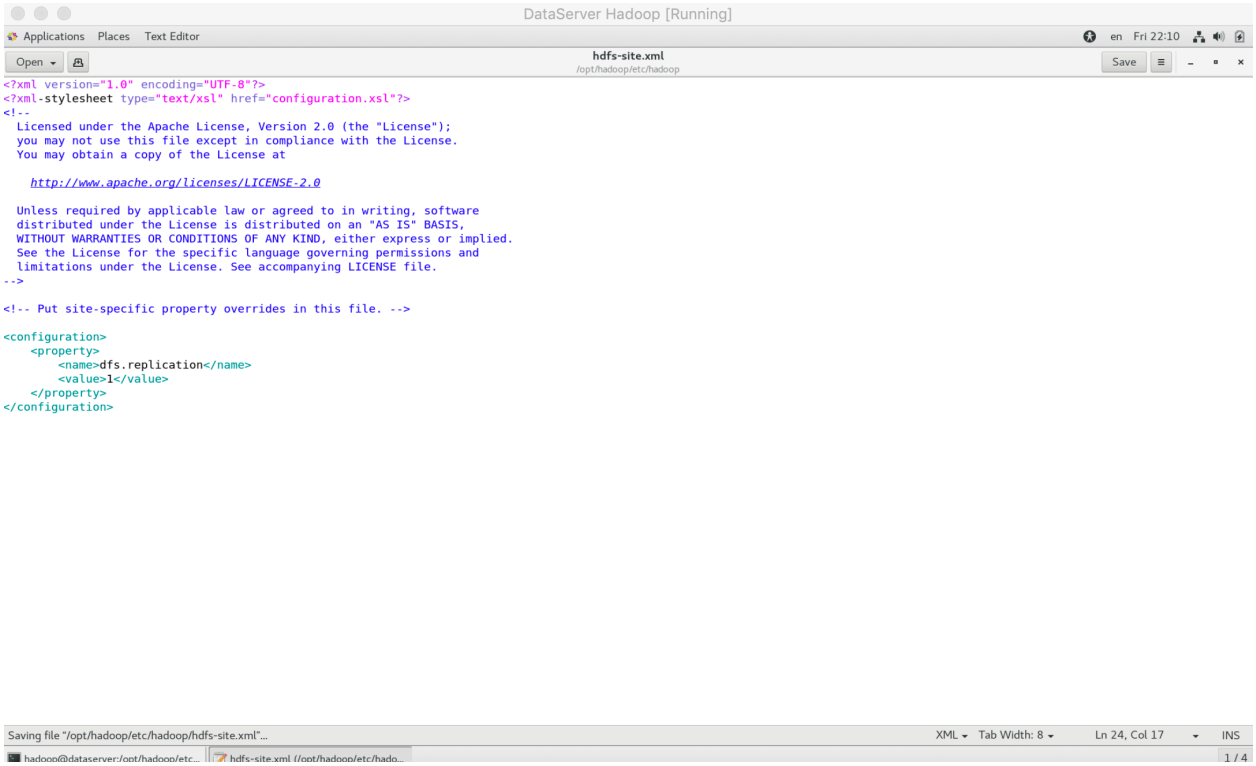

Acrescentar as propriedades conforme acima e salvar o arquivo.

5.4.2. Formatando o Namenode

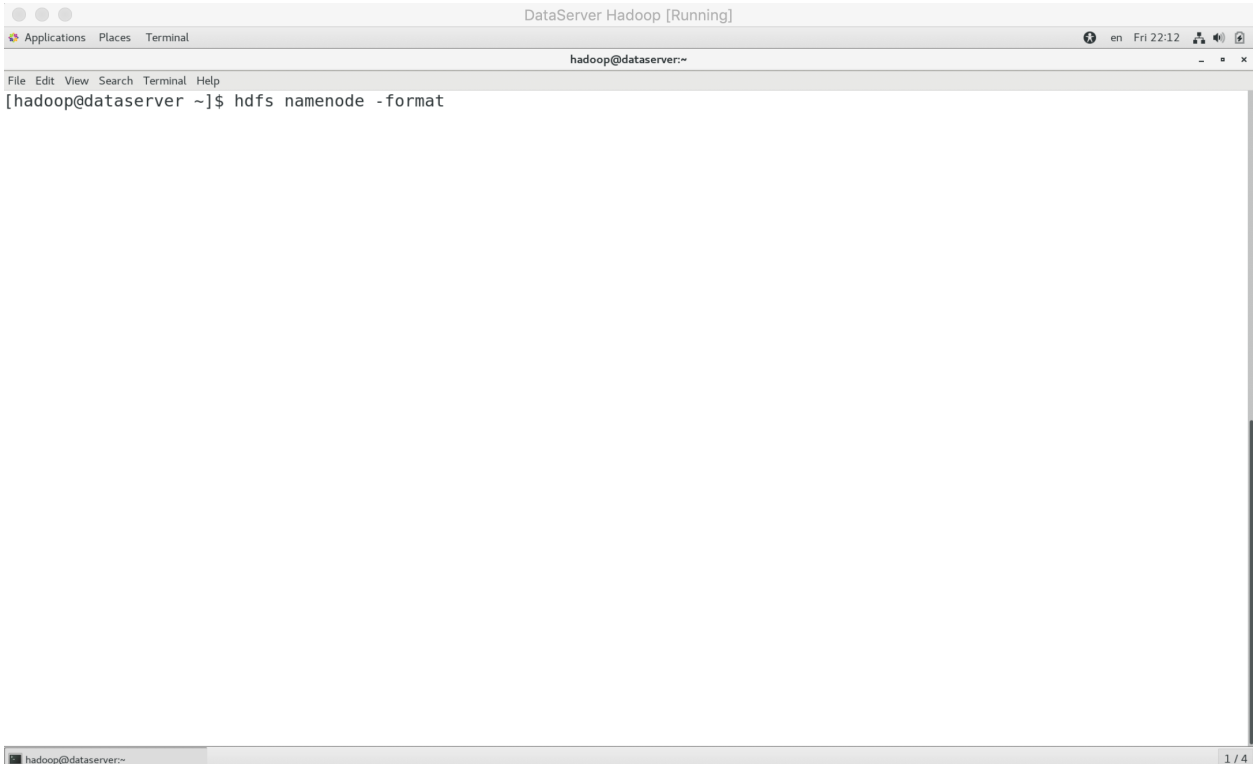

hdfs namenode –format

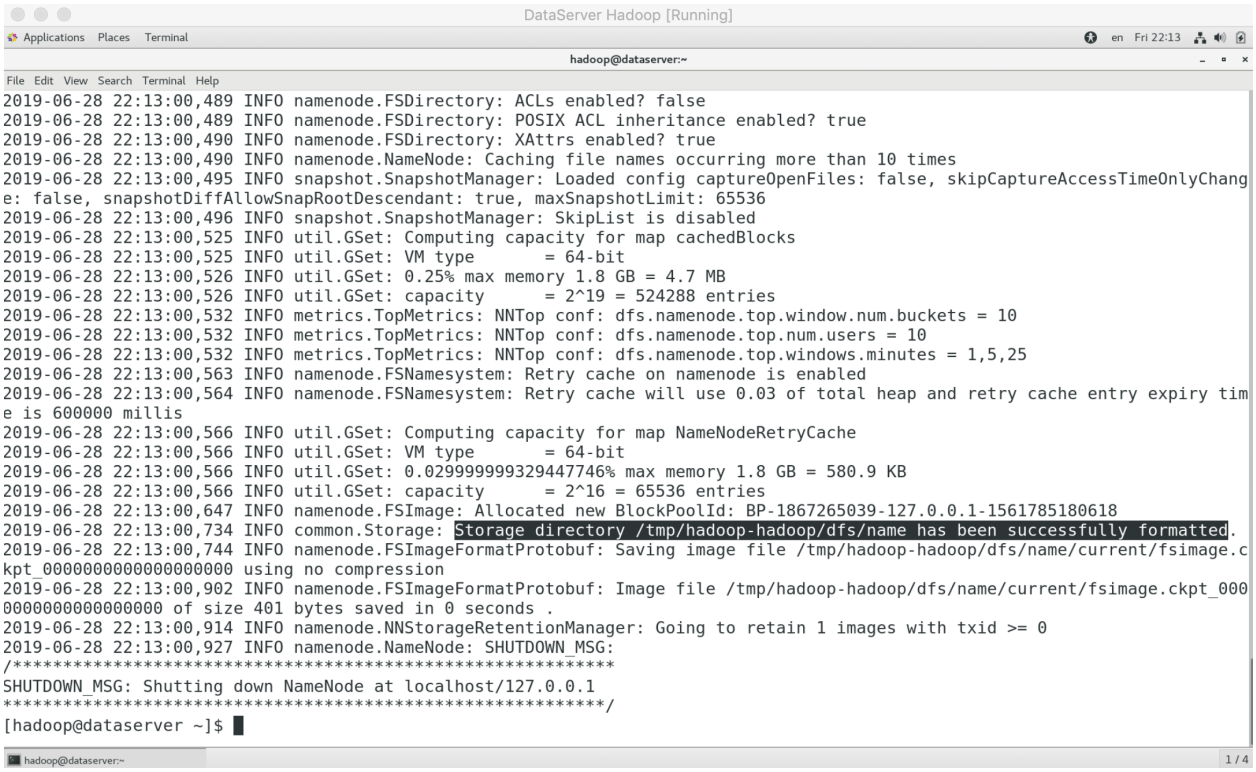

Formatação realizada com sucesso

5.4.3. Iniciando o Hadoop

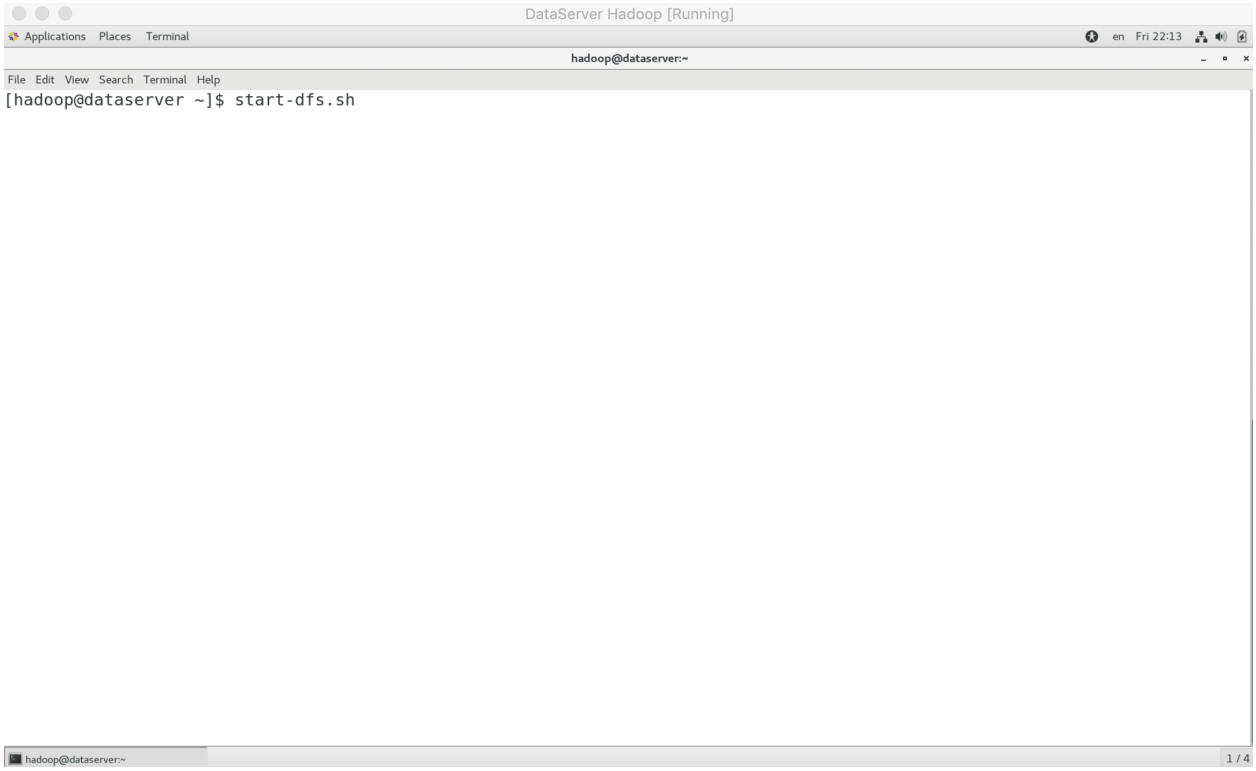

start-dfs.sh

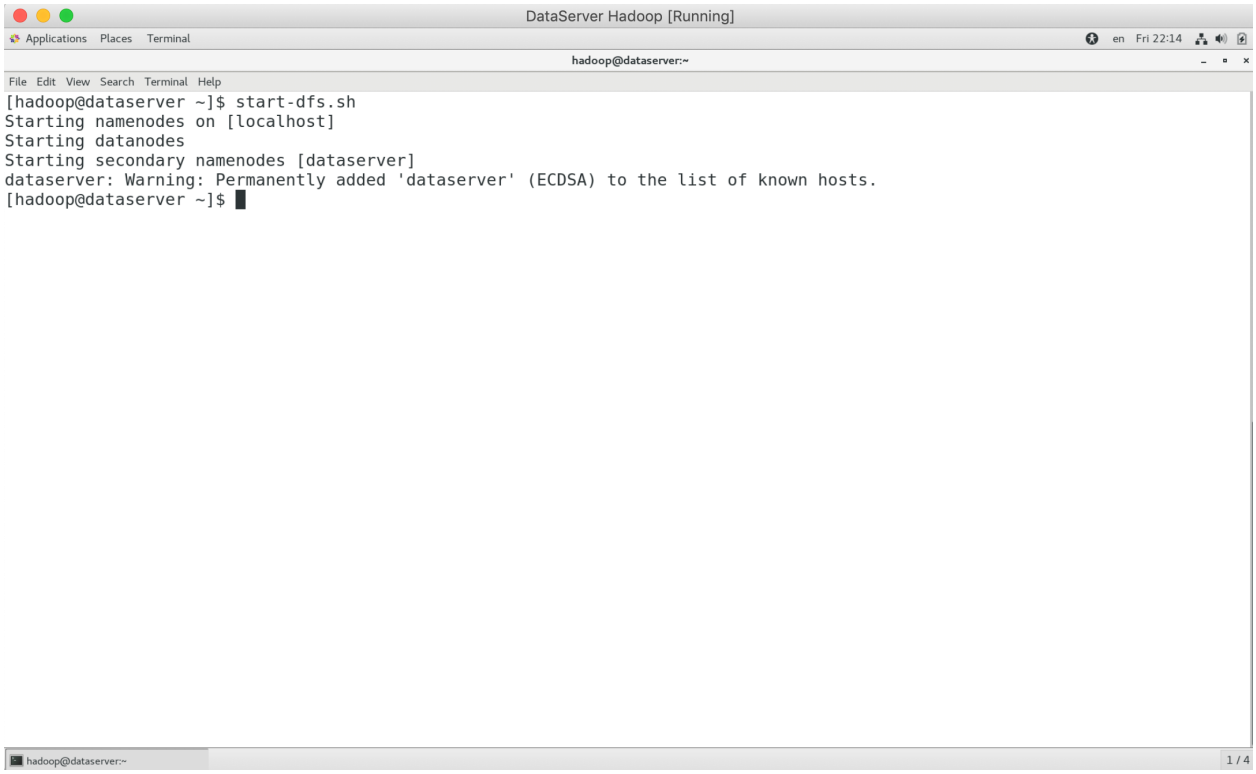

Hadoop iniciado

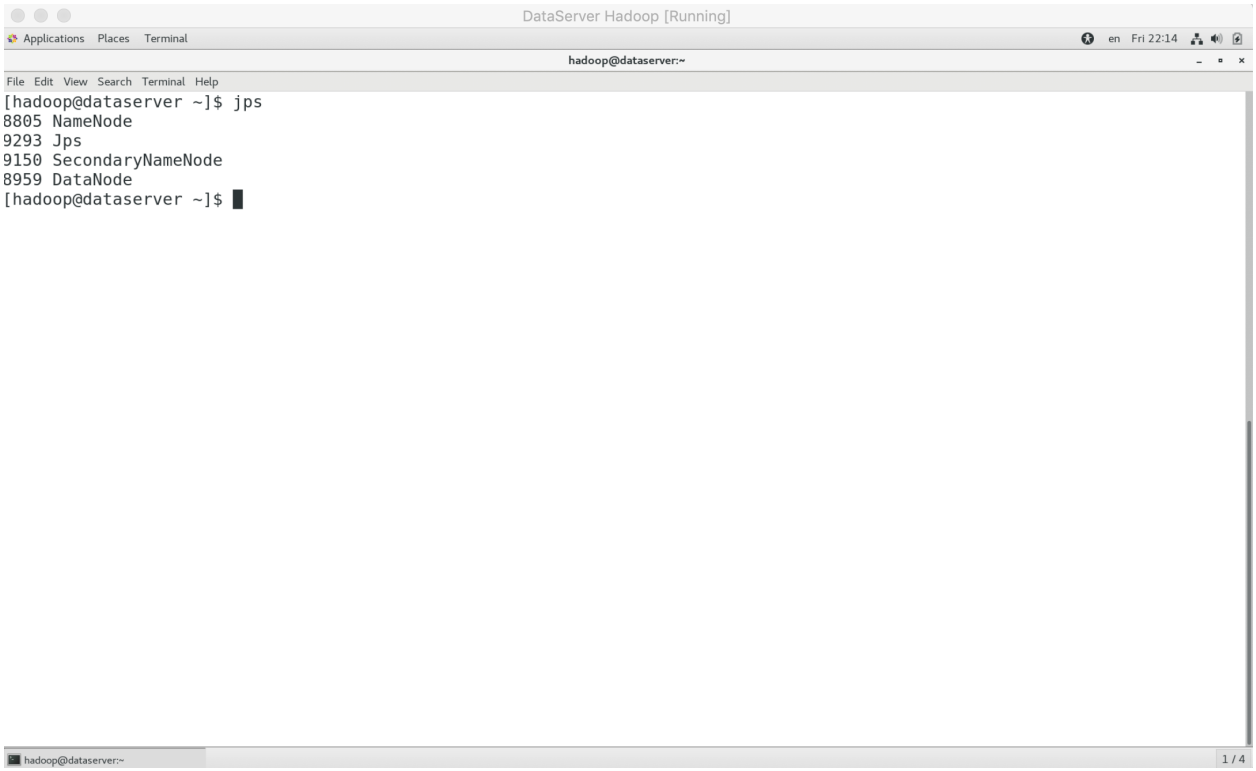

Checando os serviços inicializados com o comando **jps**

5.4.4. Iniciando o Yarn

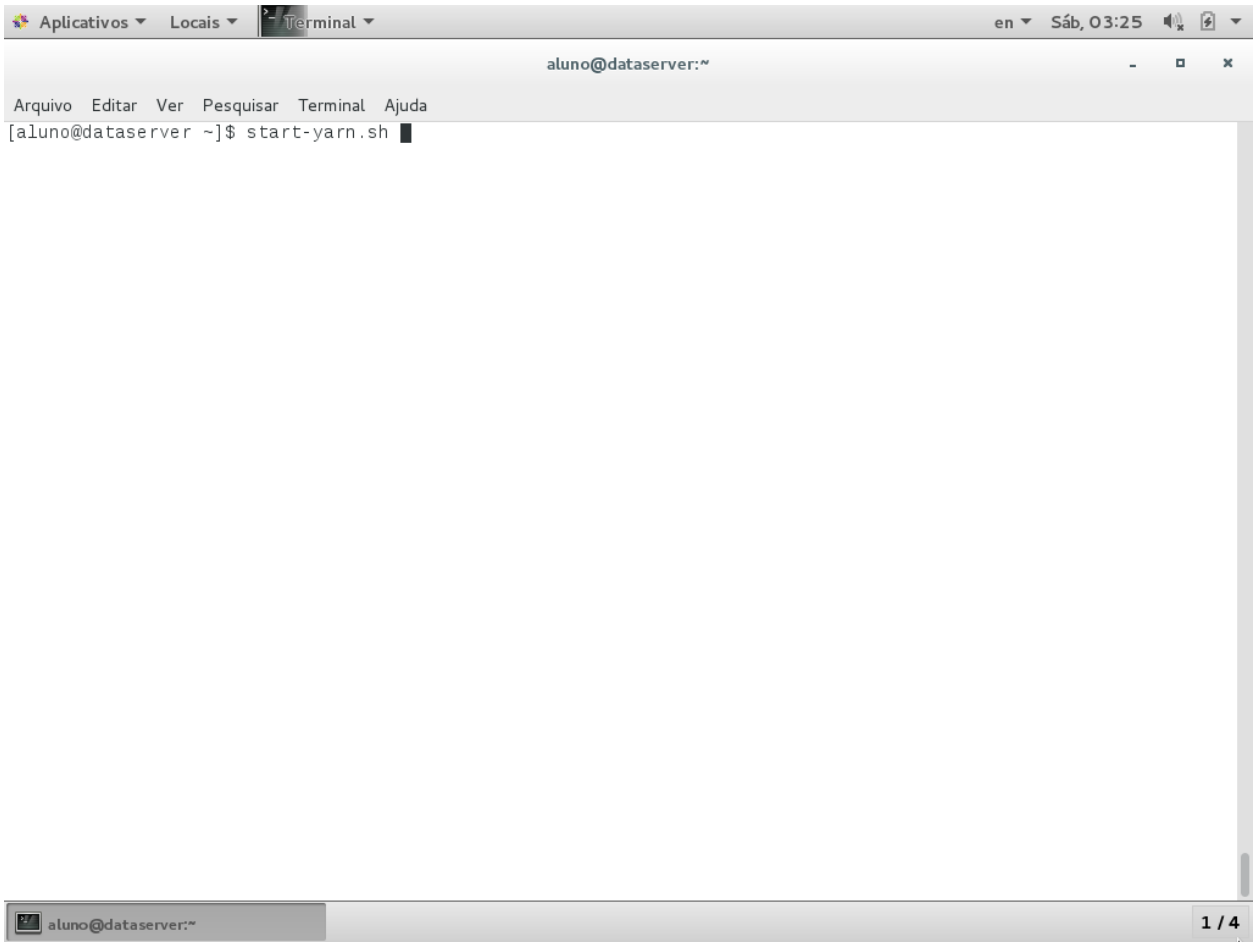

start-yarn.sh

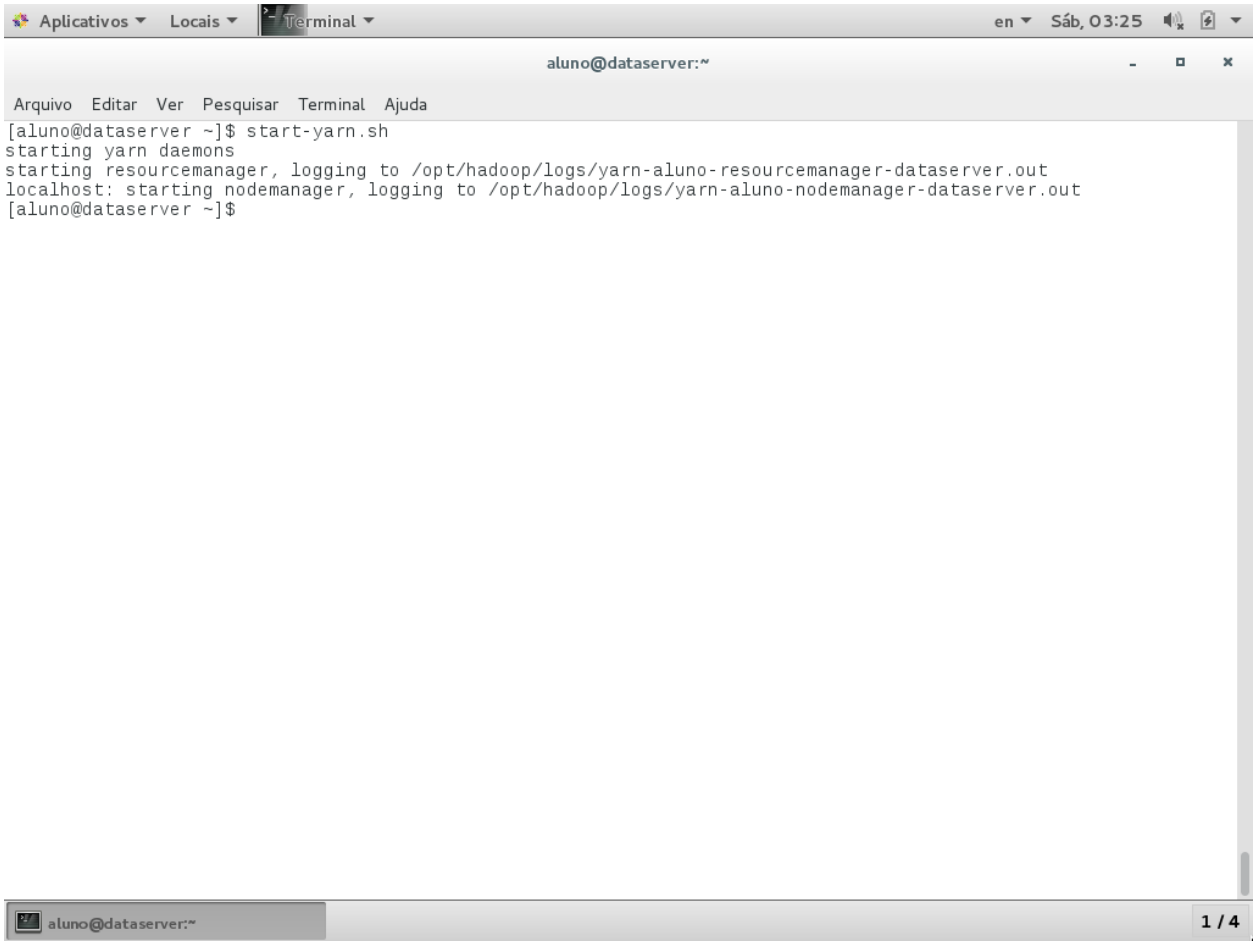

Yarn iniciado

Checando os serviços com o comando **jps**

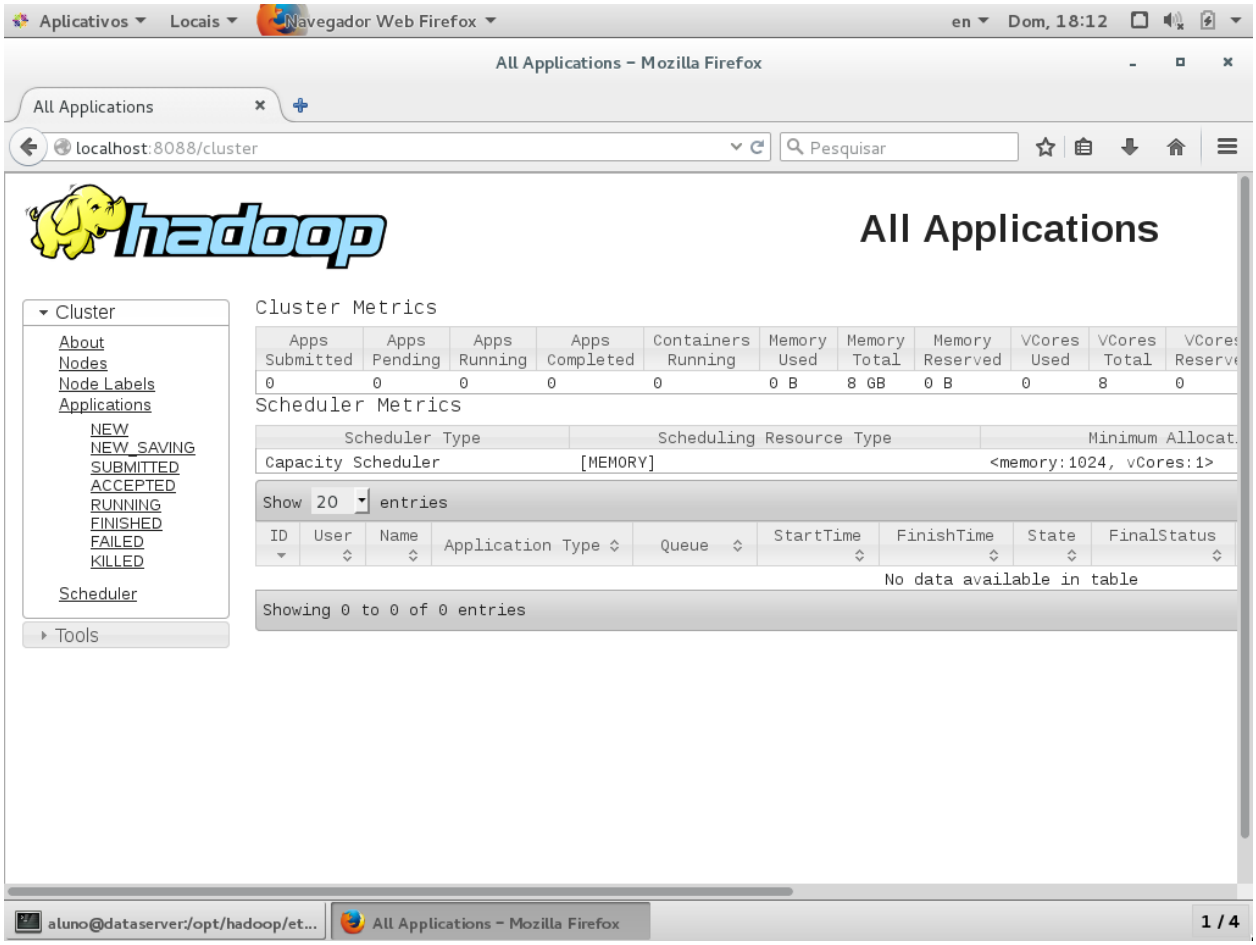

Visualizando jobs – http://localhost:8088

5.5. Processando Big Data

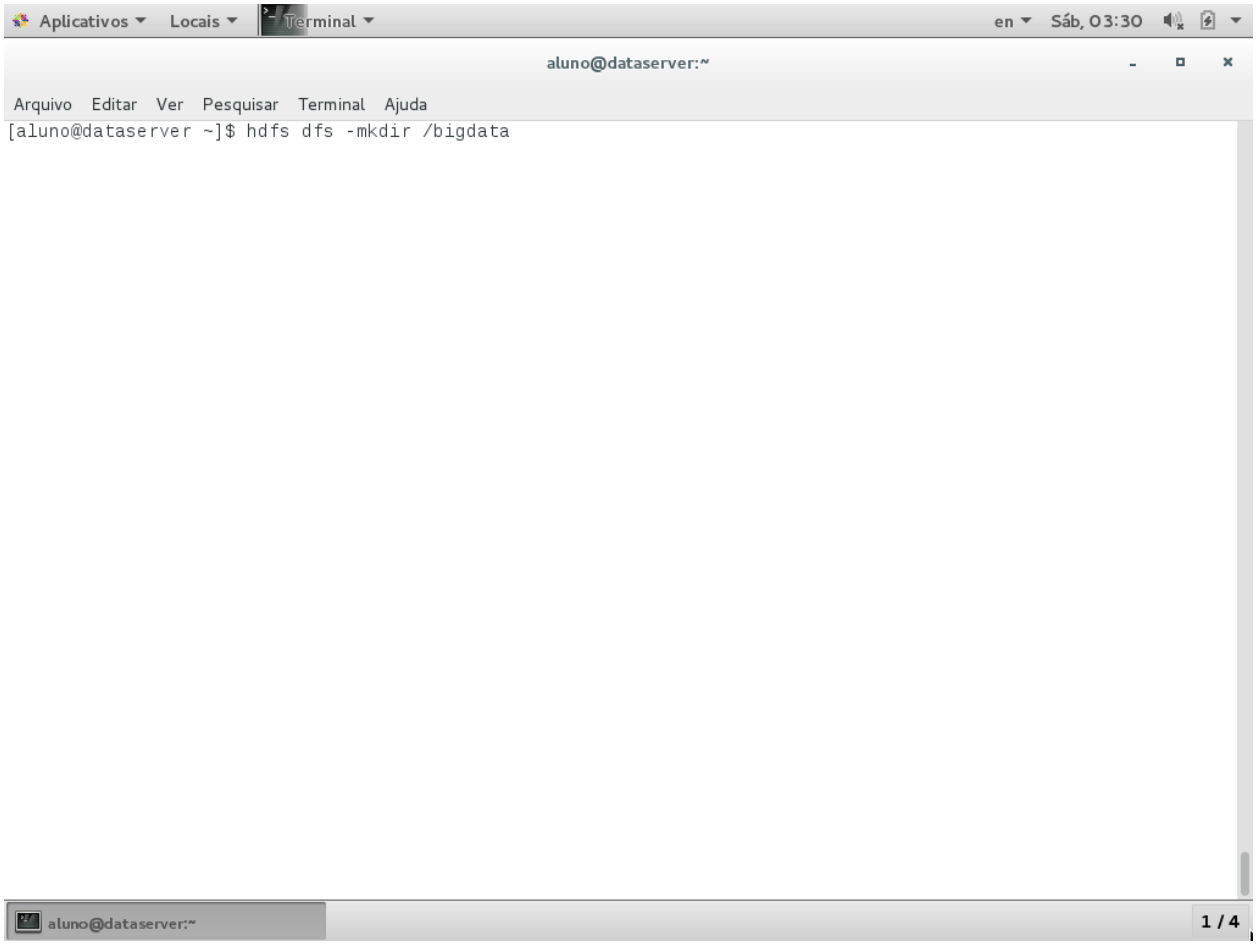

Criar o diretório **bigdata** no HDFS

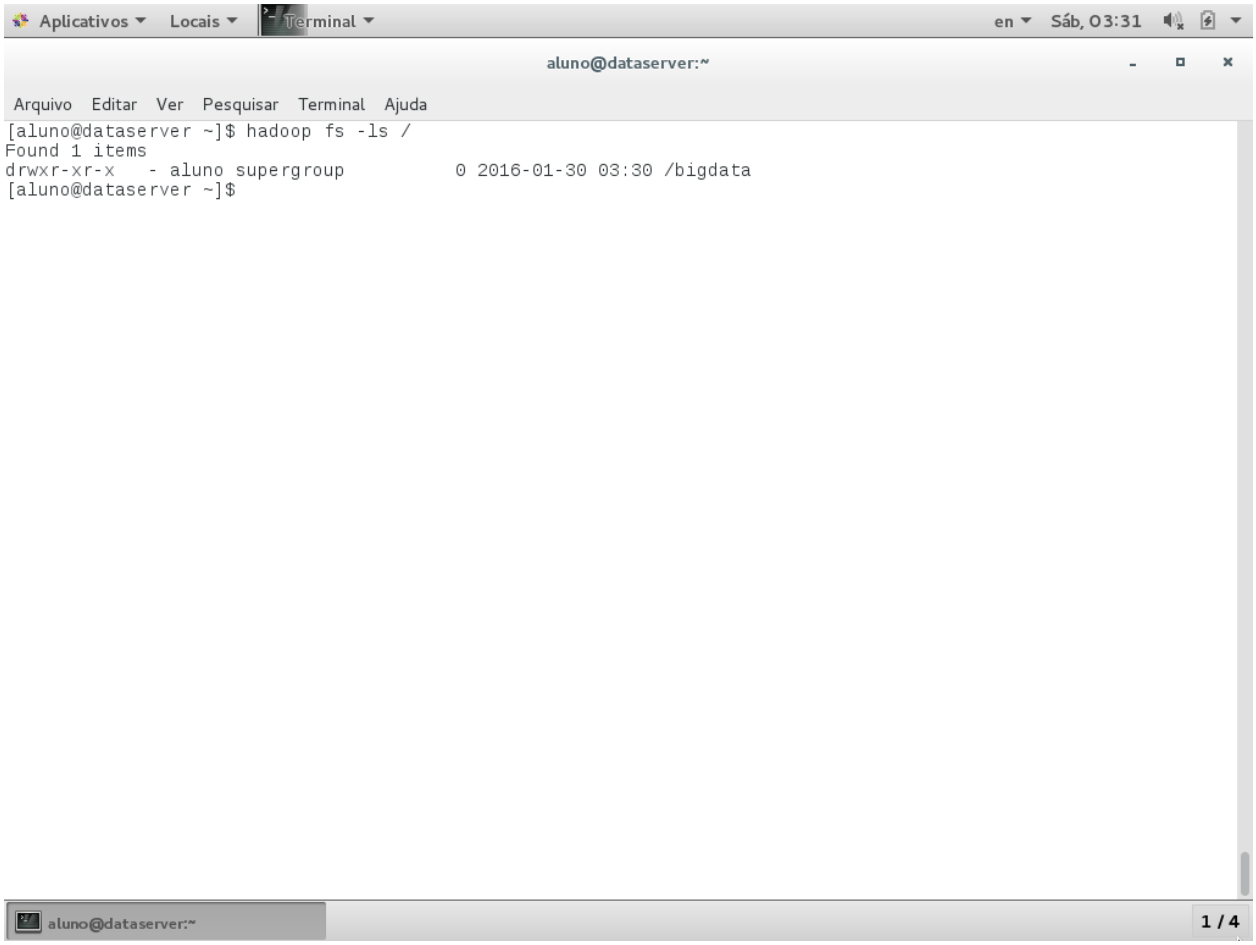

Listar o HDFS e checar o diretório criado

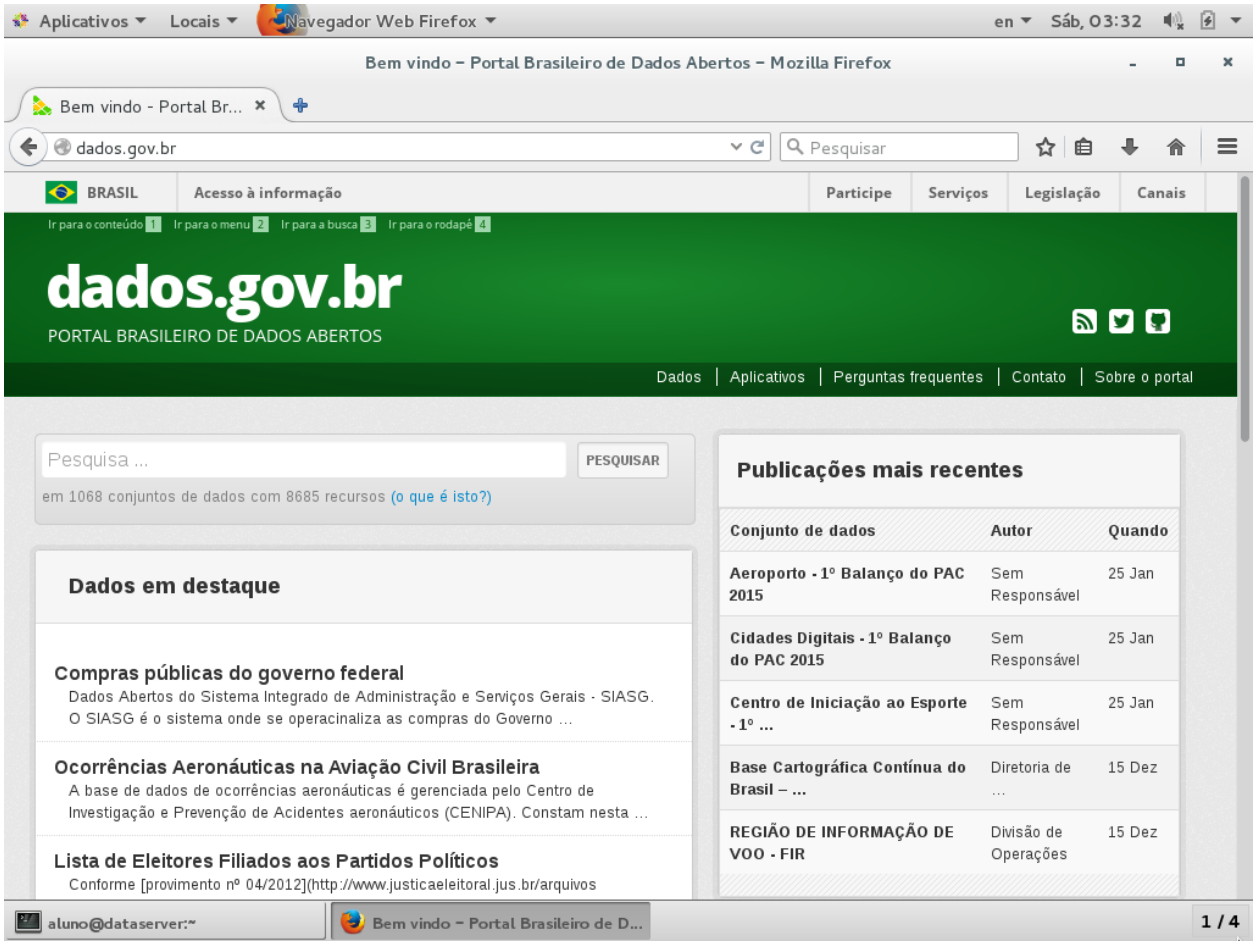

Acessar o portal de dados abertos do governo federal

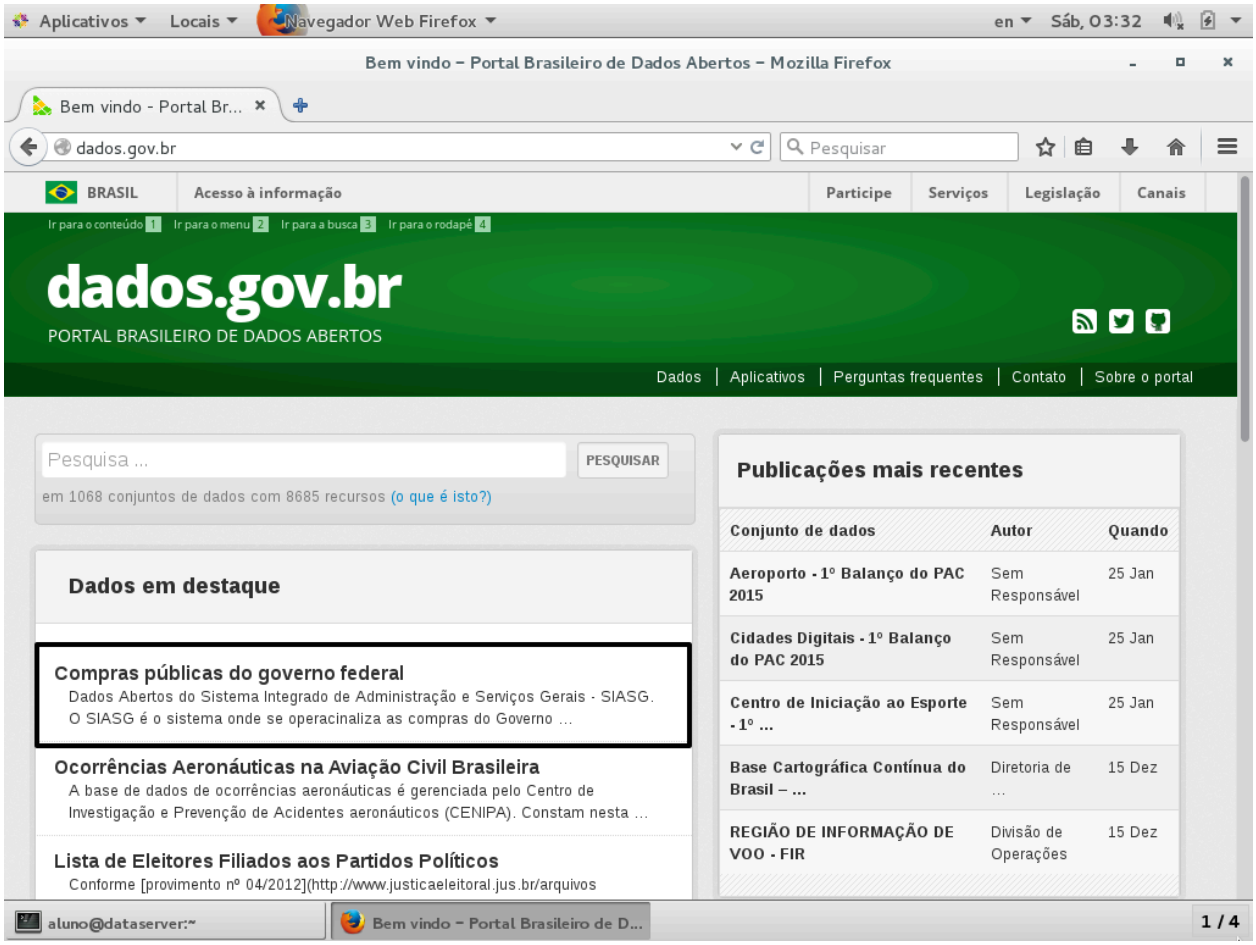

Clicar no link de compras públicas

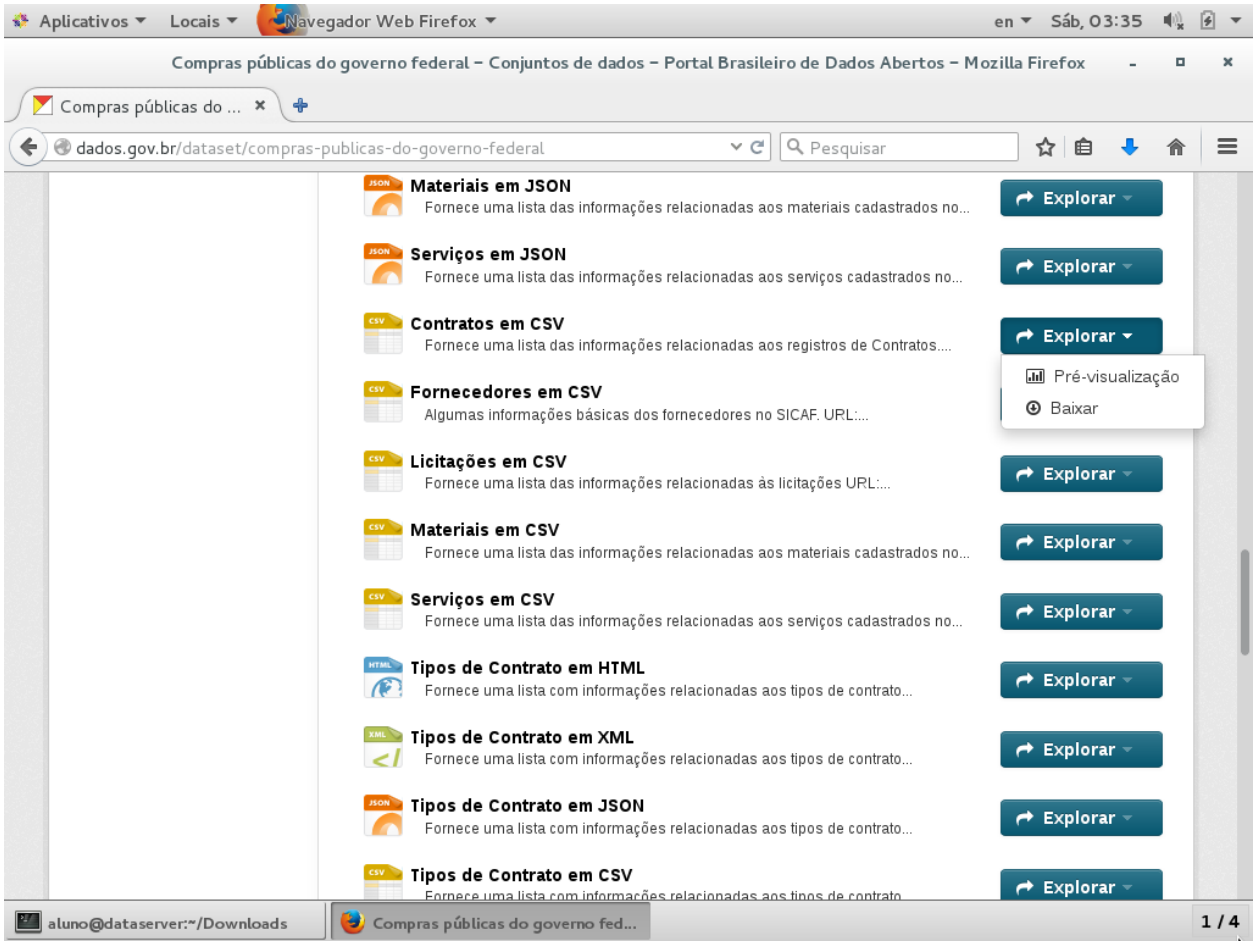

Baixar o arquivo de contratos em formato csv (pode ser qualquer um)

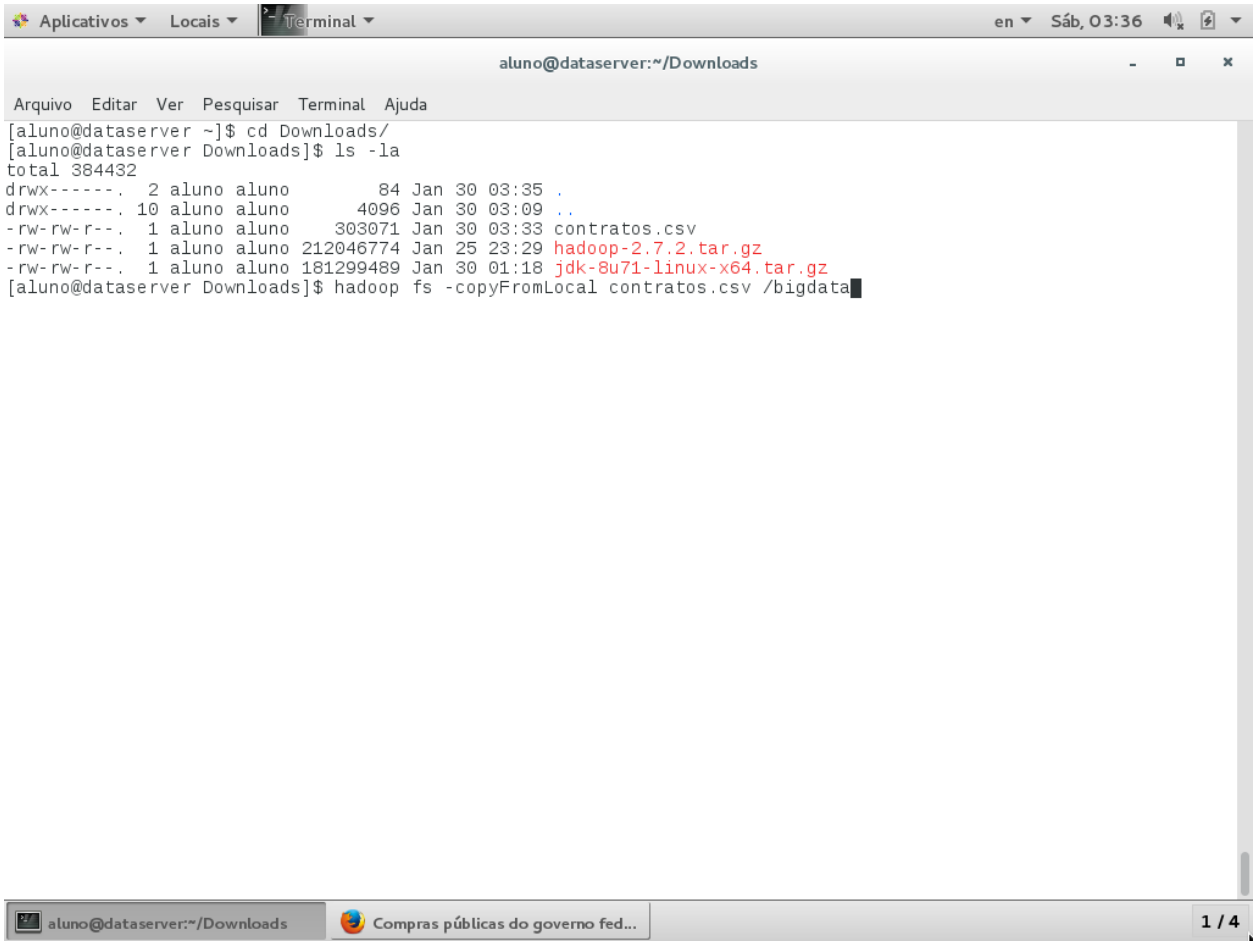

Copiar o arquivo para a pasta bigdata no HDFS

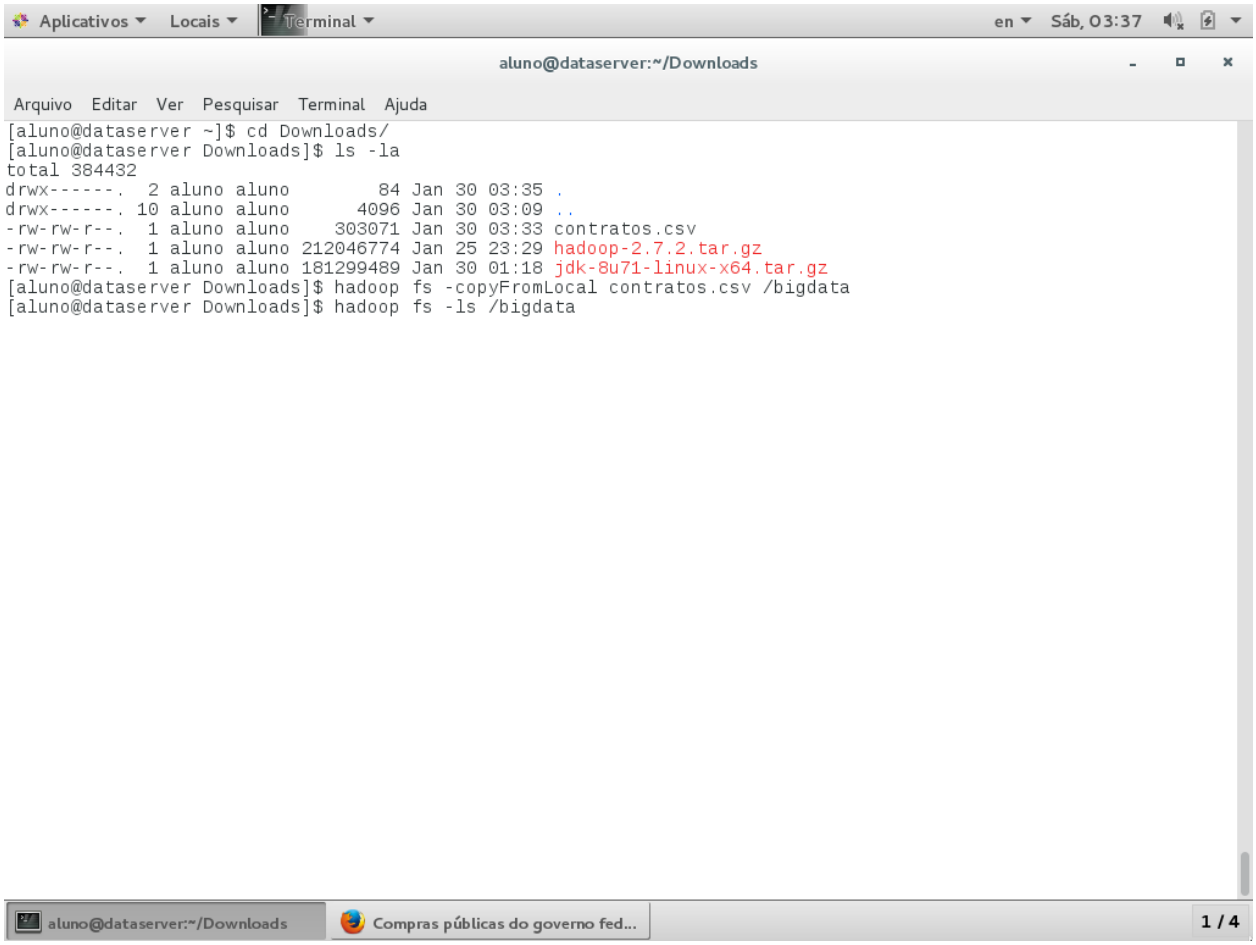

Listar o diretório bigdata

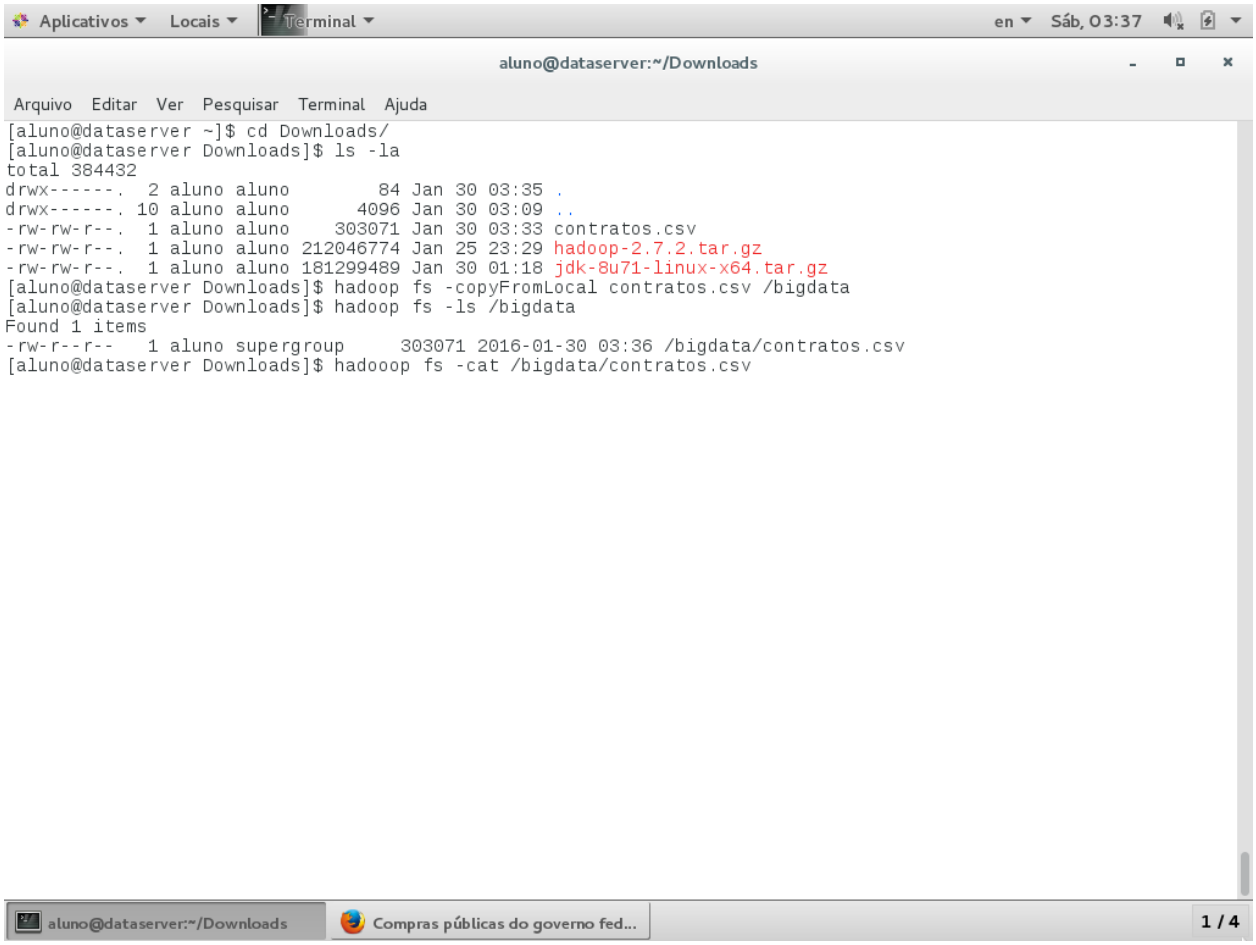

Ver o conteúdo do arquivo

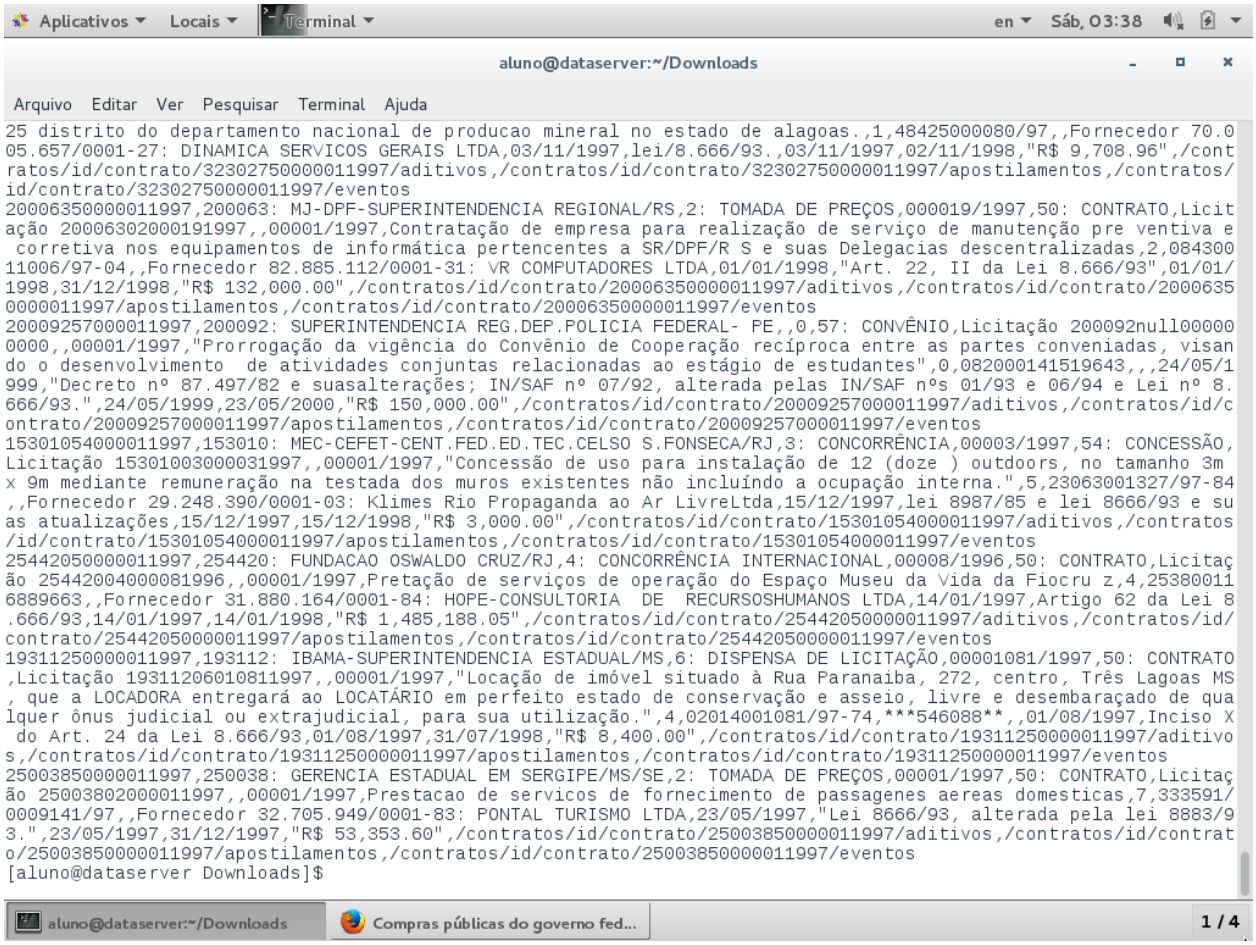

Conteúdo do arquivo já gravado no HDFS

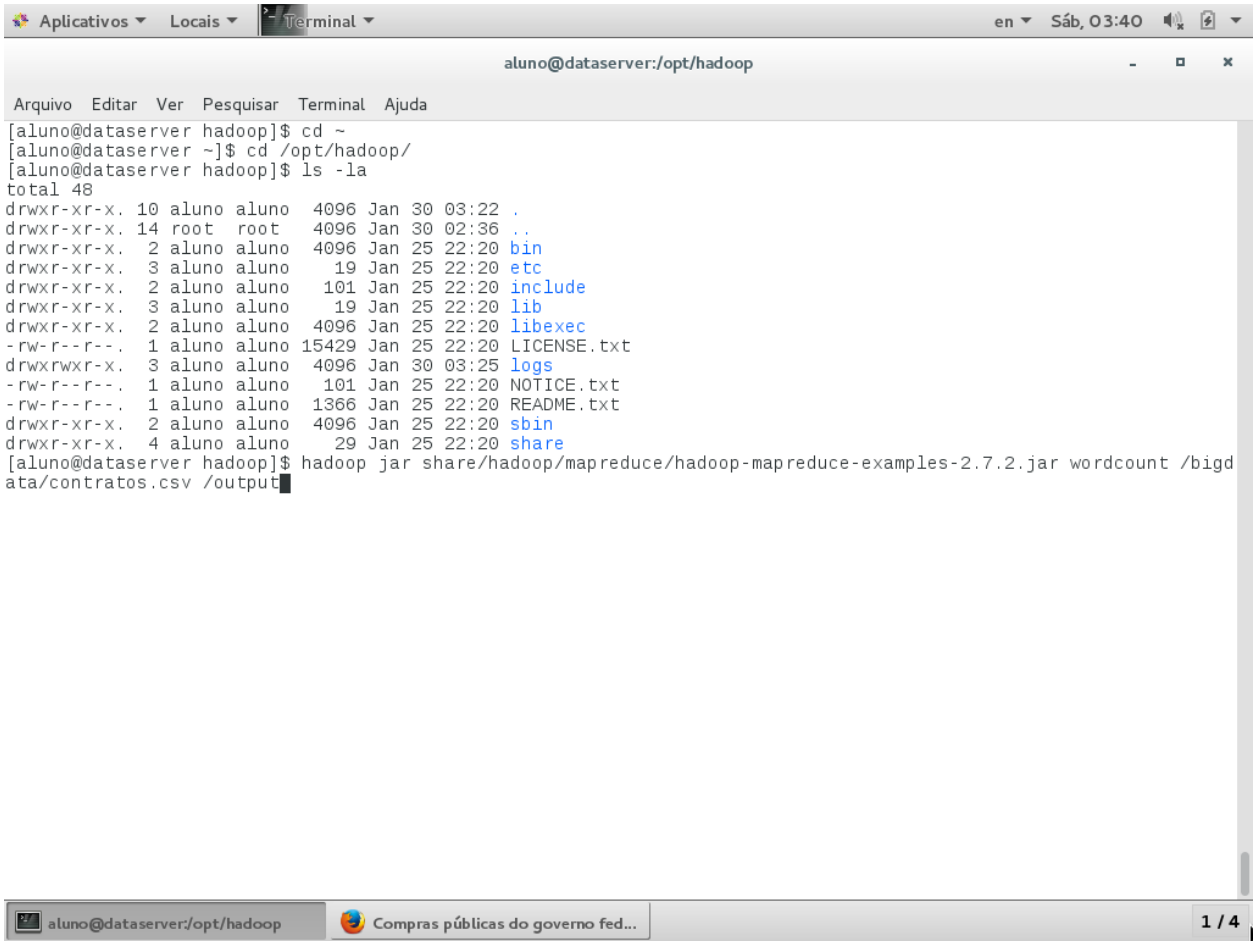

A instalação do Hadoop possui um job chamado wordcount, que pode ser usado como exemplo para processamento de Big Data. Basicamente, o job conta a ocorrência de cada palavra no arquivo. Vamos executar com o comando acima (a versão do arquivo .jar é a mesma versão do Hadoop que você instalou).

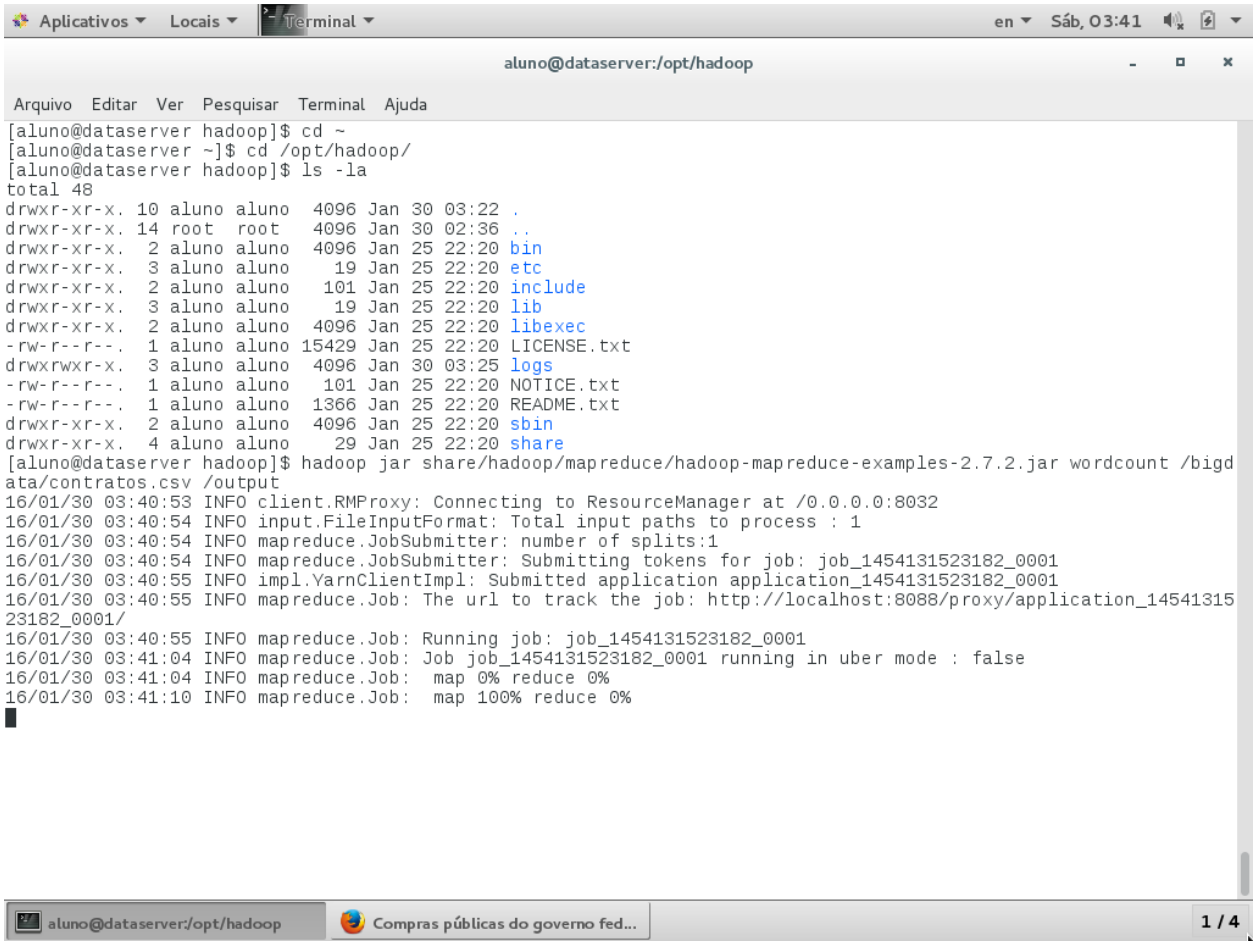

Job sendo processado

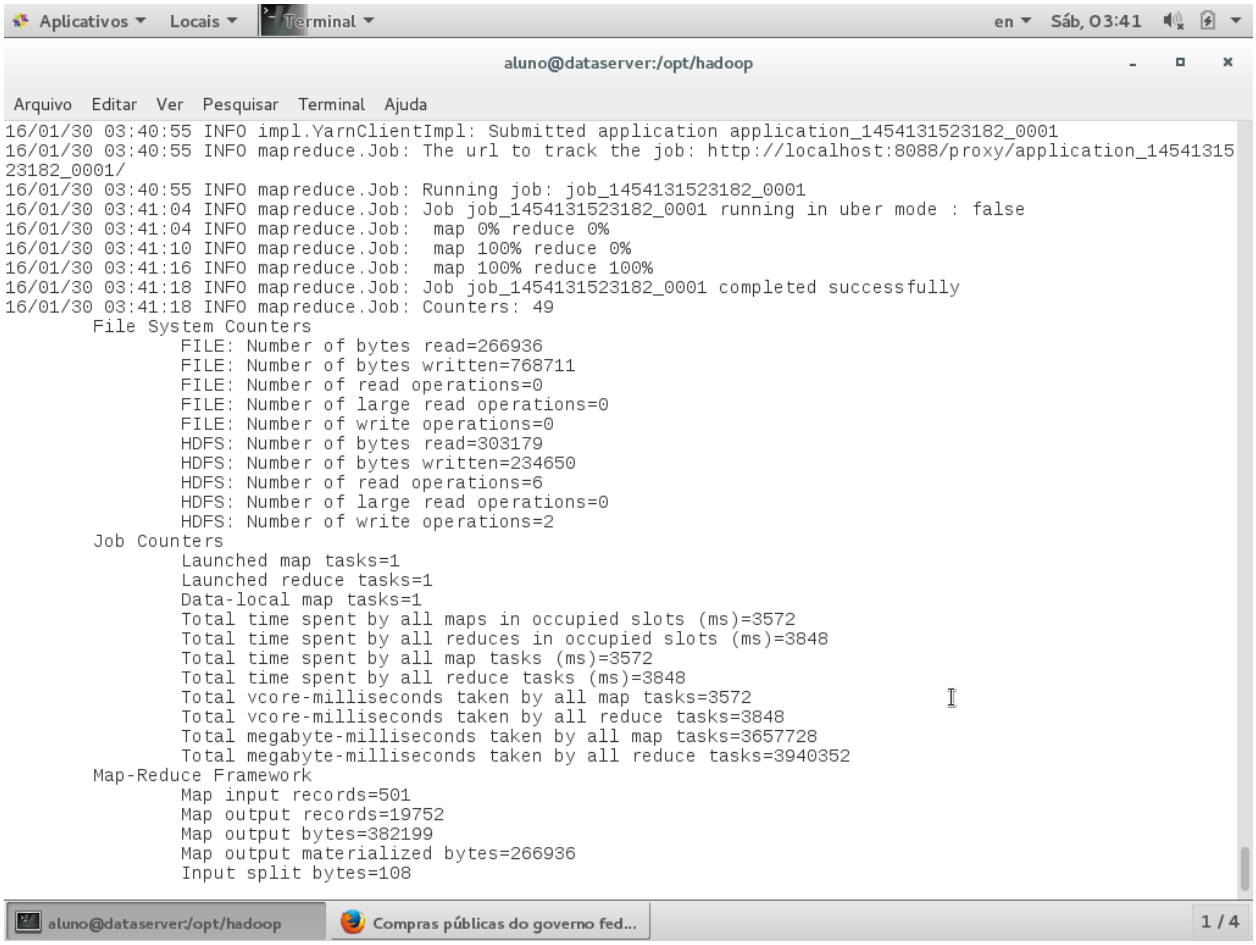

job processado com sucesso

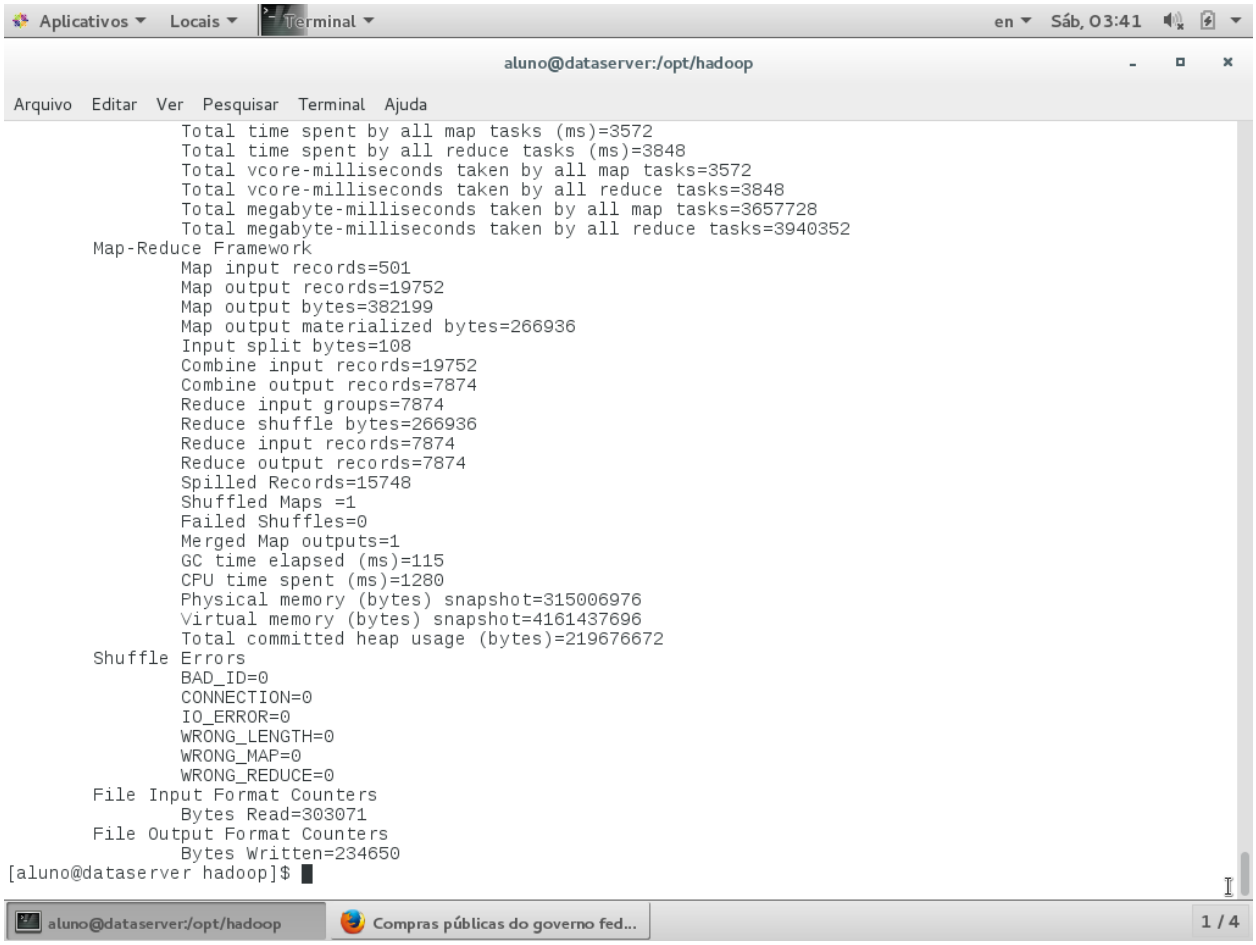

job processado com sucesso

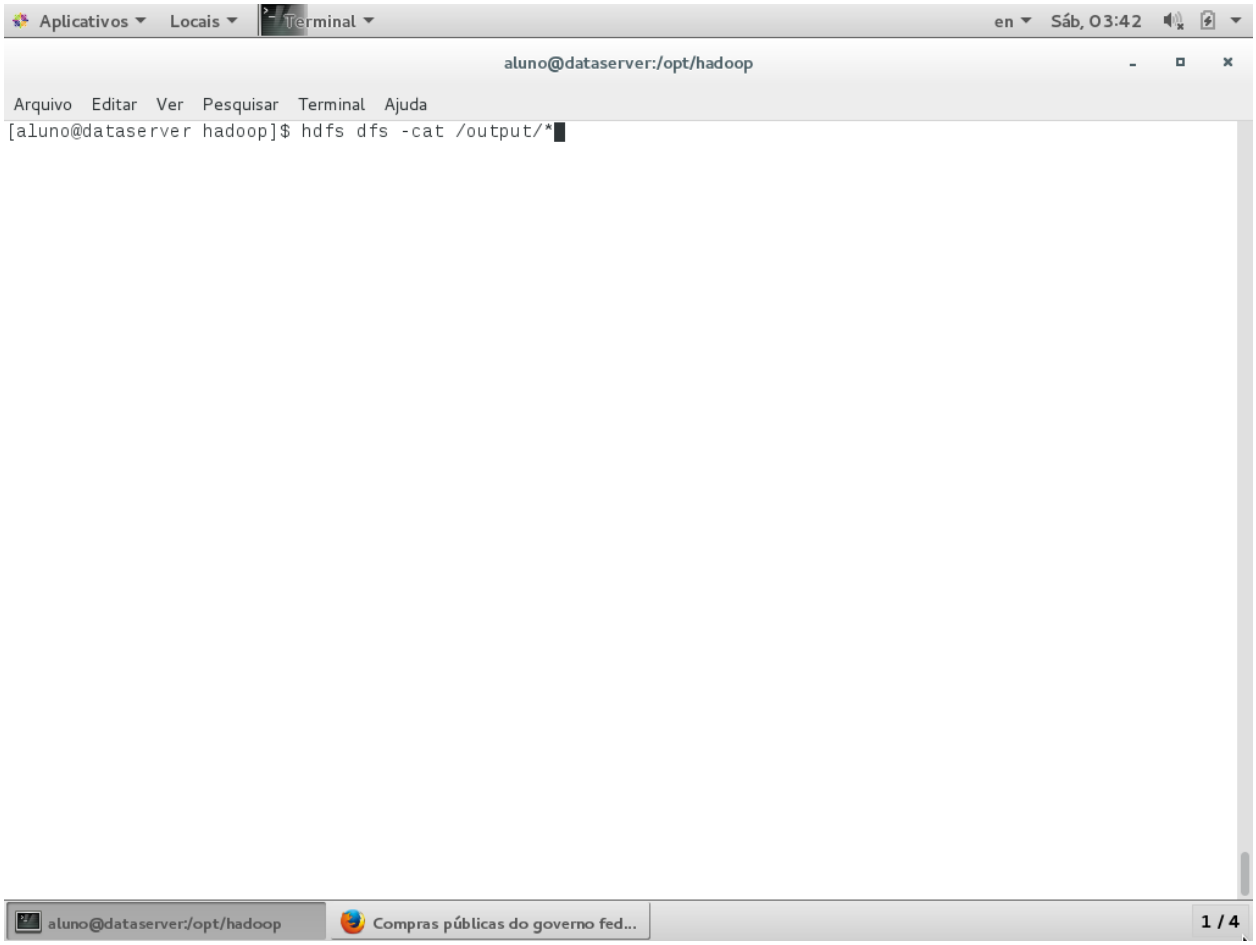

Vamos ver o resultado do processamento

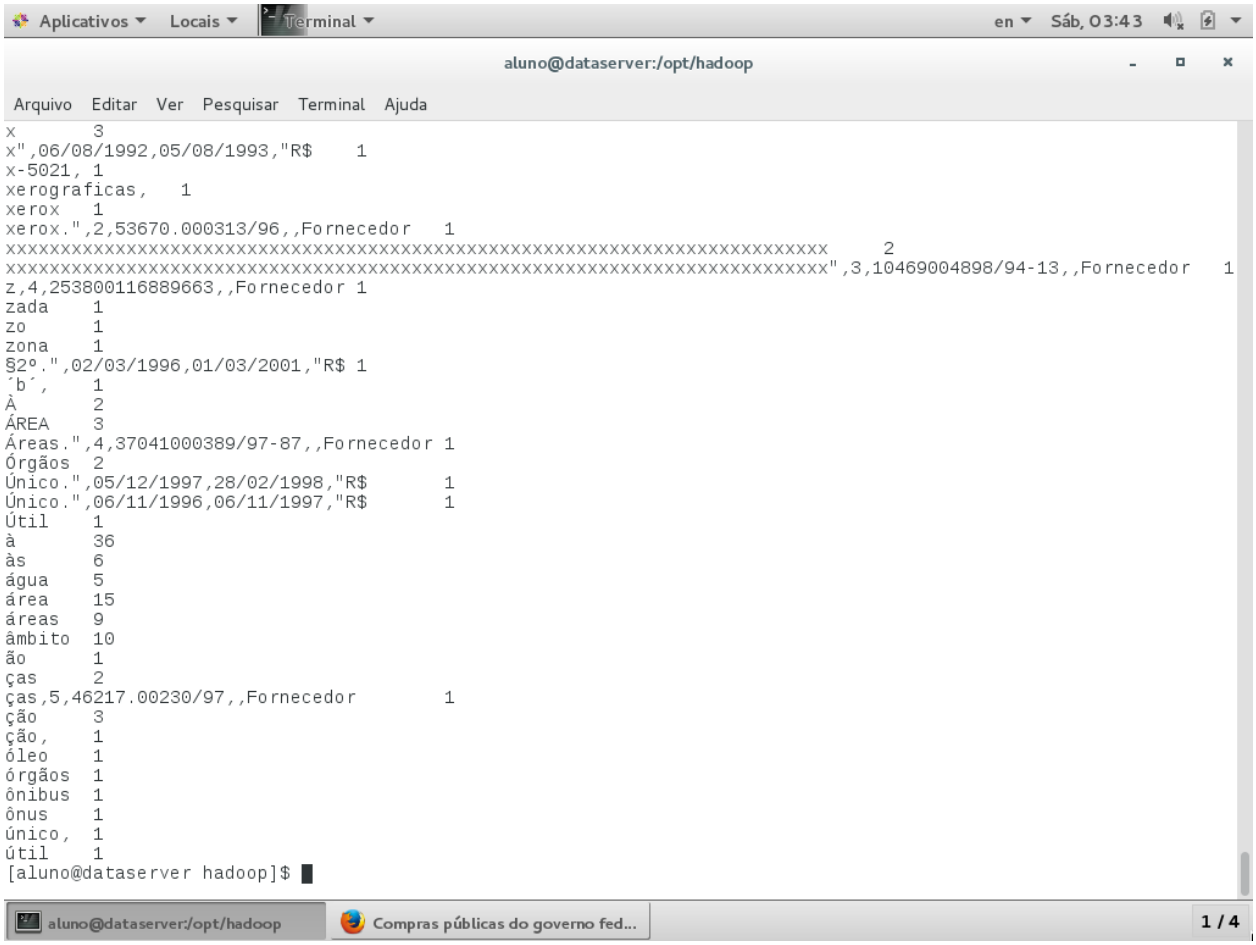

Arquivo processado. Número de ocorrência de cada palavra/termo no arquivo.

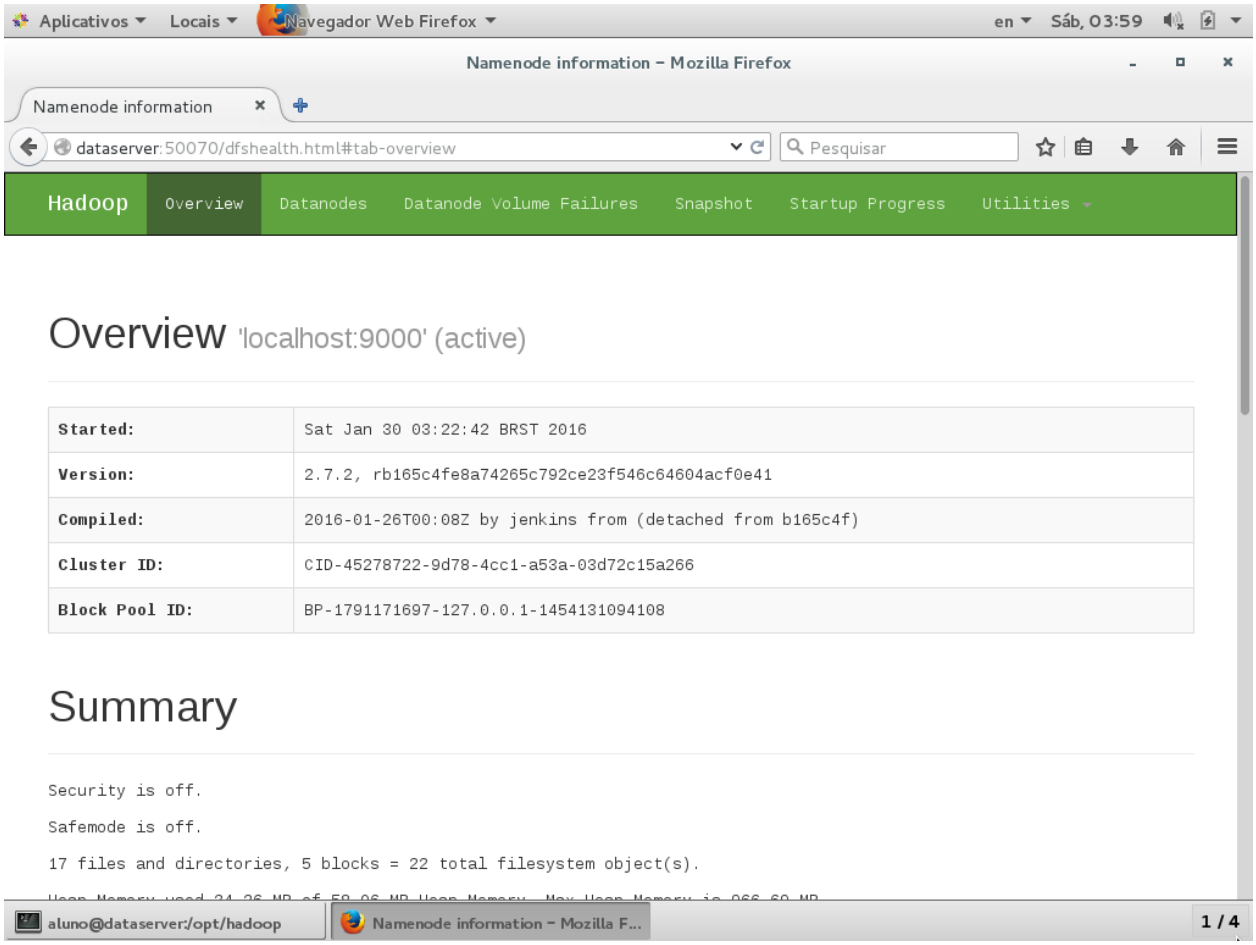

Acesso ao Hadoop pelo browser: http://dataserver:50070

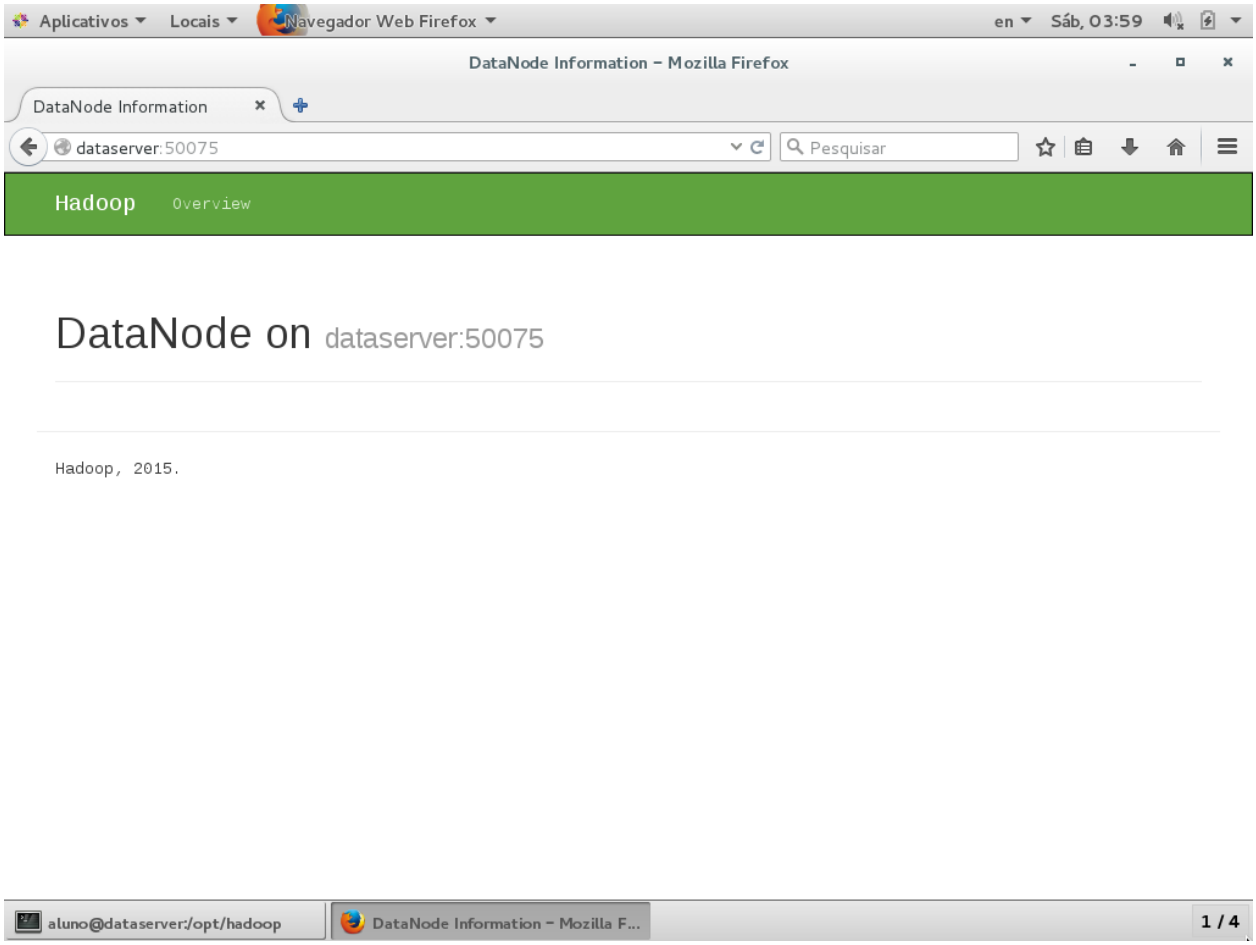

Acesso ao Hadoop pelo browser: http://dataserver:50075

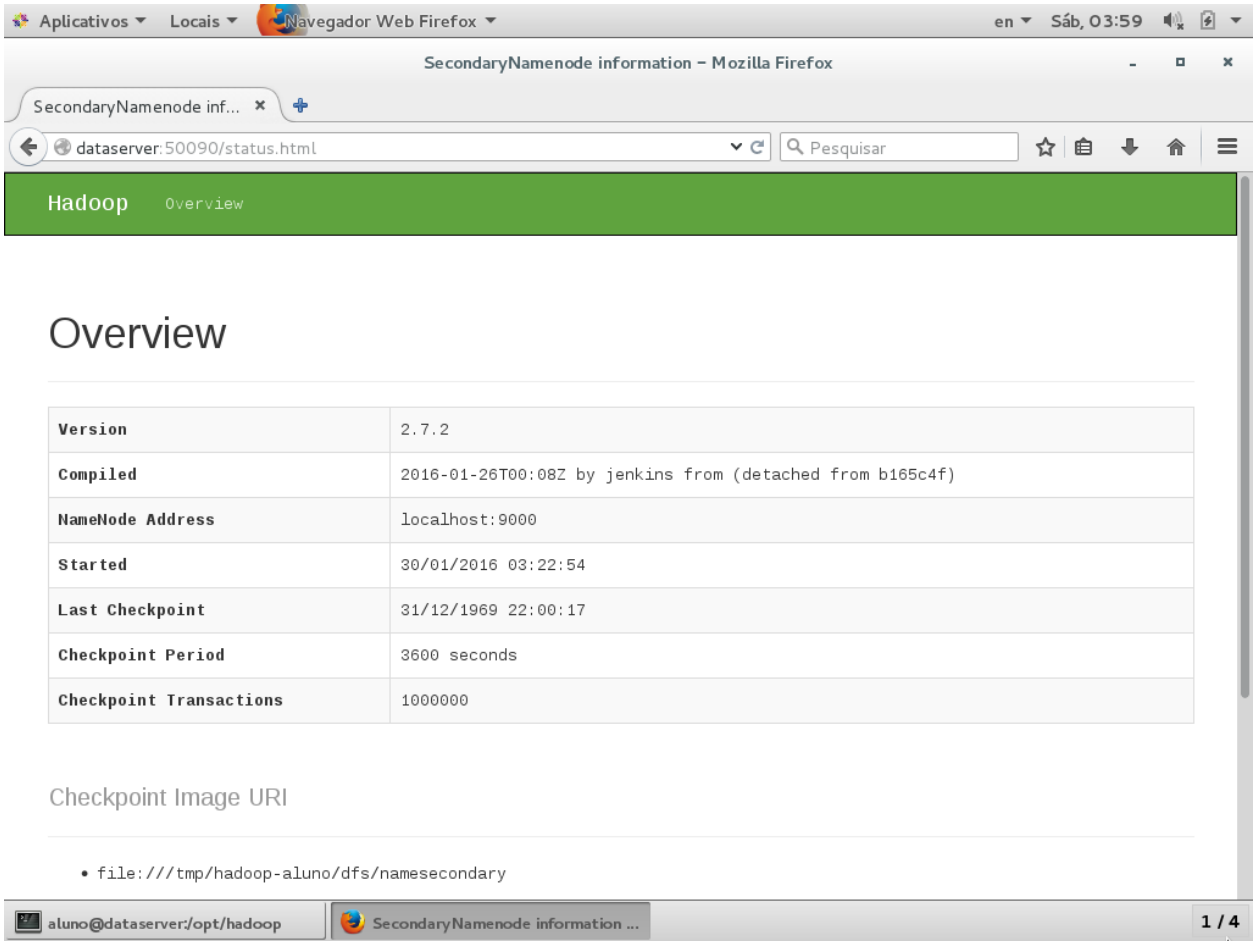

Acesso ao Hadoop pelo browser: http://dataserver:50090

Terceiro checkpoint:

Clique no meu File – Export Appliance.
Será gerada uma cópia de segurança da sua máquina virtual.

→ VM: DataServer-3.0.ova (Hadoop)

6. Instalação e Configuração do Zookeeper

6.1. Download e Instalação do Zookeeper

Download do Zookeeper – Versão 3.5.5

Faça o download, descompacte o arquivo e mova o diretório para /opt/zookeeper da mesma forma como você fez com o Java JDK e com o Hadoop.

6.2. Configurando do Zookeeper

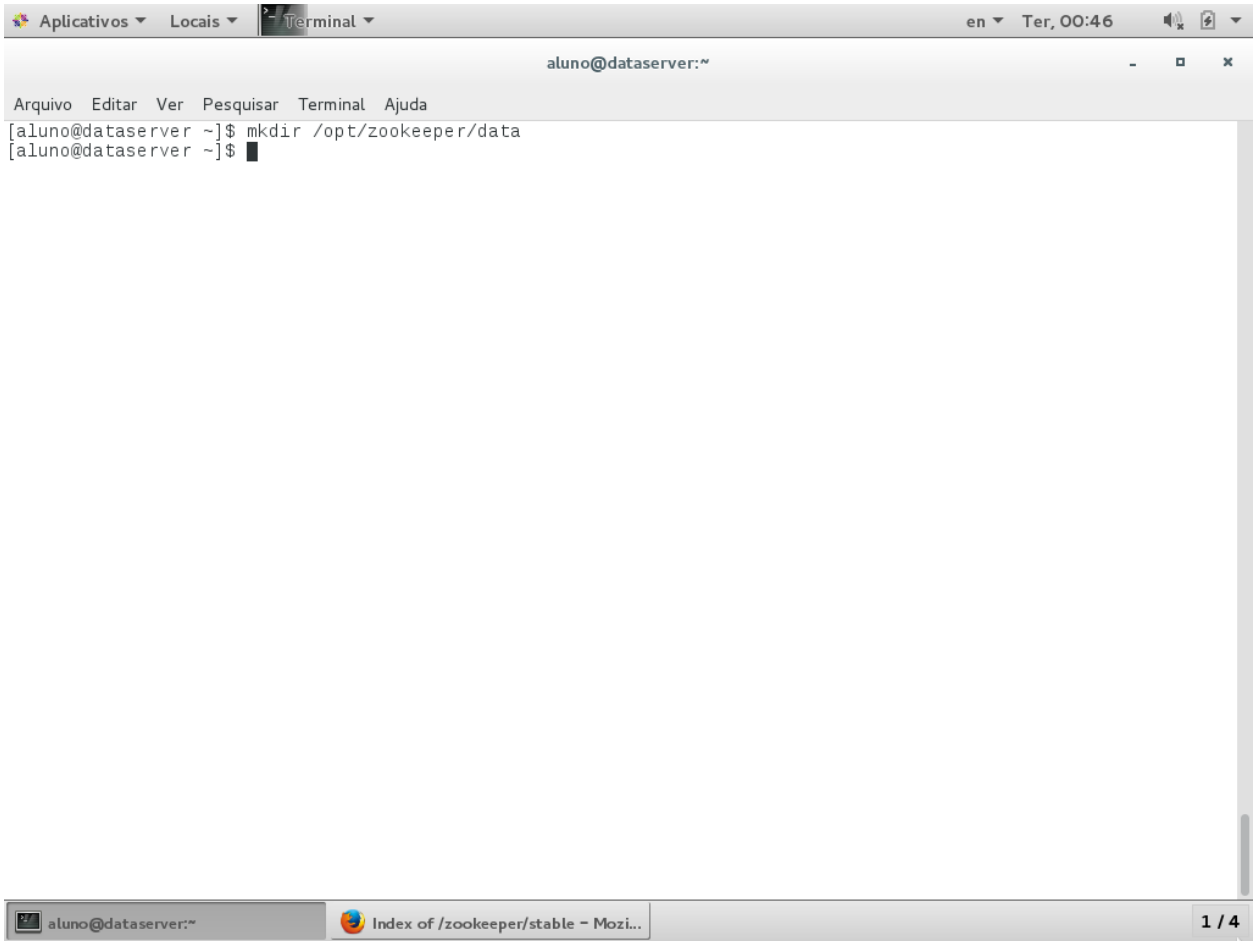

Criar o diretório **data** dentro de /opt/zookeeper

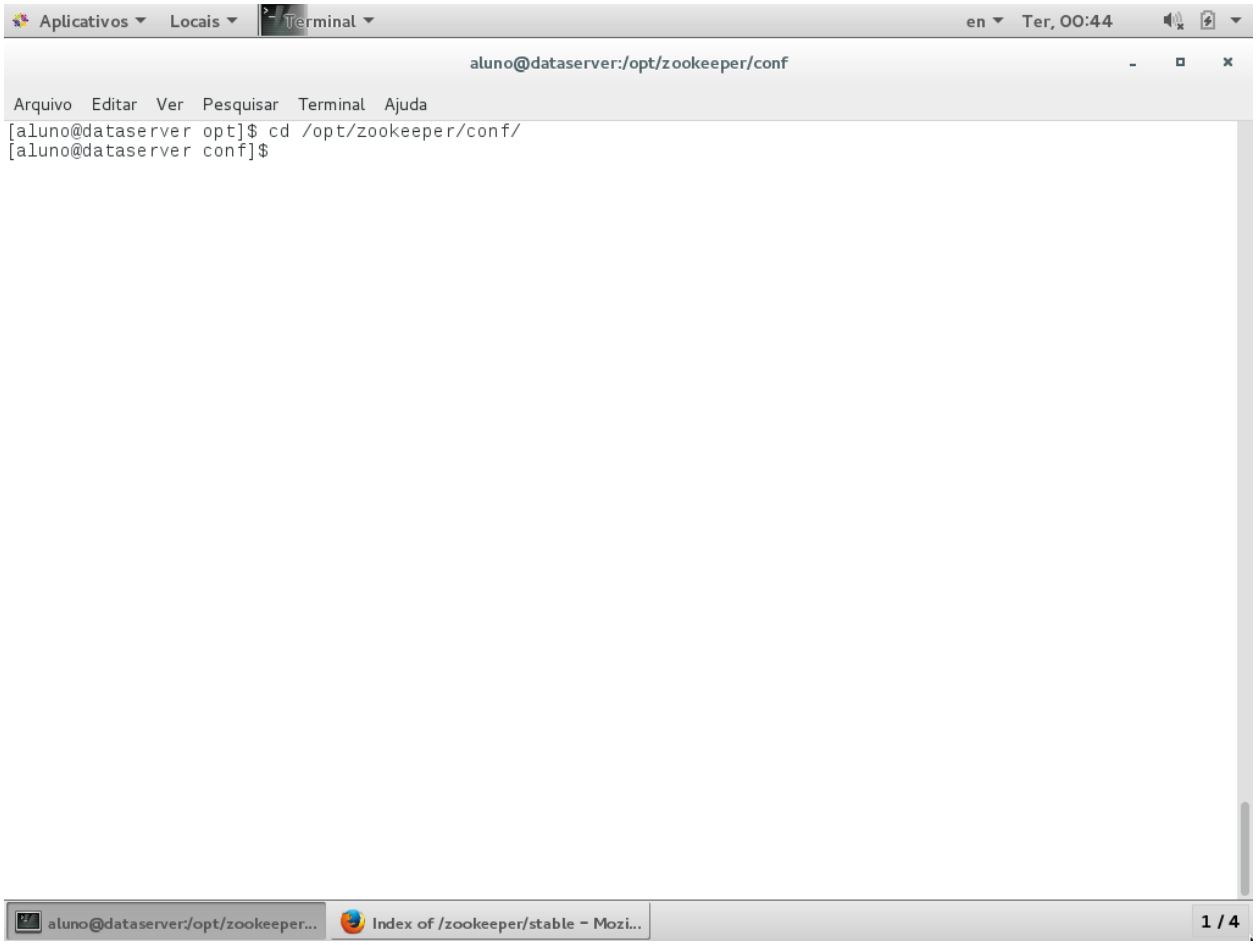

Acessar o diretório /opt/zookeeper/conf

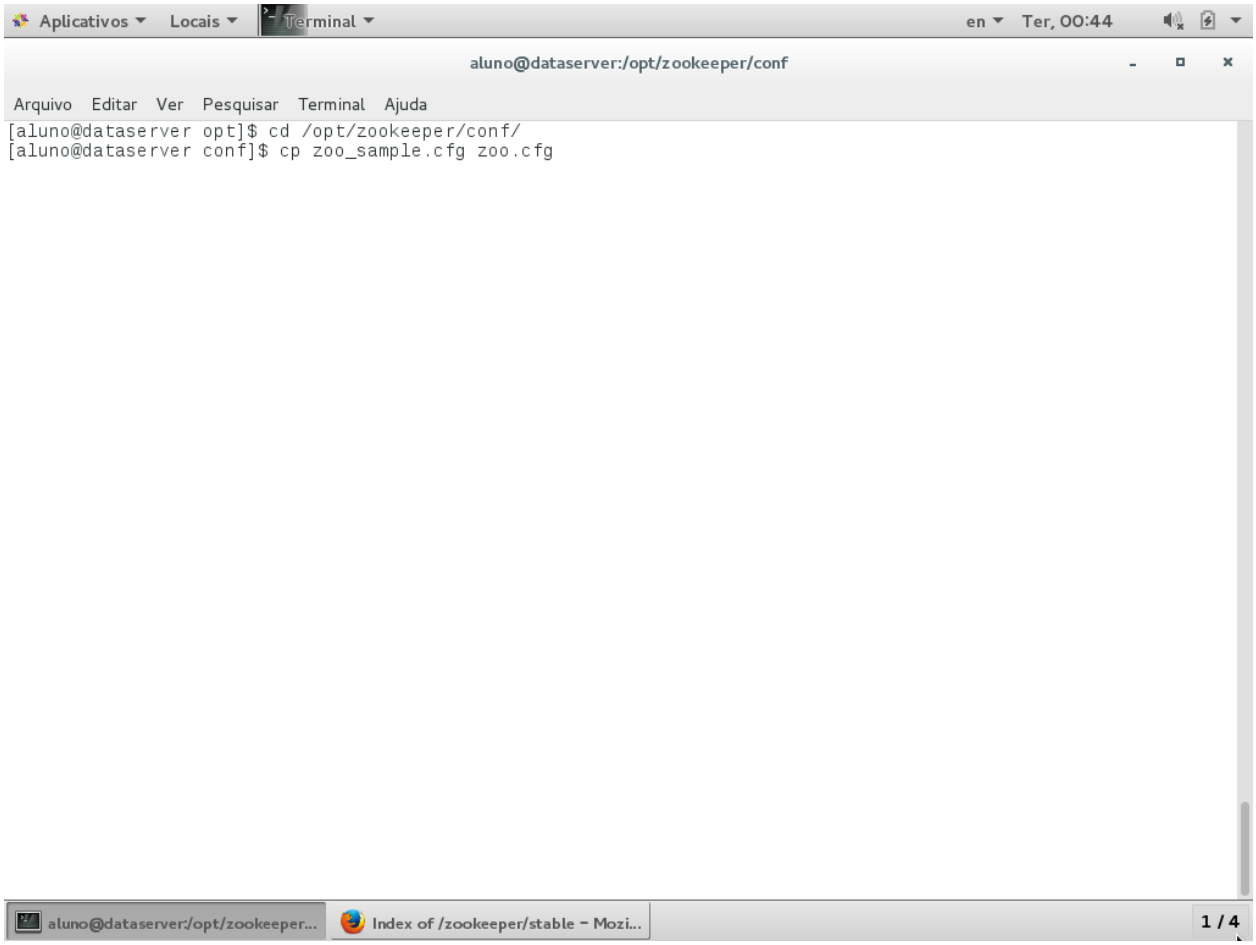

A partir do arquivo template, gerar o arquivo zoo.cfg

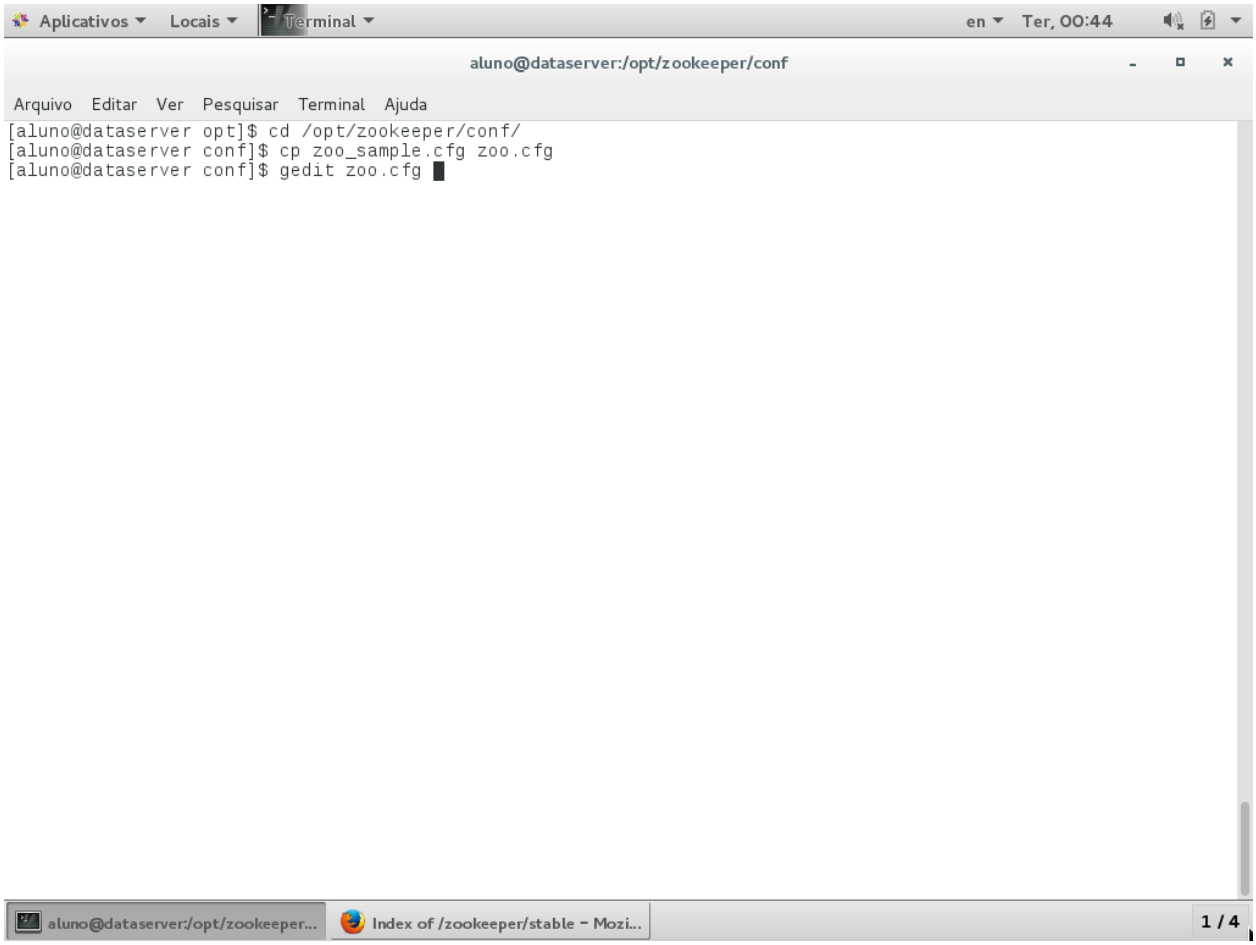

Editar o arquivo zoo.cfg

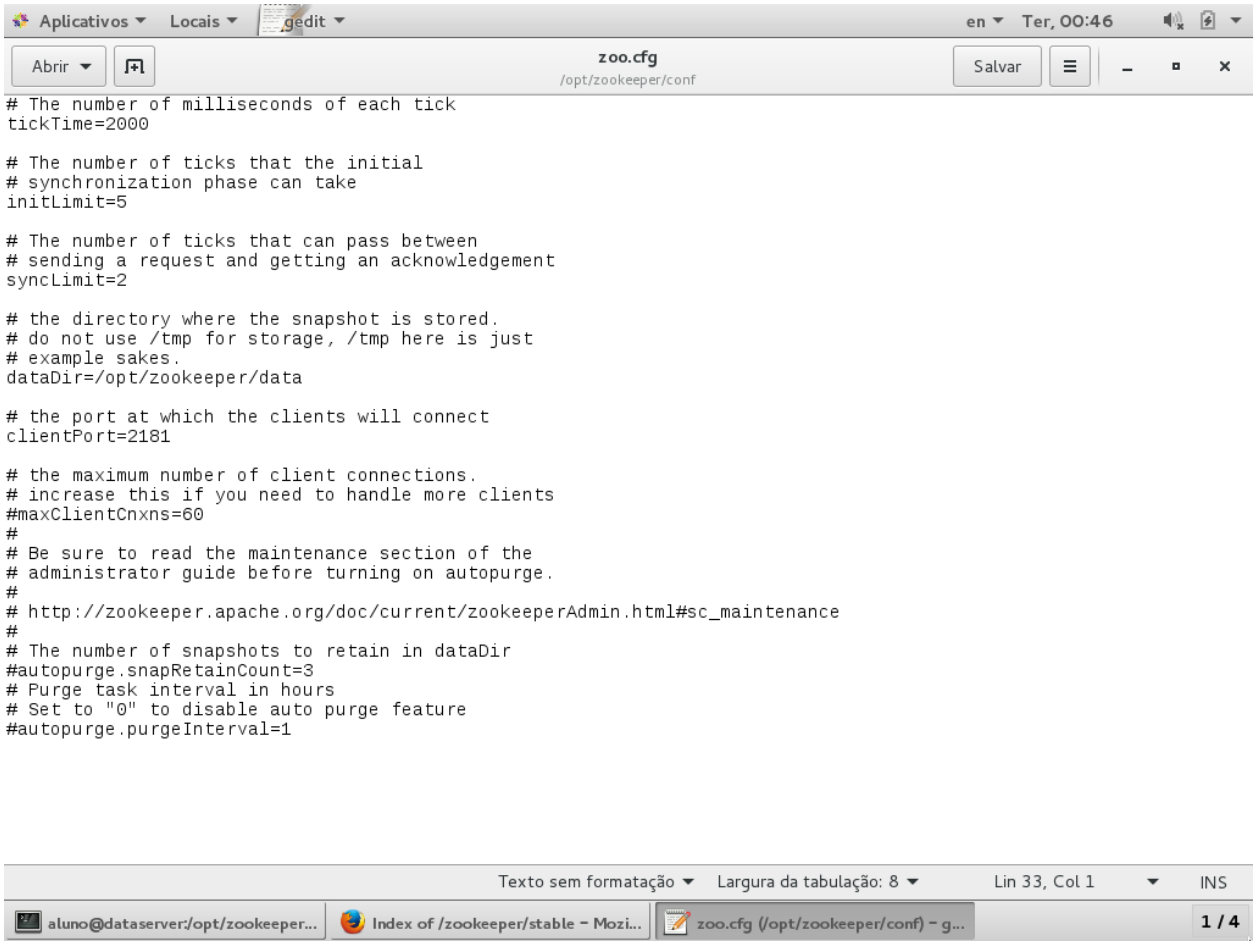

Editar o arquivo conforme tela acima

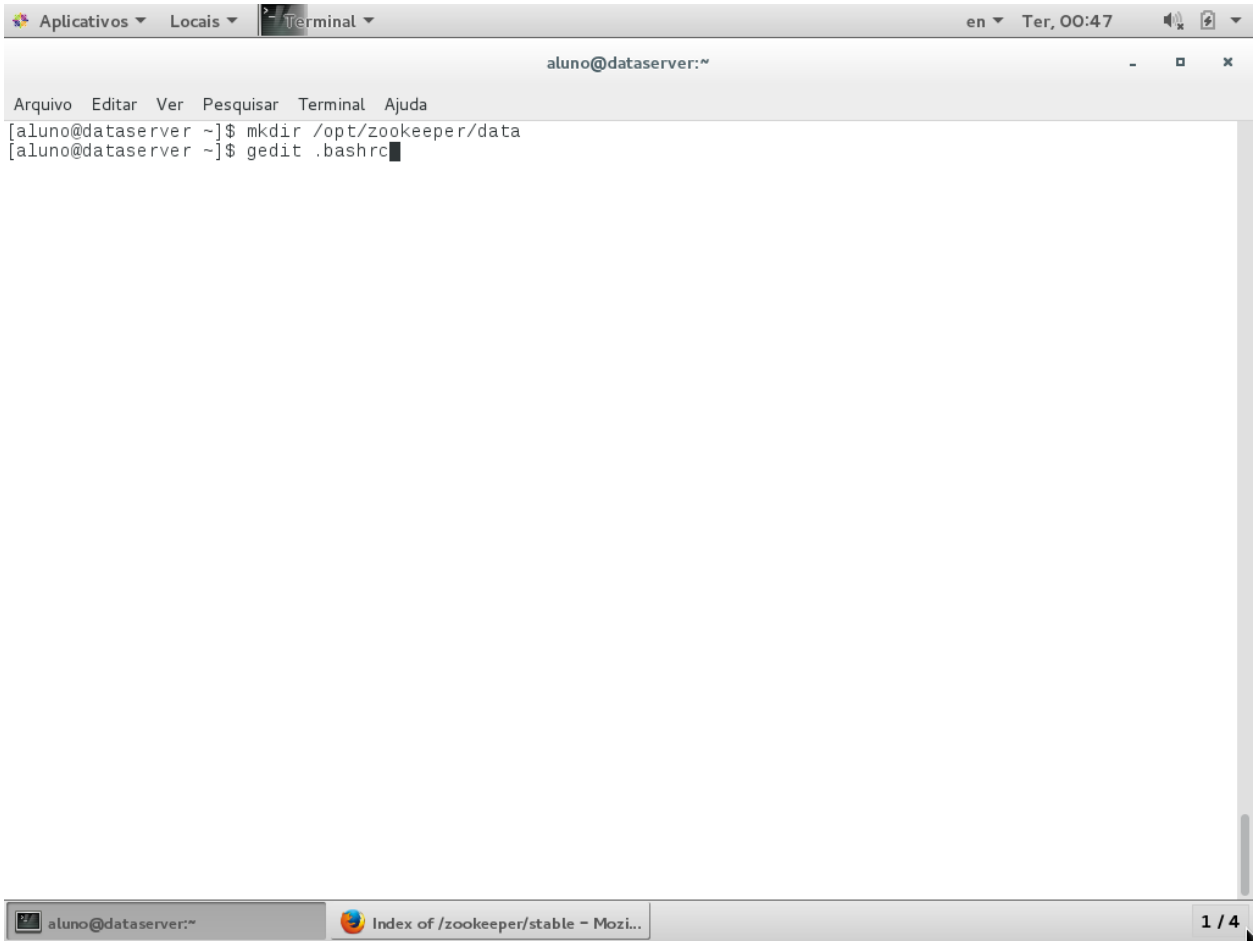

Incluir variáveis Zookeeper no /home/hadoop/.bashrc

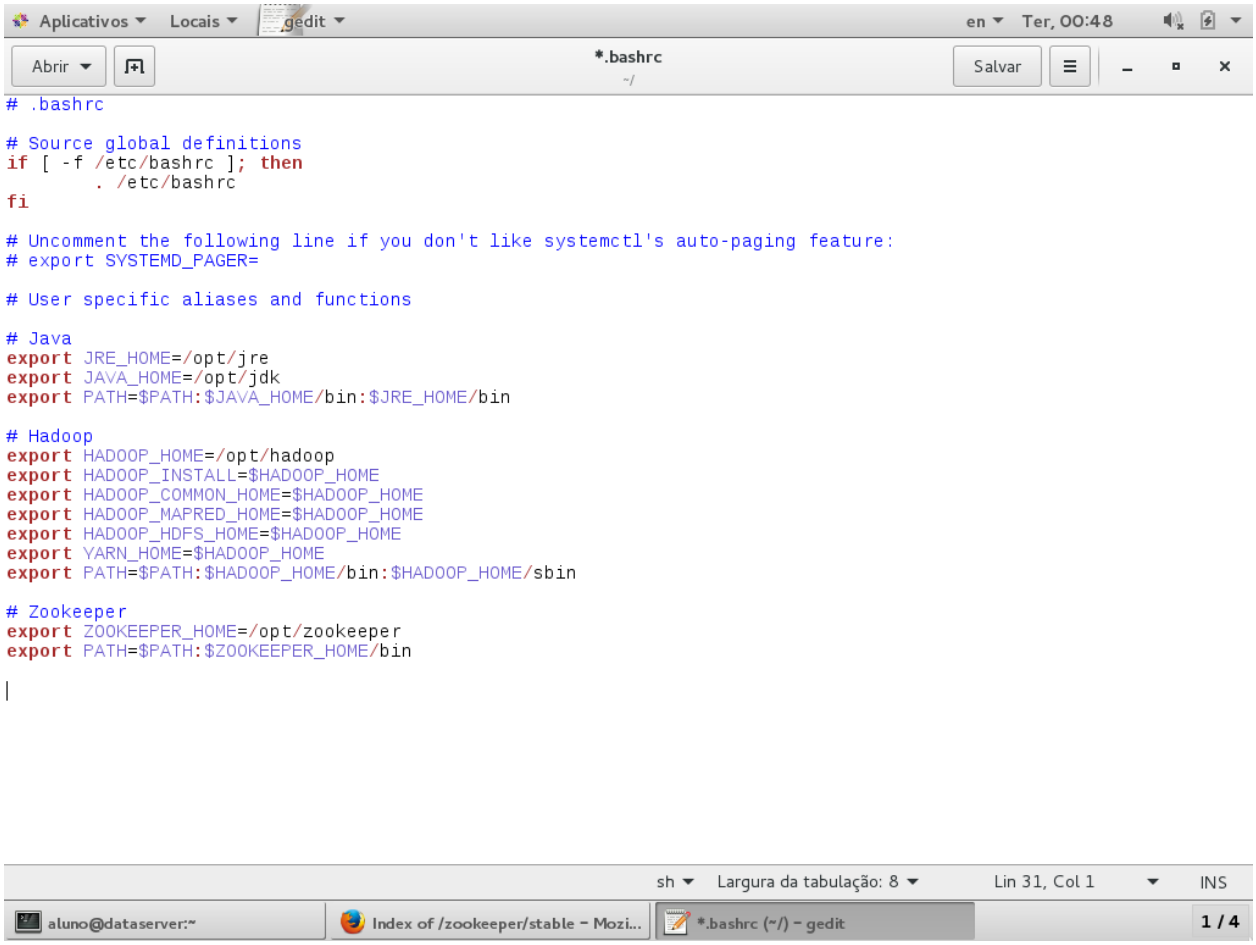

Variáveis Zookeeper

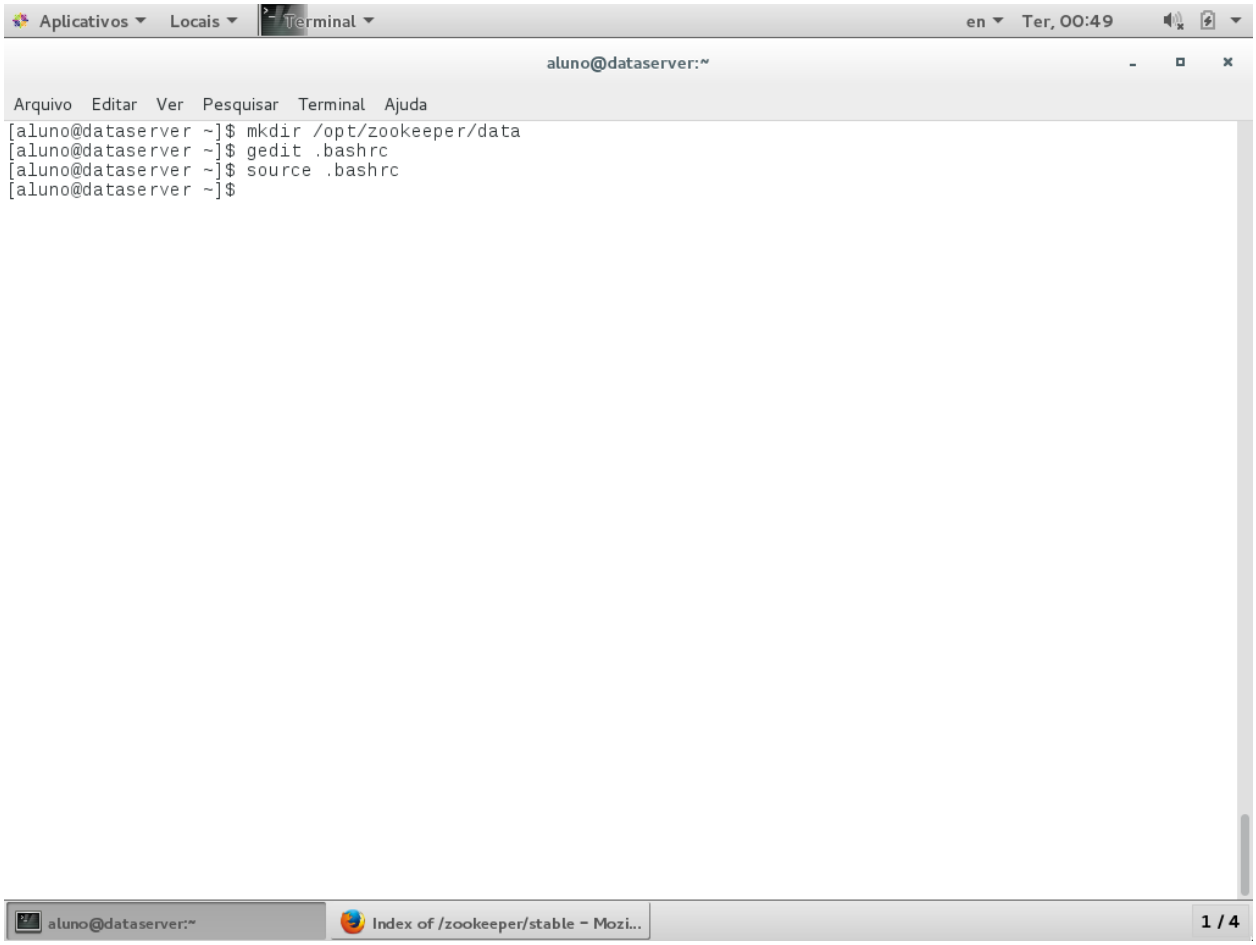

source .bashrc

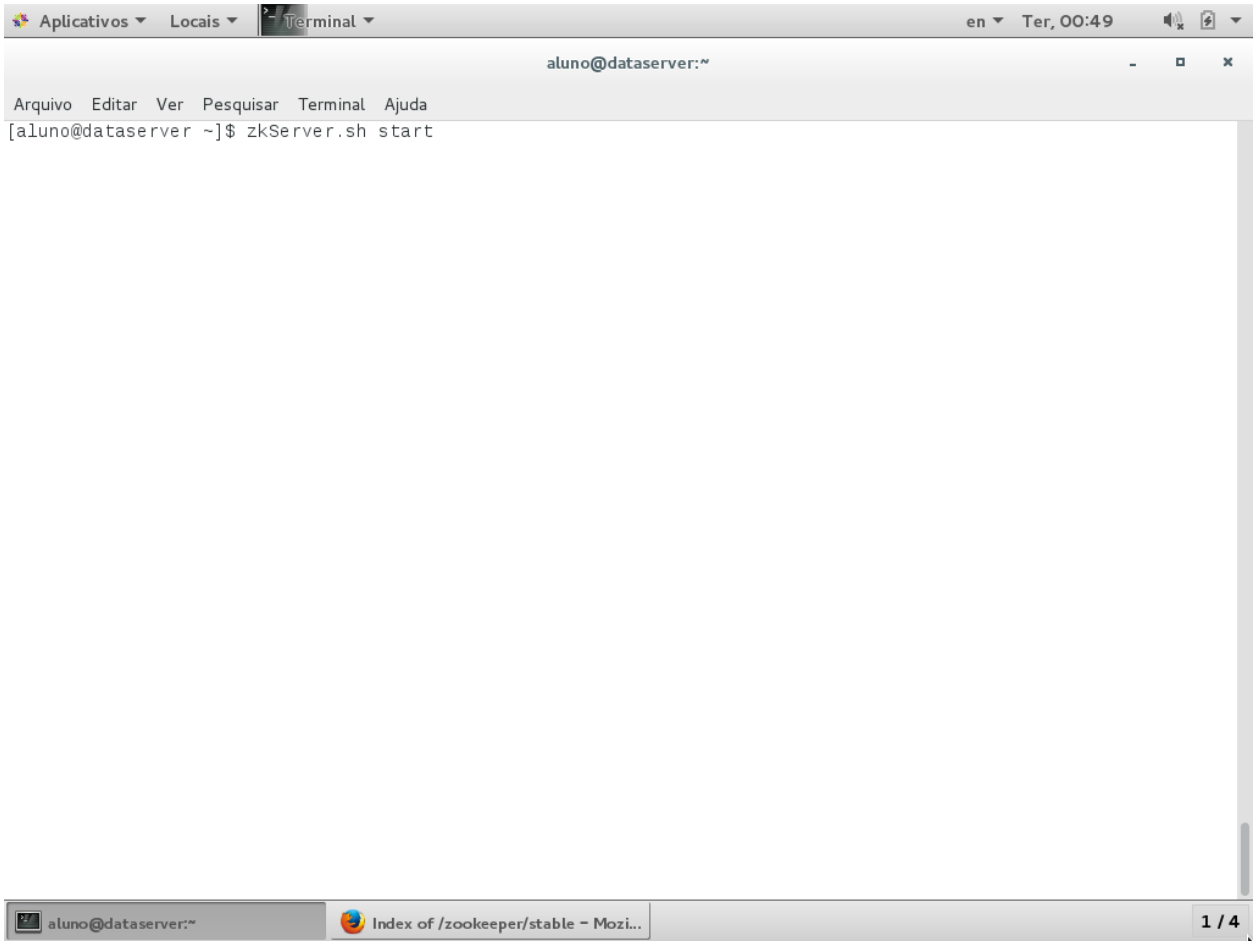

Iniciar o zookeeper

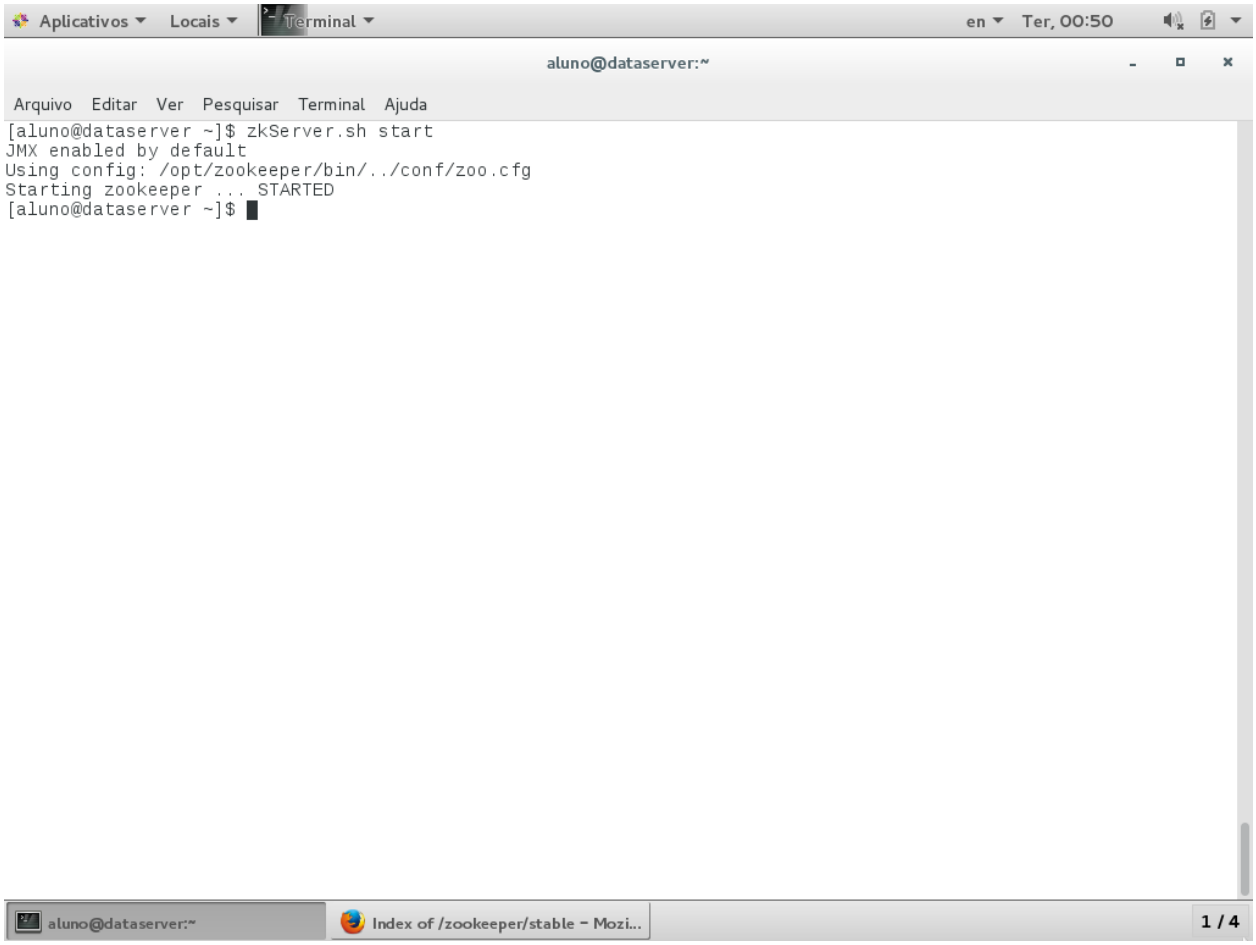

Serviço iniciado

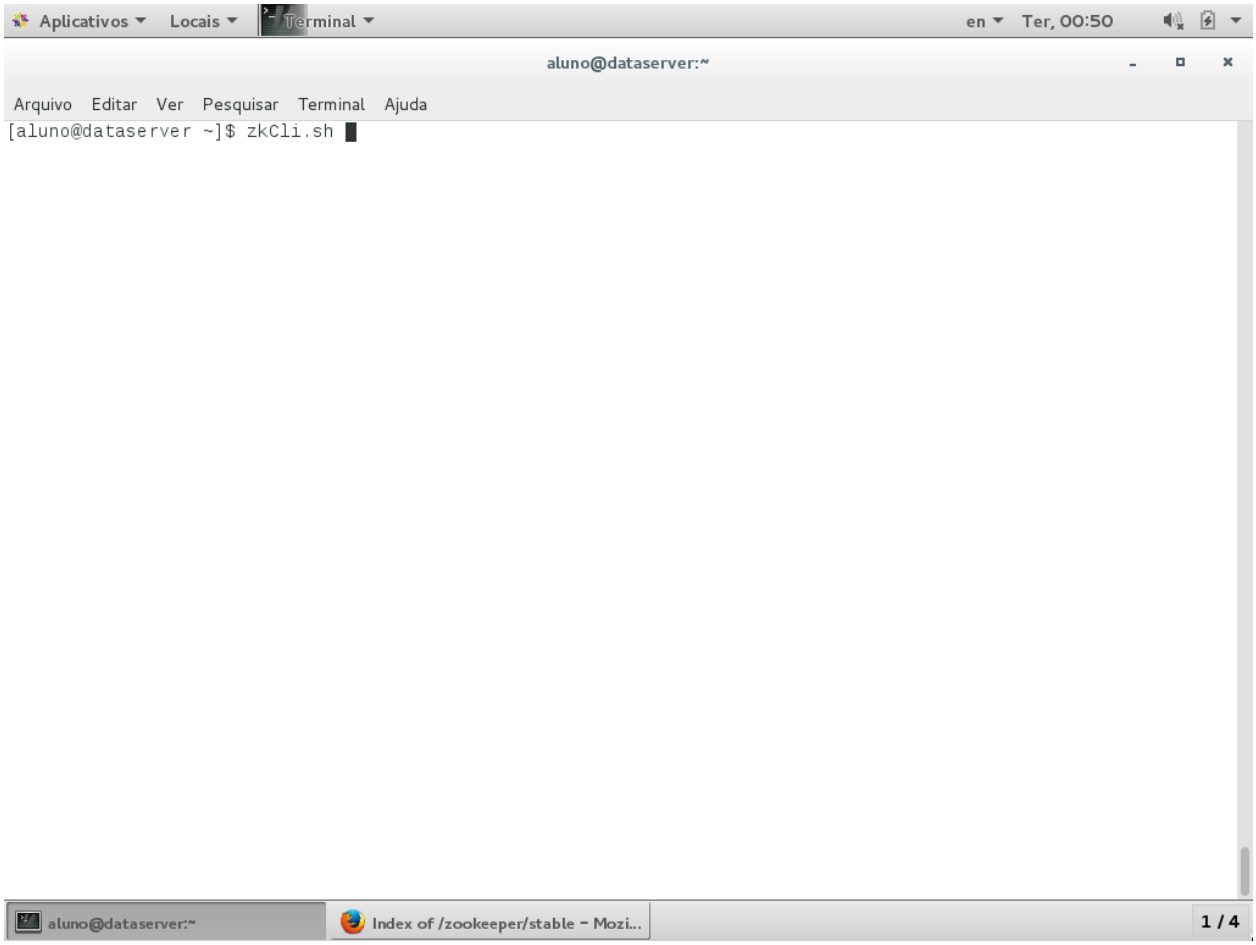

Iniciar o Zookeeper Command Line Interface (CLI)

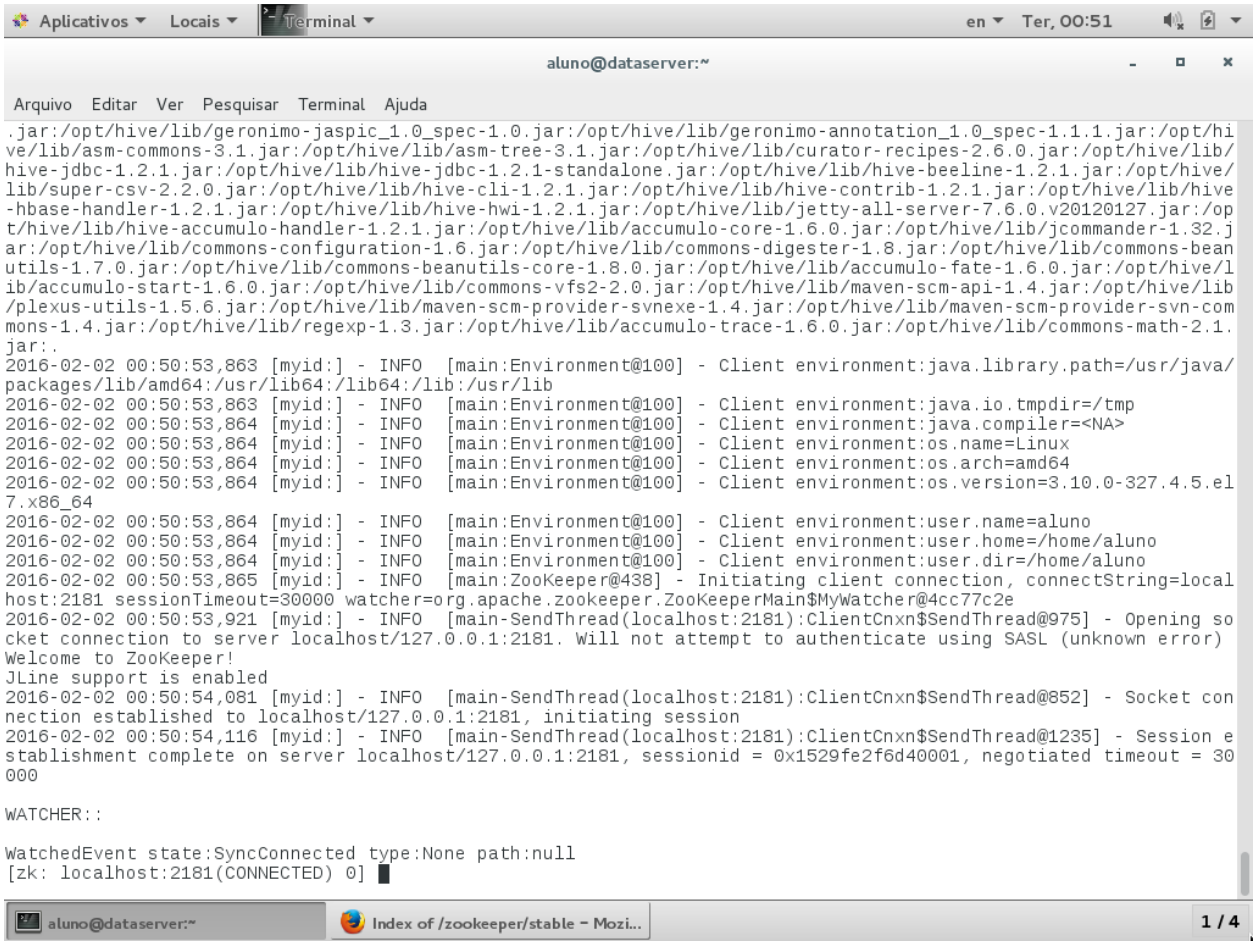

CLI iniciado

7. Instalação e Configuração do HBase

Podemos instalar HBase em qualquer um dos três modos: Standalone mode, Pseudo Distributed mode e Fully Distributed mode.

7.1. Download e Instalação do HBase

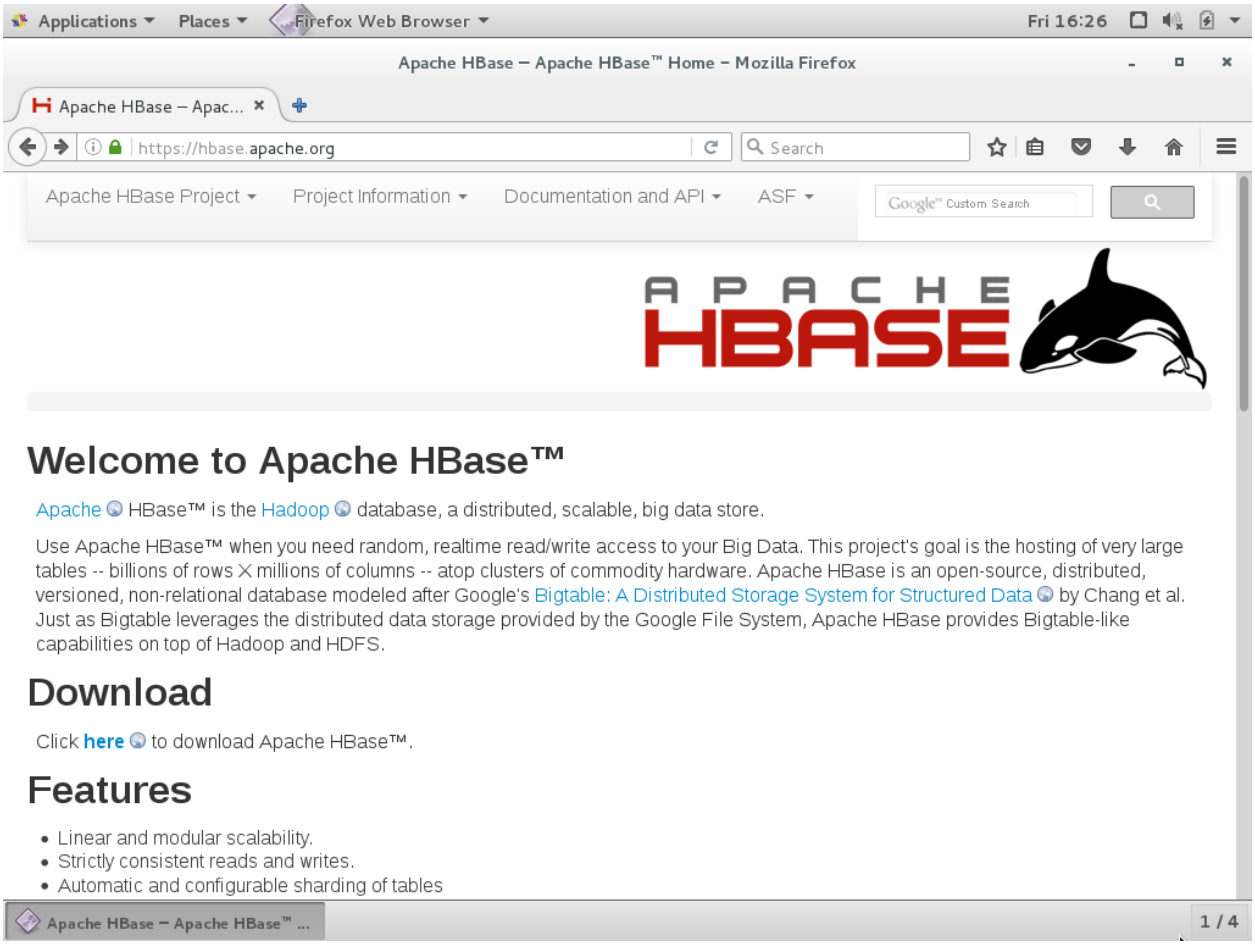

Download do Hbase – Versão 2.2.0

Faça o download, descompacte o arquivo e mova o diretório para /opt/hbase da mesma forma como você fez com o Java JDK e com o Hadoop.

7.2. Configurando o HBase

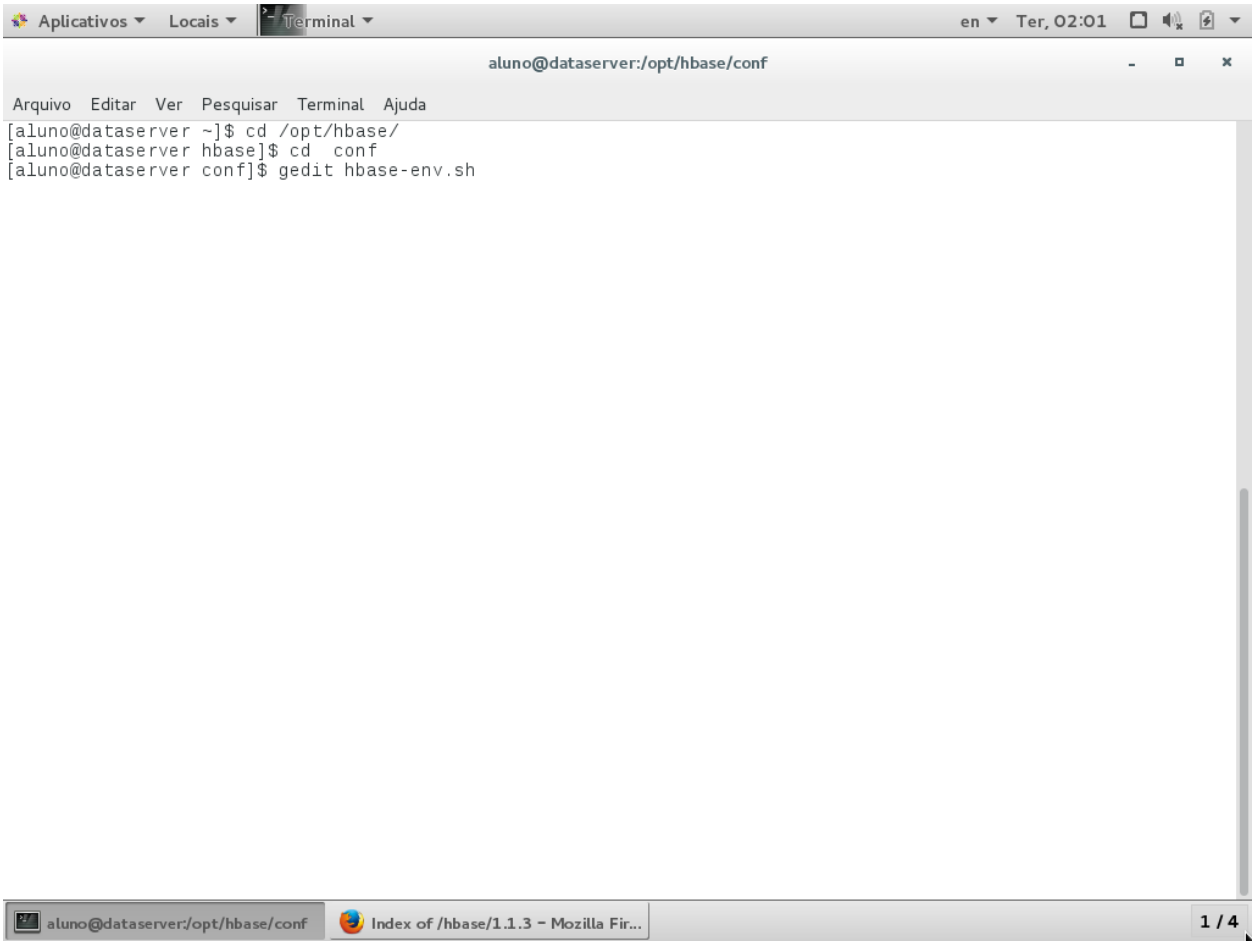

No diretório /opt/hbase/conf, editar o arquivo hbase-env.sh

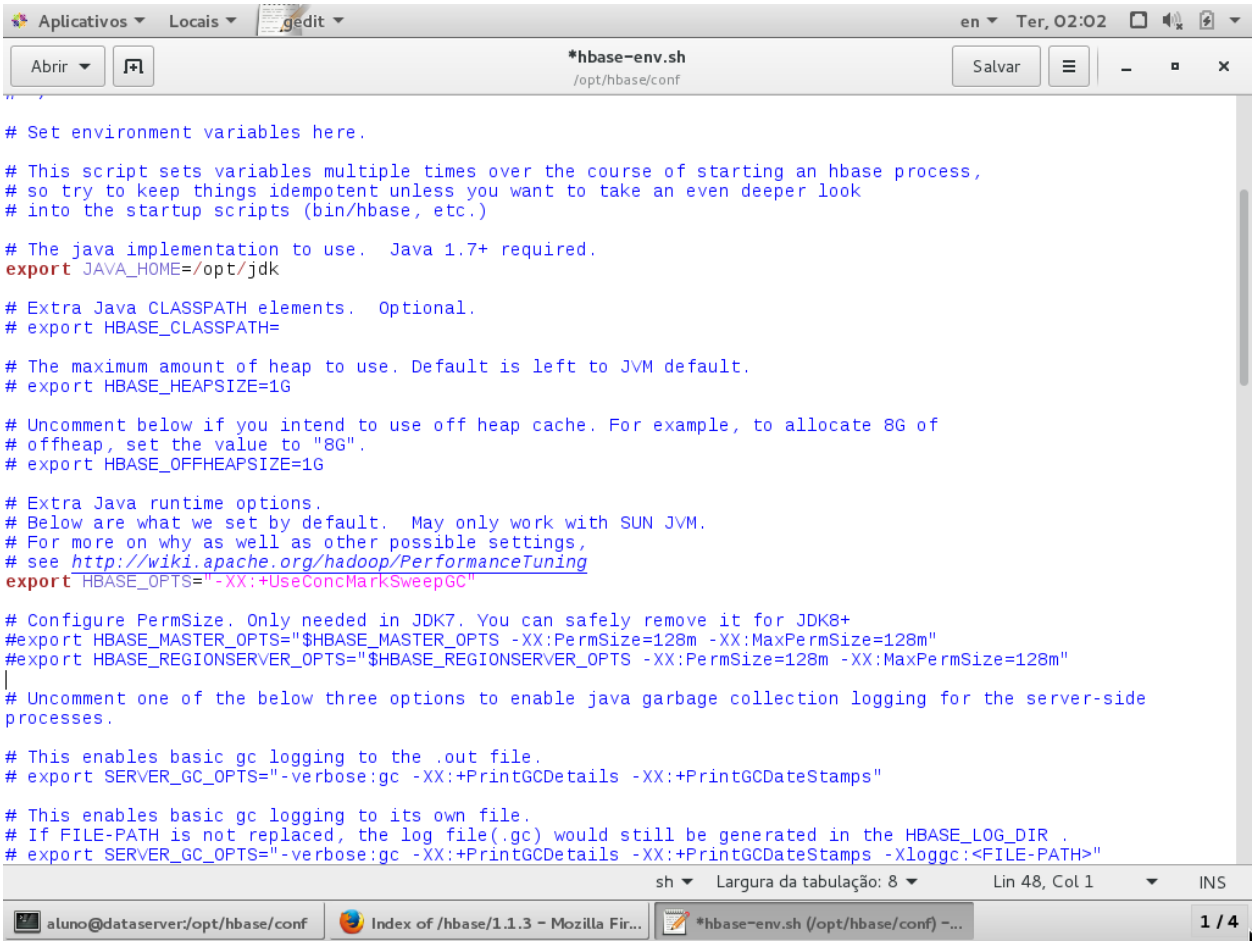

Editar o PATH do Java e comentar as linhas do PermSize

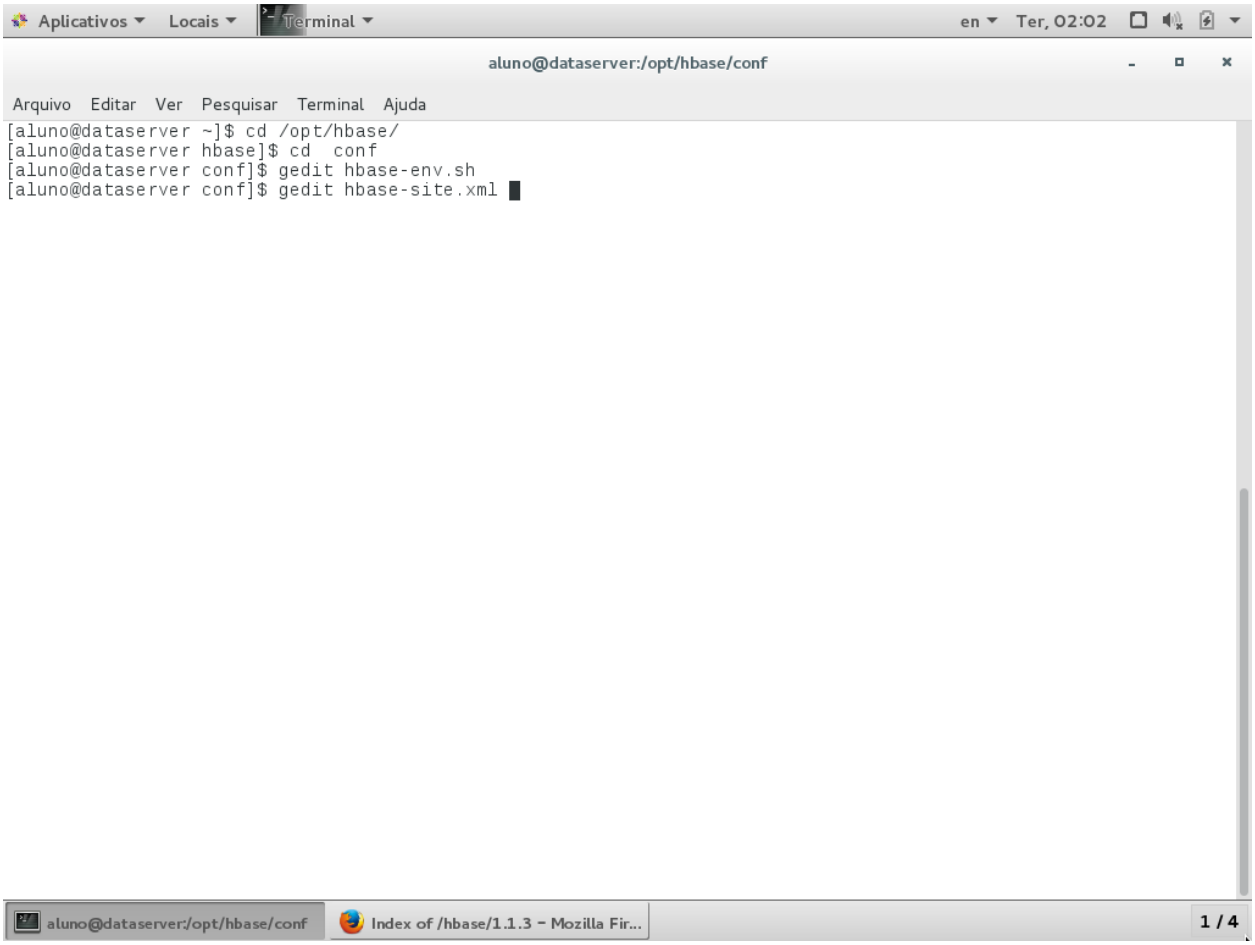

No mesmo diretório conf, editar o arquivo hbase-site.xml

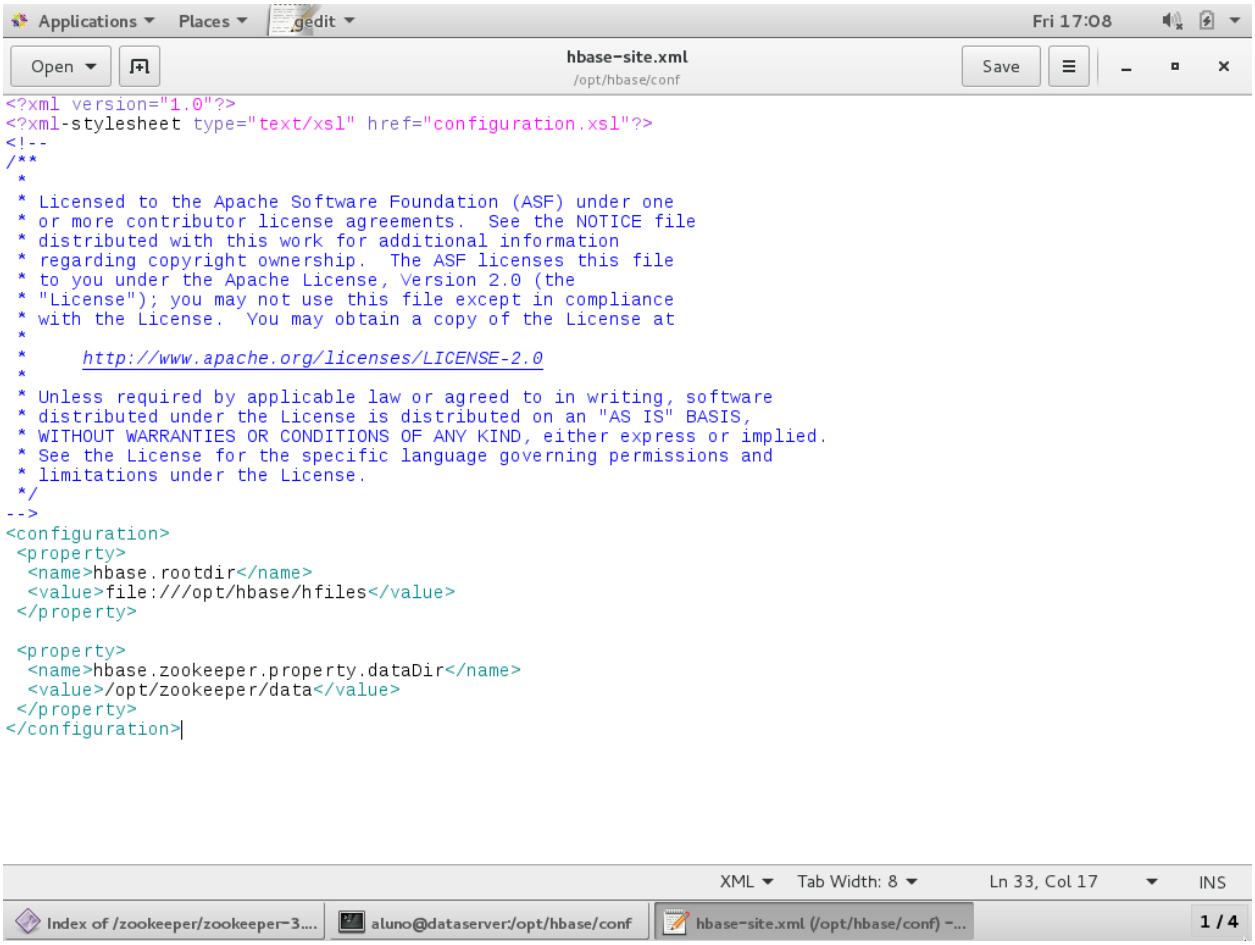

Incluir as linhas entre as tags <configuration>

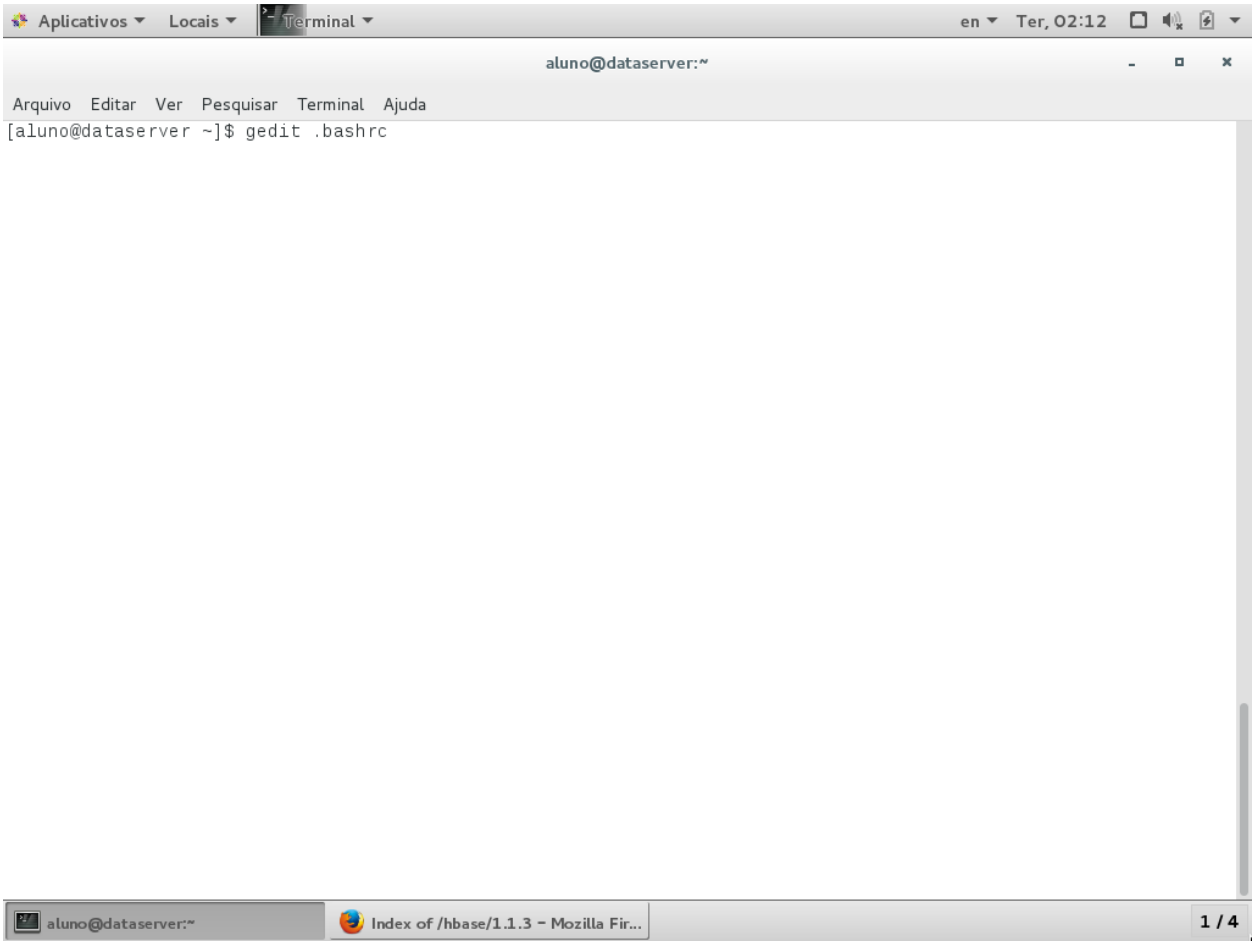

Editar o arquivo .bashrc

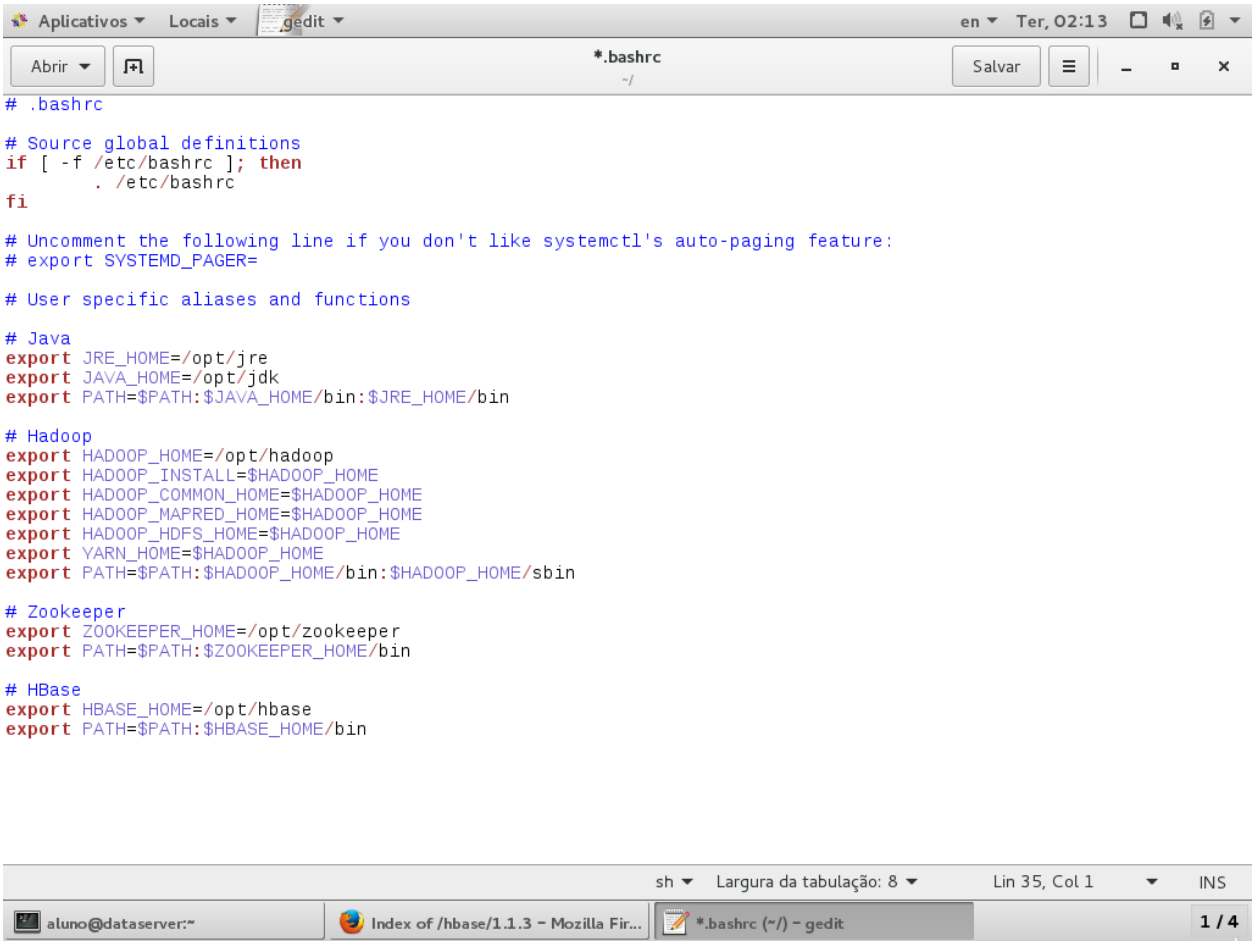

Variáveis HBase

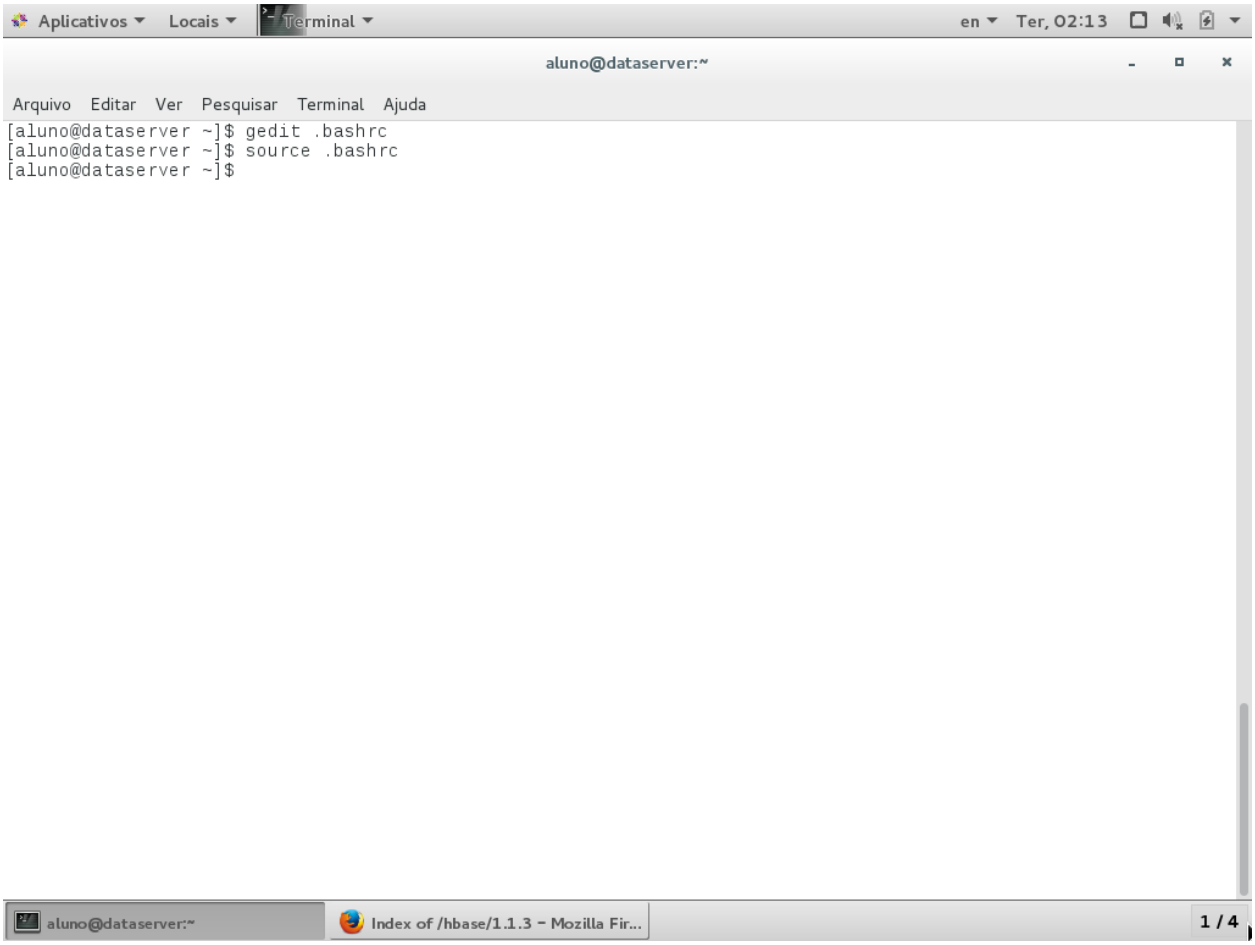

source .bashrc

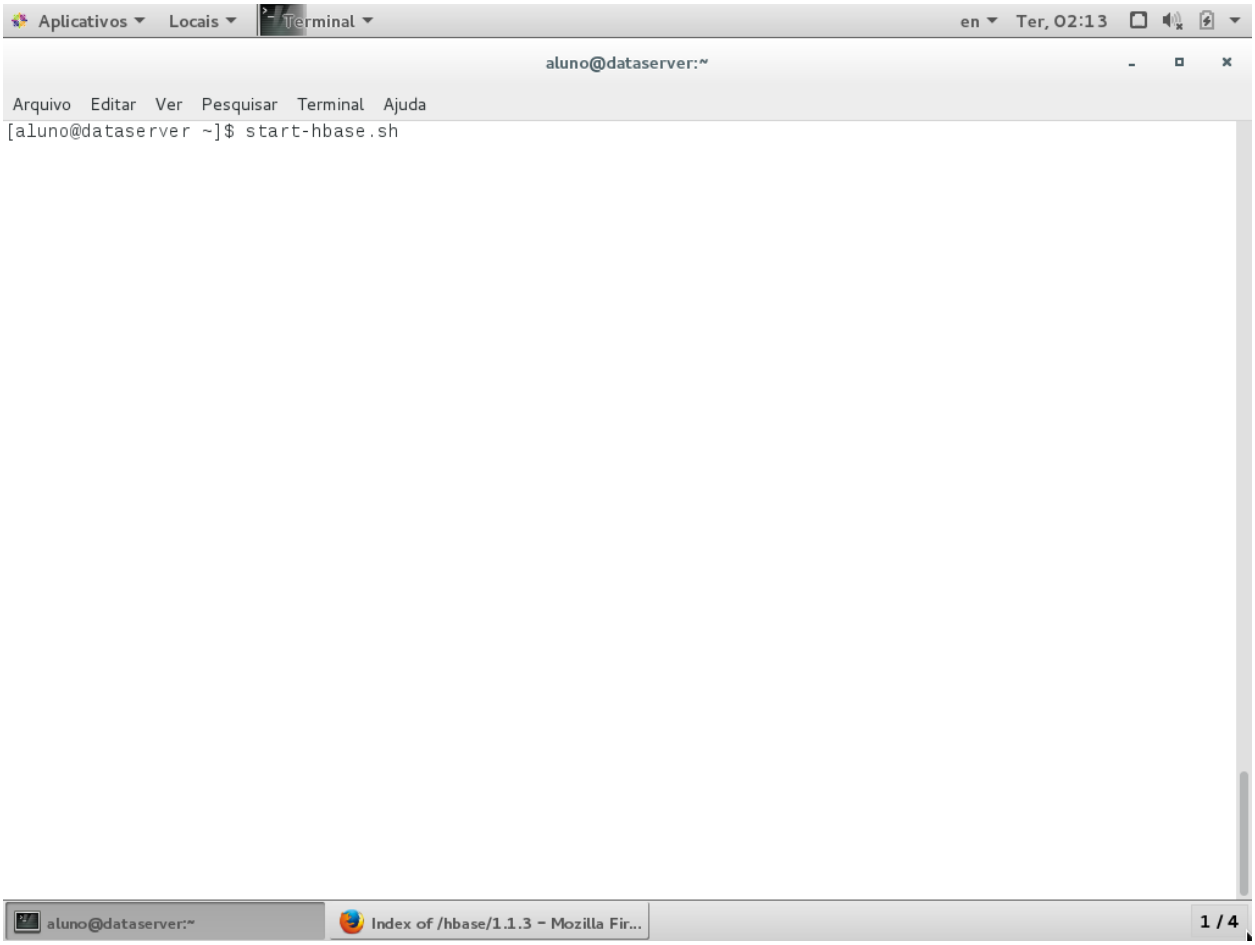

Iniciar o Hbase - start-hbase.sh

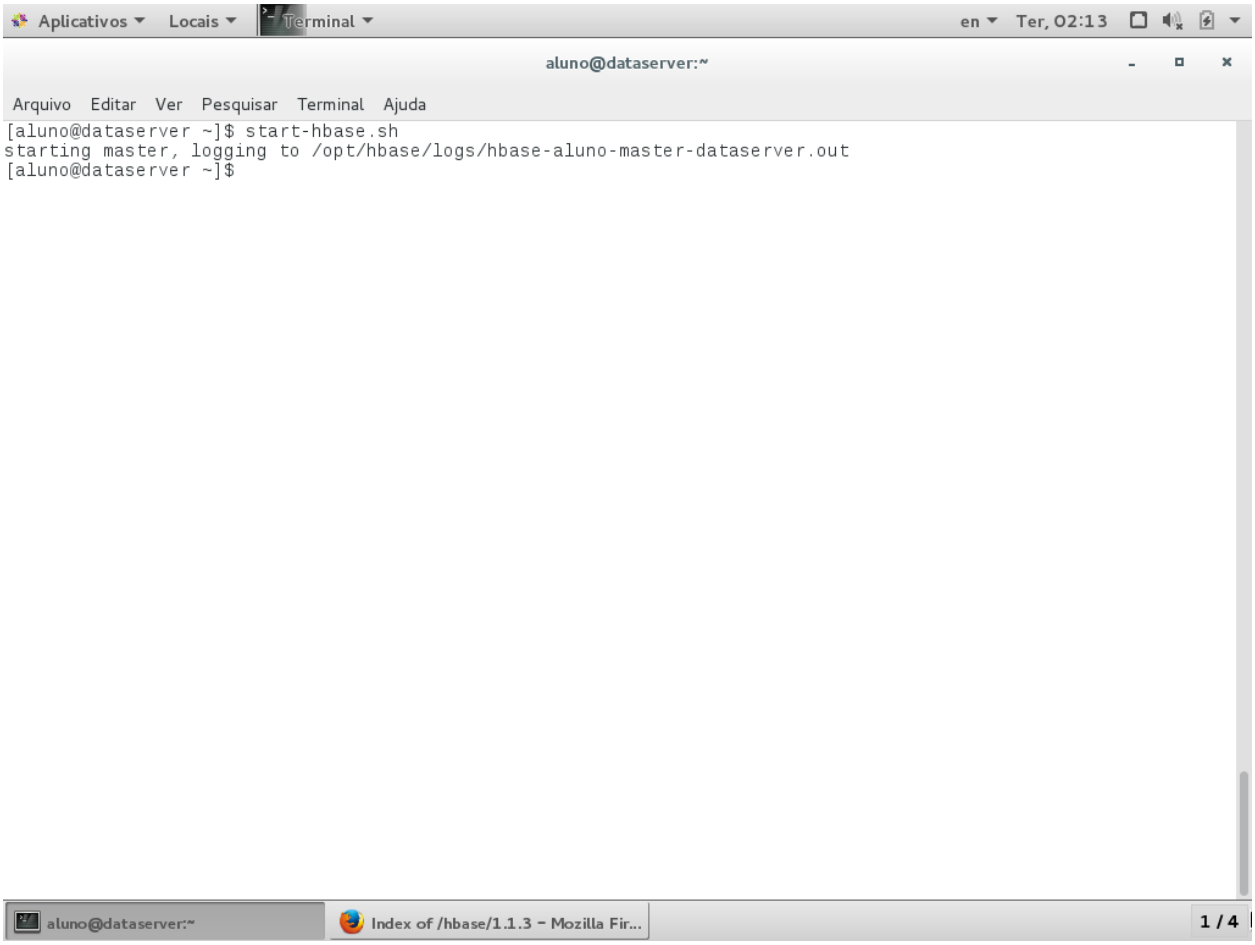

Hbase iniciado

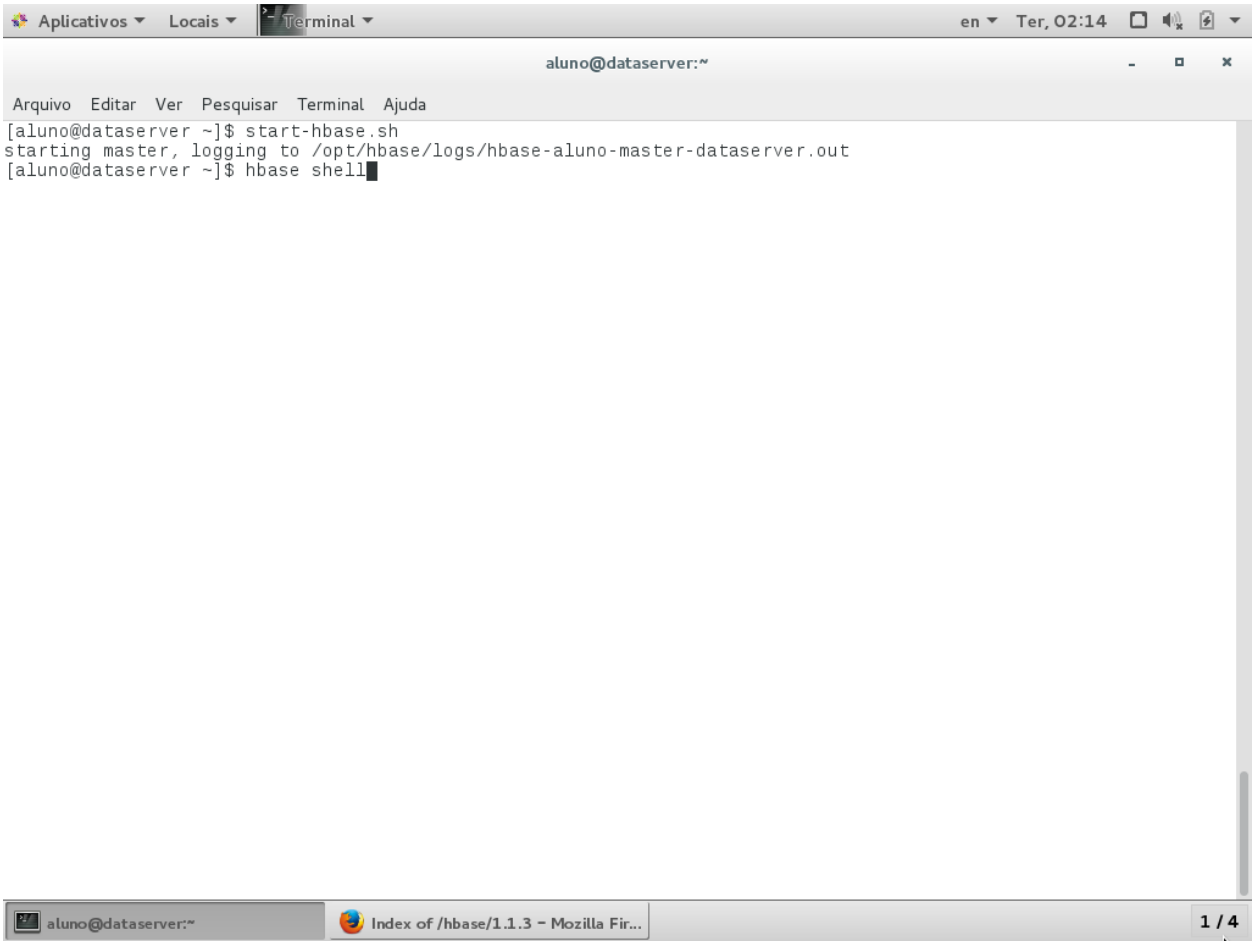

Abrir o shell do Hbase

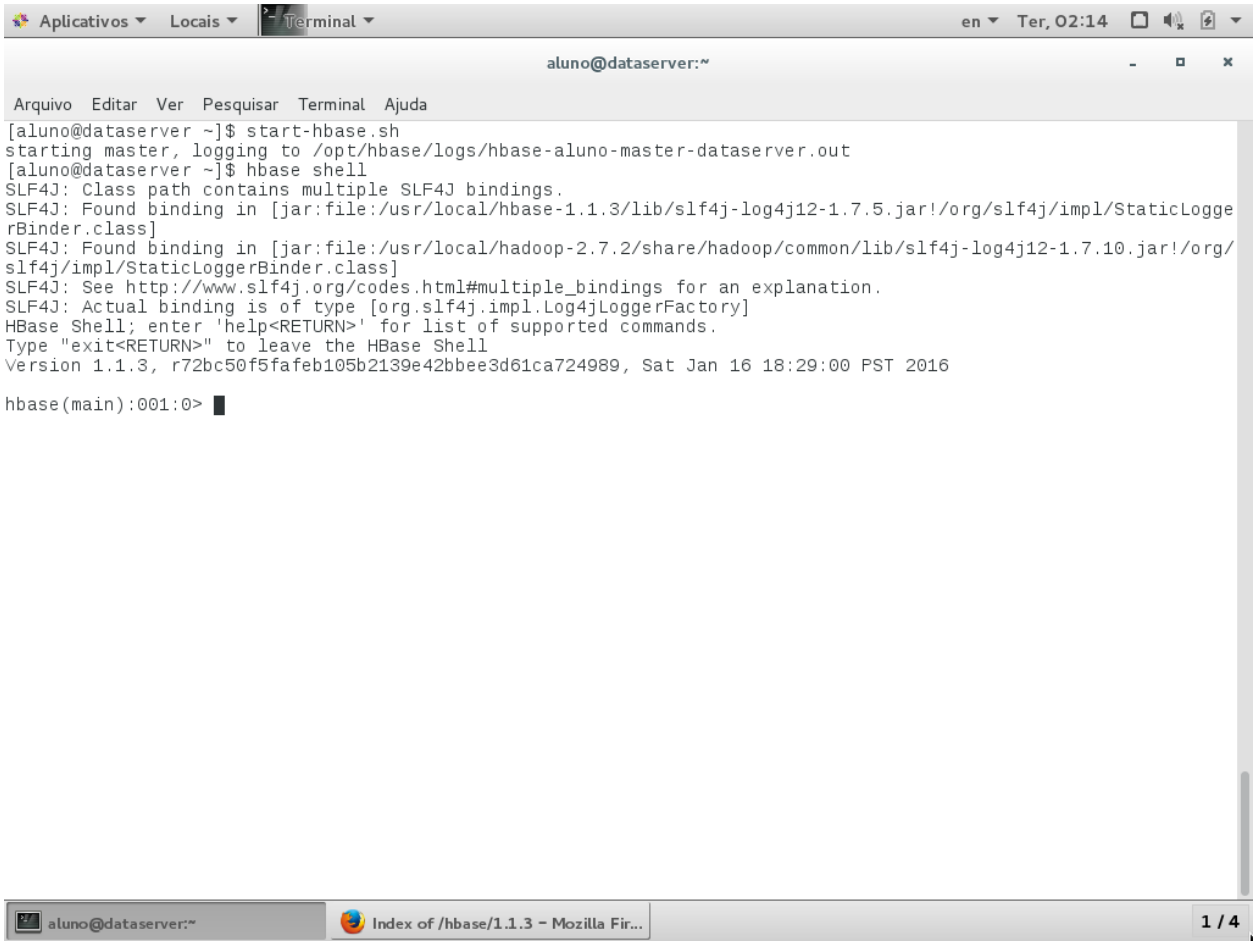

Shell iniciado

8. Instalação e Configuração do Hive

8.1. Download e Instalação do Hive

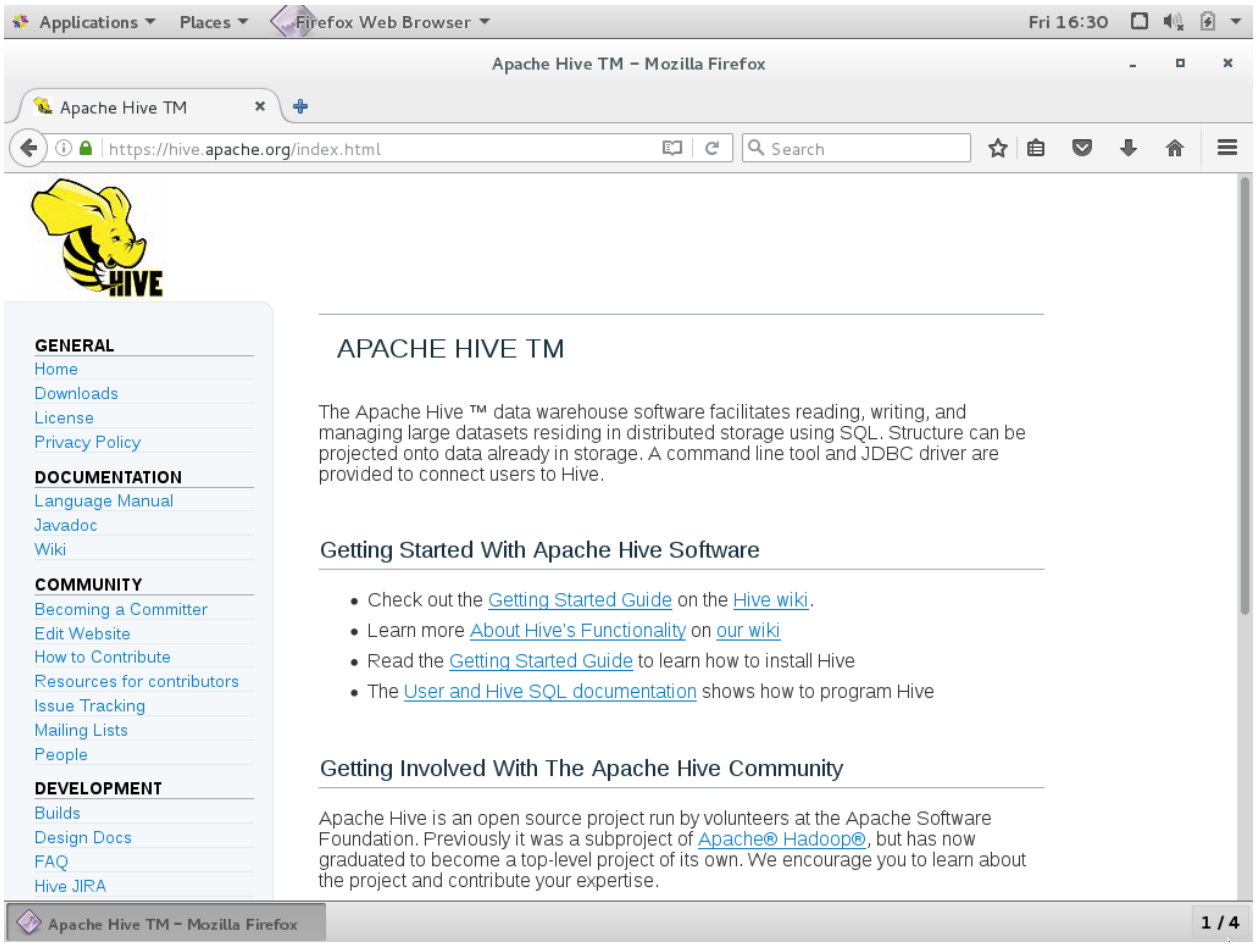

Download do Hive – Versão 3.1.1

Faça o download, descompacte o arquivo e mova o diretório para /opt/hive da mesma forma como você fez com o Java JDK e com o Hadoop.

8.2. Configurando o Hive

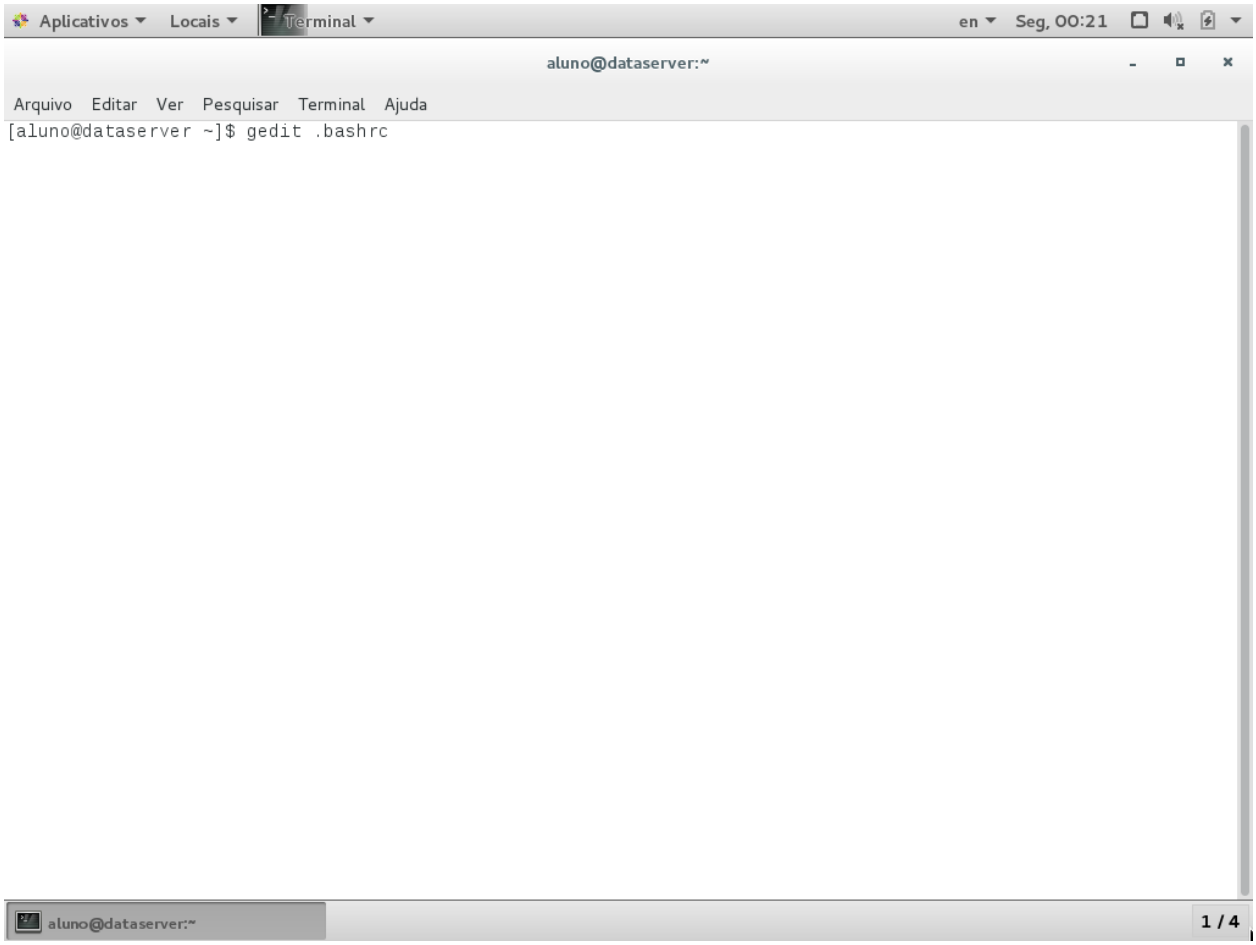

Editando o arquivo .bashrc

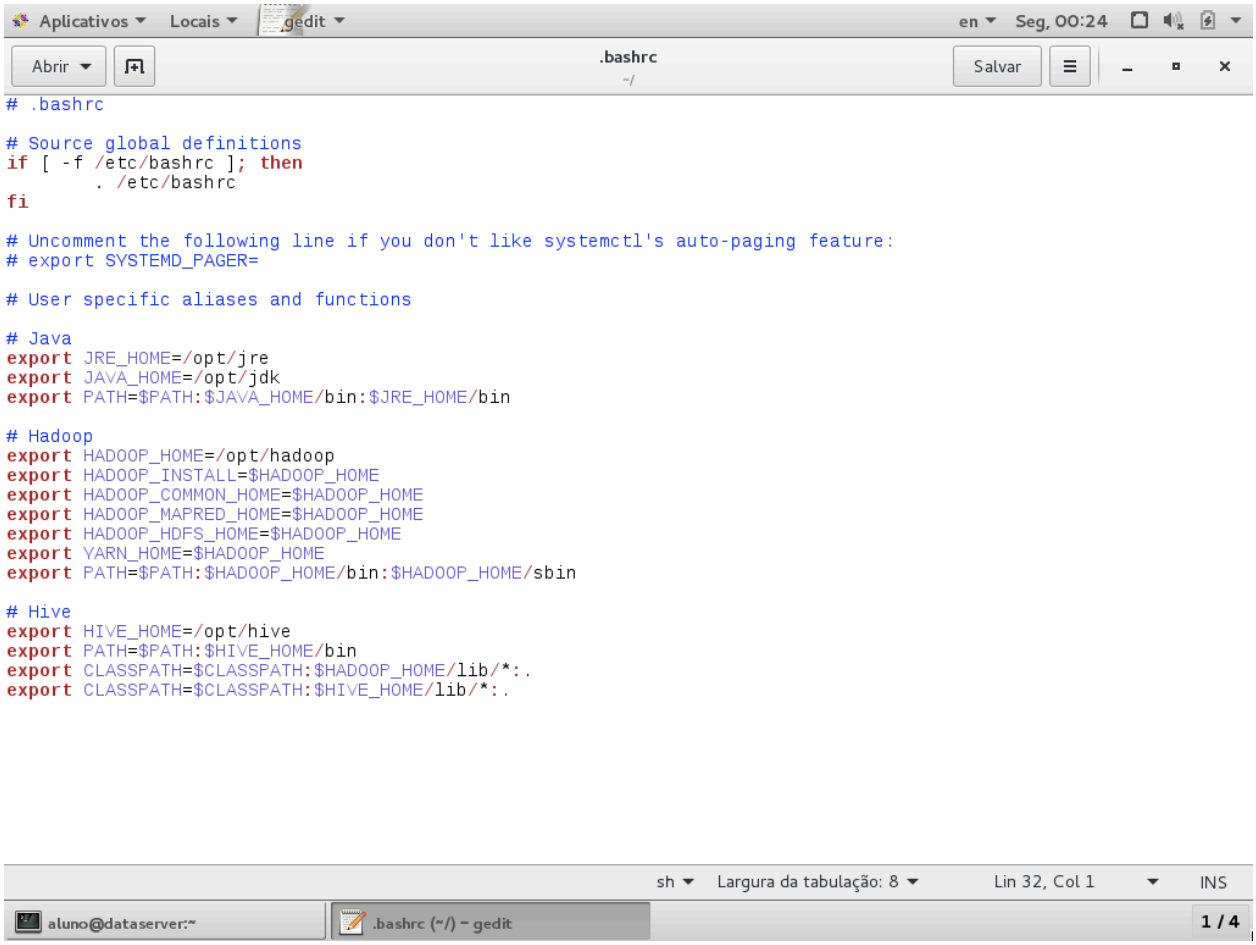

Variáveis de ambiente do Hive

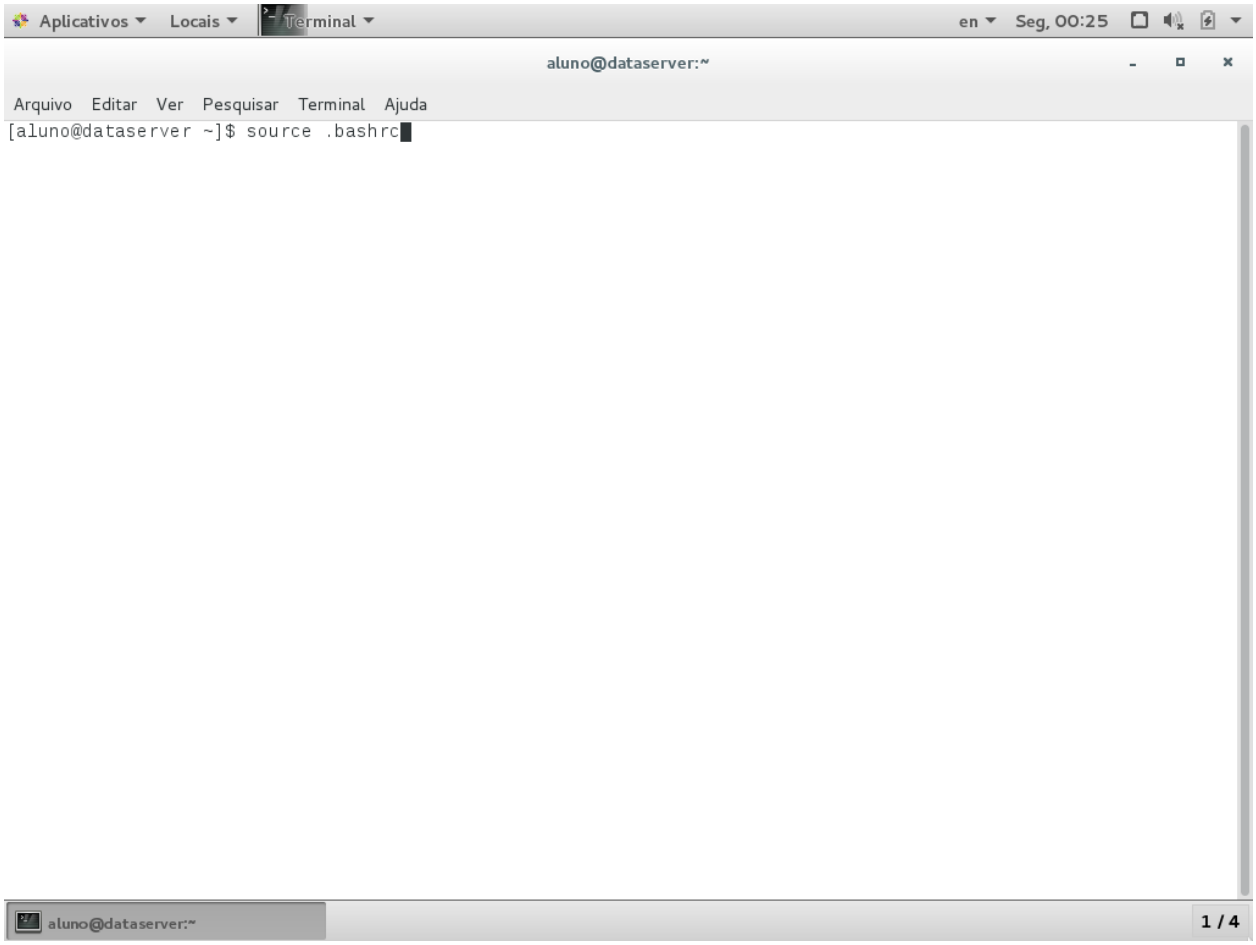

source .bashrc

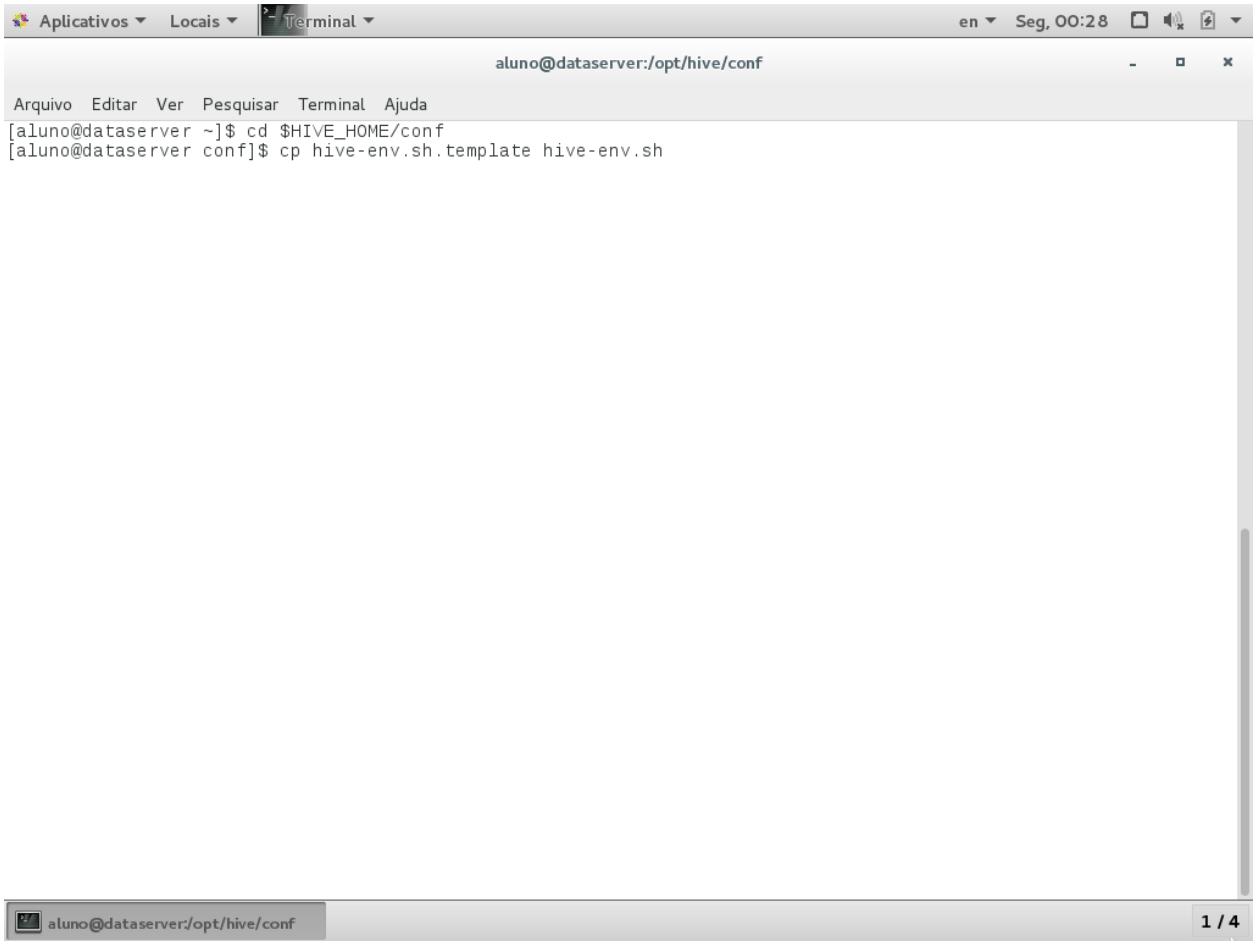

A partir do arquivo template, gerar o arquivo hive-env.sh

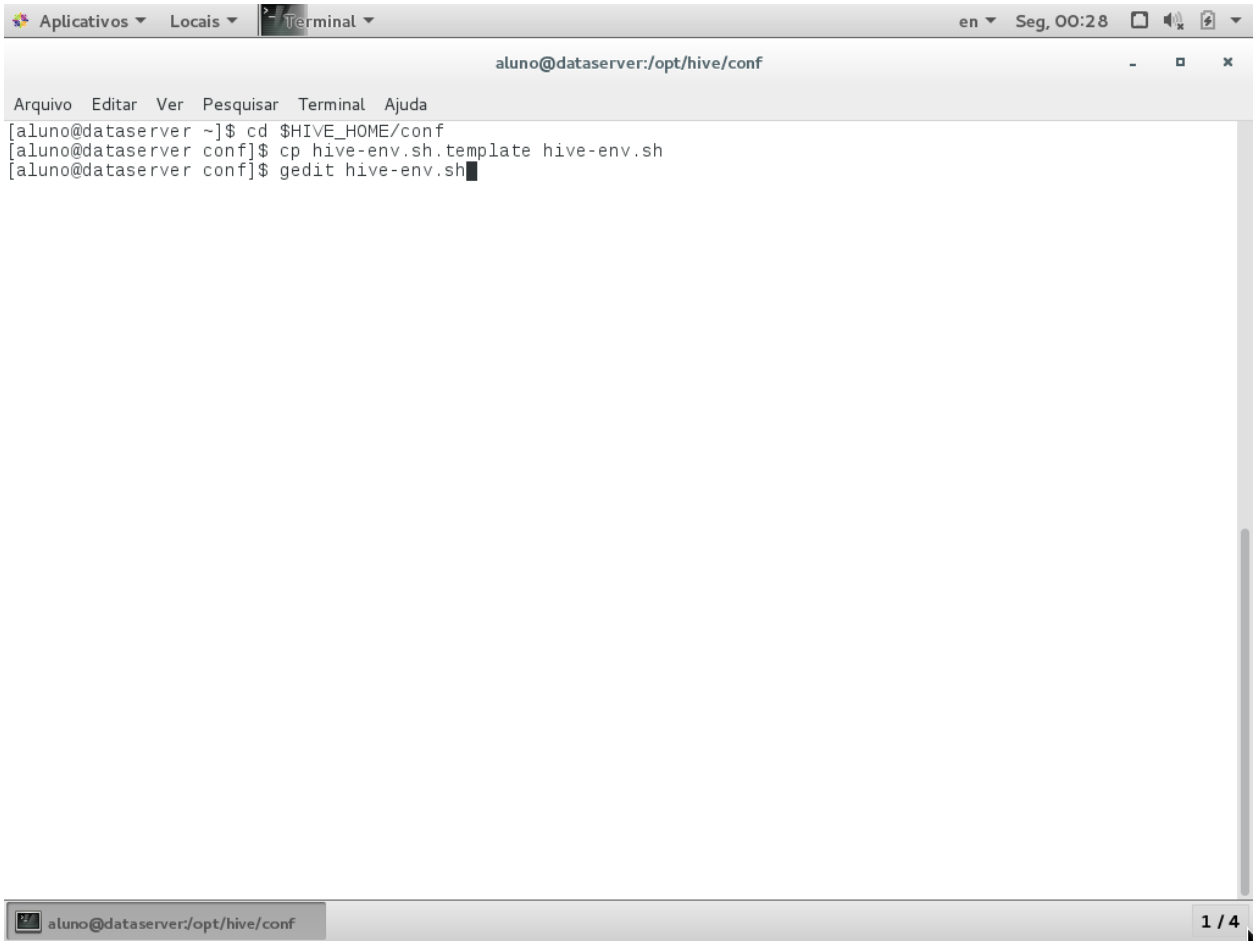

Editar o arquivo

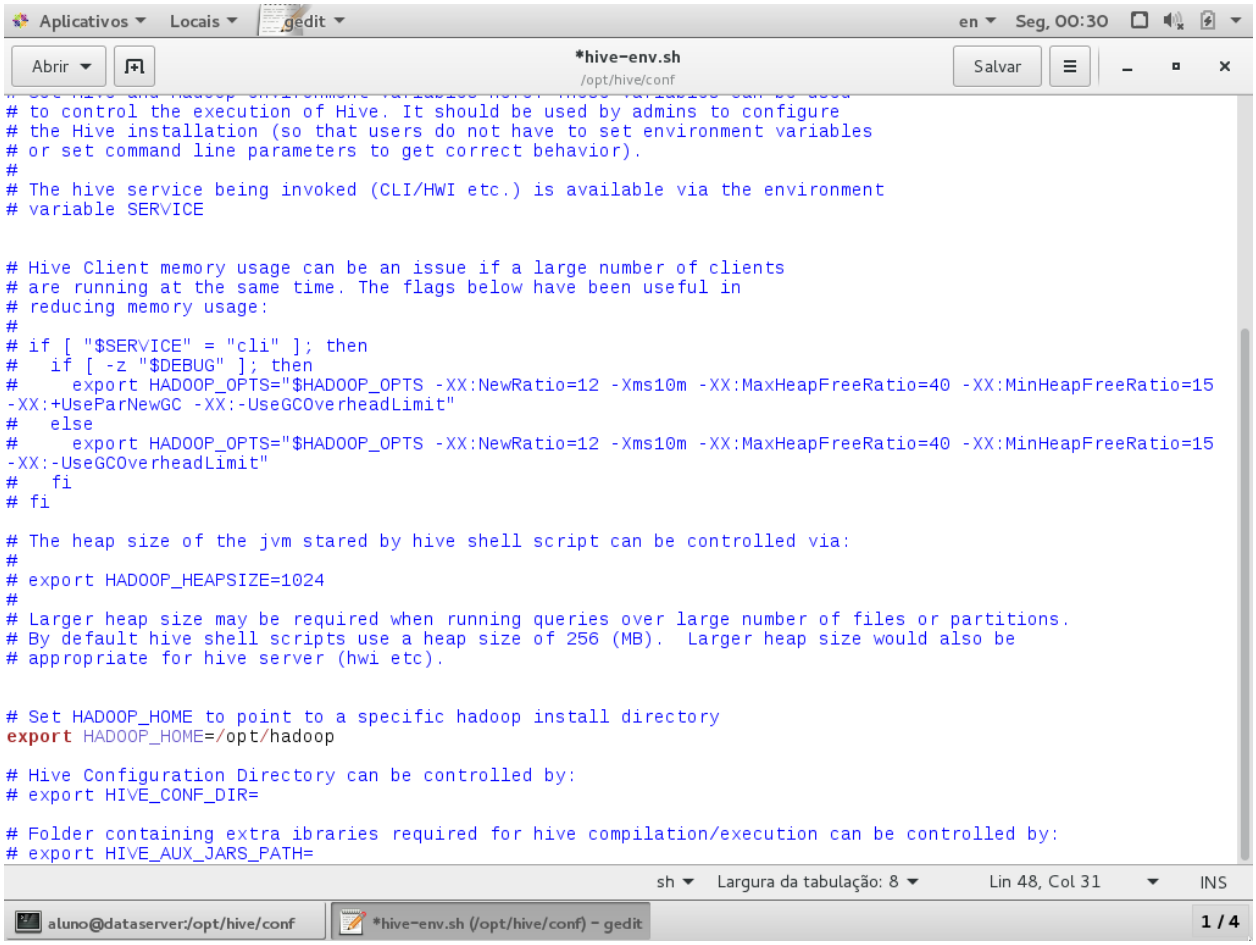

Incluir PATH do Hadoop, conforme tela acima

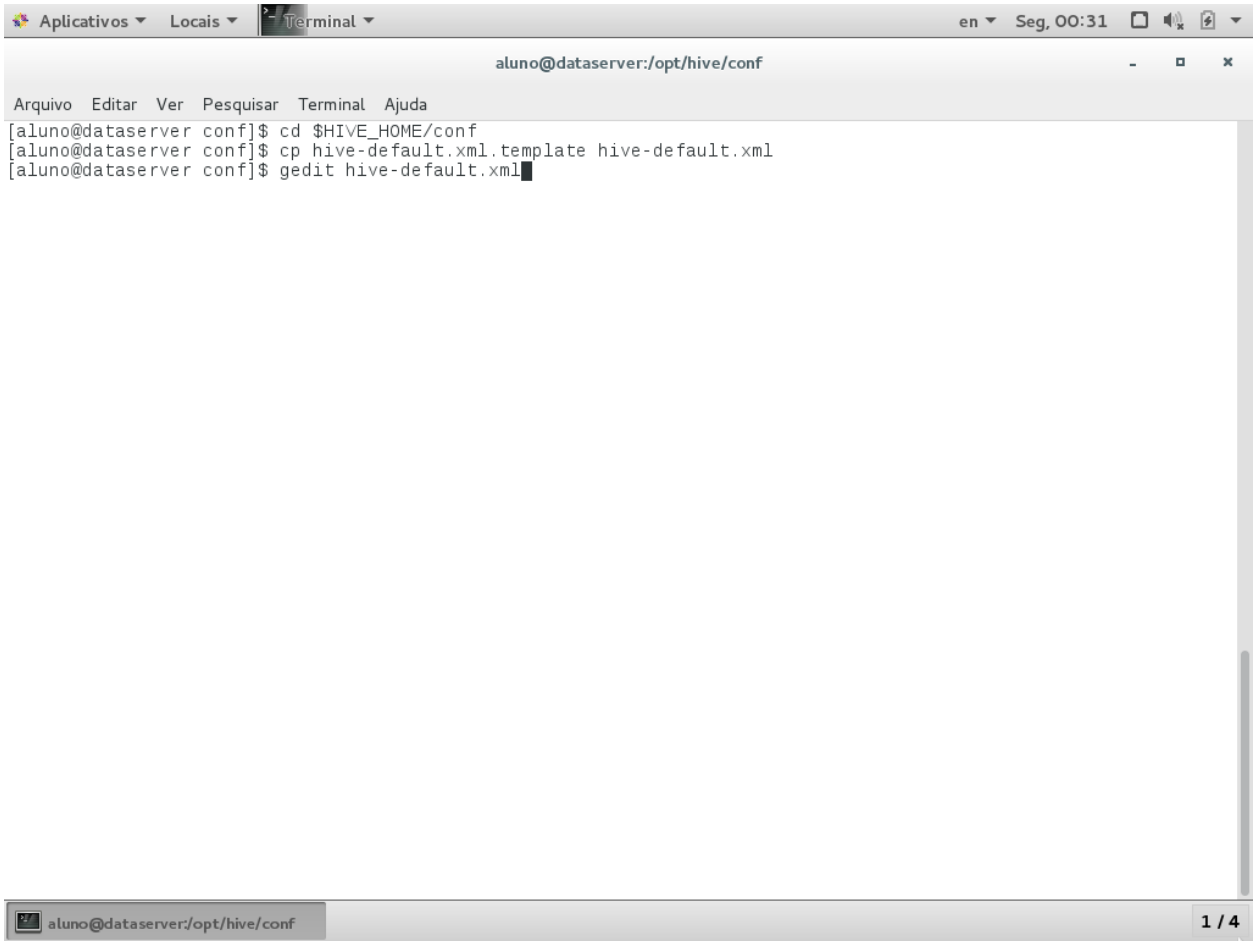

A partir do template, gerar o arquivo hive-site.xml

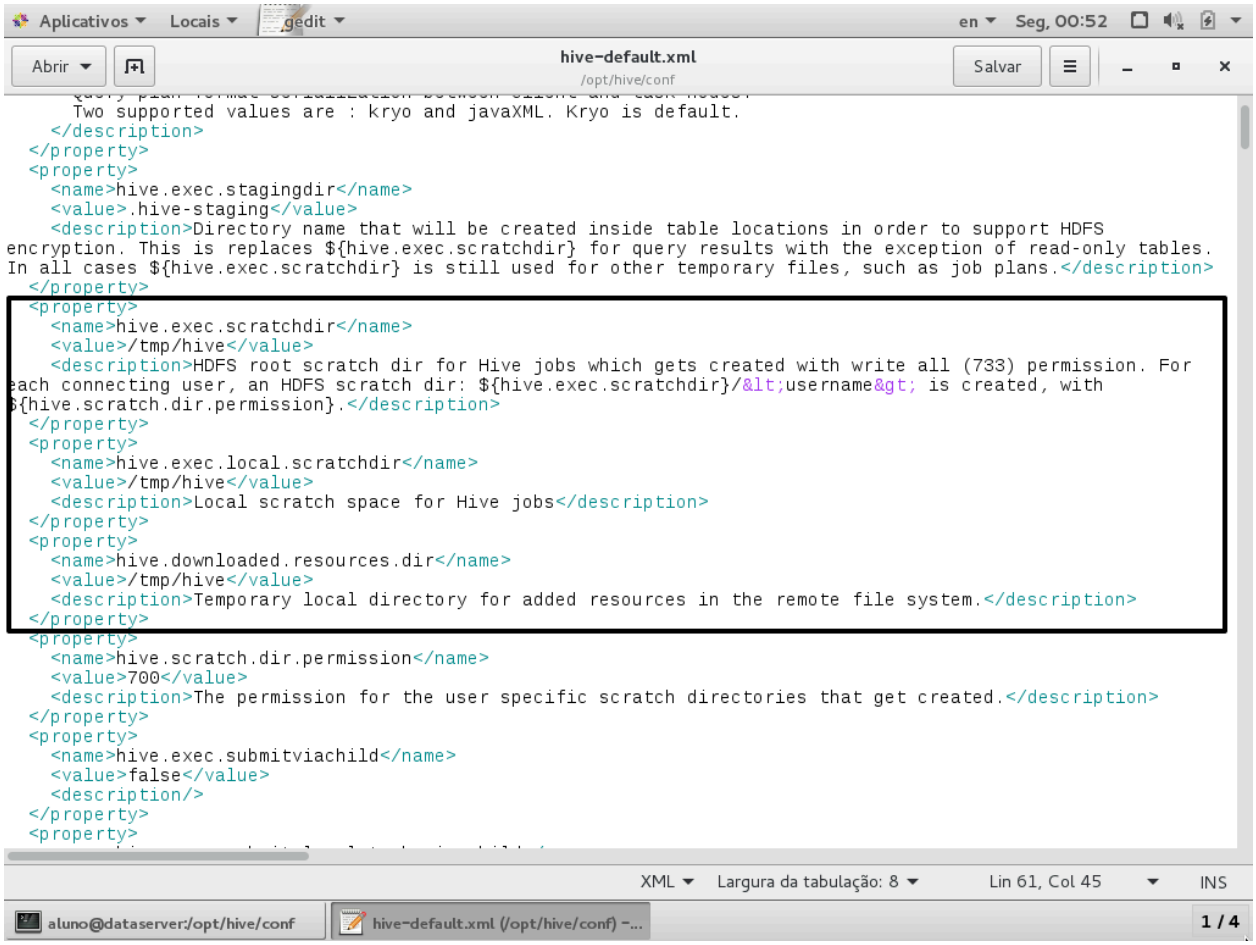

Editar as linhas conforme cima

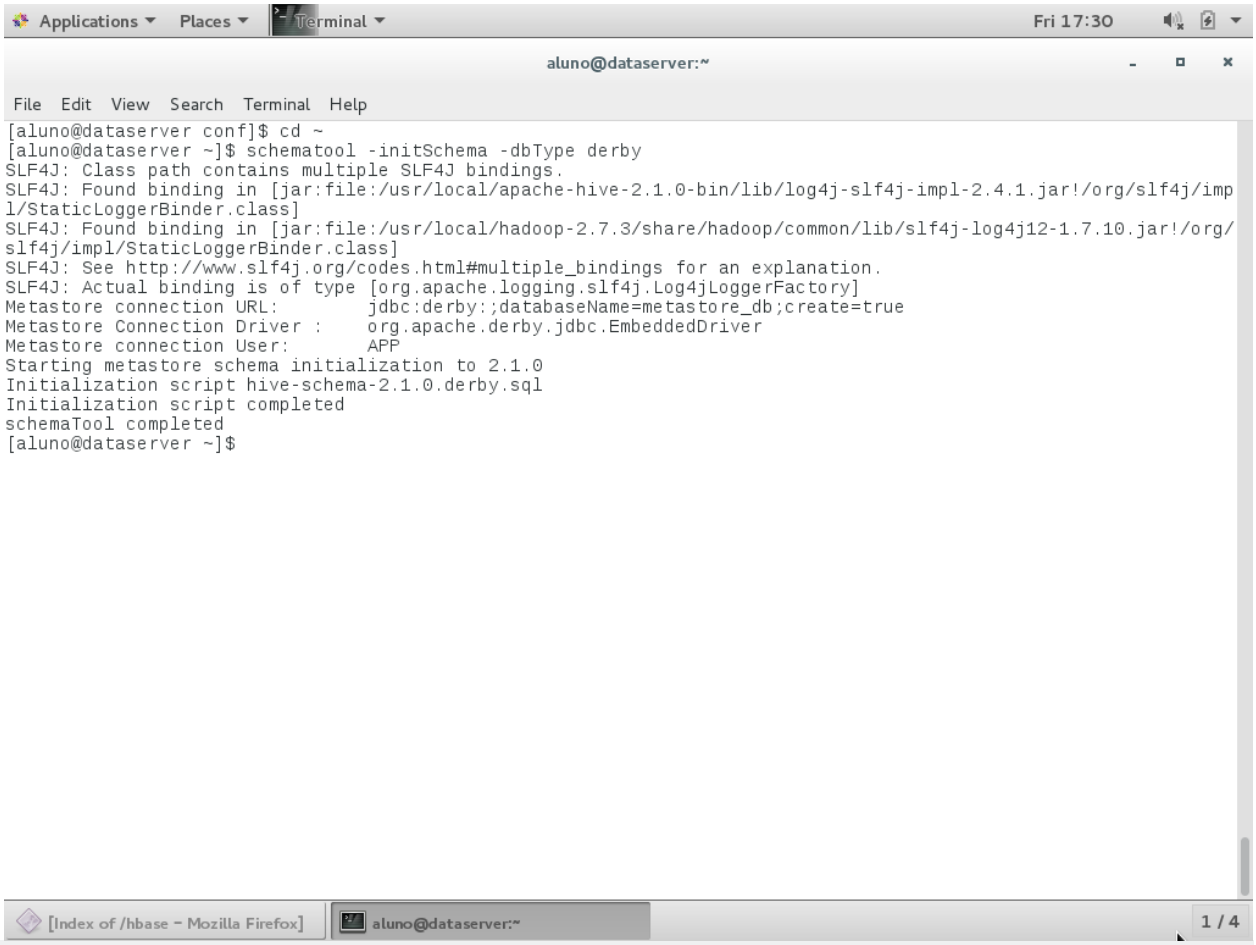

Inicializar o schema do Hive
schematool -initSchema -dbType derby

Executando o Hive (execute o comando jps para se certificar que o Hadoop está ativo)

O comando “show tables;” demonstra que o Hive foi instalado com sucesso

9. Instalação e Configuração do Pig

9.1. Download e Instalação do Pig

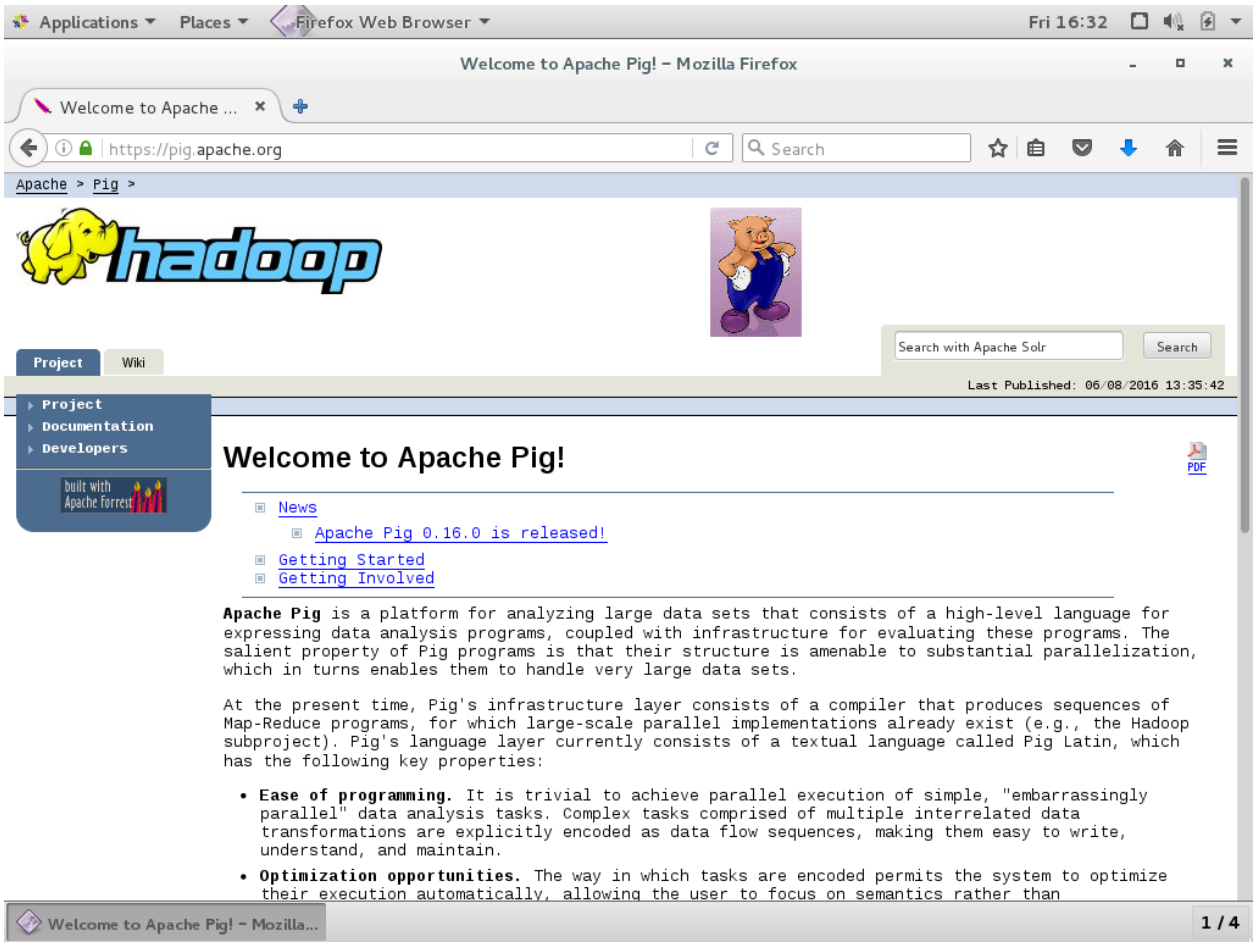

Download do Pig – Versão 0.17.0

Faça o download, descompacte o arquivo e mova o diretório para /opt/pig da mesma forma como você fez com o Java JDK e com o Hadoop.

9.2. Configurando do Pig

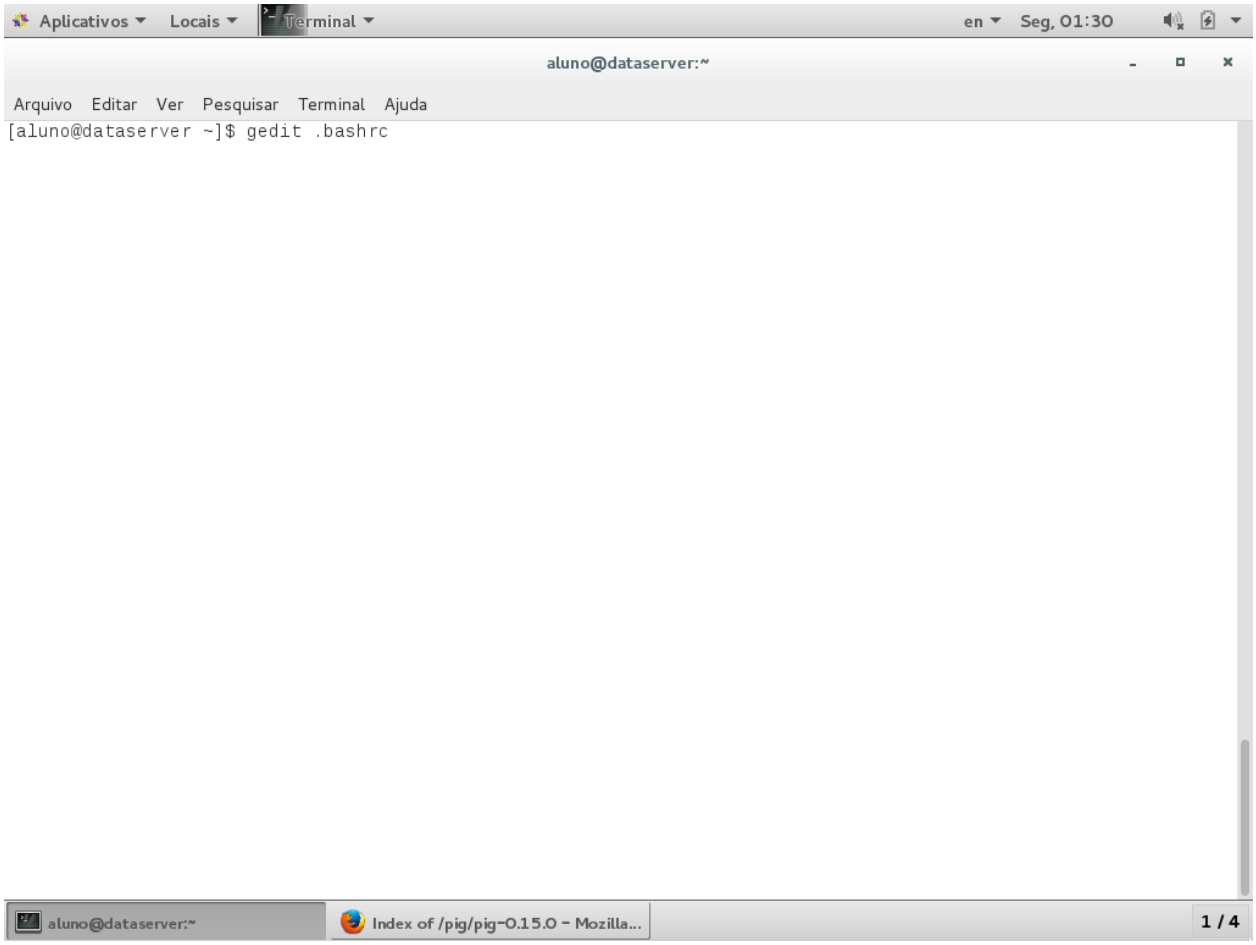

Editando o arquivo .bashrc

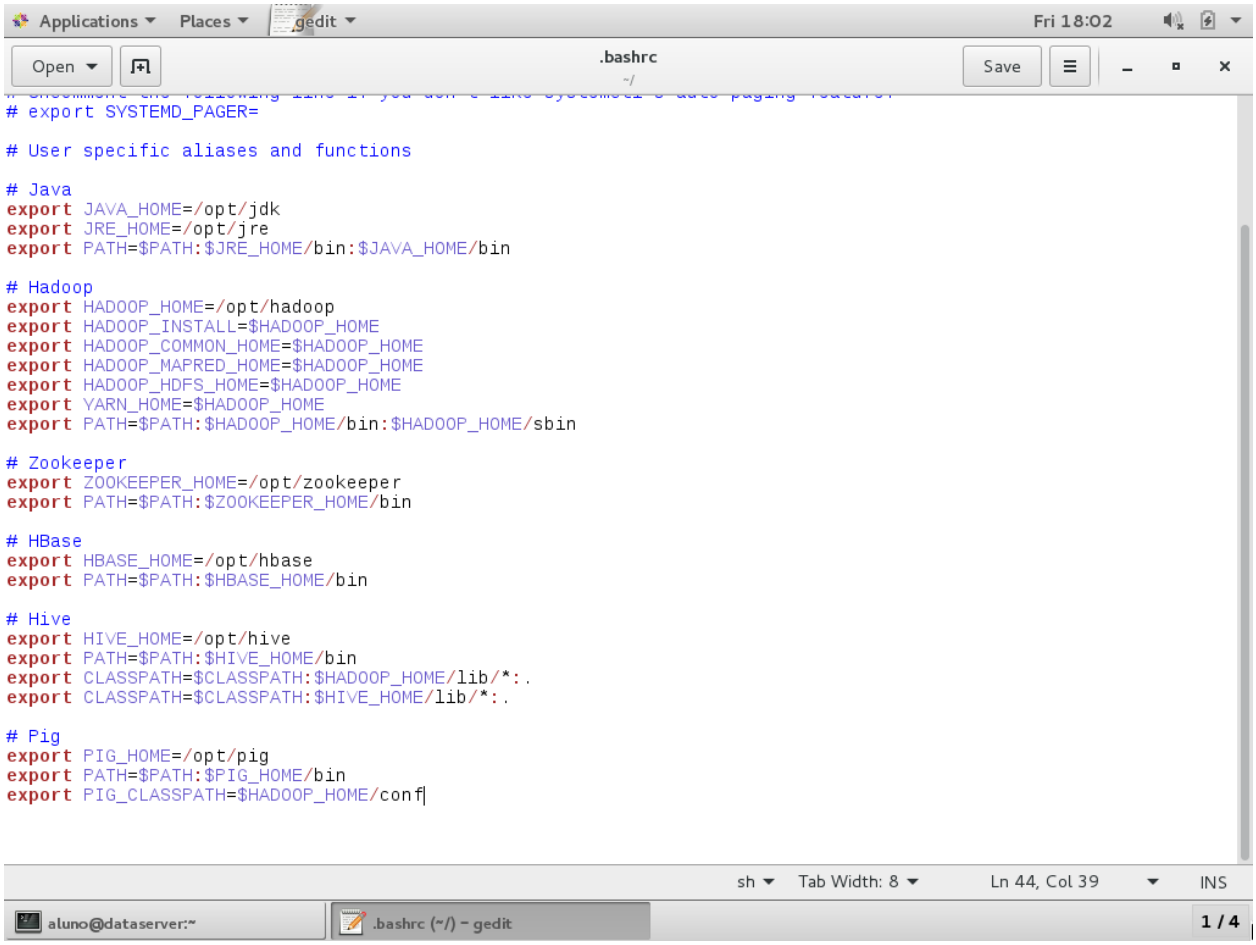

Inserir variáveis de ambiente do Pig

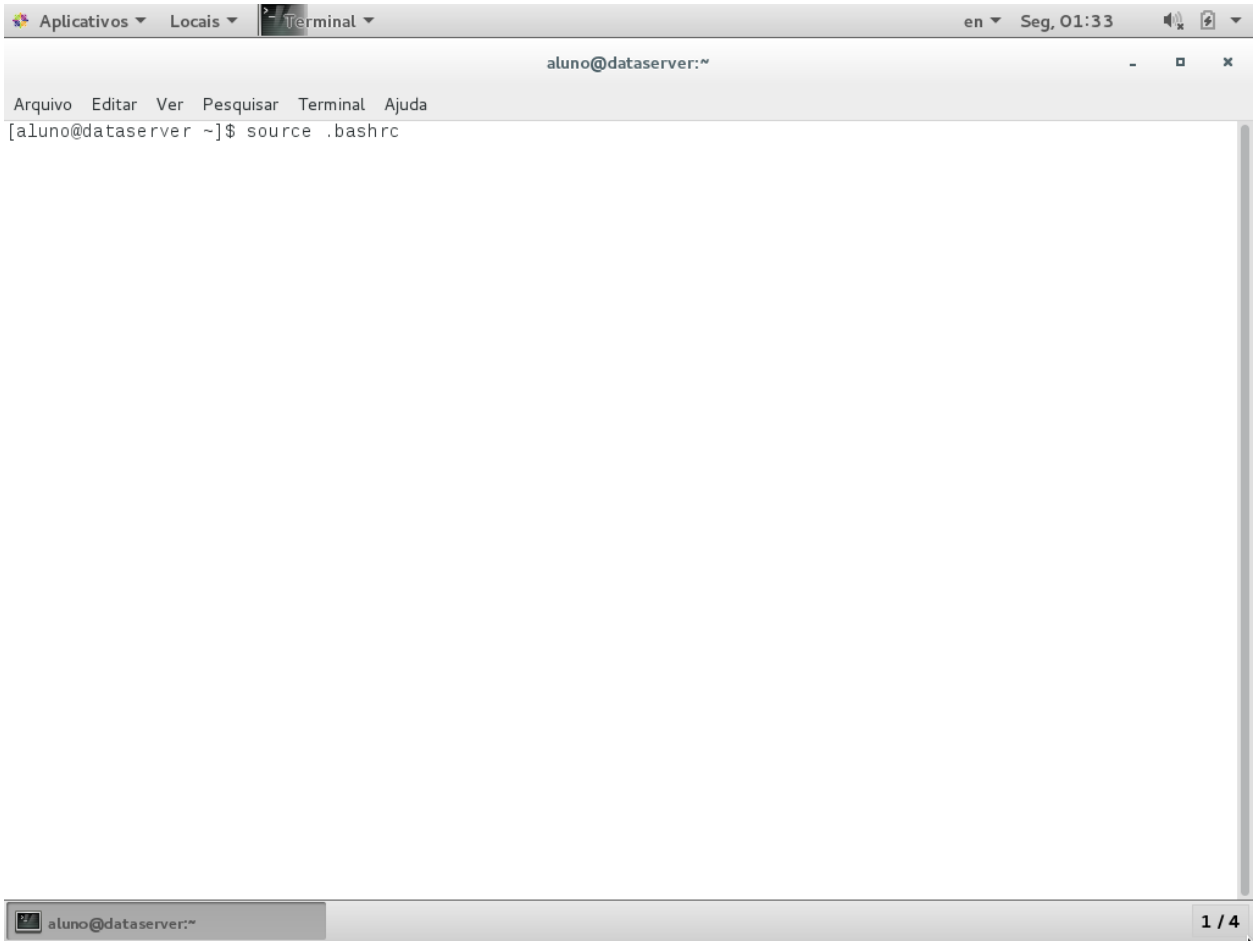

source .bashrc

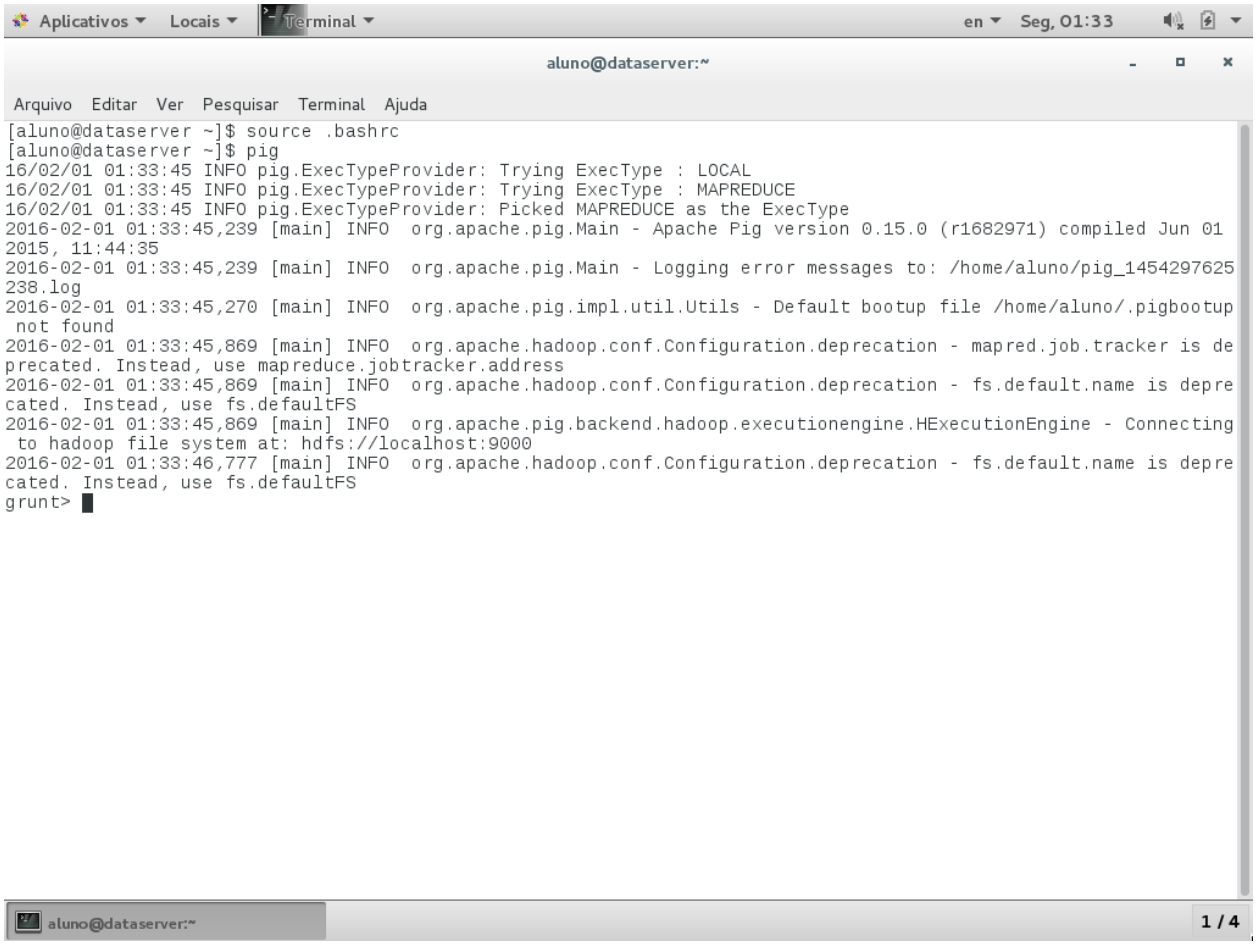

Pig instalado com sucesso

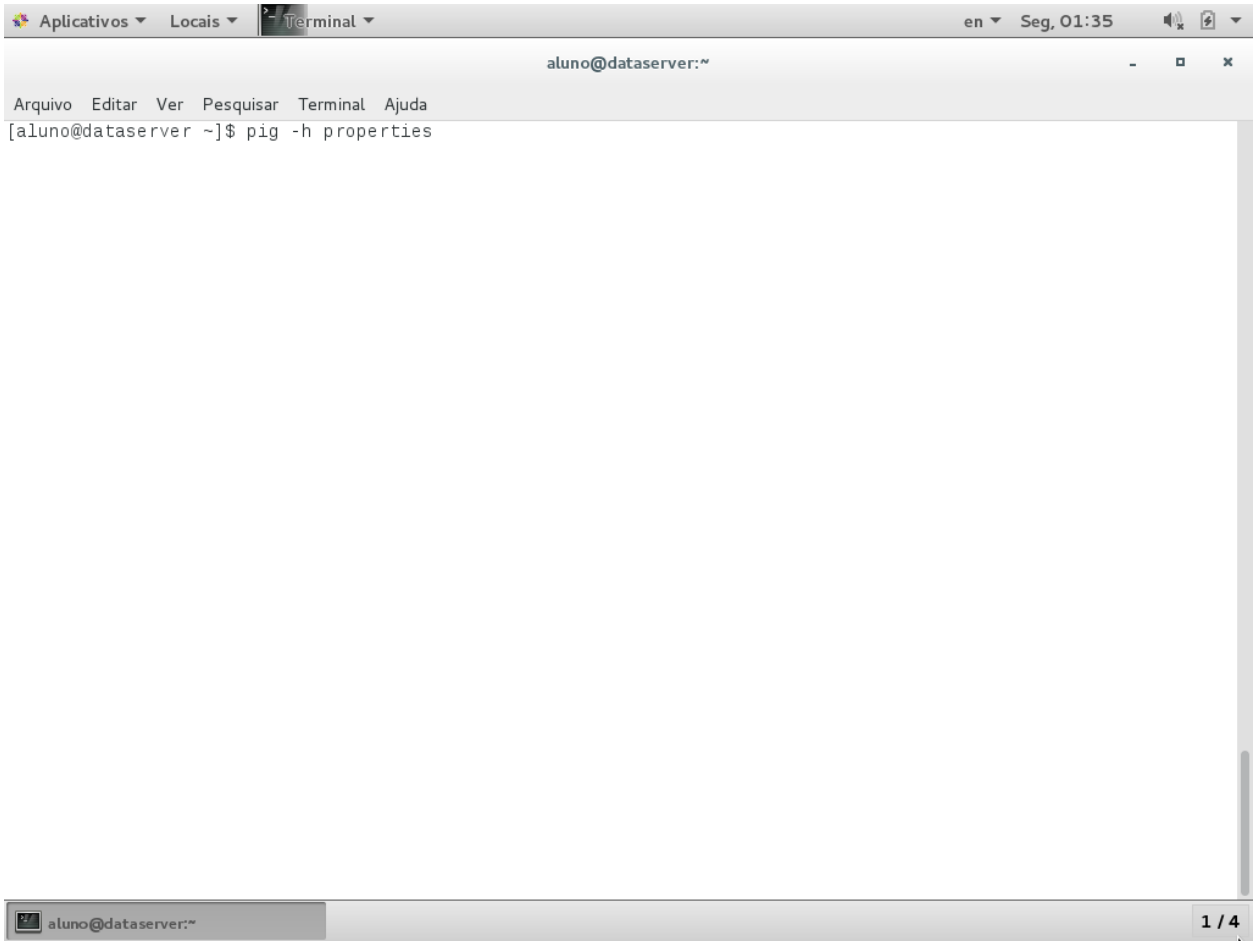

O comando **pig –h properties** lista as variáveis configuradas

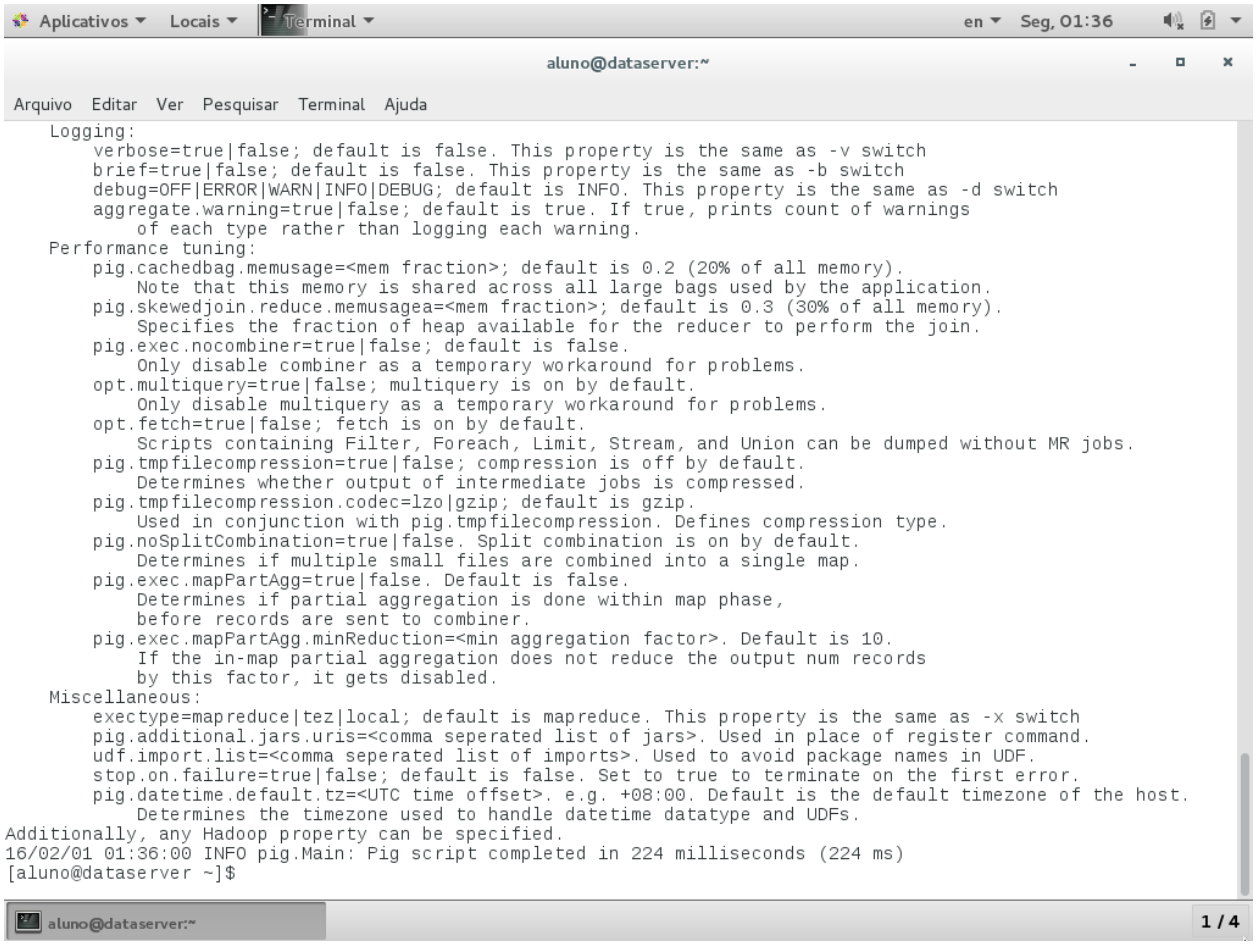

Variáveis Pig

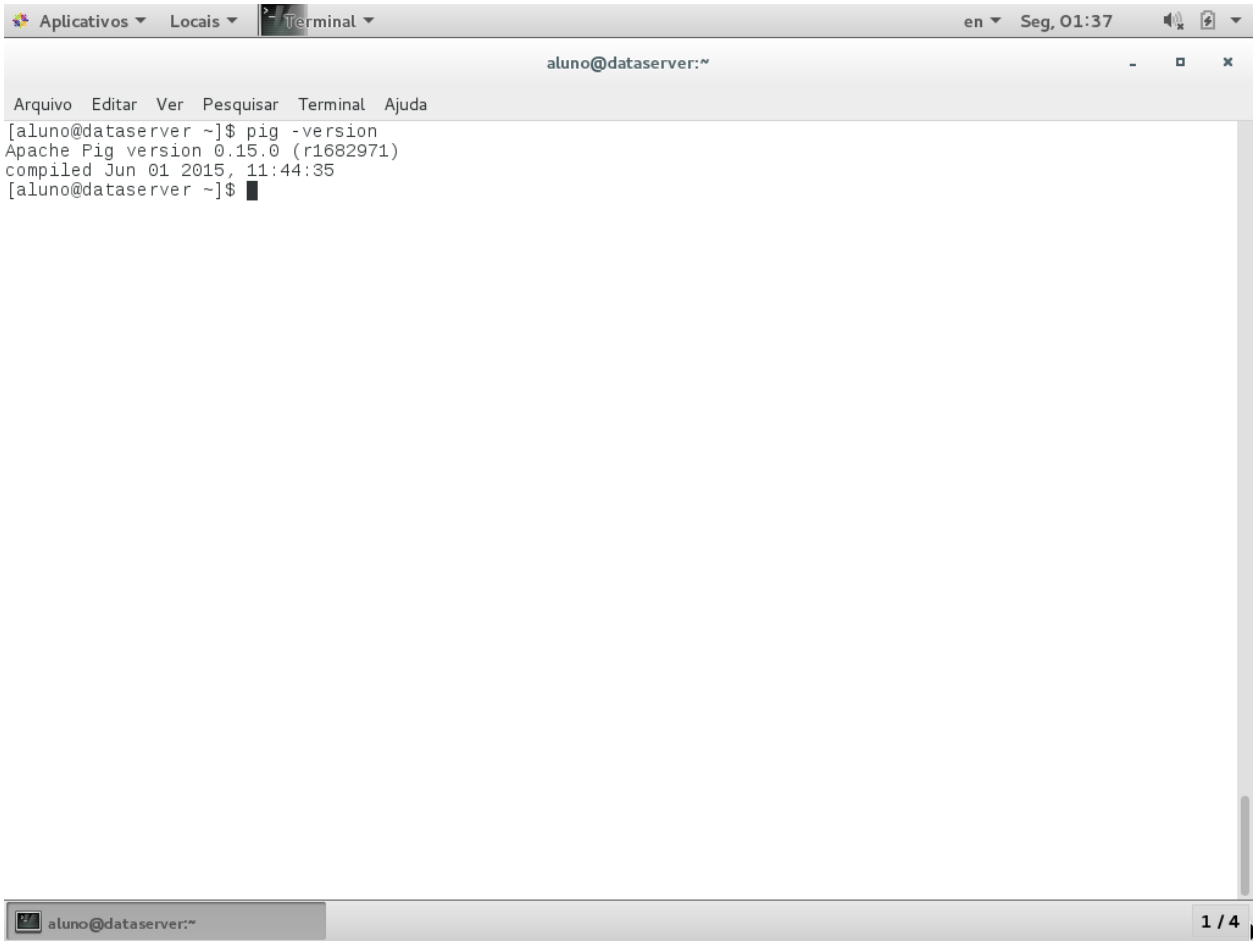

Verificar a versão do Pig

10. Instalação e Configuração do Spark

10.1. Download e Instalação do Spark

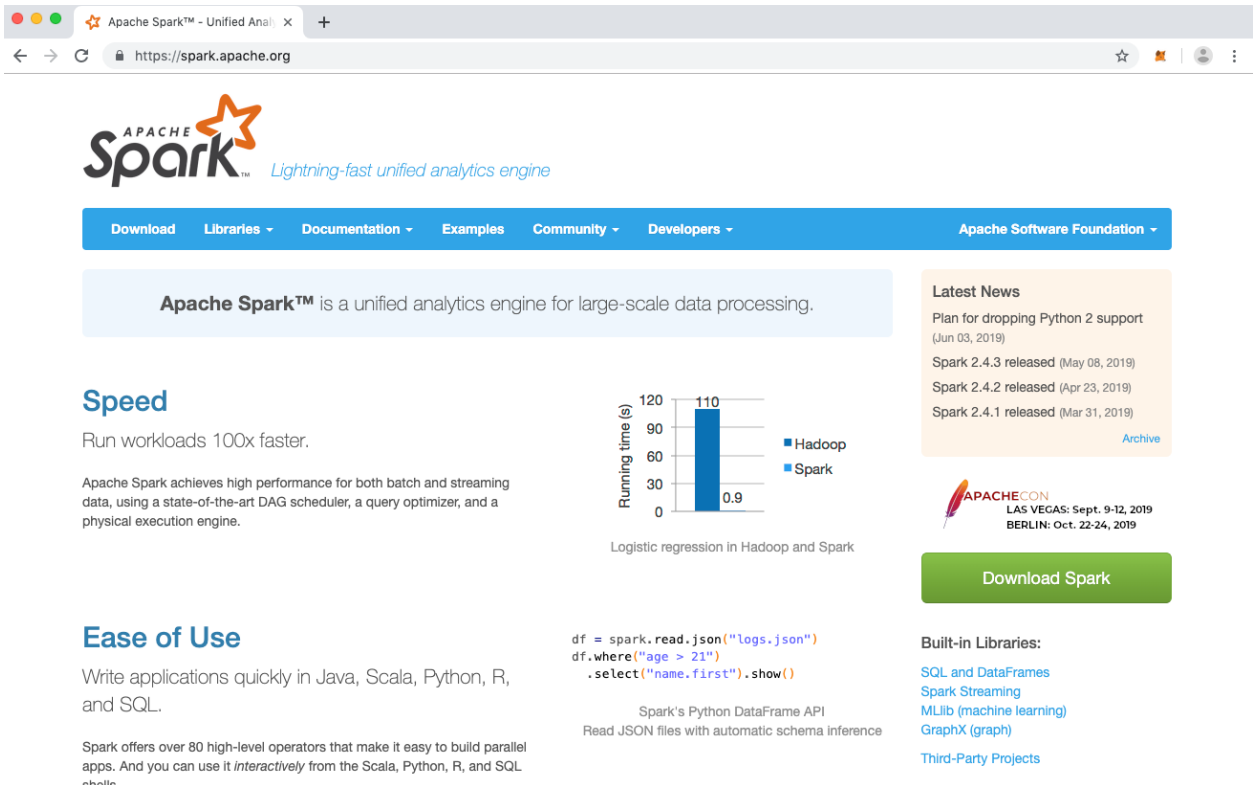

Download do Spark – Versão 2.4.3

Faça o download, descompacte o arquivo e mova o diretório para /opt/spark da mesma forma como você fez com o Java JDK e com o Hadoop.

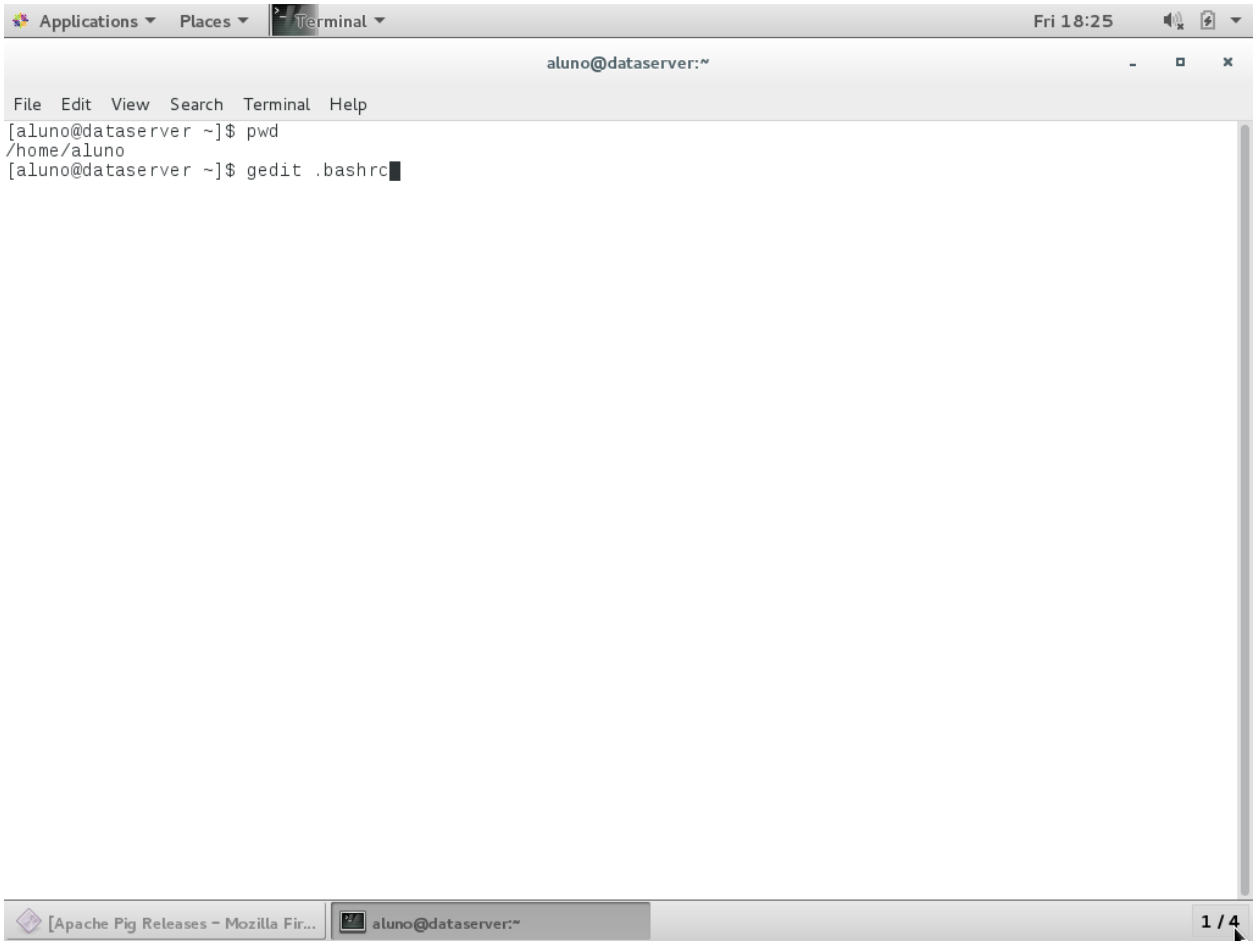

Editando o arquivo .bashrc

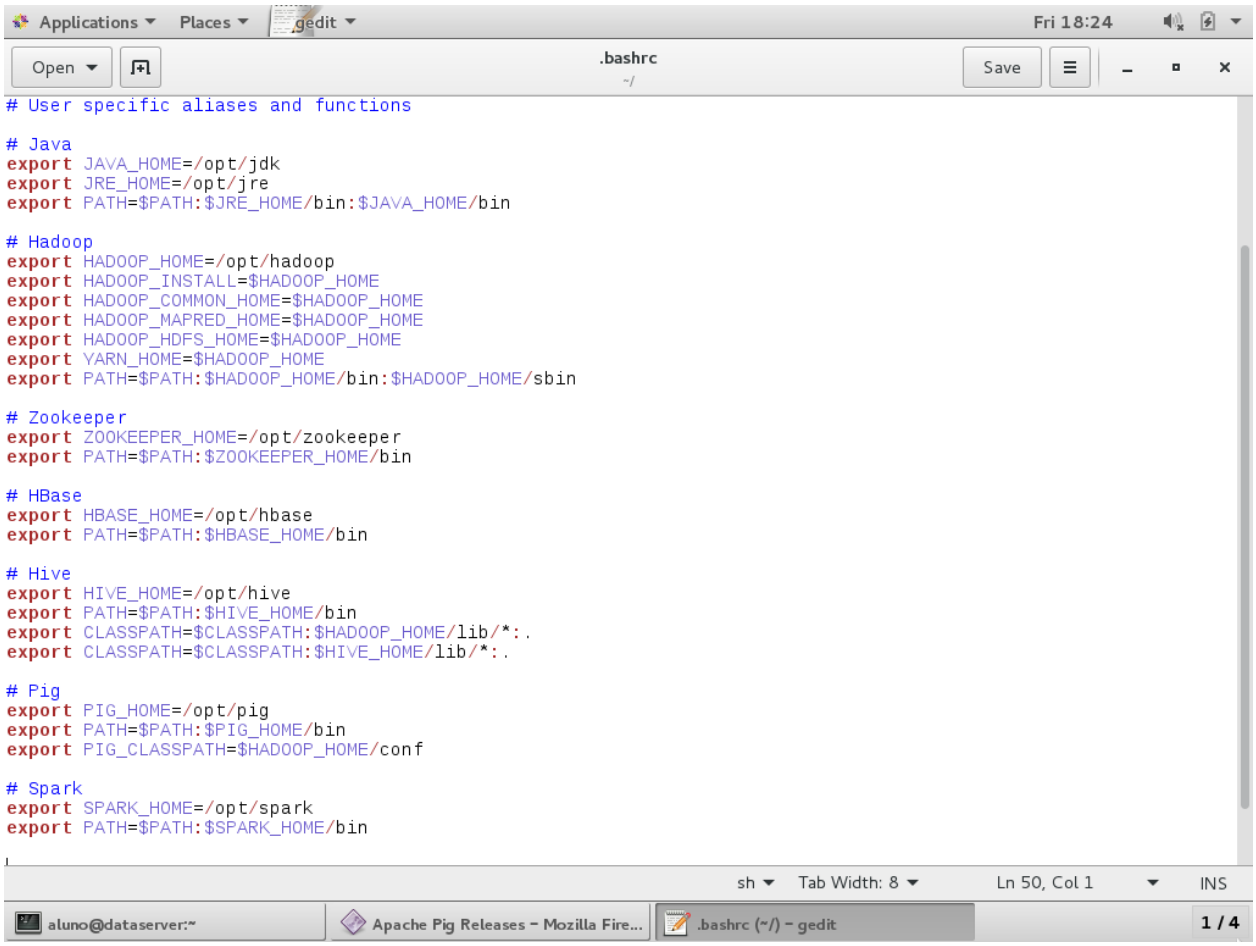

Incluir variáveis Spark

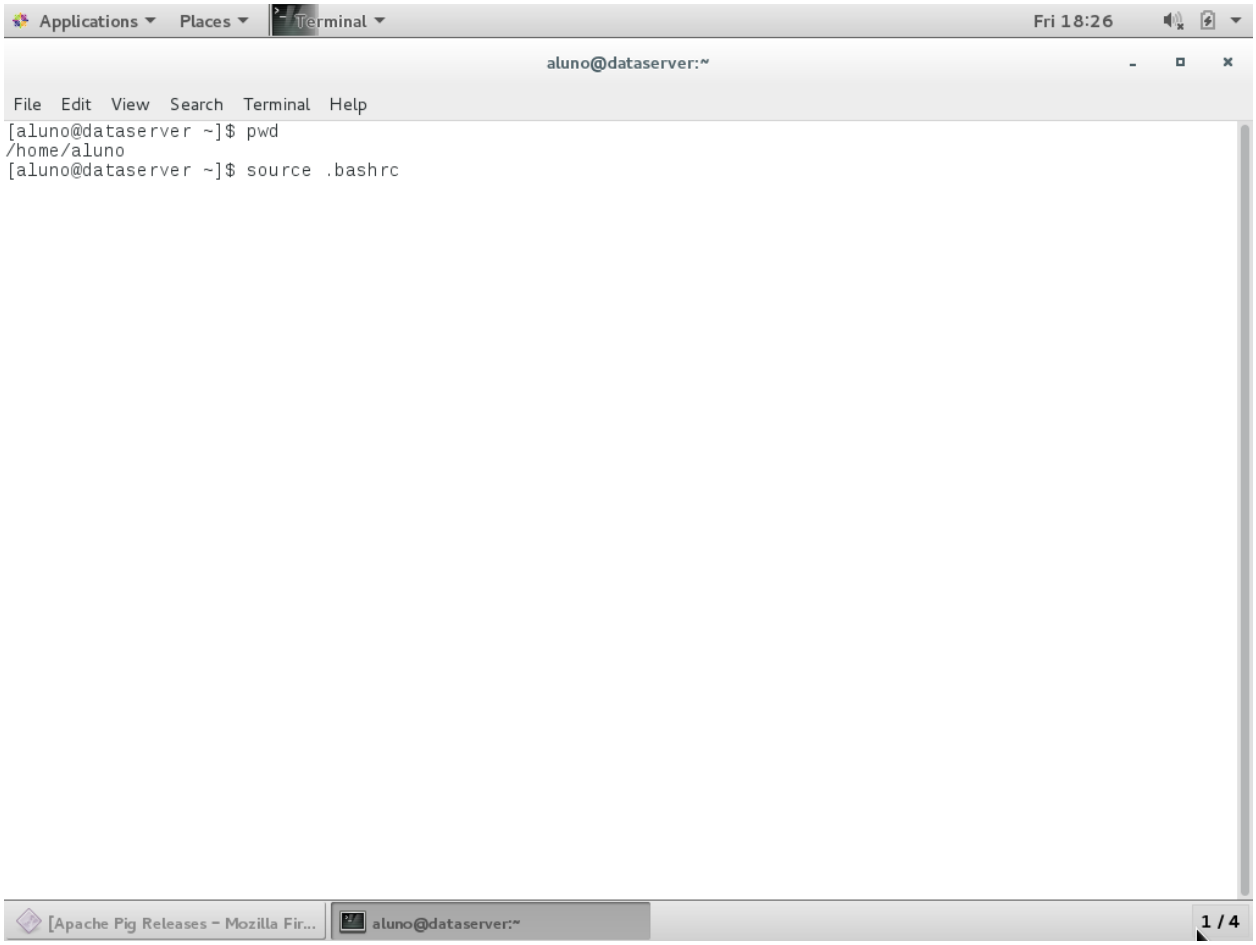

source .bashrc

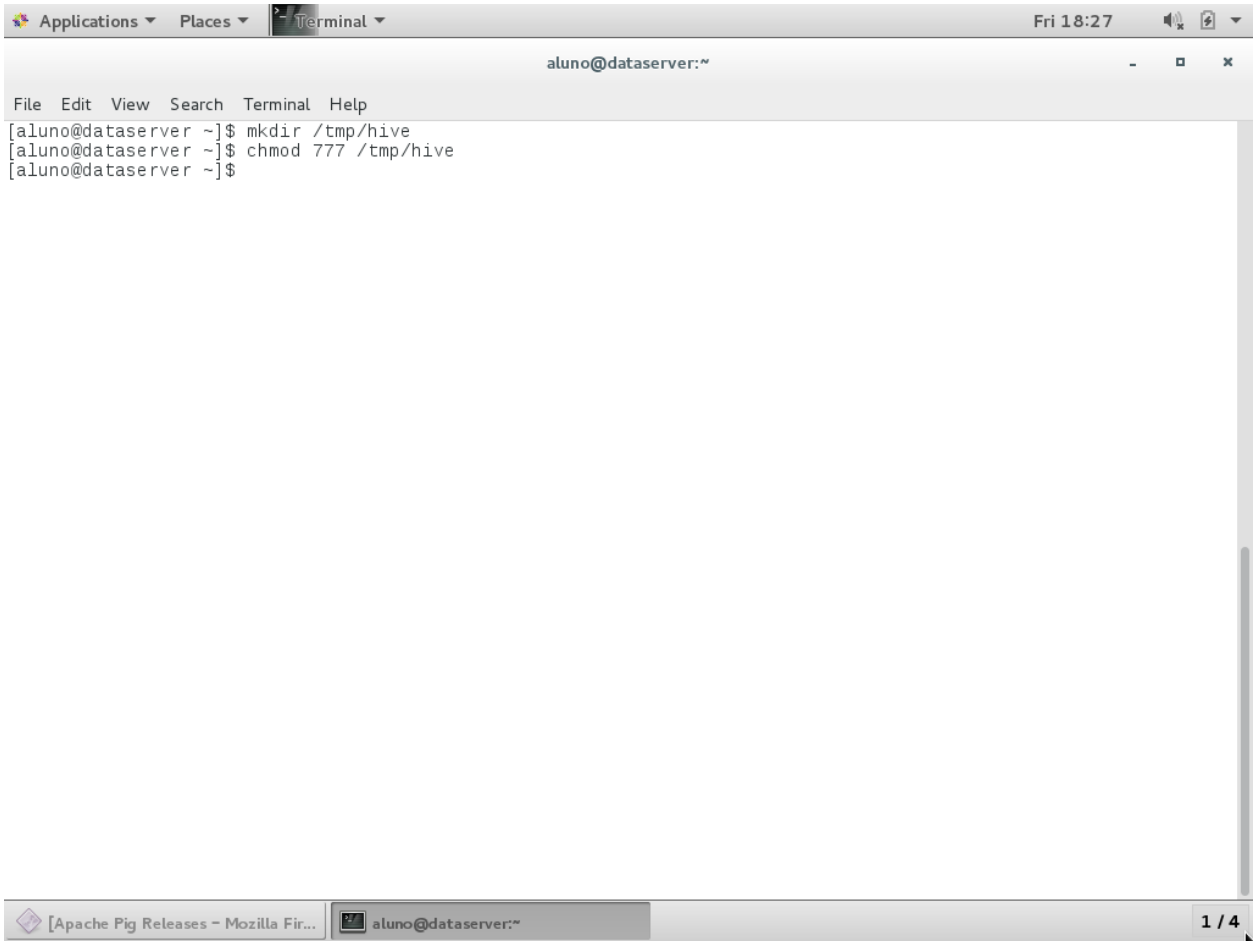

Se necessário chmod 777 /tmp/hive

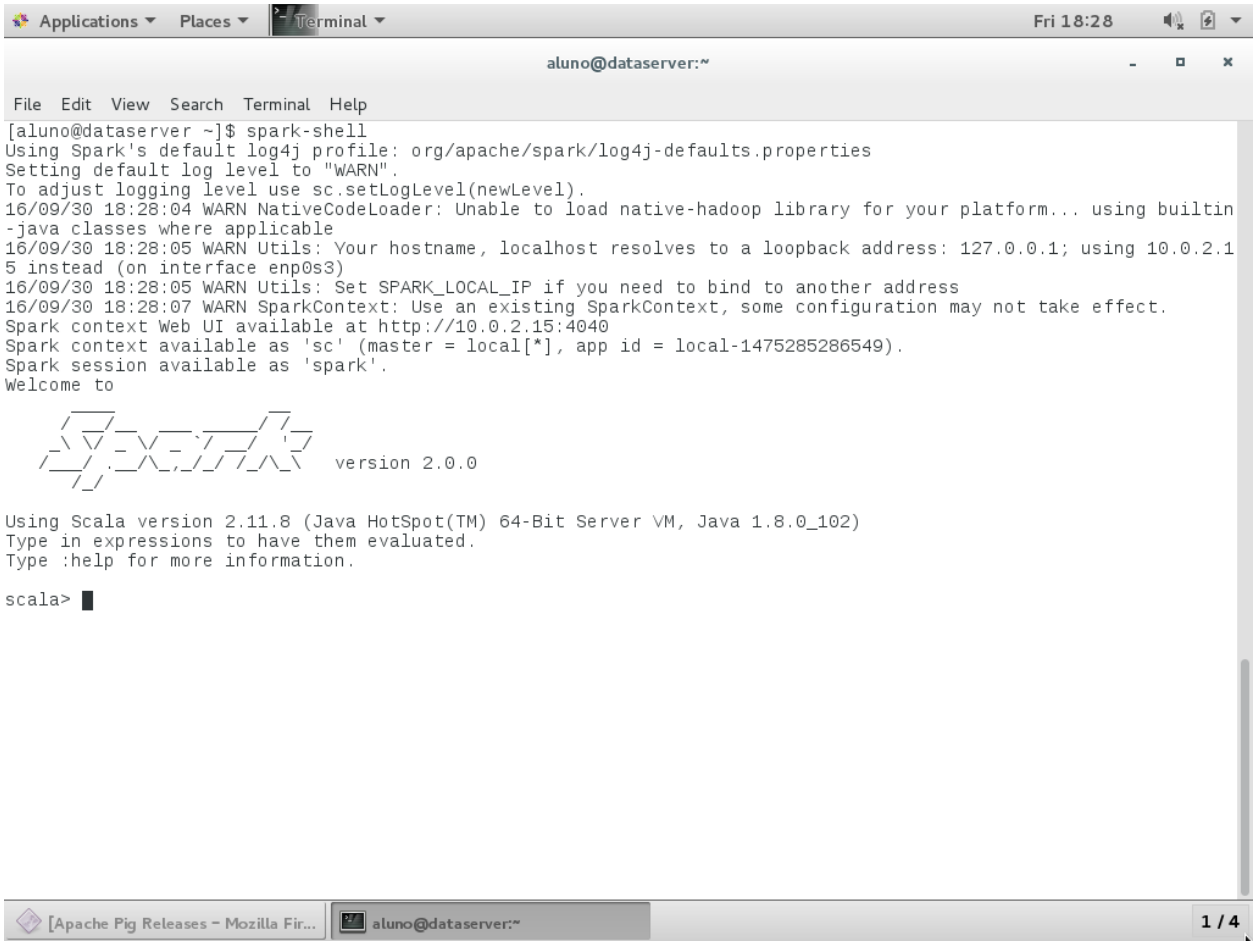
Applications
Places
Terminal
Fri 18:28
aluno@dataserver:~
File Edit View Search Terminal Help
[aluno@dataserver ~]$ spark-shell
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel).
16/09/30 18:28:04 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin
-java classes where applicable
16/09/30 18:28:05 WARN Utils: Your hostname, localhost resolves to a loopback address: 127.0.0.1; using 10.0.2.1
5 instead (on interface enp0s3)
16/09/30 18:28:05 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
16/09/30 18:28:07 WARN SparkContext: Use an existing SparkContext, some configuration may not take effect.
Spark context Web UI available at http://10.0.2.15:4040
Spark context available as 'sc' (master = local[*], app id = local-1475285286549).
Spark session available as 'spark'.
Welcome to
version 2.0.0
Using Scala version 2.11.8 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_102)
Type in expressions to have them evaluated.
Type :help for more information.
scala>
[Apache Pig Releases – Mozilla Fir...
aluno@dataserver:~
1 / 4

Spark shell

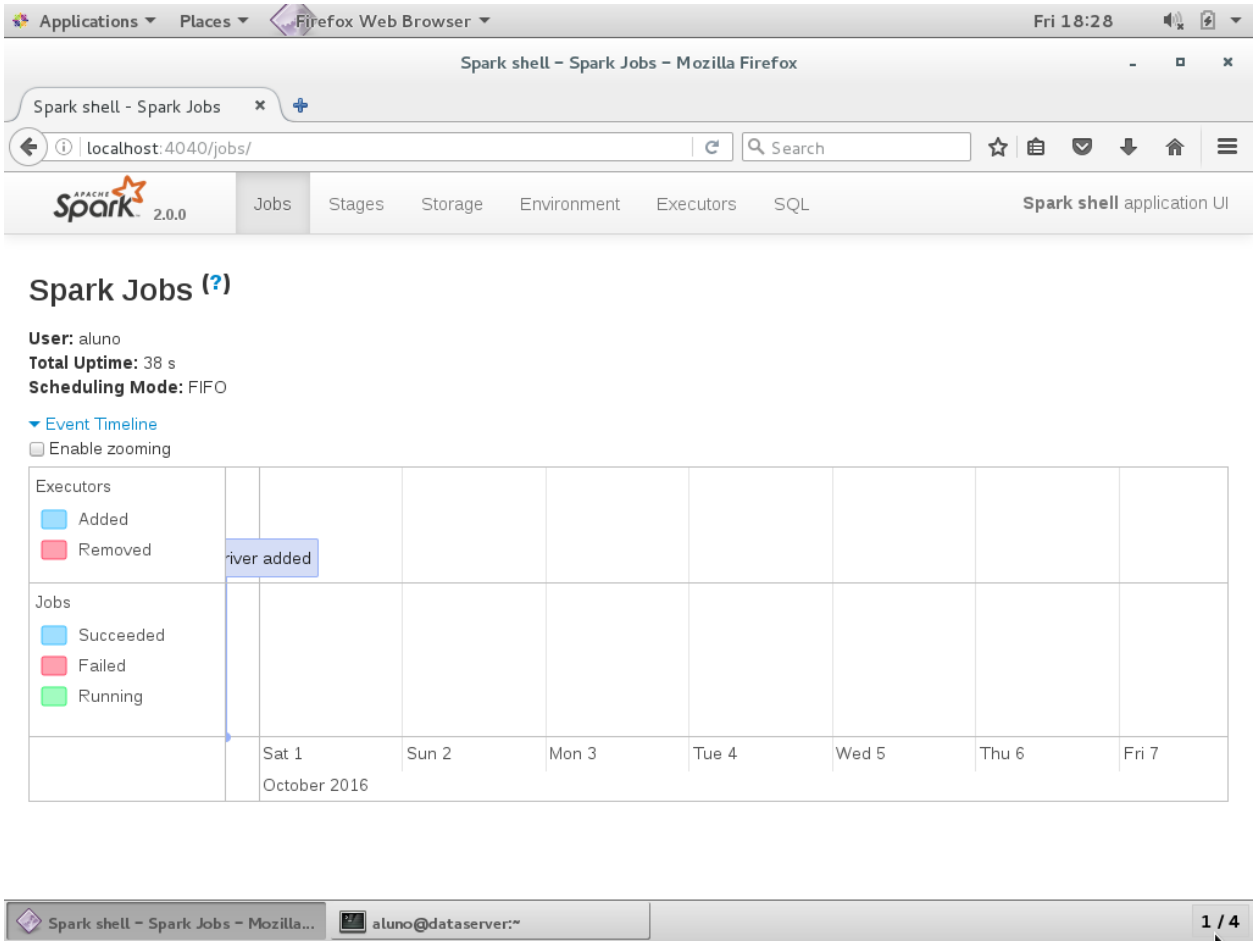

Acessando o Apache Spark pelo browser em http://localhost:4040

11. Instalação e Configuração do Sqoop

11.1. Download do Sqoop

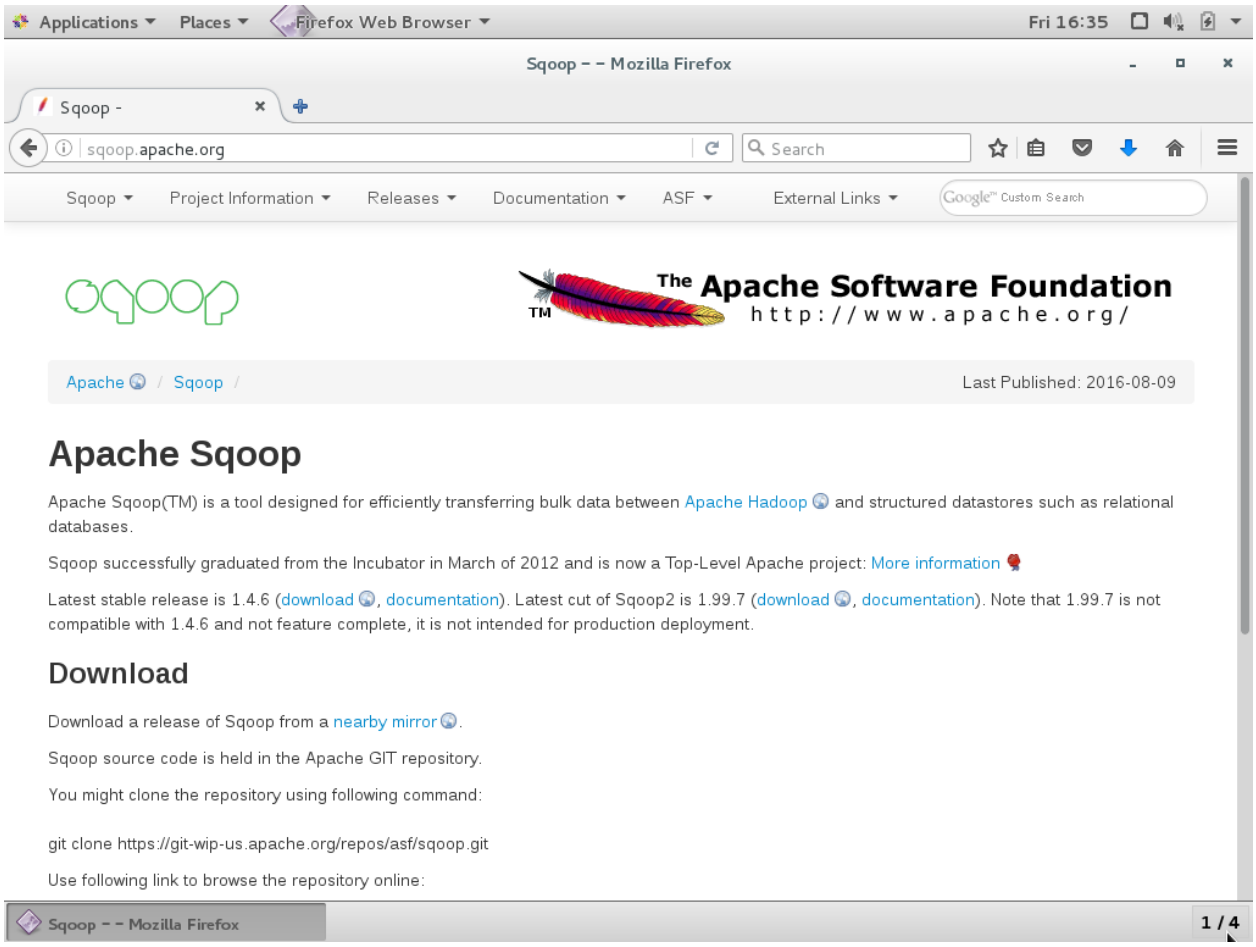

Download do Sqoop – Versão 1.4.7-hadoop-2.0.4-alpha

Faça o download, descompacte o arquivo e mova o diretório para /opt/sqoop da mesma forma como você fez com o Java JDK e com o Hadoop.

11.2. Configuração do Sqoop

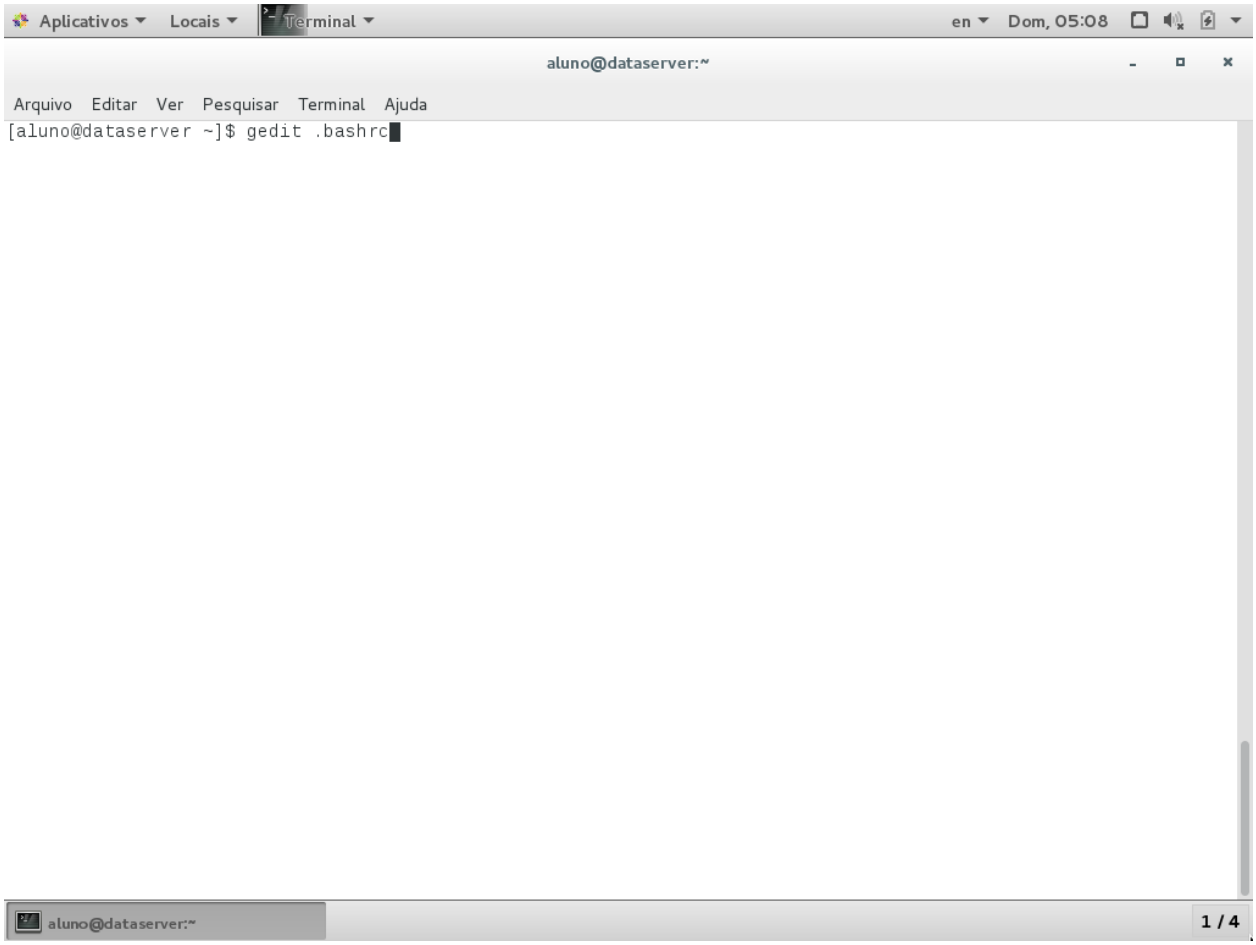

Editar arquivo .bashrc

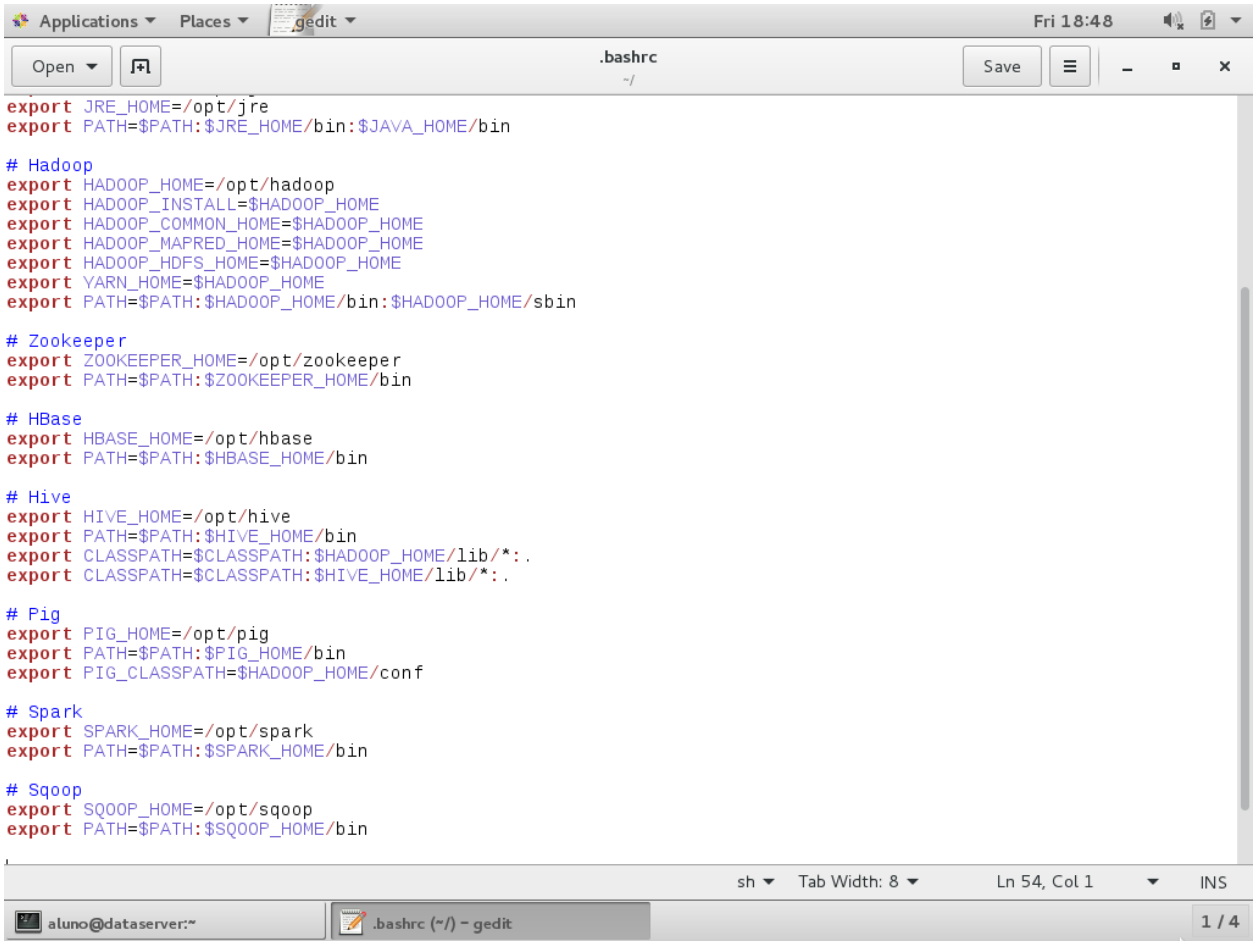

Incluir variáveis Sqoop

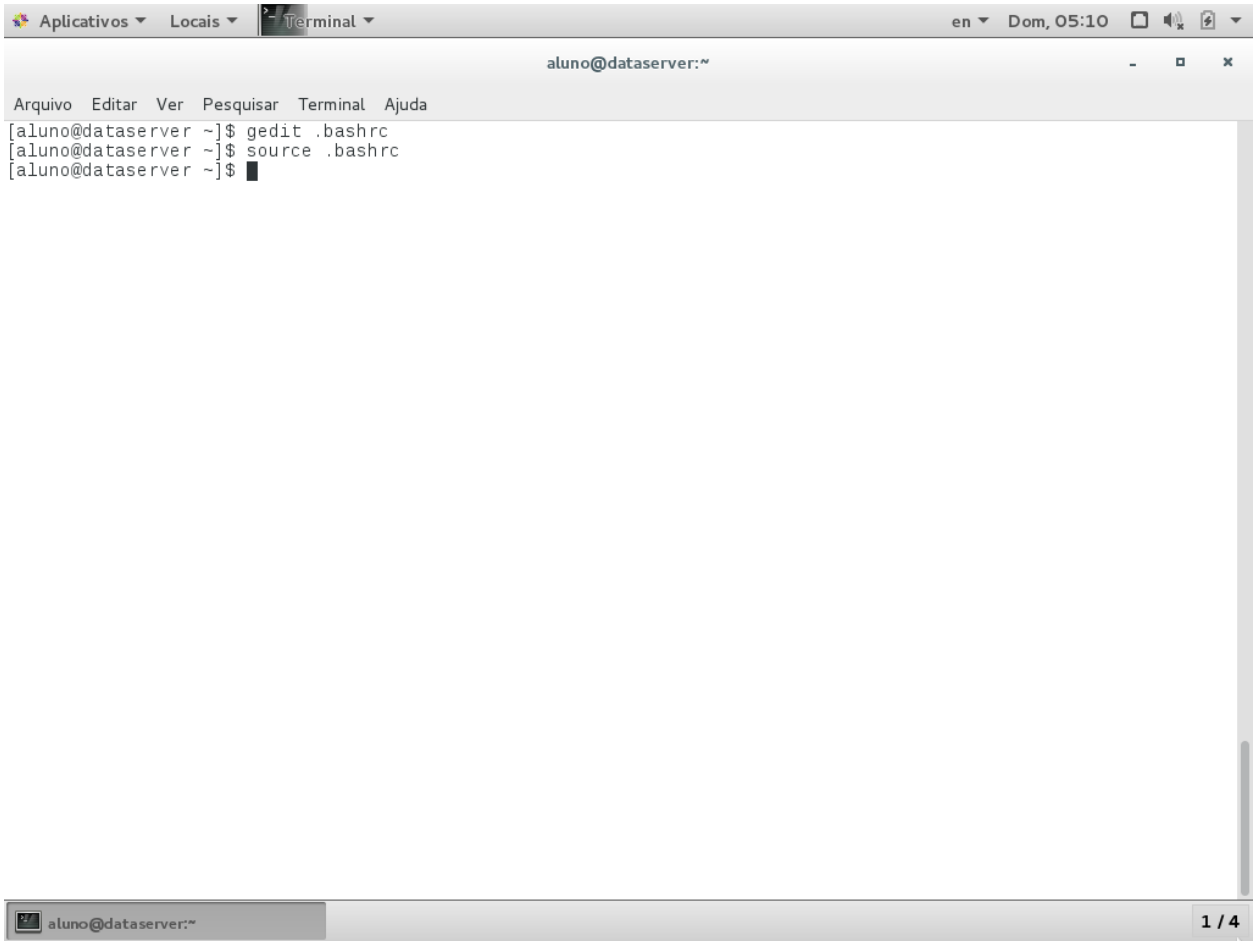

source .bashrc

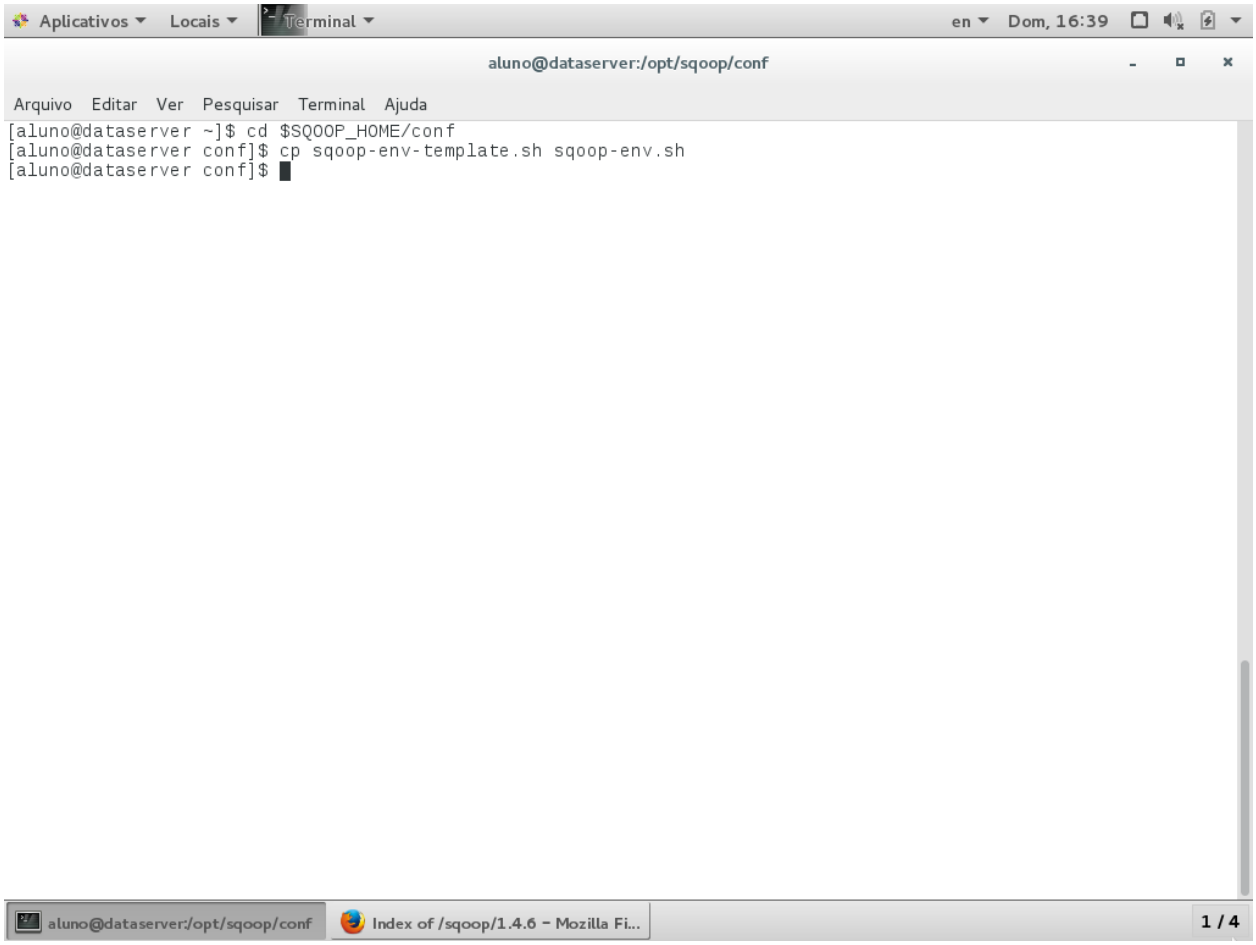

A partir do template, criar o arquivo sqoop-env.sh e editá-lo

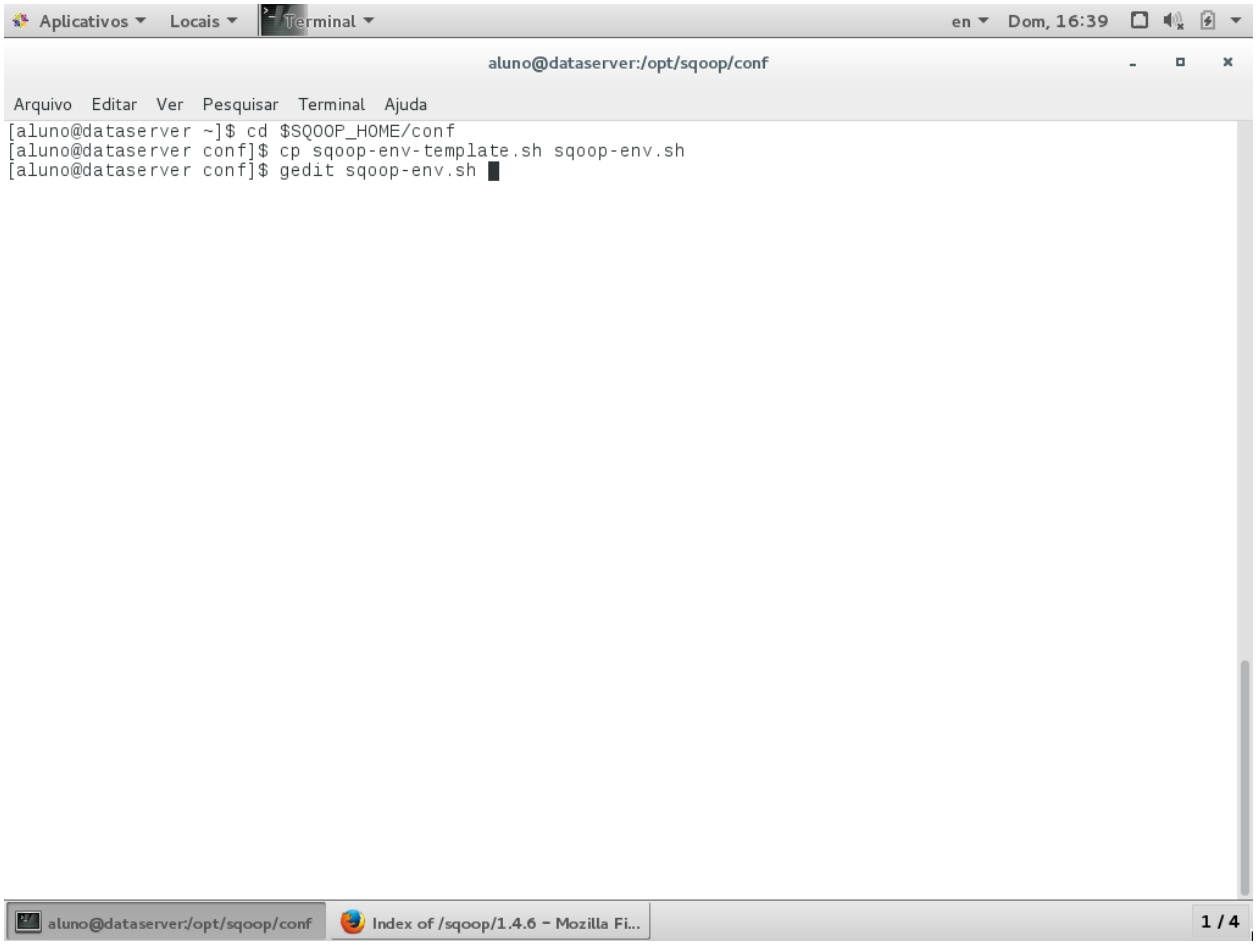

Editar o arquivo

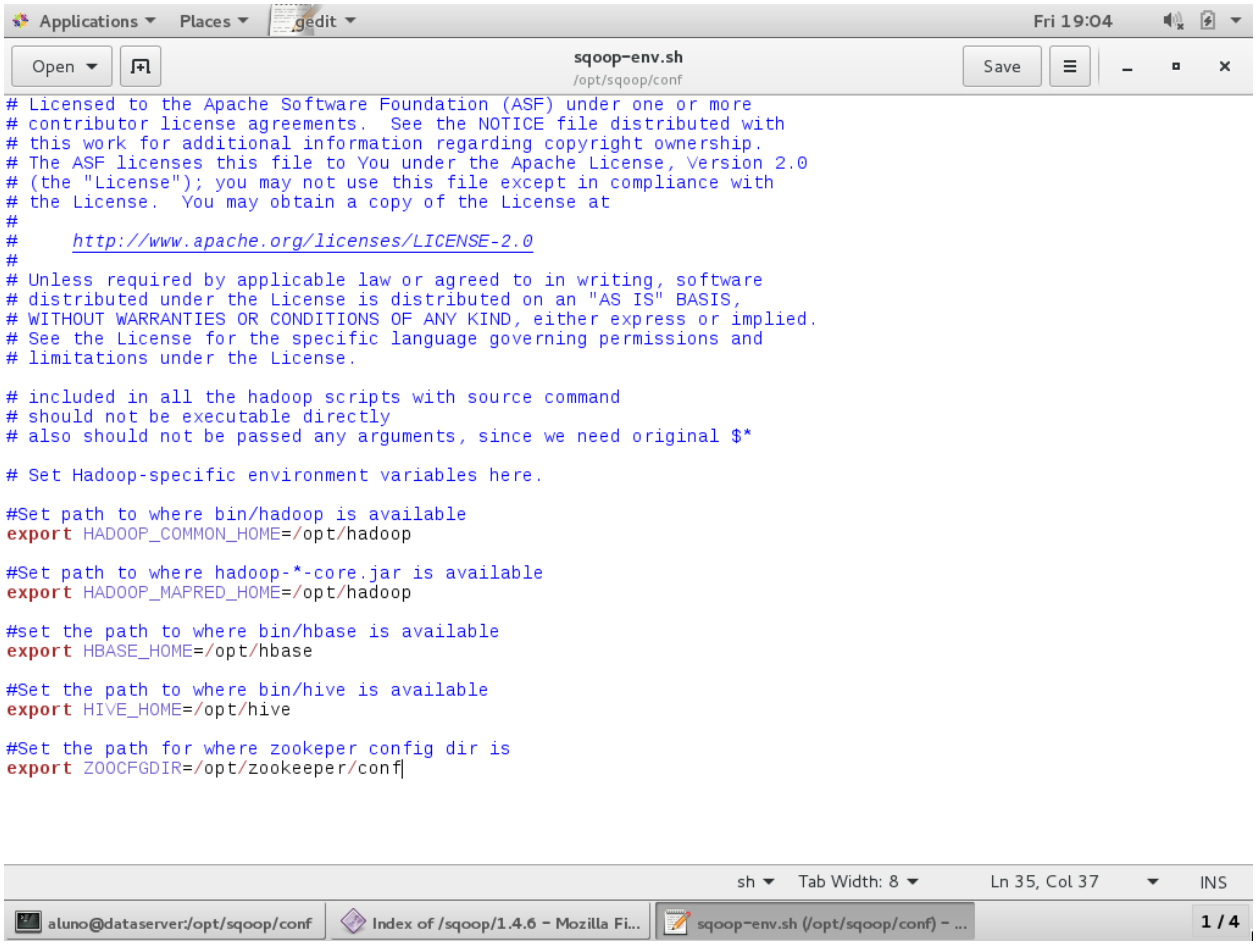

Editar variáveis conforme tela acima

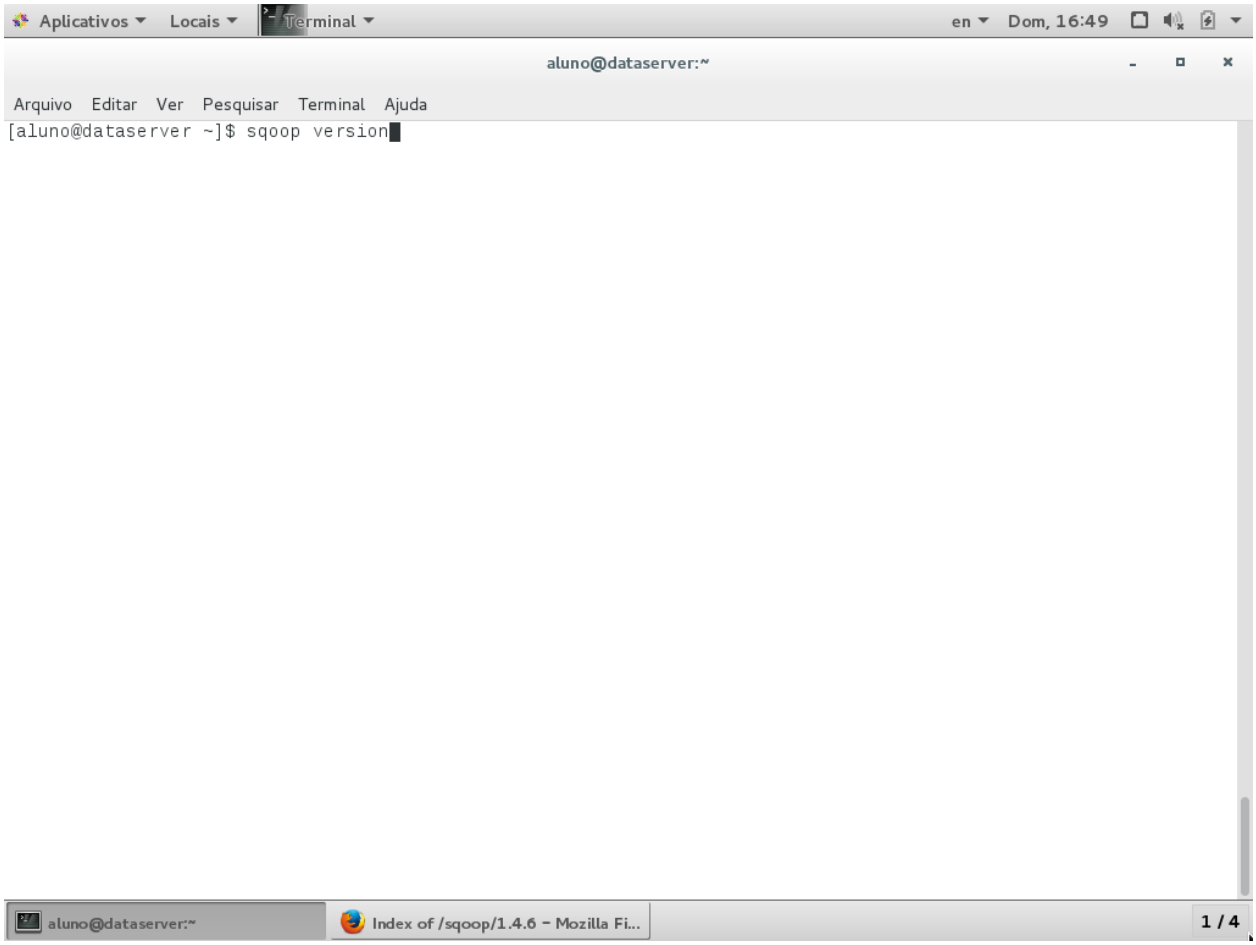

sqoop version

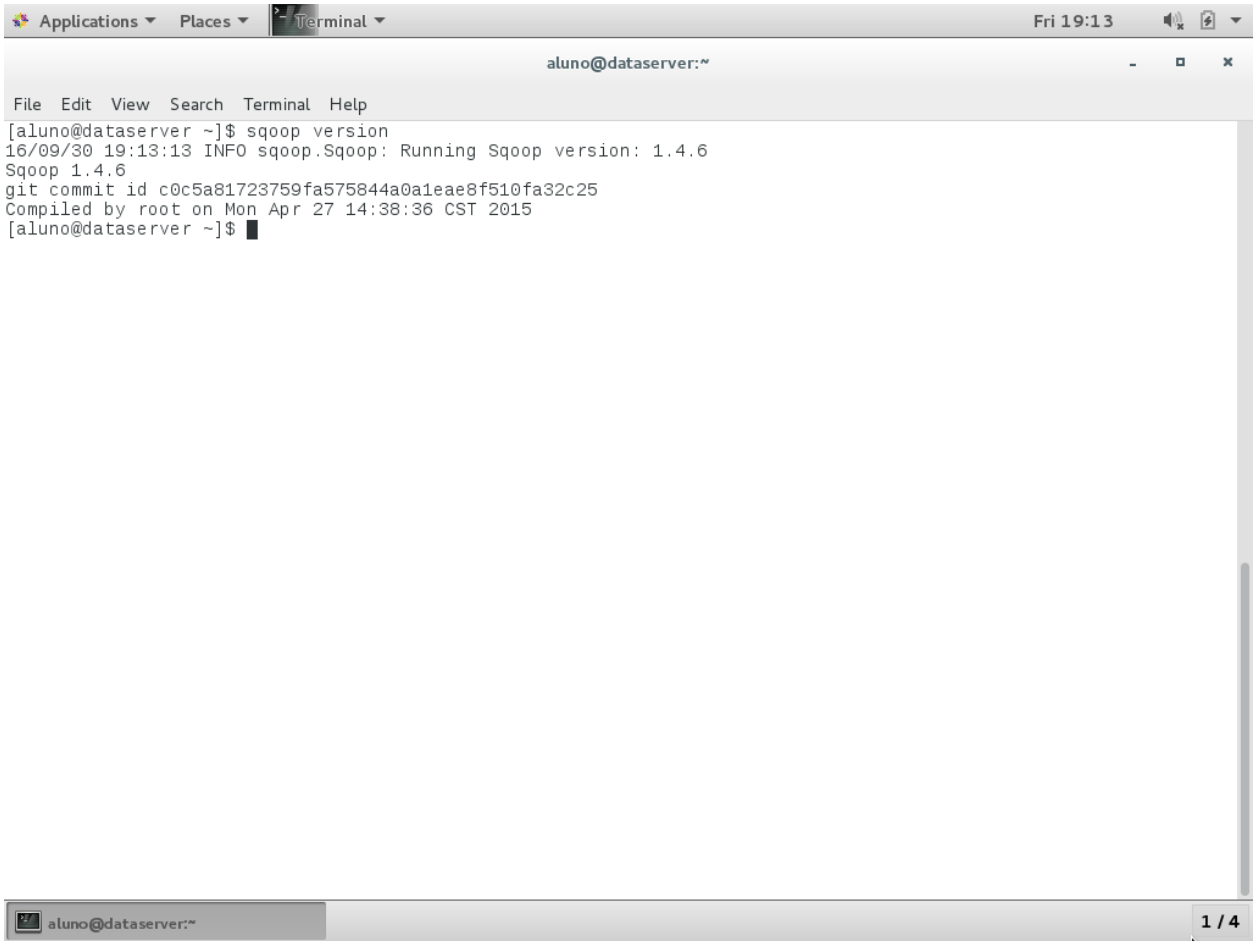

Sqoop version

12. Instalação e Configuração do Apache Flume

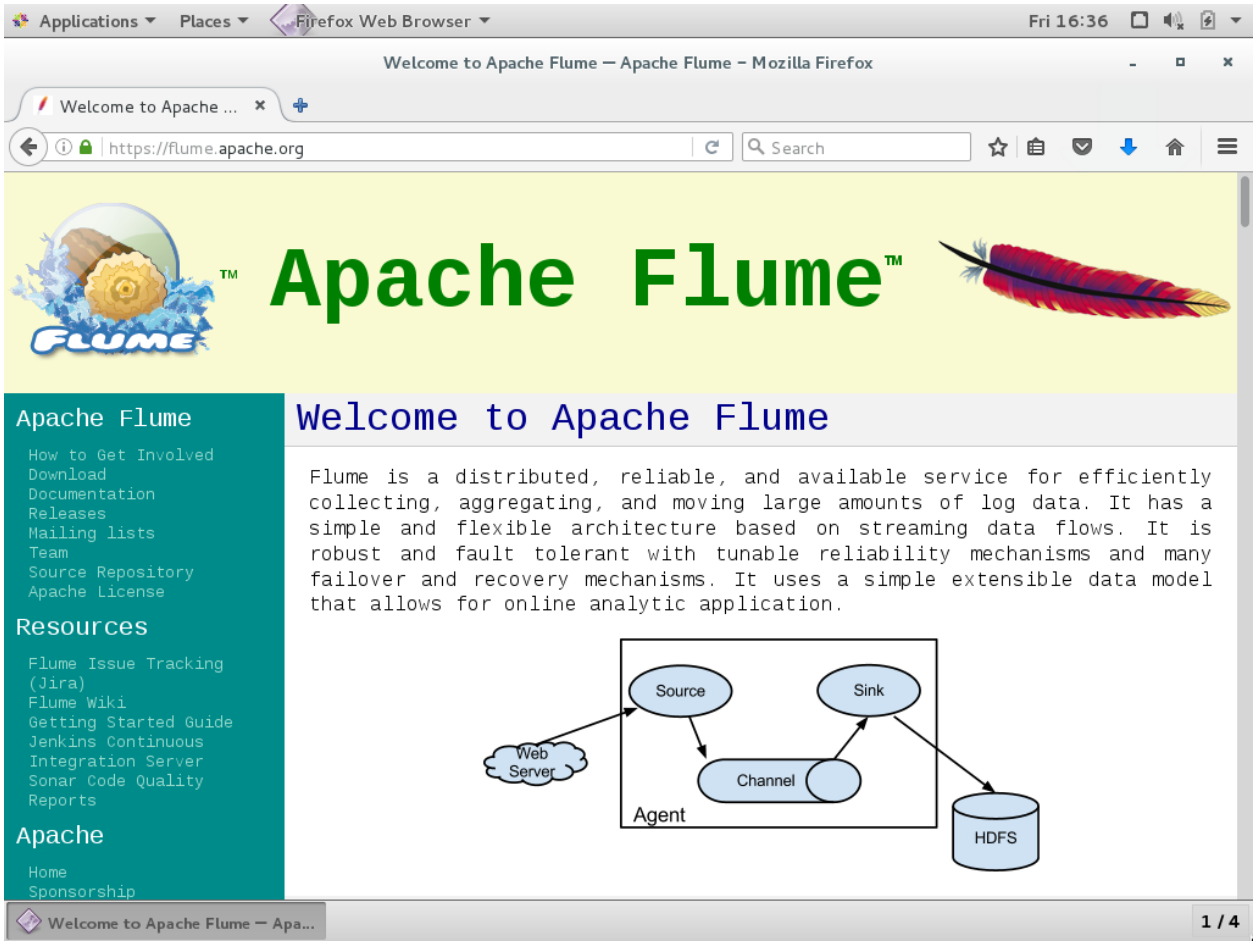

Download do Apache Flume – Versão 1.9

Faça o download, descompacte o arquivo e mova o diretório para /opt/flume da mesma forma como você fez com o Java JDK e com o Hadoop.

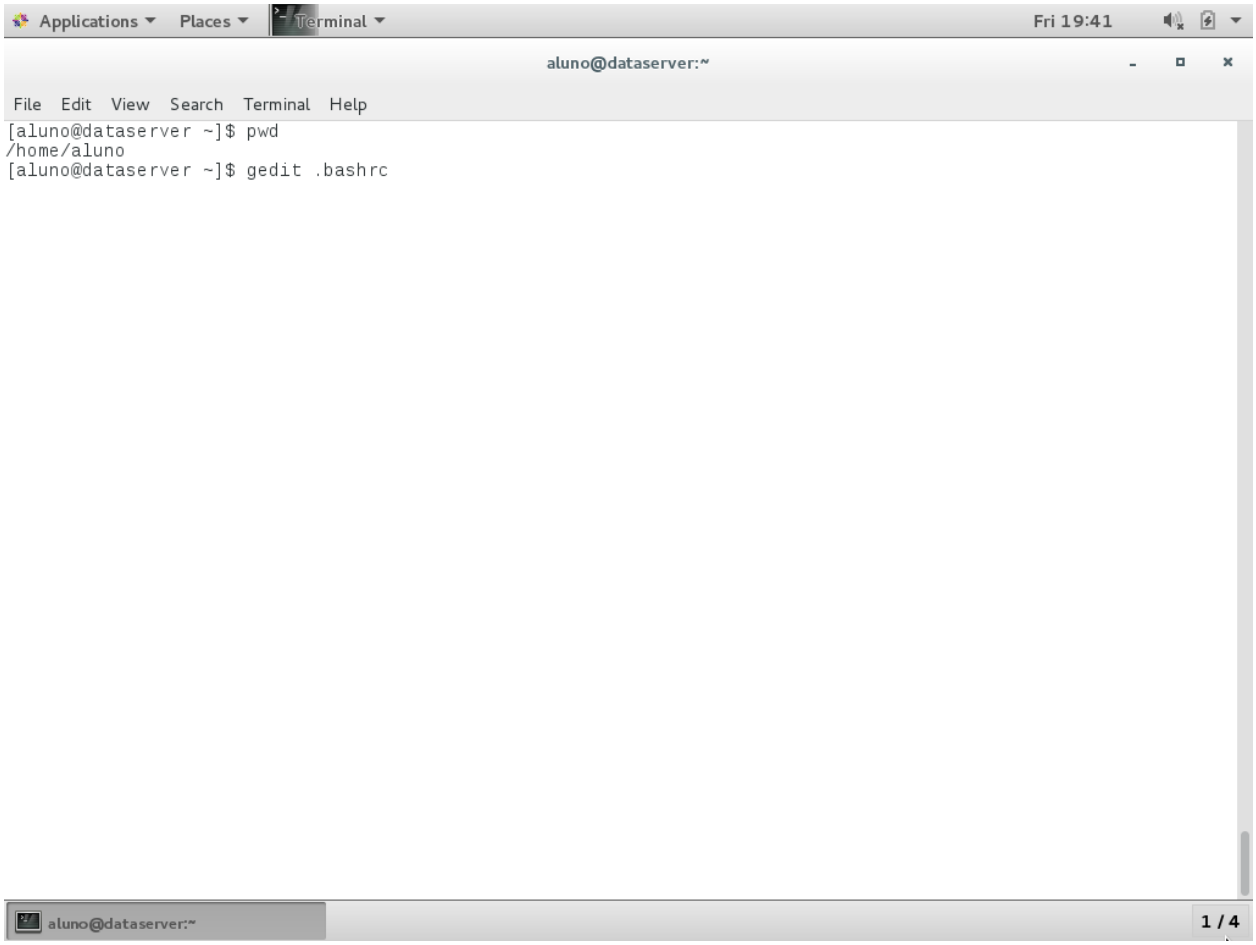

Editar as variáveis de ambiente

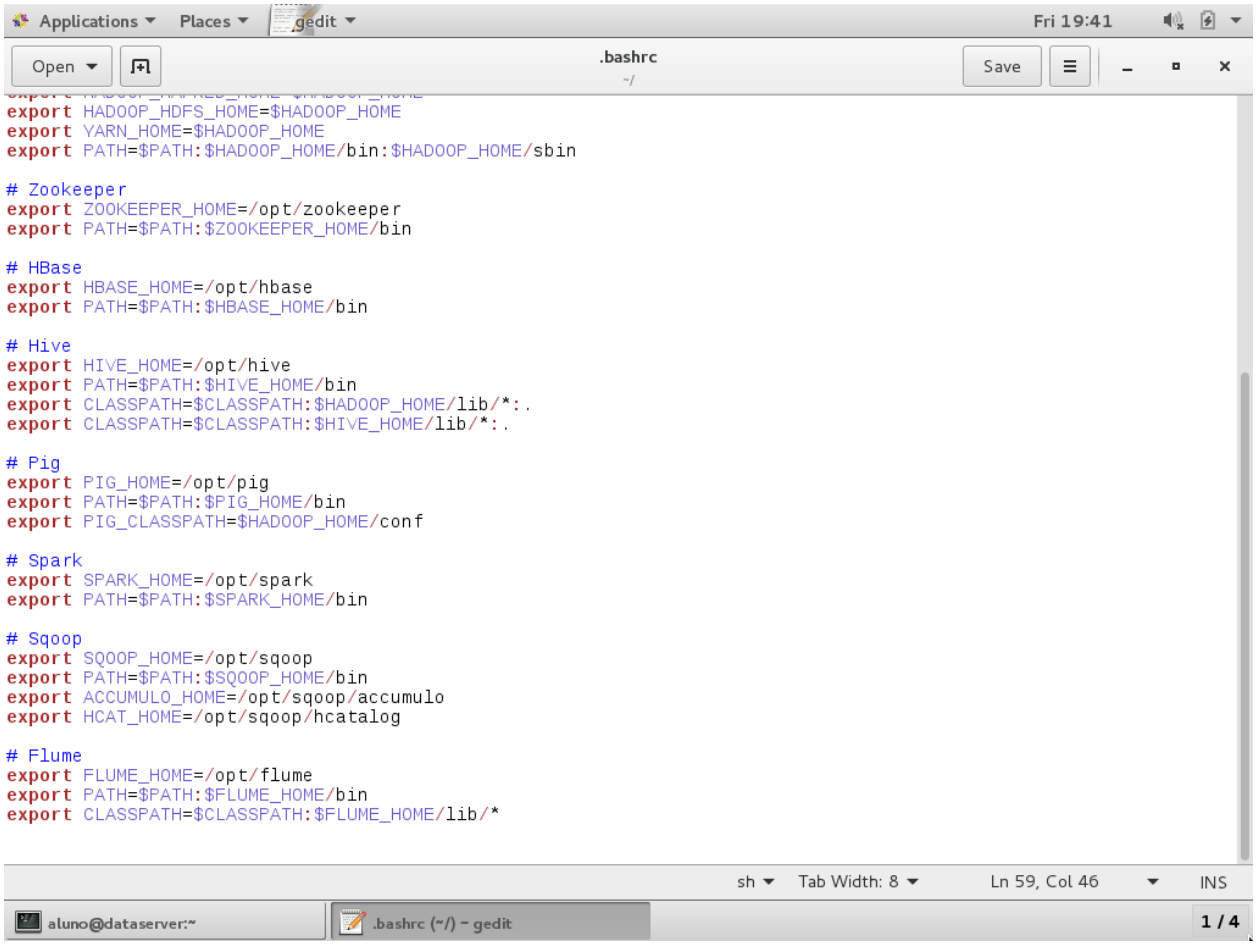

Variáveis de ambiente para o Flume

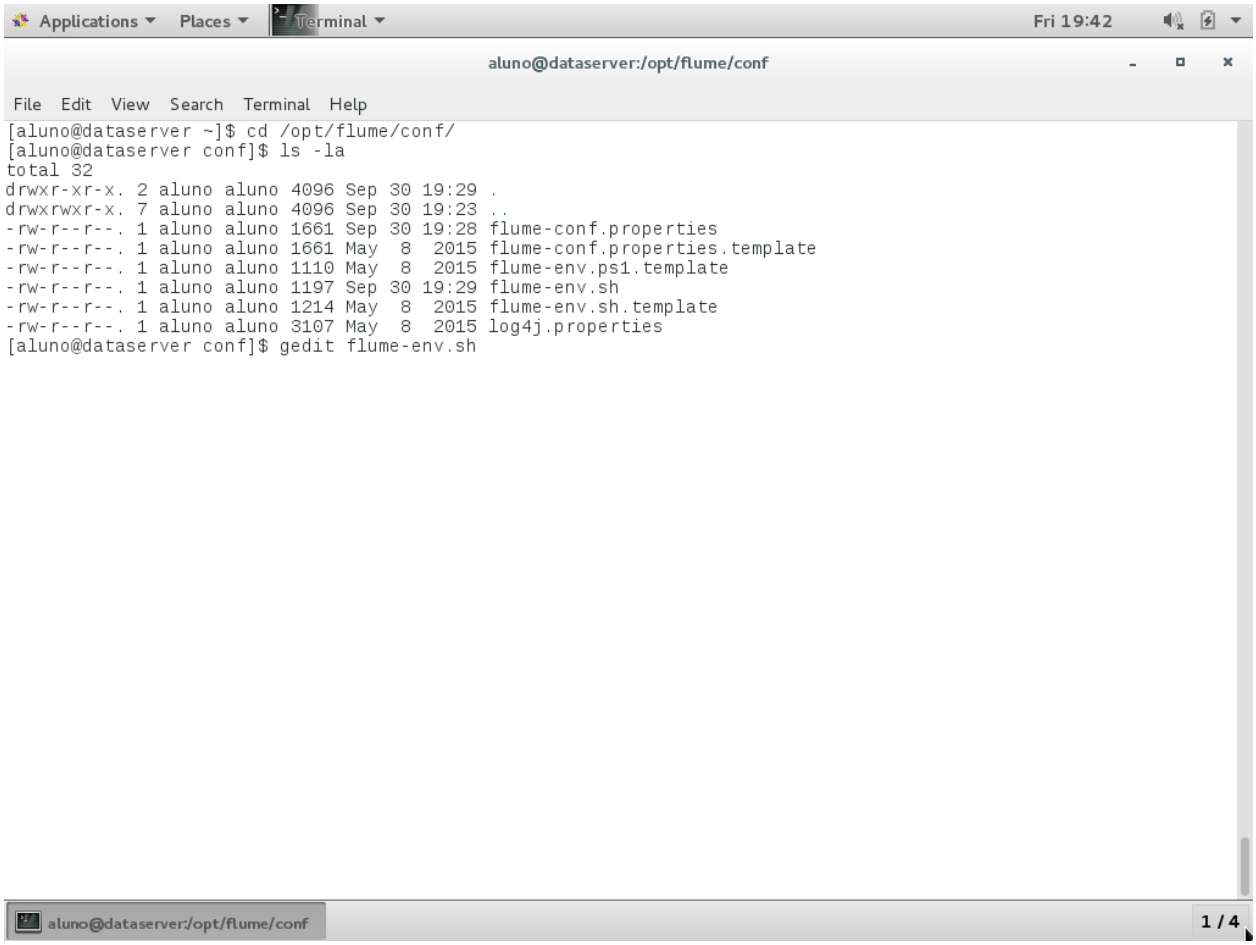

Editar o arquivo flume-env.sh

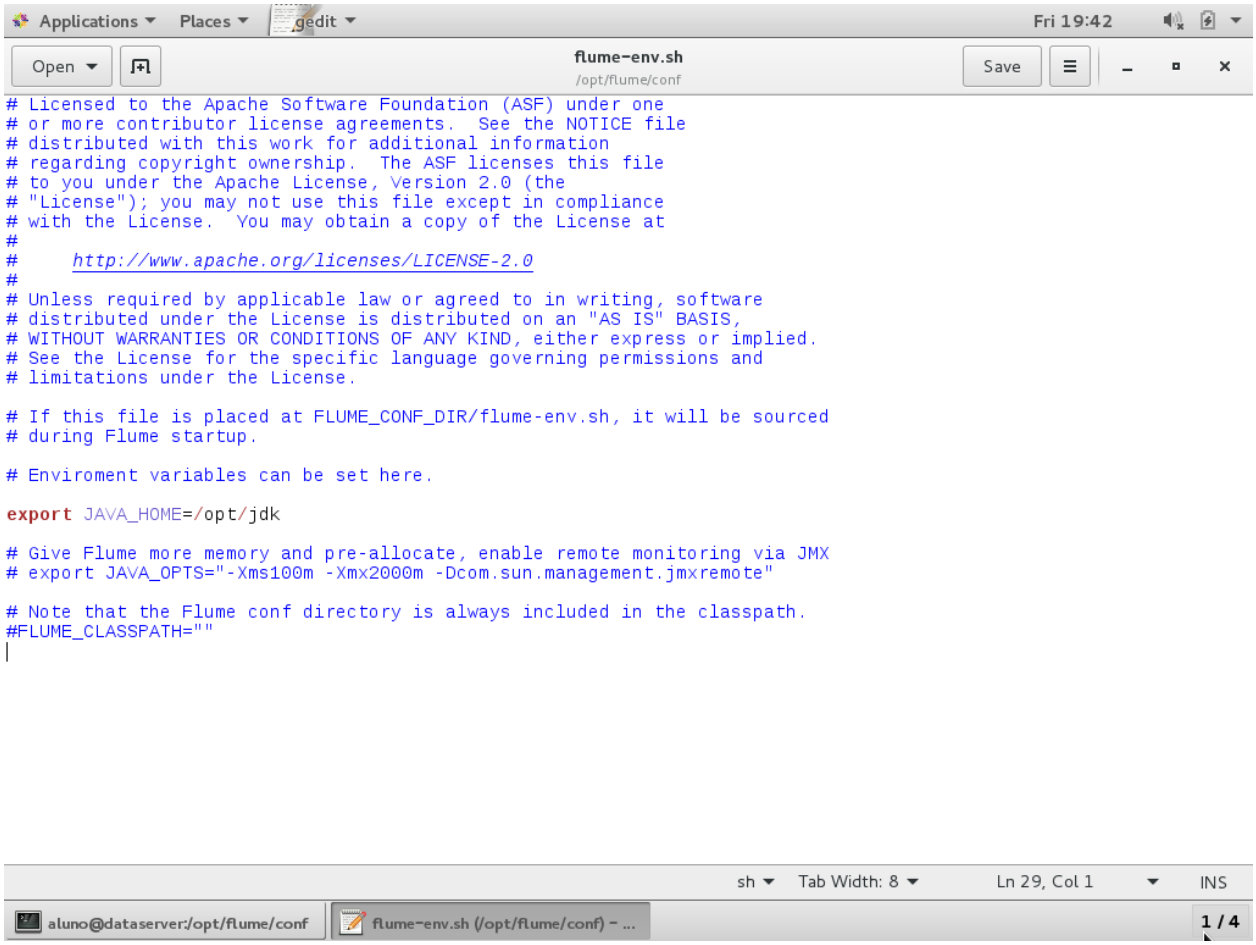

Acrescentar o JAVA_HOME

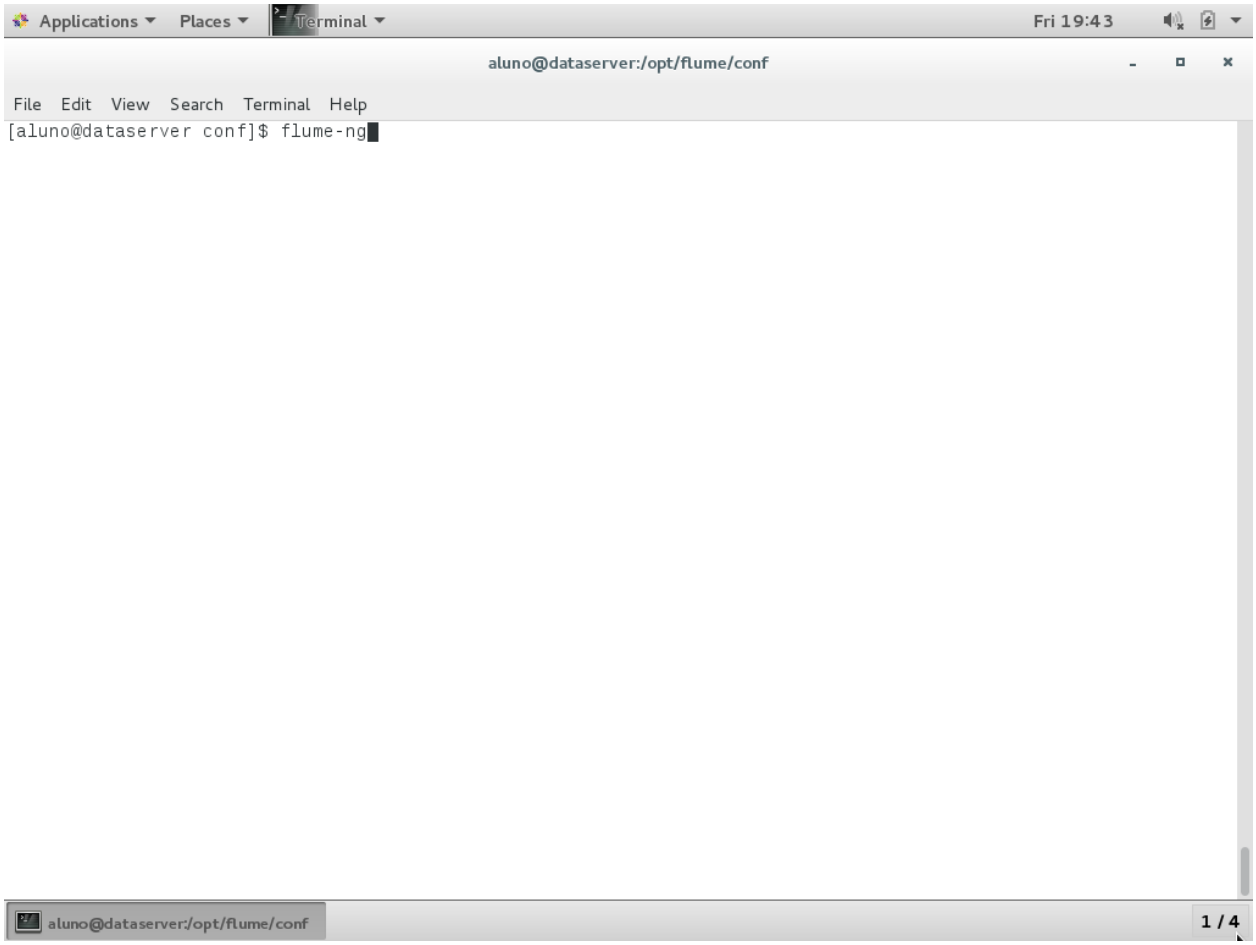

Testar a instalação

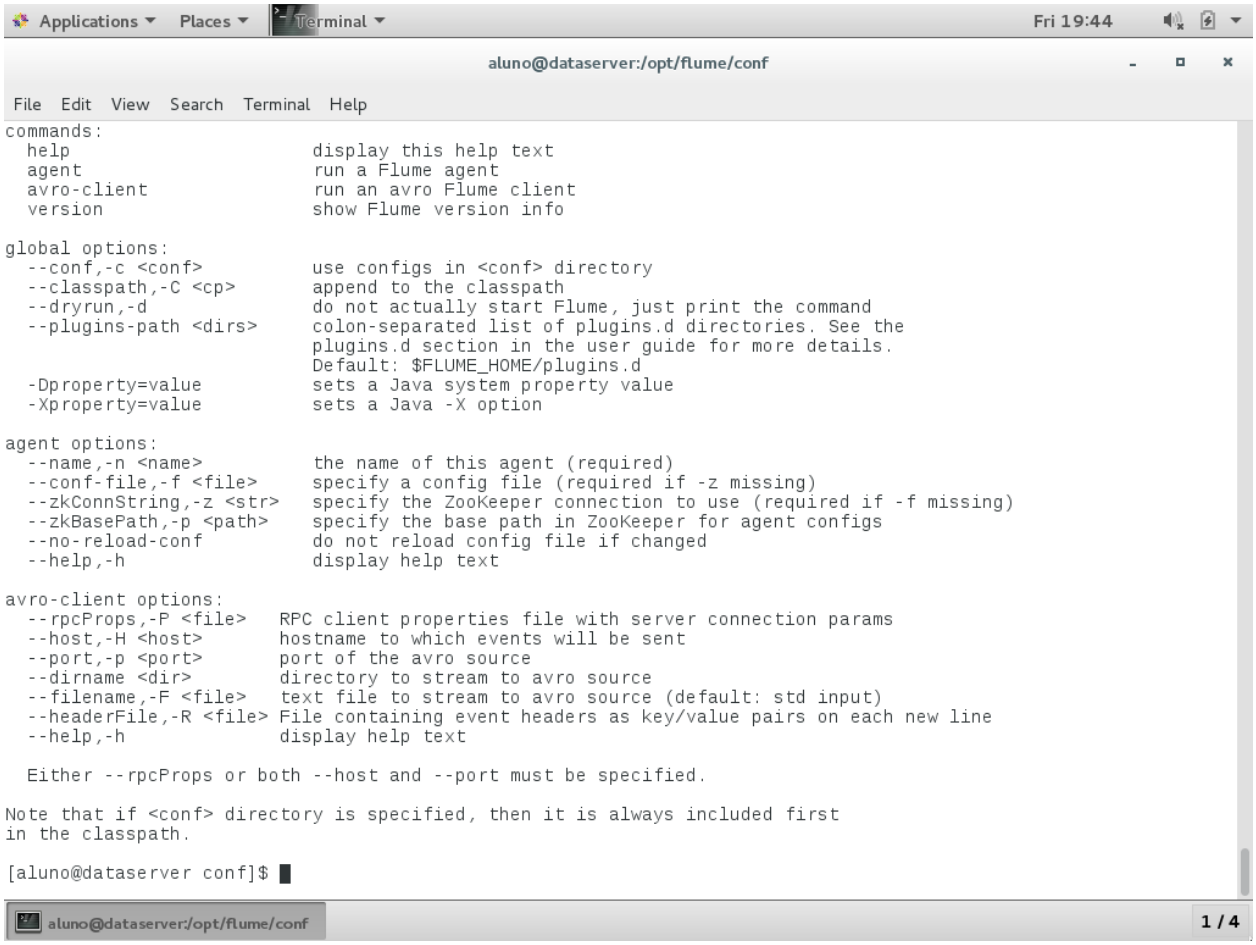

Flume instalado com sucesso

13. Instalação e Configuração do Ambari (Opcional)

Nota: No CentOS, o Ambari pode ser instalado mais facilmente através do gerenciador de pacotes yum.

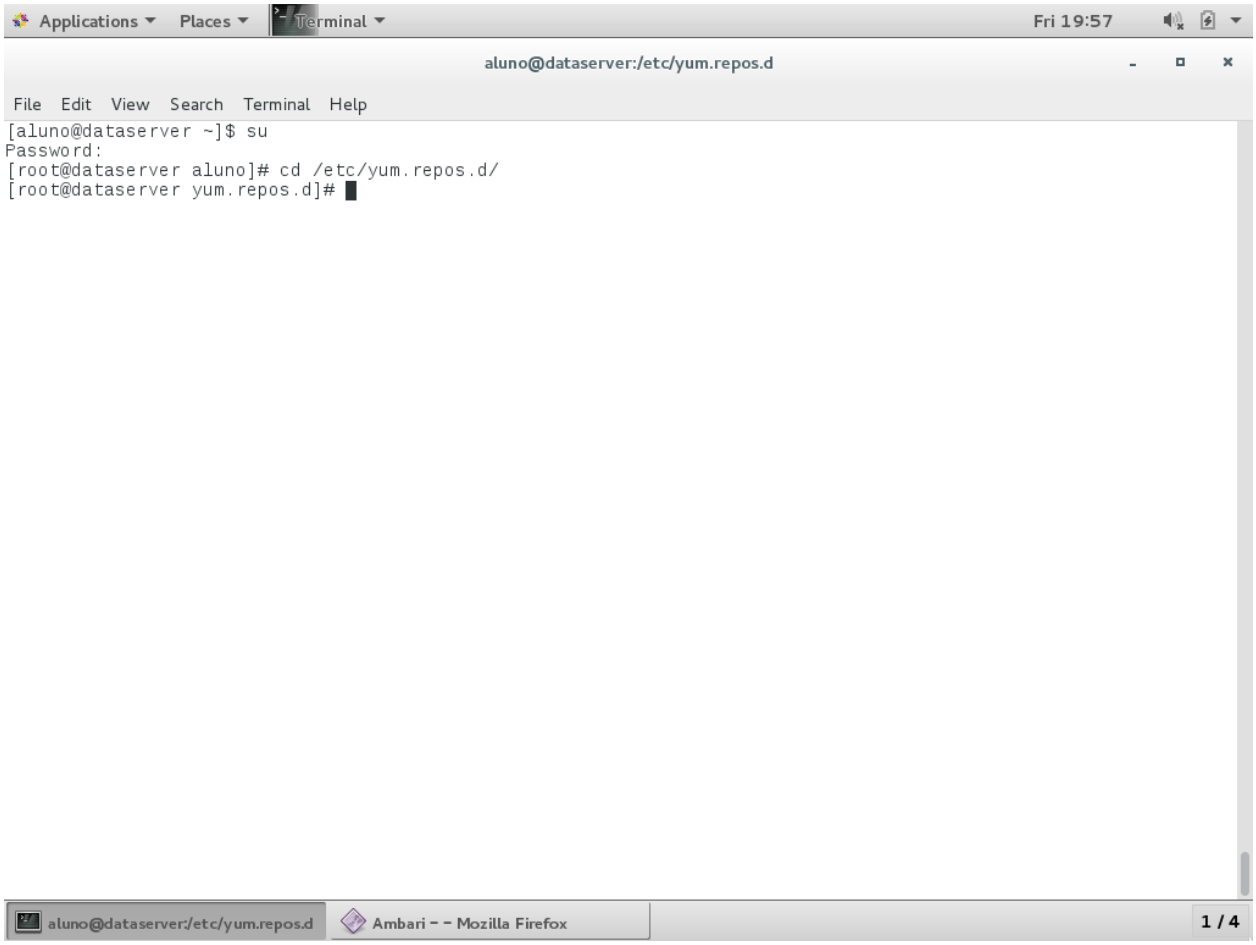

Conectado como root, acessar o diretório de repositórios do CentOS

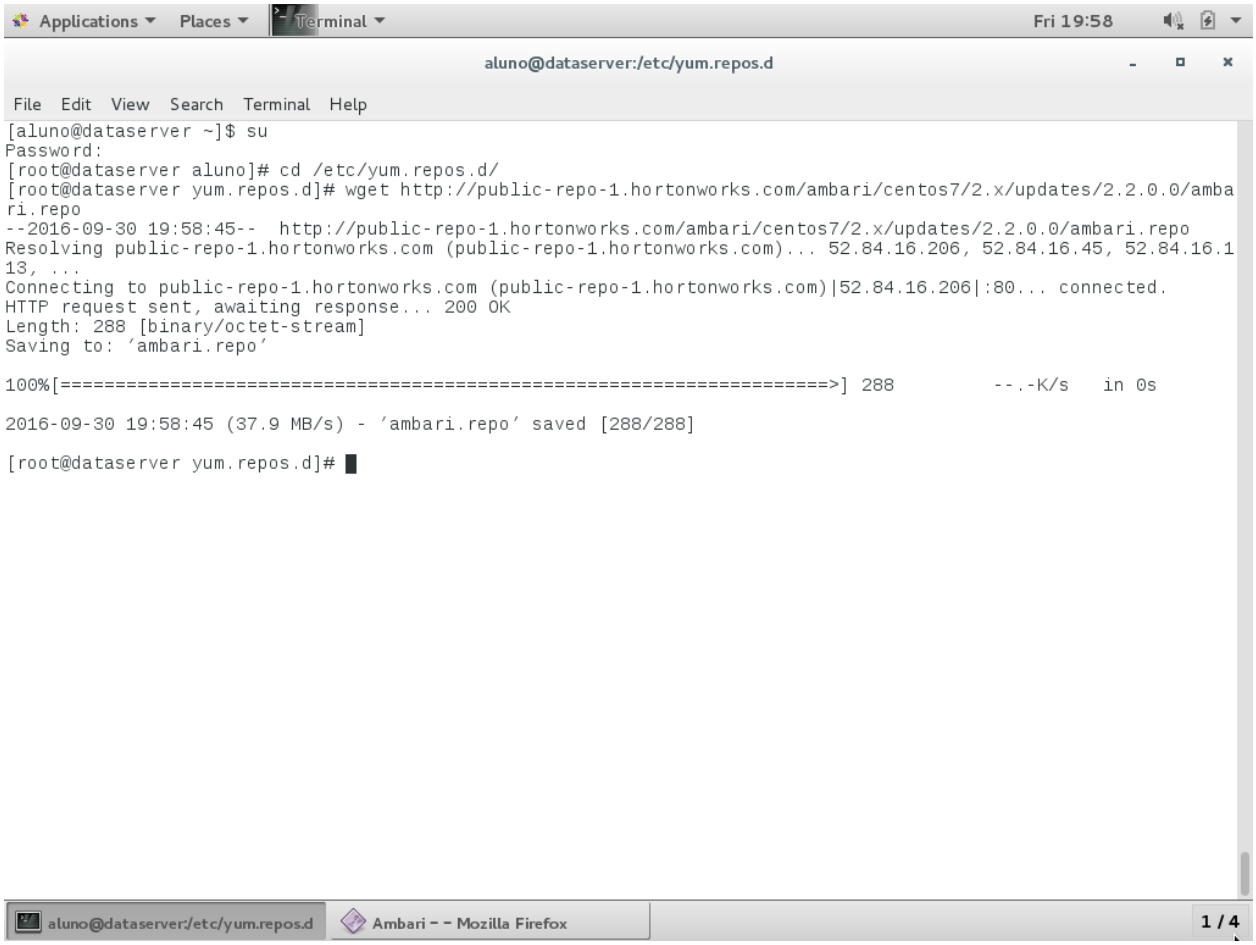

Download do arquivo de repositório do Ambari

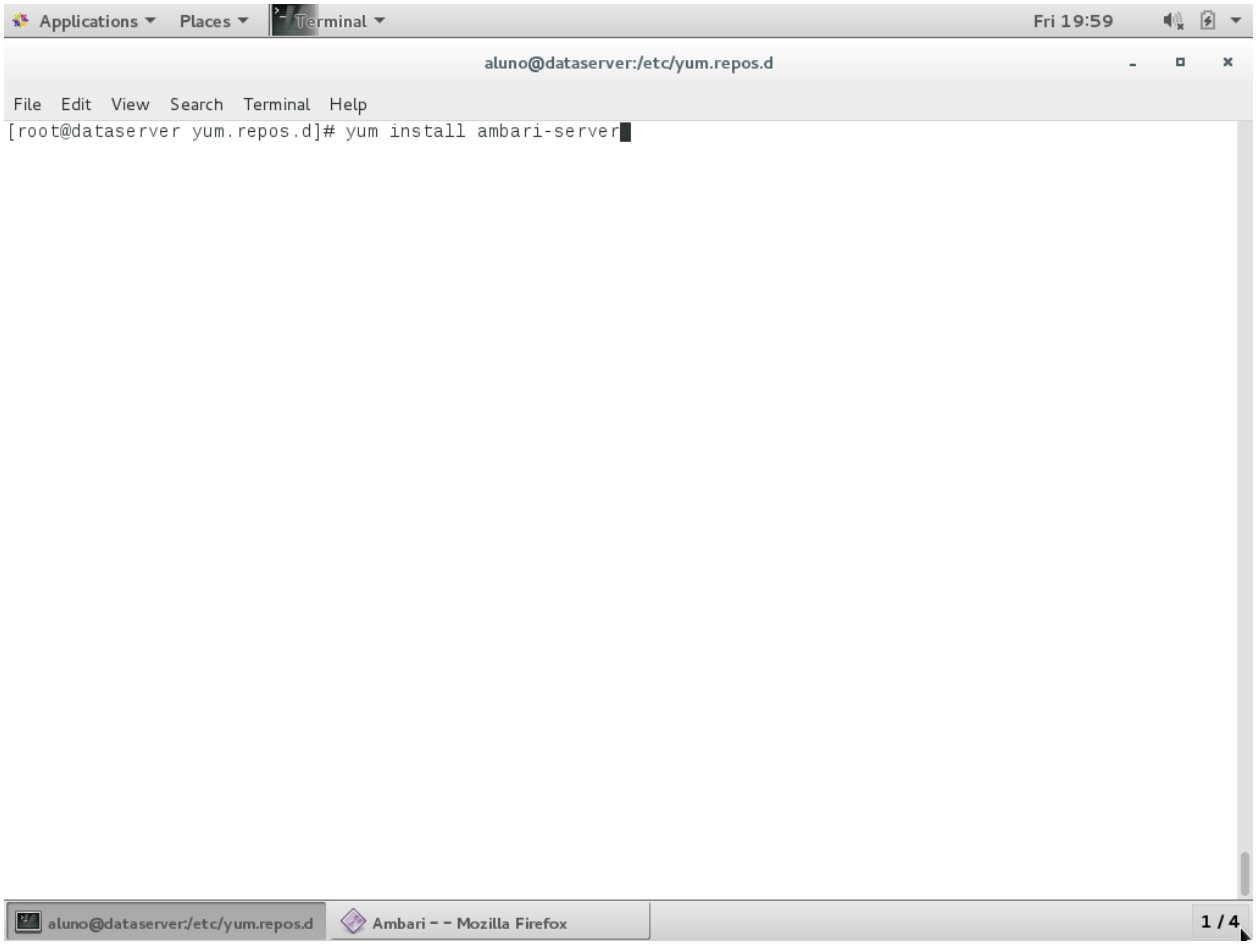

Como root, executar: yum install ambari-server

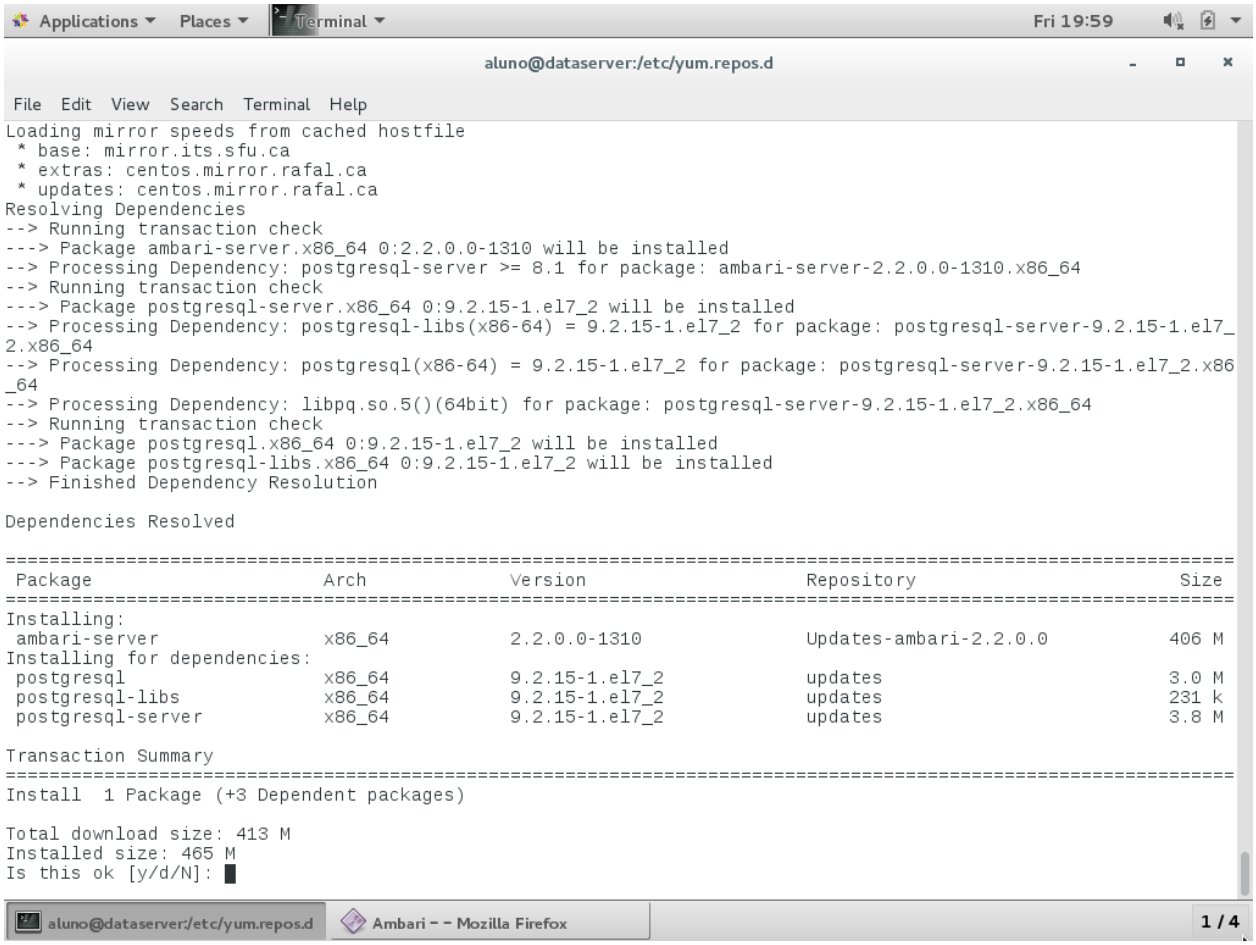

Instalação do Ambari

Instalação concluída com sucesso

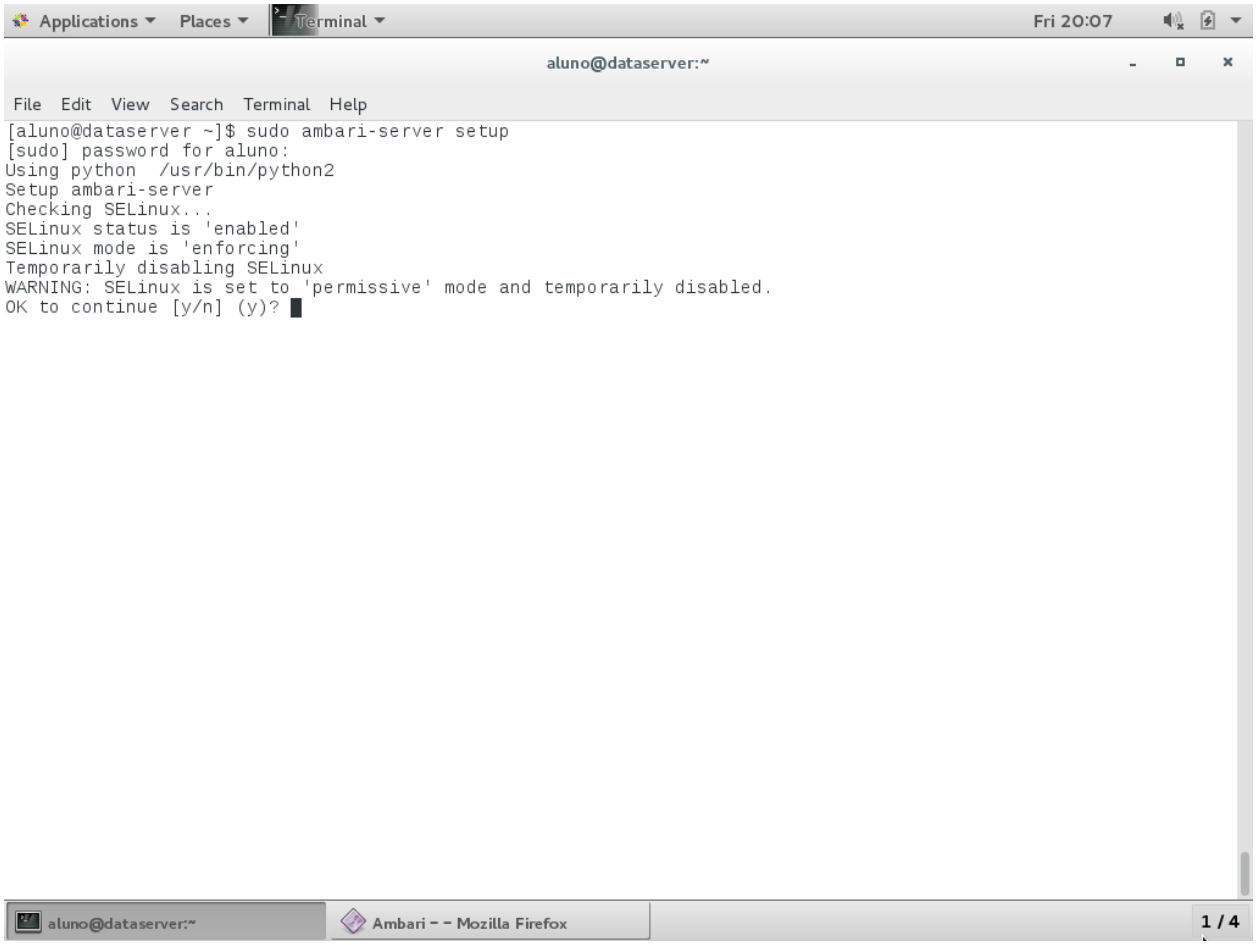

Configuração do Ambari

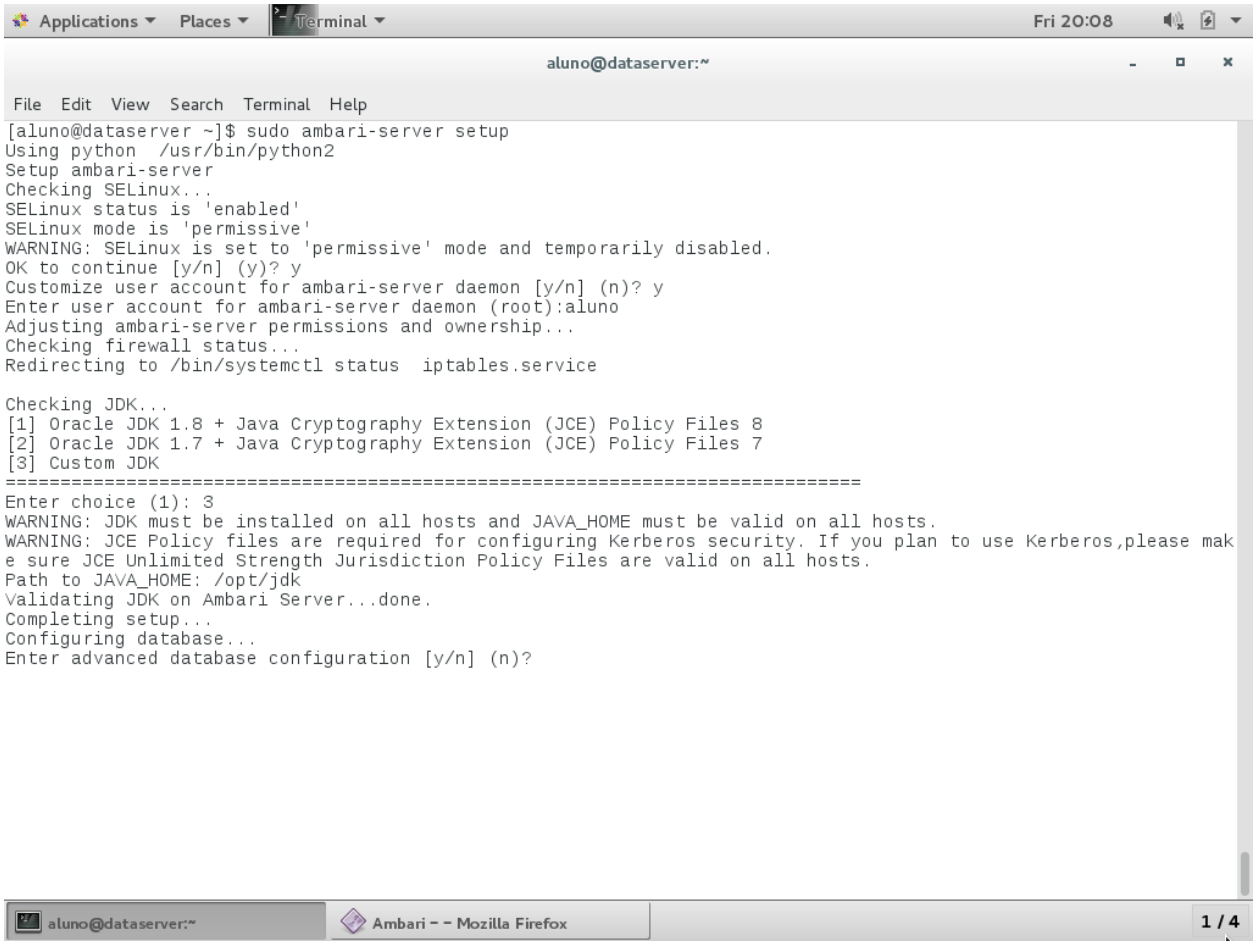

Configuração do Ambari

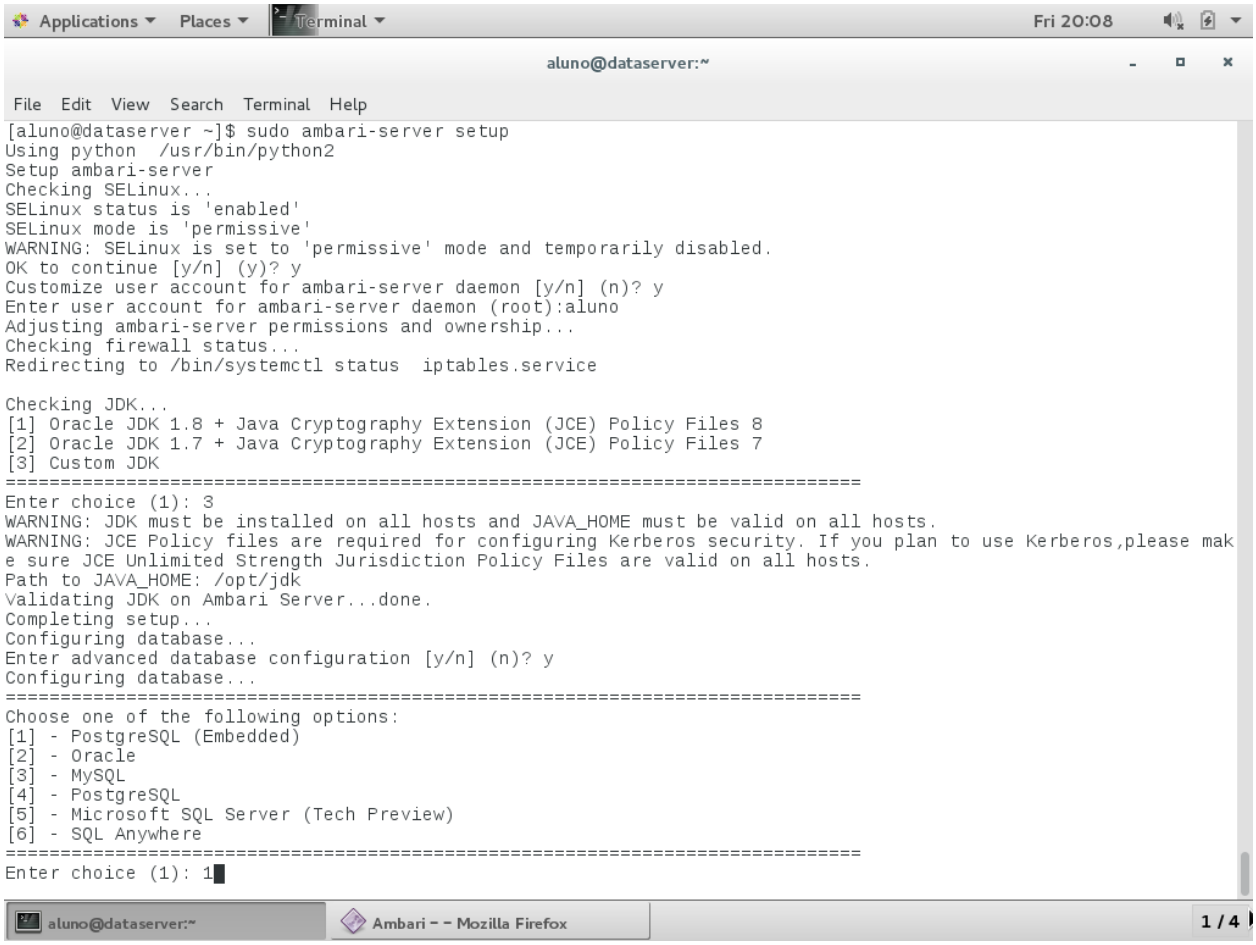

Configuração do Ambari

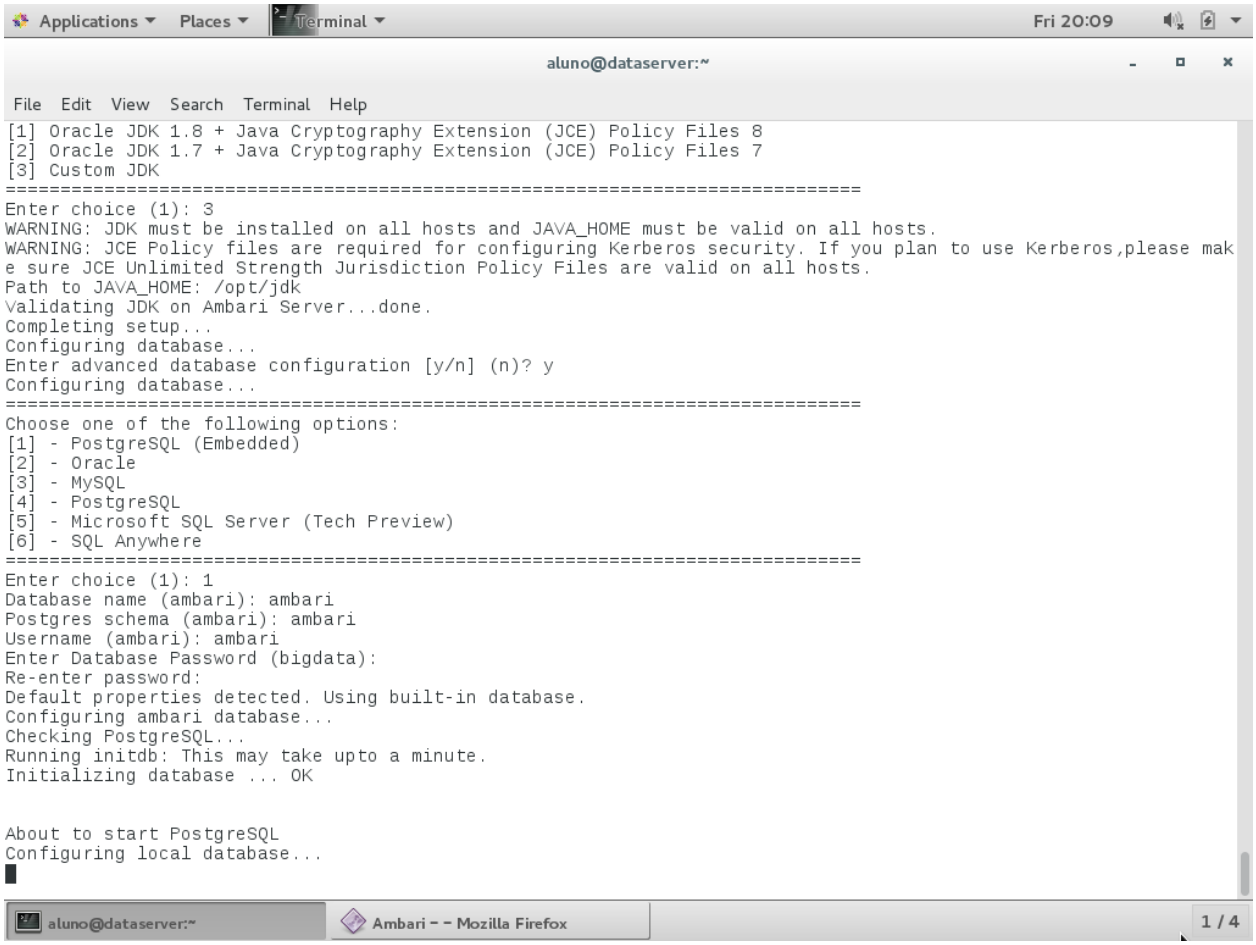

Configuração em andamento

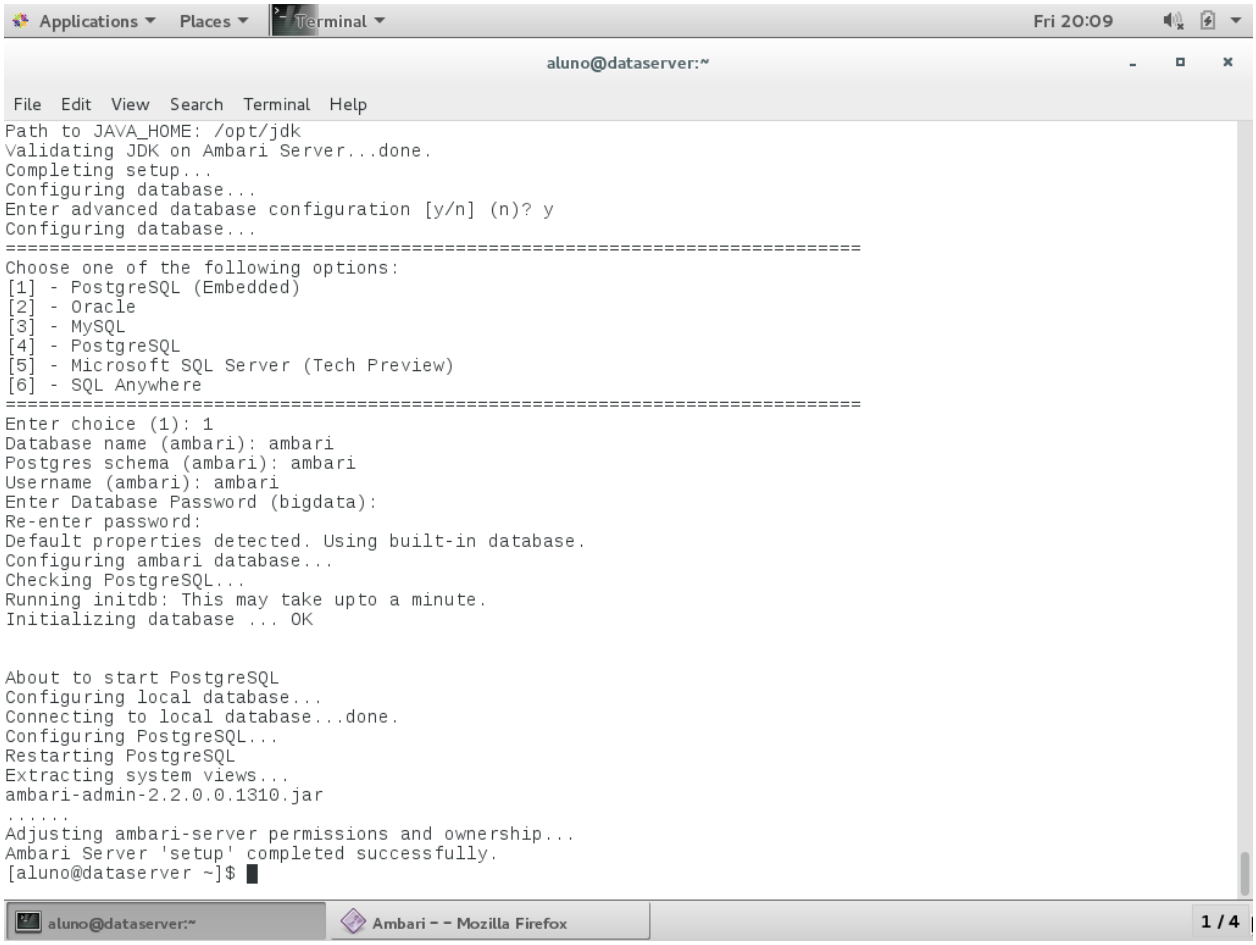

Configuração concluída

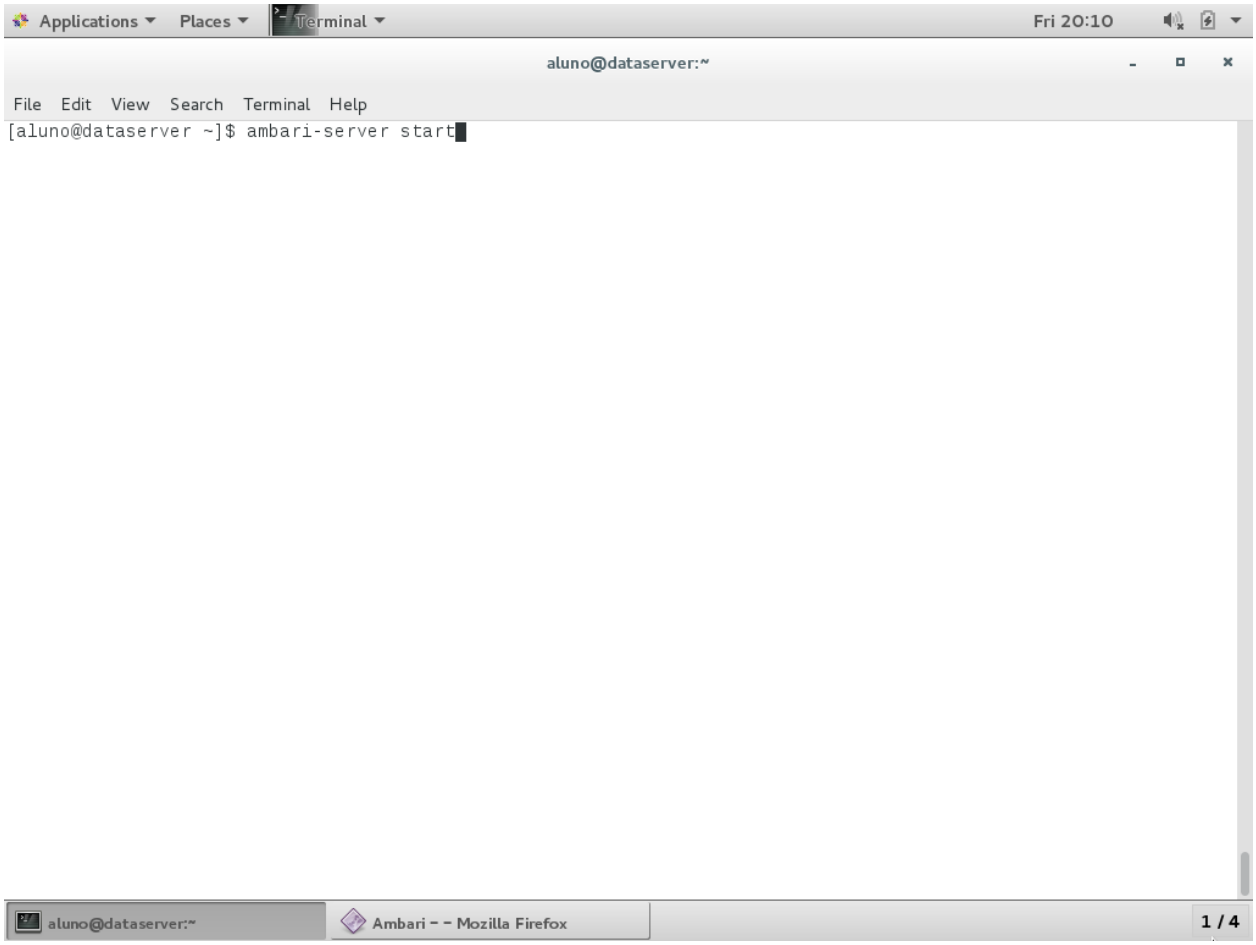

Inicializar o Ambari

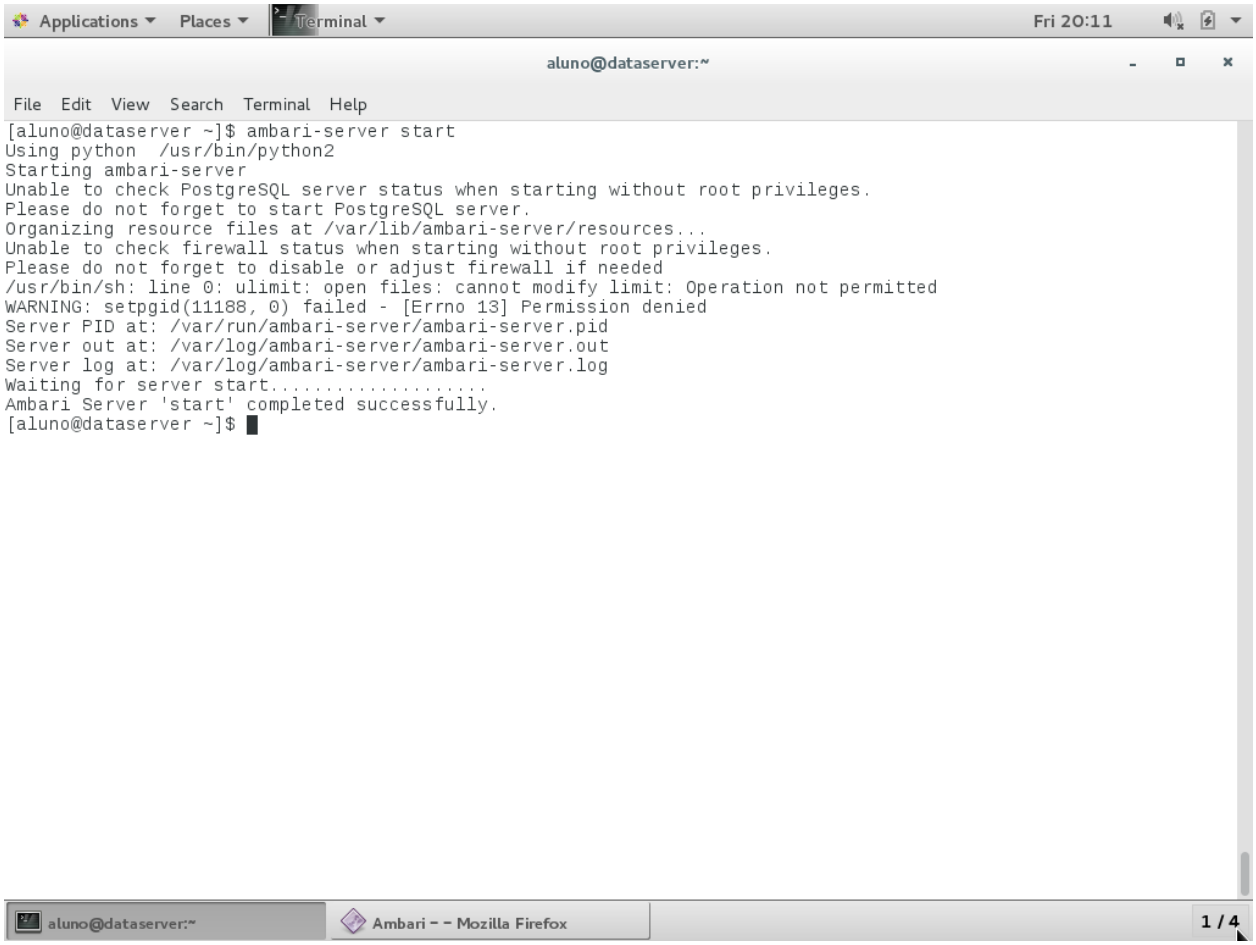

Inicializado

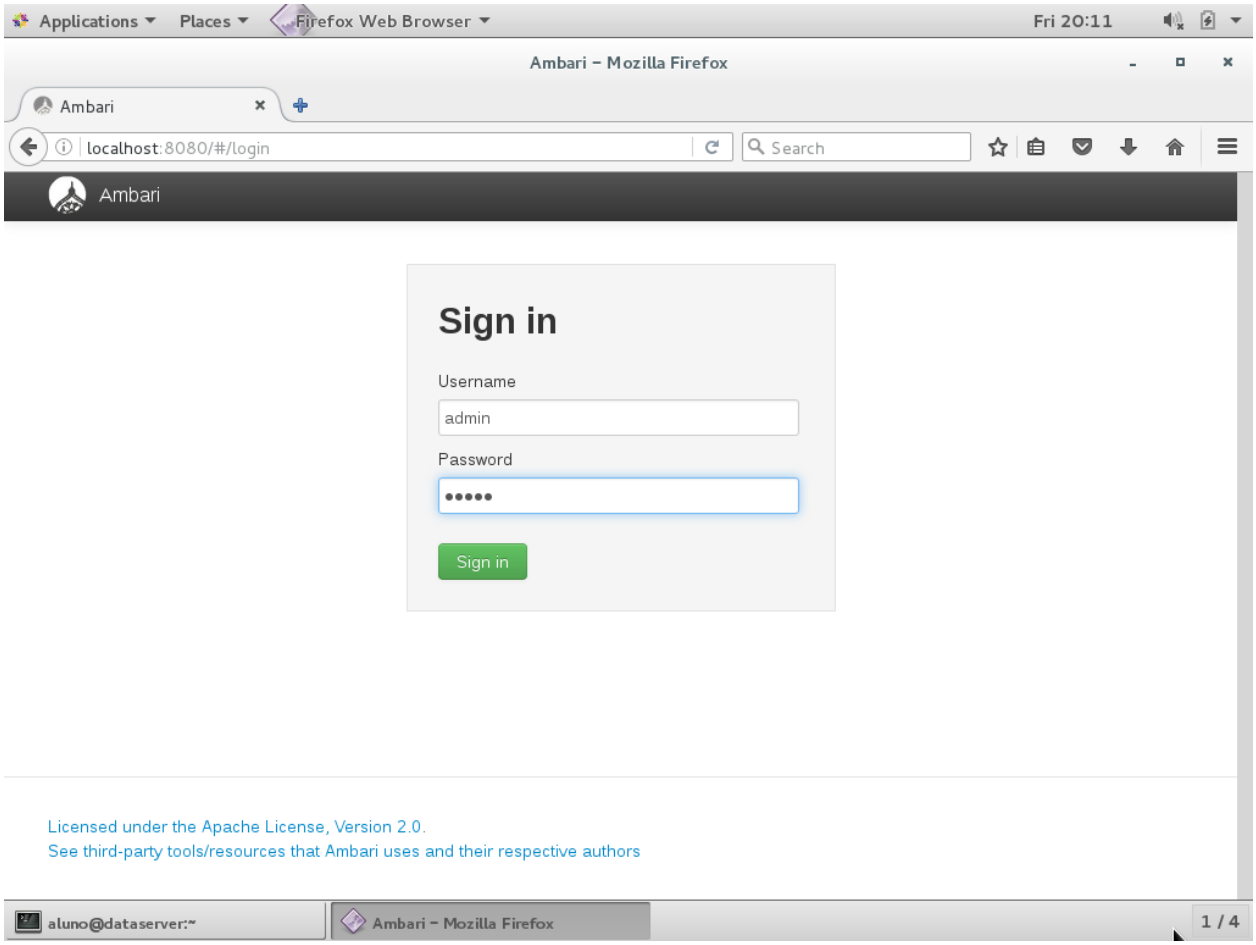

Acessar o browser – http://dataserver:8080 - usuário: admin / senha: admin

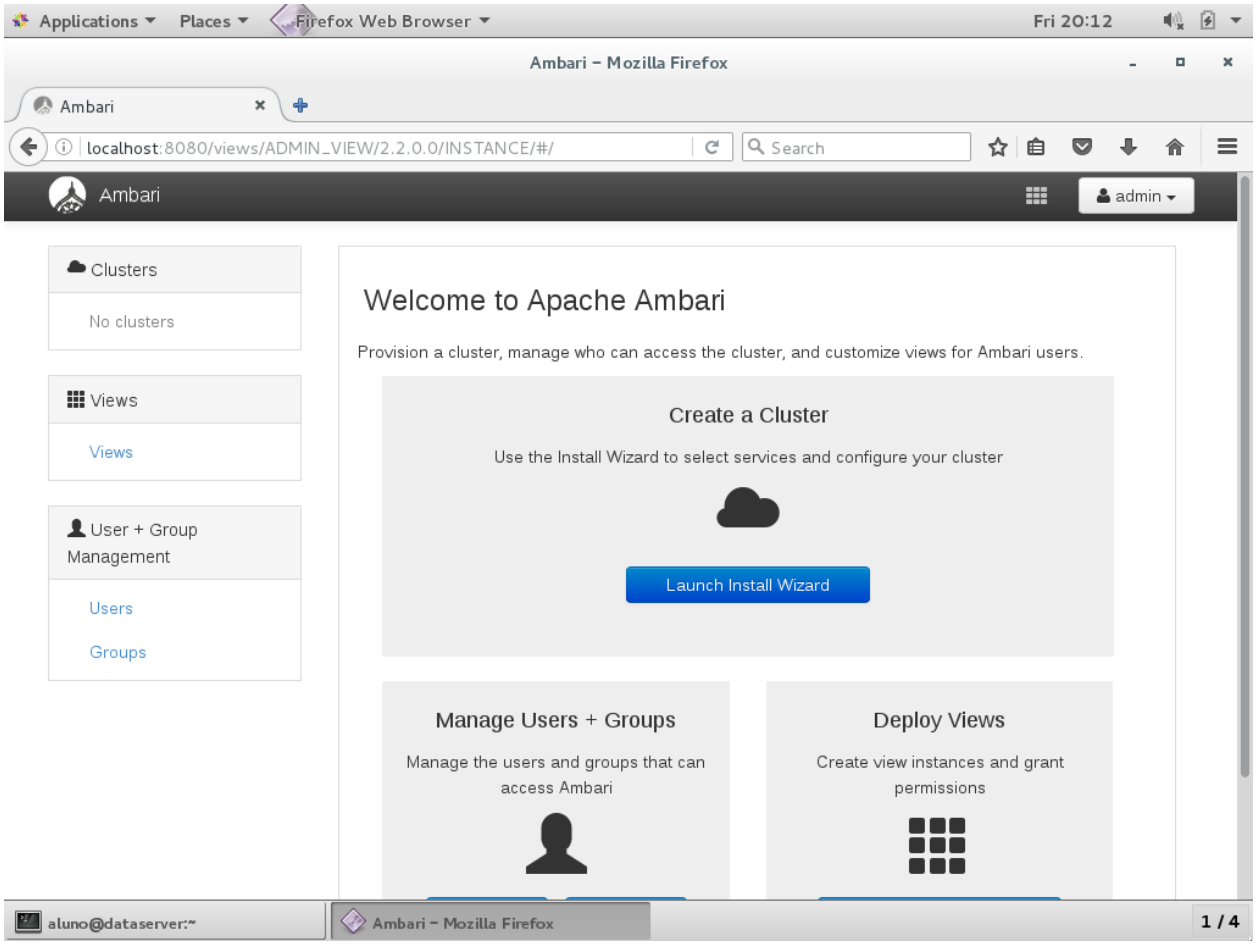

Pronto para configuração do cluster

Quarto checkpoint:

Clique no meu File – Export Appliance.
Será gerada uma cópia de segurança da sua máquina virtual.

→ VM: DataServer-vFinal.ova (Completa)

Parabéns!

Você tem um ambiente de testes para armazenar e processar Big Data!